

COMP0173-CW2-NOTEBOOK-Data-P2_2

December 10, 2025

1 COMP0173: Coursework 2

The paper HEARTS: A Holistic Framework for Explainable, Sustainable, and Robust Text Stereotype Detection by Theo King, Zekun Wu et al. (2024) presents a comprehensive approach to analysing and detecting stereotypes in text [1]. The authors introduce the HEARTS framework, which integrates model explainability, carbon-efficient training, and accurate evaluation across multiple bias-sensitive datasets. By using transformer-based models such as ALBERT-V2, BERT, and DistilBERT, this research project demonstrates that stereotype detection performance varies significantly across dataset sources, underlining the need for diverse evaluation benchmarks. The paper provides publicly available datasets and code [2], allowing full reproducibility and offering a standardised methodology for future research on bias and stereotype detection in Natural Language Processing (NLP).

2 Instructions

All figures produced during this notebook are stored in the project's `/COMP0173_Figures` directory. The corresponding LaTeX-formatted performance comparison tables, including ALBERT-V2, BERT, and DistilBERT are stored in `/COMP0173_PDF`, with the compiled document available as `COMP0173-CW2-TABLES.pdf`.

3 Technical Implementation (70%)

```
[1]: # %%capture
# pip install -r requirements.txt
# pip install transformers
# pip install --upgrade transformers
# pip install --upgrade tokenizers
# pip install -U sentence-transformers
# pip install natasha
# pip install datasets
# pip install --user -U nltk
# conda install -c anaconda nltk
# pip install --upgrade openai pandas tqdm
# pip install dotenv
```

```
[2]: # pip install -U pip setuptools wheel
# pip install -U spacy
# python -m spacy download en_core_web_trf
# python -m spacy download en_core_web_sm
# python -m spacy download ru_core_news_lg

# # GPU
# pip install -U 'spacy[cuda12x]'
# # GPU - Train Models
# pip install -U 'spacy[cuda12x,transformers,lookups]'
```

```
[3]: # Import the libraries
import random, numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
sns.set(color_codes=True)
plt.style.use('seaborn-v0_8')

# To ignore warnings
import warnings
warnings.filterwarnings('ignore')
np.random.seed(23)

warnings.filterwarnings(
    "ignore",
    message="pkg_resources is deprecated as an API"
)
```

```
[4]: # Import libraries
import pandas as pd
import os
import sys
import importlib.util, pathlib
from pathlib import Path
import warnings
from importlib import reload
from importlib.machinery import SourceFileLoader
from IPython.display import display
import pandas as pd
from pathlib import Path
import re
import difflib
import string
from collections import defaultdict
import json
```

```
[5]: import torch
import transformers
from transformers import AutoModelForMaskedLM, XLMWithLMHeadModel
from transformers import AutoTokenizer, AutoConfig
from transformers import TrainingArguments, Trainer
from sentence_transformers import SentenceTransformer, util
import platform
from datasets import load_dataset
import spacy
import requests
from tqdm import tqdm
import yaml

import os
os.environ["TOKENIZERS_PARALLELISM"] = "false"

import natasha
from natasha import (
    Segmenter,
    MorphVocab,
    NewsEmbedding,
    NewsMorphTagger,
    NewsSyntaxParser,
    NewsNERTagger,
    PER,
    NamesExtractor,
    Doc
)
segmenter = Segmenter()
morph_vocab = MorphVocab()

segmenter = Segmenter()
morph_vocab = MorphVocab()

emb = NewsEmbedding()
morph_tagger = NewsMorphTagger(emb)
```

```
[6]: # # Check the GPU host (UCL access)
# print("CUDA available:", torch.cuda.is_available())
# print("Device:", torch.cuda.get_device_name(0))

# # Path
# import os
```

```
# os.chdir("/tmp/HEARTS-Text-Stereotype-Detection")
# os.getcwd()
```

3.1 Part 2: Identify a contextually relevant challenge in your country or region of your choice that can be addressed using the same AI approach

Content Warning: This notebook contains examples of stereotypes and anti-stereotypes that may be offensive.

3.1.1 *Question 1 : Problem and SDG alignment*

This coursework supports Sustainable Development Goal (SDG) 5: Gender Equality - *Achieve gender equality and empower all women and girls*, SDG 9: Industry, Innovation, and Infrastructure - *Build resilient infrastructure, promote inclusive and sustainable industrialization and foster innovation*, SDG 10: Reduced Inequalities - *Reduce inequality within and among countries*, and SDG 16: Peace, Justice, and Strong Institutions: - *Promote peaceful and inclusive societies for sustainable development, provide access to justice for all and build effective, accountable and inclusive institutions at all levels* [5].

The specific targets covered by this coursework are:

- SDG 5.1: *End all forms of discrimination against all women and girls everywhere*
- SDG 5.b: *Enhance the use of enabling technology, in particular information and communications technology, to promote the empowerment of women*
- SDG 10.2: *By 2030, empower and promote the social, economic and political inclusion of all, irrespective of age, sex, disability, race, ethnicity, origin, religion or economic or other status*
- SDG 10.3: *Ensure equal opportunity and reduce inequalities of outcome, including by eliminating discriminatory laws, policies and practices and promoting appropriate legislation, policies and action in this regard*
- SDG 16.1: *Significantly reduce all forms of violence and related death rates everywhere*
- SDG 16.6: *Develop effective, accountable and transparent institutions at all levels*
- SDG 16.10: *Ensure public access to information and protect fundamental freedoms, in accordance with national legislation and international agreements*
- SDG 16.b: *Promote and enforce non-discriminatory laws and policies for sustainable development*

3.1.2 *Question 2 : Limitations and ethical considerations*

3.1.3 *Question 3 : Scalability and sustainability analysis*

3.2 Part 3: Curate or identify an alternative dataset appropriate for your context

3.2.1 *Question 1 : Identify contextually appropriate dataset*

1. RuBias
2. Kaggle

3. RuHateBe

3.2.2 Question 2 : Document data collection/access process and ethical considerations

Mention where you got these datasets - provide refs and what should be cleaned from these datasets

```
[7]: # Load dataset in its raw format
# RuBias
rubias = pd.read_csv("COMP0173_Data/rubias.tsv", sep="\t", encoding="utf-8")
```

```
[8]: # Rename column
rubias = rubias.rename(columns={"domain": "stereotype_type"})

# Change the level name
rubias["stereotype_type"] = rubias["stereotype_type"].replace("class", "profession")
```

```
[9]: # Load dataset in its raw format
# RuSter
ruster = pd.read_json("COMP0173_Stereotypes/stereotypes.json")
```

```
[10]: # Save
ruster.to_csv("COMP0173_Data/ruster.csv", index=False)
```

Helper Functions

```
[11]: def pie_chart_domain(df, column, name):

    """
    Plot the percentage distribution of social-group domains as a styled pie chart.

    Parameters
    -----
    df : pandas.DataFrame
        Input dataframe containing a categorical column representing social domains.
    column : str, optional
        Name of the column in `df` holding domain labels.
    name : str, optional
        Name of the dataset.

    Returns
    -----
    None
        Displays a pie chart visualising the proportional distribution of categories.
    """

    # Your code here to generate the pie chart visualization.
```

Notes

The function applies a custom colour palette tailored for the RuBias dataset (gender, class, nationality, LGBTQ). Any unseen categories default to grey.

```
"""
```

```
# Compute relative frequency (%) of categories
domain_counts = df[column].value_counts(normalize=True) * 100
labels = domain_counts.index
sizes = domain_counts.values

# Predefined colour palette
color_map = {
    'gender': "#CA5353",
    'profession': "#F1A72F",
    'nationality': "#559A67",
    'lgbtq': "#527BCD",
}
# Assign colours; fallback to grey for unknown labels
colors = [color_map.get(lbl, 'grey') for lbl in labels]

# Create compact, high-resolution figure
plt.figure(figsize=(5.5, 4), dpi=155)

# Draw pie chart with formatted percentages
wedges, texts, autotexts = plt.pie(
    sizes,
    labels=None,
    autopct='%1.1f%%',
    pctdistance=0.55,
    startangle=90,
    colors=colors,
    wedgeprops={'linewidth': 2, 'edgecolor': 'white'}
)

# Style displayed percentage numbers
for t in autotexts:
    t.set_fontsize(10)
    t.set_color("black")

# Title
plt.title(f"Social Group Distribution: {name}", fontsize=16)

# Legend placed to the right of the figure
plt.legend()
```

```

        wedges,
        labels,
        title="Domain",
        loc="center left",
        bbox_to_anchor=(1.02, 0.5),
        fontsize=11,
        title_fontsize=12
    )

plt.tight_layout()
plt.show()

```

[12]: `def format_string(texts: pd.Series) -> pd.Series:`

```

    """
    Normalise Russian stereotype strings.

    Operations
    -----
    - lowercase
    - remove punctuation (except comma, hyphen, underscore)
    - replace '-' and '--' with spaces
    - collapse multiple spaces
    - normalise ' ' → ''

```

```

    Parameters
    -----
    texts : pd.Series
        Series of raw text strings.

```

```

    Returns
    -----
    pd.Series
        Normalised text strings.
    """

```

```

# keep comma, hyphen, underscore
punc = ''.join(ch for ch in string.punctuation if ch not in ',--_')

trans_table = str.maketrans('--', ' ', punc)

def _norm(s: str) -> str:
    s = str(s).lower().translate(trans_table)
    s = " ".join(s.split())
    s = s.replace(' ', '')
    return s

```

```

    return texts.apply(_norm)

[13]: def data_prep(df: pd.DataFrame) -> pd.DataFrame:

    """
    Preprocess the RUBIAS dataset into a unified stereotype format.

    Removes index-like columns, anti-trope content and irrelevant
    task types, standardises column names and stereotype-type labels,
    cleans the text field, and removes empty/duplicate rows.

    Output schema:
        * text
        * category          (fixed to 'stereotype')
        * stereotype_type   (e.g. gender, profession, nationality)

    Parameters
    -----
    df : pd.DataFrame
        Raw RUBIAS dataframe.

    Returns
    -----
    pd.DataFrame
        Cleaned dataframe ready for manual curation or augmentation.
    """

# Drop any index-like columns such as 'Unnamed: 0'
unnamed_cols = [c for c in df.columns if c.startswith("Unnamed")]
if unnamed_cols:
    df = df.drop(columns=unnamed_cols)

# Remove anti-stereotype variants
if "anti-trope" in df.columns:
    df = df.drop(columns=["anti-trope"])

# Remove non-relevant generation templates
irrelevant = {"template_hetpos", "freeform_repres"}
if "task_type" in df.columns:
    df = df[~df["task_type"].isin(irrelevant)]
    df = df.drop(columns=["task_type"])

# Standardise schema
df = df.rename(columns={"pro-trope": "text"})

# Keep only relevant columns
df = df[["text", "stereotype_type"]]

```

```

# Assign fixed category label
df["category"] = "stereotype"

# Format strings
df["text"] = format_string(df["text"])

# Optional: drop duplicates and empties
df = df[df["text"].notna() & (df["text"].str.len() > 0)]
df = df.drop_duplicates(subset="text")

# Order columns
df = df[["text", "category", "stereotype_type"]]

return df

```

[14]: def drop_semantic_duplicates(
 df: pd.DataFrame,
 text_col: str = "text",
 group_col: str = "stereotype_type",
 model_name: str = "DeepPavlov/rubert-base-cased-sentence",
 border_sim: float = 0.98,
):

"""
 Remove semantically near-duplicate text entries from a dataframe.

 This function computes sentence embeddings using a SentenceTransformer
 model and identifies near-duplicate sentences based on cosine similarity.
 Only sentences belonging to the same group (e.g., same stereotype type)
 are compared. For each pair of sentences that exceed the similarity
 threshold, the later-indexed entry is removed. Detected duplicates
 are printed to stdout.

 Parameters

 df : pandas.DataFrame
 Input dataframe containing at least the text column and optionally a
 grouping column.
 text_col : str, default "text"
 Name of the column containing raw text to evaluate for duplicates.
 group_col : str, default "stereotype_type"
 Column name determining groups within which similarity comparisons
 are performed. Sentences from different groups are never compared.
 model_name : str, default "DeepPavlov/rubert-base-cased-sentence"
 Identifier of a SentenceTransformer model used to compute embeddings.
 border_sim : float, default 0.98

Cosine similarity threshold above which two sentences are considered near-duplicates. Must be in the range [0, 1].

Returns

pandas.DataFrame

A cleaned dataframe with near-duplicate rows removed and the index reset.

Notes

- The function prints each detected near-duplicate pair, including the kept sentence, removed sentence, and similarity score.
- Duplicate detection is greedy: the earliest occurrence is preserved, and any later duplicates are removed.
- Performance may degrade for very large datasets due to $O(n^2)$ pairwise similarity comparisons.

Examples

```
>>> df_clean = drop_semantic_duplicates(  
...     df,  
...     text_col="text",  
...     group_col="stereotype_type",  
...     border_sim=0.90,  
... )  
>>> df_clean.head()  
"""  
  
df = df.reset_index(drop=True).copy()  
  
sent_encoder = SentenceTransformer(model_name)  
texts = df[text_col].tolist()  
embeddings = sent_encoder.encode(texts, convert_to_tensor=True)  
  
to_remove = set()  
n = len(df)  
  
for i in range(n):  
    if i in to_remove:  
        continue  
    for j in range(i + 1, n):  
        if j in to_remove:  
            continue  
  
        if df.loc[i, group_col] != df.loc[j, group_col]:  
            continue
```

```

sim = util.pytorch_cos_sim(embeddings[i], embeddings[j]).item()

if sim > border_sim:
    print("-" * 80)
    print(f"Duplicates Found (Similarity = {sim:.3f})")
    print(f"Saved [{i}]: {df.loc[i, text_col]}")
    print(f"Removed [{j}]: {df.loc[j, text_col]}")
    print("-" * 80)

    to_remove.add(j)

print(f"\nTotal near-duplicates removed: {len(to_remove)}\n")

return df.drop(index=list(to_remove)).reset_index(drop=True)

```

[15]: def augment_sentence_claude(sentence: str,
 stereotype_type: str,
 temperature: float = 0.5) -> dict:

"""
Generate neutral and unrelated (nonsensical) augmentations for a given Russian stereotype sentence using the Bedrock Proxy API.

This function embeds the entire instruction prompt and examples inside a single user message, because the proxy does not support the `system` role. The output is validated via a strict JSON schema.

Parameters

sentence : str
The original stereotype sentence in Russian.

stereotype_type : str
The associated stereotype group (e.g., 'gender', 'profession').

temperature : float, optional
Sampling temperature for the LLM. Default is 0.7.

Returns

dict
A dictionary containing:

- *'neutral'*: str
A neutralised version of the input sentence.
- *'unrelated'*: str
A nonsensical, unrelated version of the input sentence.

Raises

```

-----
RuntimeError
    If the API returns a non-200 status code.
ValueError
    If JSON parsing fails or required keys are missing.
"""

# Build full user prompt: instructions + examples + concrete task
user_content = (
    SYSTEM_PROMPT_RU.strip()
    + "\n\n                .\n"
    + "                :\n"
    + f"\n{sentence}\n"
    + f"        : {stereotype_type}\n\n"
    + "                ."
    "            JSON      :\n"
    "{ \"neutral\": \"...\", \"unrelated\": \"...\" }"
)

# Message container for API
messages = [{  

    "role": "user",  

    "content": user_content,  

}]

# Request payload (API requires team_id, api_token, model inside JSON)
payload = {  

    "team_id": TEAM_ID,  

    "api_token": API_TOKEN,  

    "model": MODEL_ID,  

    "messages": messages,  

    "max_tokens": 300,  

    "temperature": temperature,  

    "response_format": {  

        "type": "json_schema",  

        "json_schema": {  

            "name": "rubist_augmentation",  

            "strict": True,  

            "schema": AUG_SCHEMA,  

        },
    },
}

# Execute POST request
response = requests.post(  

    API_ENDPOINT,  

    headers={
```

```

        "Content-Type": "application/json",
        "X-Team-ID": TEAM_ID,
        "X-API-Token": API_TOKEN,
    },
    json=payload,
    timeout=60,
)

# Validate HTTP layer
if response.status_code != 200:
    raise RuntimeError(
        f"API error {response.status_code}: {response.text[:500]}"
    )

# Parse API response
result = response.json()

# Quota reporting
if "metadata" in result and "remaining_quota" in result["metadata"]:
    quota = result["metadata"]["remaining_quota"]
    print(
        f"[Quota] LLM={quota['llm_cost']} | GPU={quota['gpu_cost']} | "
        f"Used={quota['total_cost']}/{quota['budget_limit']} | "
        f"Remaining={quota['remaining_budget']} | "
        f"Usage={quota['budget_usage_percent']}%"
    )
)

# Extract model-generated JSON text
try:
    raw_text = result["content"][0]["text"]
except Exception as exc:
    raise ValueError(f"Malformed response structure: {result}") from exc

# Parse JSON output from the model
try:
    data = json.loads(raw_text)
except json.JSONDecodeError:
    raise ValueError(
        f"Could not parse JSON from model output:\n{raw_text}"
    )

# Validate required keys
if not all(k in data for k in ("neutral", "unrelated")):
    raise ValueError(f"Missing required keys in response: {data}")

# Clean and return output
return {

```

```
        "neutral": data["neutral"].strip(),
        "unrelated": data["unrelated"].strip(),
    }
```

Exploratory Data Analysis

```
[16]: # Print the shape
print(rubias.columns)
print(rubias.shape)

rubias.head()
```

```
Index(['Unnamed: 0', 'pro-trope', 'anti-trope', 'stereotype_type',
       'task_type'],
      dtype='object')
(2221, 5)
```

```
[16]:   Unnamed: 0                               pro-trope \
0          0                               ...
1          1                               -
2          2                               -
3          3                               -
4          4                               .

                                         anti-trope stereotype_type \
0                               ...                               gender
1                               -                               gender
2                               -                               gender
3                               -                               gender
4                               .                               gender

                           task_type
0  freeform_generic
1  freeform_generic
2  freeform_generic
3  freeform_generic
4  freeform_generic
```

```
[17]: # Print the shape
print(ruster.columns)
print(ruster.shape)

ruster.head()
```

```
Index(['text', 'category', 'stereotype_type', 'notes'], dtype='object')
(408, 4)
```

```
[17]:                                text     category \
0          -     ,     -     stereotype
```

```

1                      stereotype
2                      stereotype
3 ,      -      stereotype
4          ... stereotype

    stereotype_type           notes
0     gender   A woman's happiness is in the family, a man's ...
1     gender       Men don't know how to handle children.
2     gender        Ladies can't drive cars.
3     gender   Women speak in riddles, men speak directly.
4     gender   Men must stoically endure pain and not seek help.

```

[18]: # Display the general information and variable type of the dataset
rubias.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2221 entries, 0 to 2220
Data columns (total 5 columns):
 #   Column           Non-Null Count  Dtype  
---  --  
 0   Unnamed: 0        2221 non-null   int64  
 1   pro-trope         2221 non-null   object  
 2   anti-trope        2221 non-null   object  
 3   stereotype_type   2221 non-null   object  
 4   task_type         2221 non-null   object  
dtypes: int64(1), object(4)
memory usage: 86.9+ KB

```

[19]: # Display the general information and variable type of the dataset
ruster.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 408 entries, 0 to 407
Data columns (total 4 columns):
 #   Column           Non-Null Count  Dtype  
---  --  
 0   text             408 non-null    object  
 1   category         408 non-null    object  
 2   stereotype_type  408 non-null    object  
 3   notes            408 non-null    object  
dtypes: object(4)
memory usage: 12.9+ KB

```

[20]: # Display the count of unique rows
rubias.nunique()

[20]: Unnamed: 0 2221
pro-trope 1989

```
anti-trope          2005
stereotype_type      4
task_type            19
dtype: int64
```

```
[21]: # Display the count of unique rows
ruster.nunique()
```

```
[21]: text          408
category        1
stereotype_type 4
notes           408
dtype: int64
```

```
[22]: # Print unique domains
print("Unique stereotype types: RuBias")
print(rubias['stereotype_type'].unique())

# Print unique domains
print("Unique subdomains:")
print(rubias['task_type'].unique())
```

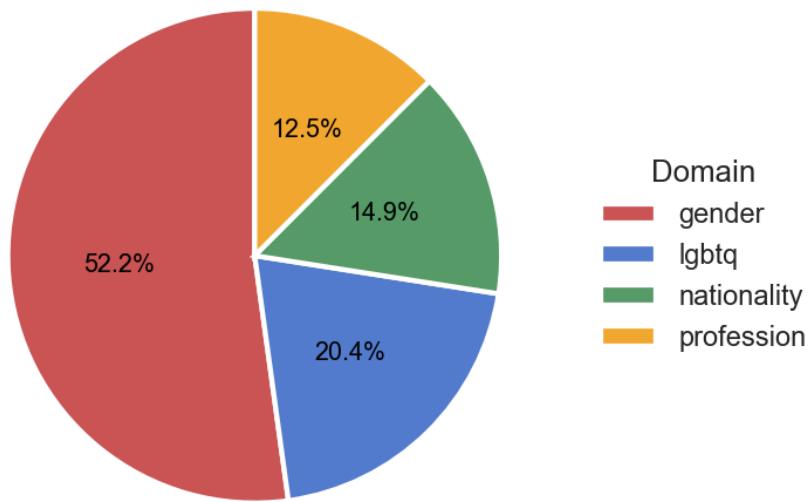
```
Unique stereotype types: RuBias
['gender' 'profession' 'nationality' 'lgbtq']
Unique subdomains:
['freeform_generic' 'freeform_prof' 'freeform_prof_full'
 'freeform_prof_stereotype' 'freeform_job' 'freeform_family_stereotype'
 'freeform_family_full' 'template_hetpos' 'freeform_enemy' 'template_rich'
 'freeform_full' 'freeform_immigrant' 'template_assoc' 'template_poor'
 'freeform_antisem' 'freeform_gendergap' 'freeform_repres' 'freeform_lgb'
 'freeform_transnb']
```

```
[23]: # Print unique domains
print("Unique stereotype types: RuSter")
print(ruster['stereotype_type'].unique())
```

```
Unique stereotype types: RuSter
['gender' 'profession' 'nationality' 'lgbtq']
```

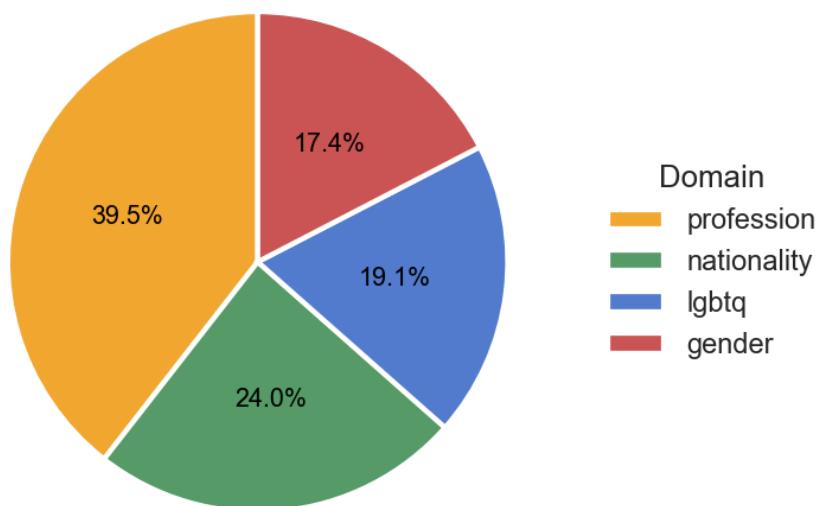
```
[24]: # Plot the pie chart
pie_chart_domain(rubias, column='stereotype_type', name = "RuBias (Original, ↴2221 samples)")
```

Social Group Distribution: RuBias (Original, 2221 samples)



```
[25]: # Plot the pie chart
pie_chart_domain(ruster, column='stereotype_type', name = "RuSter (Original, ↴408 samples)")
```

Social Group Distribution: RuSter (Original, 408 samples)



3.2.3 *Question 3 : Provide data preprocessing pipeline*

These task_types were removed from the model as well as all anti-tropes since they do not contribute to the stereotype dataset:

1. template_hetpos

1.1 pro-tope:

- Russian: ,
- English: Igor loves Anya because she is beautiful

1.2 anti-trope:

- Russian: ,
- English: Anya loves Igor because he is handsome

2. freeform_repres

2.1 pro-tope:

- Russian:
- English: He invited her to a romantic dinner

2.2 anti-trope:

- Russian:
- English: She invited her to a romantic dinner

Data Preprocessing

[26]: # Preprocess and clean

```
rubias_prep = data_prep(rubias)

# Print the shape
print(rubias_prep.columns)
print(rubias_prep.shape)

rubias_prep.head()
```

```
Index(['text', 'category', 'stereotype_type'], dtype='object')
(1855, 3)
```

[26]:

```
          text      category \
0           ... stereotype
1           ... stereotype
2           ... stereotype
3           ... stereotype
4           ... stereotype

stereotype_type
0            gender
```

```
1      gender
2      gender
3      gender
4      gender
```

```
[27]: # # Save for manual cleaning
# rubias_prep.to_excel("COMP0173_Temp_Data/rubias_prep.xlsx", index=False)
# print("Converted successfully!")
```

```
[28]: # Format strings
ruster["text"] = format_string(ruster["text"])
ruster = ruster.drop(columns=["notes"])

# Print the shape
print(ruster.columns)
print(ruster.shape)

ruster.head()
```

```
Index(['text', 'category', 'stereotype_type'], dtype='object')
(408, 3)
```

```
[28]:
```

	text	category	\
0	,	stereotype	
1		stereotype	
2		stereotype	
3	,	stereotype	
4	..	stereotype	

	stereotype_type
0	gender
1	gender
2	gender
3	gender
4	gender

Manual Cleaning

1. Since the original dataset was mostly about the biases and the texts are generated by using the “she” and “he” pronouns - I will replace these by “Woman”, “Man”, and drop the irrelevant stereotypes than are either not common stereotypes or counterfactual, over negative, duplicates - replace common slur words in russian to more formal

```
[29]: # Download the manually cleaned dataset
rubias_manual = pd.read_excel("COMP0173_Temp_Data/rubias_manual.xlsx")
rubias_manual.shape
```

```
[29]: (1087, 3)
```

```
[30]: # Create a new dataset by merging
rubist = pd.concat([rubias_manual, ruster], ignore_index=True)

# Drop duplicates
rubist = rubist.drop_duplicates(subset="text")
rubist = rubist.dropna()
```

RuBiST - New Dataset

```
[31]: # Drop examples that are similar to others
print("RuBiSt Dataset Shape - Before:", rubist.shape)

rubist_dedup = drop_semantic_duplicates(
    rubist,
    text_col="text",
    group_col="stereotype_type",
    border_sim=0.85,
)

print("RuBiSt Dataset Shape - After:", rubist_dedup.shape)
```

RuBiSt Dataset Shape - Before: (1486, 3)

Exception when trying to download <https://sbert.net/models/DeepPavlov/rubert-base-cased-sentence.zip>. Response 404
 WARNING:root:SentenceTransformer-Model
<https://sbert.net/models/DeepPavlov/rubert-base-cased-sentence.zip> not found.
 Try to create it from scratch
 WARNING:root:Try to create Transformer Model DeepPavlov/rubert-base-cased-sentence with mean pooling
 You're using a BertTokenizerFast tokenizer. Please note that with a fast tokenizer, using the `__call__` method is faster than using a method to encode the text followed by a call to the `pad` method to get a padded encoding.

Duplicates Found (Similarity = 0.882)

Saved [0]:

Removed [5]: -

Duplicates Found (Similarity = 0.865)

Saved [0]:

Removed [764]:

Duplicates Found (Similarity = 0.851)

Saved [1]:

Removed [89]: ,

Duplicates Found (Similarity = 0.868)

Saved [2] : ,

Removed [3] :

Duplicates Found (Similarity = 0.881)

Saved [2] : ,

Removed [17] : ,

Duplicates Found (Similarity = 0.900)

Saved [2] : ,

Removed [134] :

Duplicates Found (Similarity = 0.869)

Saved [2] : ,

Removed [265] : ,

Duplicates Found (Similarity = 0.876)

Saved [2] : ,

Removed [287] : ,

Duplicates Found (Similarity = 0.868)

Saved [2] : ,

Removed [384] :

Duplicates Found (Similarity = 0.891)

Saved [2] : ,

Removed [530] :

Duplicates Found (Similarity = 0.879)

Saved [2] : ,

Removed [538] : ,

Duplicates Found (Similarity = 0.916)

Saved [2] : ,

Removed [681] :

Duplicates Found (Similarity = 0.866)

Saved [2] : ,

Removed [726] :

Duplicates Found (Similarity = 0.852)

Saved [2] : ,

Removed [738] :

Duplicates Found (Similarity = 0.864)

Saved [2] : ,

Removed [827] :

Duplicates Found (Similarity = 0.903)

Saved [2] : ,

Removed [1009] :

Duplicates Found (Similarity = 0.856)

Saved [2] : ,

Removed [1027] : ,

Duplicates Found (Similarity = 0.893)

Saved [2] : ,

Removed [1413] :

Duplicates Found (Similarity = 0.855)

Saved [2] : ,

Removed [1433] :

Duplicates Found (Similarity = 0.919)

Saved [2] : ,

Removed [1467] :

Duplicates Found (Similarity = 0.852)

Saved [2] : ,

Removed [1472] :

Duplicates Found (Similarity = 0.878)
Saved [4] : ,
Removed [19] :

Duplicates Found (Similarity = 0.864)
Saved [4] : ,
Removed [188] :

Duplicates Found (Similarity = 0.906)
Saved [4] : ,
Removed [250] :

Duplicates Found (Similarity = 0.896)
Saved [4] : ,
Removed [260] :

Duplicates Found (Similarity = 0.857)
Saved [4] : ,
Removed [293] :

Duplicates Found (Similarity = 0.869)
Saved [4] : ,
Removed [294] :

Duplicates Found (Similarity = 0.898)
Saved [4] : ,
Removed [377] :

Duplicates Found (Similarity = 0.880)
Saved [4] : ,
Removed [381] :

Duplicates Found (Similarity = 0.877)
Saved [4] : ,

Removed [542] :

Duplicates Found (Similarity = 0.876)

Saved [4] : ,

Removed [671] :

Duplicates Found (Similarity = 0.861)

Saved [4] : ,

Removed [810] :

Duplicates Found (Similarity = 0.891)

Saved [4] : ,

Removed [831] :

Duplicates Found (Similarity = 0.900)

Saved [4] : ,

Removed [903] : ,

Duplicates Found (Similarity = 0.874)

Saved [4] : ,

Removed [923] :

Duplicates Found (Similarity = 0.918)

Saved [4] : ,

Removed [1029] :

Duplicates Found (Similarity = 0.856)

Saved [4] : ,

Removed [1034] :

Duplicates Found (Similarity = 0.916)

Saved [4] : ,

Removed [1036] :

Duplicates Found (Similarity = 0.874)

Saved [4] : ,

Removed [1037] :

Duplicates Found (Similarity = 0.875)

Saved [4] : ,

Removed [1042] :

Duplicates Found (Similarity = 0.853)

Saved [4] : ,

Removed [1044] :

Duplicates Found (Similarity = 0.873)

Saved [4] : ,

Removed [1410] :

Duplicates Found (Similarity = 0.892)

Saved [4] : ,

Removed [1421] : ,

Duplicates Found (Similarity = 0.884)

Saved [4] : ,

Removed [1459] : ,

Duplicates Found (Similarity = 0.989)

Saved [6] :

Removed [909] :

Duplicates Found (Similarity = 0.871)

Saved [8] :

Removed [878] : ,

Duplicates Found (Similarity = 0.851)

Saved [10] :

Removed [745] : ,

Duplicates Found (Similarity = 0.854)

Saved [10] :

Removed [834] :

Duplicates Found (Similarity = 0.923)

Saved [10] :

Removed [1172] :

Duplicates Found (Similarity = 0.874)

Saved [12] :

Removed [591] :

Duplicates Found (Similarity = 0.862)

Saved [13] :

Removed [159] :

Duplicates Found (Similarity = 0.856)

Saved [13] :

Removed [524] :

Duplicates Found (Similarity = 0.928)

Saved [13] :

Removed [525] :

Duplicates Found (Similarity = 0.915)

Saved [13] :

Removed [617] : ,

Duplicates Found (Similarity = 0.898)

Saved [13] :

Removed [621] : ,

Duplicates Found (Similarity = 0.872)

Saved [13] :

Removed [781] : ,

Duplicates Found (Similarity = 0.866)

Saved [13] :

Removed [946] :

Duplicates Found (Similarity = 0.851)

Saved [13] :
Removed [1005] :

Duplicates Found (Similarity = 0.874)
Saved [13] :
Removed [1045] :

Duplicates Found (Similarity = 0.901)
Saved [13] :
Removed [1047] :

Duplicates Found (Similarity = 0.867)
Saved [13] :
Removed [1054] :

Duplicates Found (Similarity = 0.893)
Saved [13] :
Removed [1064] :

Duplicates Found (Similarity = 0.859)
Saved [13] :
Removed [1162] : ,

Duplicates Found (Similarity = 0.892)
Saved [13] :
Removed [1175] :

Duplicates Found (Similarity = 0.895)
Saved [13] :
Removed [1181] : , , ,
,

Duplicates Found (Similarity = 0.878)
Saved [13] :
Removed [1189] : ,

Duplicates Found (Similarity = 0.867)
Saved [13] :

Removed [1193] :

Duplicates Found (Similarity = 0.861)

Saved [13] :

Removed [1212] :

Duplicates Found (Similarity = 0.889)

Saved [13] :

Removed [1213] :

Duplicates Found (Similarity = 0.895)

Saved [13] :

Removed [1215] :

Duplicates Found (Similarity = 0.879)

Saved [13] :

Removed [1216] :

Duplicates Found (Similarity = 0.893)

Saved [13] :

Removed [1226] :

Duplicates Found (Similarity = 0.856)

Saved [13] :

Removed [1230] :

Duplicates Found (Similarity = 0.872)

Saved [13] :

Removed [1240] :

Duplicates Found (Similarity = 0.908)

Saved [13] :

Removed [1269] :

Duplicates Found (Similarity = 0.859)

Saved [13] :

Removed [1285] : , ,

Duplicates Found (Similarity = 0.924)

Saved [18] :

Removed [62] :

Duplicates Found (Similarity = 0.882)

Saved [20] :

Removed [297] : ,

Duplicates Found (Similarity = 0.885)

Saved [21] :

Removed [38] : ,

Duplicates Found (Similarity = 0.925)

Saved [21] :

Removed [197] :

Duplicates Found (Similarity = 0.875)

Saved [21] :

Removed [205] :

Duplicates Found (Similarity = 0.903)

Saved [21] :

Removed [228] :

Duplicates Found (Similarity = 0.896)

Saved [21] :

Removed [230] :

Duplicates Found (Similarity = 0.893)

Saved [21] :

Removed [426] :

Duplicates Found (Similarity = 0.871)

Saved [21] :

Removed [429] :

Duplicates Found (Similarity = 0.893)

Saved [21] :
Removed [652] :

Duplicates Found (Similarity = 0.887)
Saved [21] :
Removed [679] :

Duplicates Found (Similarity = 0.924)
Saved [21] :
Removed [1192] :

Duplicates Found (Similarity = 0.858)
Saved [25] : ,

Removed [1203] :

Duplicates Found (Similarity = 0.891)
Saved [29] : ,
Removed [37] : ,

Duplicates Found (Similarity = 0.892)
Saved [30] : ,
Removed [31] : ,

Duplicates Found (Similarity = 0.872)
Saved [33] : ,
Removed [1001] : ,

Duplicates Found (Similarity = 0.859)
Saved [33] : ,
Removed [1311] :

Duplicates Found (Similarity = 0.857)
Saved [34] : ,

Removed [55] :

Duplicates Found (Similarity = 0.861)

Saved [40] : ,
Removed [876] :

Duplicates Found (Similarity = 0.853)
Saved [45] :
Removed [180] :

Duplicates Found (Similarity = 0.853)
Saved [45] :
Removed [280] : ,

Duplicates Found (Similarity = 0.871)
Saved [45] :
Removed [528] :

Duplicates Found (Similarity = 0.880)
Saved [45] :
Removed [963] :

Duplicates Found (Similarity = 0.883)
Saved [45] :
Removed [981] :

Duplicates Found (Similarity = 0.858)
Saved [45] :
Removed [986] :

Duplicates Found (Similarity = 0.892)
Saved [45] :
Removed [1056] :

Duplicates Found (Similarity = 0.889)
Saved [45] :
Removed [1057] :

Duplicates Found (Similarity = 0.878)
Saved [45] :
Removed [1169] :

Duplicates Found (Similarity = 0.874)

Saved [45] :

Removed [1195] : ,

Duplicates Found (Similarity = 0.863)

Saved [45] :

Removed [1198] :

Duplicates Found (Similarity = 0.862)

Saved [45] :

Removed [1271] :

Duplicates Found (Similarity = 0.880)

Saved [48] :

Removed [1166] :

Duplicates Found (Similarity = 0.858)

Saved [50] :

Removed [72] :

Duplicates Found (Similarity = 0.868)

Saved [50] :

Removed [160] :

Duplicates Found (Similarity = 0.850)

Saved [53] :

Removed [67] :

Duplicates Found (Similarity = 0.878)

Saved [53] :

Removed [1004] :

Duplicates Found (Similarity = 0.886)

Saved [53] :

Removed [1170] :

Duplicates Found (Similarity = 0.893)

Saved [59] :

Removed [216] :

Duplicates Found (Similarity = 0.851)

Saved [59] :

Removed [217] :

Duplicates Found (Similarity = 0.879)

Saved [65] :

Removed [1341] :

Duplicates Found (Similarity = 0.885)

Saved [66] :

Removed [1339] :

Duplicates Found (Similarity = 0.899)

Saved [70] :

Removed [80] :

Duplicates Found (Similarity = 0.864)

Saved [71] :

Removed [189] :

Duplicates Found (Similarity = 0.879)

Saved [73] :

Removed [226] :

Duplicates Found (Similarity = 0.861)

Saved [74] : , ,

Removed [84] :

Duplicates Found (Similarity = 0.899)

Saved [74] : , ,

Removed [182] :

Duplicates Found (Similarity = 0.876)

Saved [74] : , ,

Removed [628] : ,

Duplicates Found (Similarity = 0.881)

Saved [74]: , ,

Removed [1070]: ,

Duplicates Found (Similarity = 0.852)

Saved [74]: , ,

Removed [1077]:

Duplicates Found (Similarity = 0.864)

Saved [74]: , ,

Removed [1345]:

Duplicates Found (Similarity = 0.902)

Saved [74]: , ,

Removed [1361]:

Duplicates Found (Similarity = 0.874)

Saved [74]: , ,

Removed [1395]:

Duplicates Found (Similarity = 0.875)

Saved [74]: , ,

Removed [1408]:

Duplicates Found (Similarity = 0.904)

Saved [75]:

Removed [1199]:

Duplicates Found (Similarity = 0.881)

Saved [79]: ,

Removed [693]:

Duplicates Found (Similarity = 0.857)

Saved [79]: ,

Removed [710]: , , ,

Duplicates Found (Similarity = 0.880)

Saved [79]: ,

Removed [725] :

Duplicates Found (Similarity = 0.861)

Saved [79] : ,

Removed [1326] :

Duplicates Found (Similarity = 0.858)

Saved [79] : ,

Removed [1383] :

Duplicates Found (Similarity = 0.859)

Saved [83] : ,

Removed [1030] : ,

,

Duplicates Found (Similarity = 0.861)

Saved [86] : ,

Removed [766] :

Duplicates Found (Similarity = 0.909)

Saved [90] :

Removed [128] :

Duplicates Found (Similarity = 0.908)

Saved [90] :

Removed [135] :

Duplicates Found (Similarity = 0.923)

Saved [90] :

Removed [138] :

Duplicates Found (Similarity = 0.892)

Saved [90] :

Removed [713] :

Duplicates Found (Similarity = 0.859)

Saved [90] :

Removed [787] :

Duplicates Found (Similarity = 0.913)

Saved [90] :

Removed [788] :

Duplicates Found (Similarity = 0.889)

Saved [90] :

Removed [1417] : ,

Duplicates Found (Similarity = 0.851)

Saved [90] :

Removed [1424] :

Duplicates Found (Similarity = 0.851)

Saved [94] : ,

Removed [373] :

Duplicates Found (Similarity = 0.854)

Saved [94] : ,

Removed [678] : ,

Duplicates Found (Similarity = 0.861)

Saved [94] : ,

Removed [1046] :

Duplicates Found (Similarity = 0.890)

Saved [95] :

Removed [118] :

Duplicates Found (Similarity = 0.852)

Saved [95] :

Removed [639] :

Duplicates Found (Similarity = 0.870)

Saved [96] :

Removed [626] :

Duplicates Found (Similarity = 0.876)

Saved [99] :

Removed [100] :

Duplicates Found (Similarity = 0.871)

Saved [103] : ,

Removed [379] :

Duplicates Found (Similarity = 0.862)

Saved [103] : ,

Removed [452] :

Duplicates Found (Similarity = 0.860)

Saved [103] : ,

Removed [533] :

Duplicates Found (Similarity = 0.897)

Saved [103] : ,

Removed [1093] :

Duplicates Found (Similarity = 0.903)

Saved [103] : ,

Removed [1108] : ,

Duplicates Found (Similarity = 0.870)

Saved [103] : ,

Removed [1129] :

Duplicates Found (Similarity = 0.885)

Saved [103] : ,

Removed [1134] : ,

Duplicates Found (Similarity = 0.864)

Saved [105] :

Removed [532] :

Duplicates Found (Similarity = 0.862)

Saved [105] :

Removed [854] :

Duplicates Found (Similarity = 0.851)

Saved [107] :

Removed [191] :

Duplicates Found (Similarity = 0.921)

Saved [107] :

Removed [264] :

Duplicates Found (Similarity = 0.857)

Saved [107] :

Removed [321] :

Duplicates Found (Similarity = 0.896)

Saved [107] :

Removed [575] :

Duplicates Found (Similarity = 0.856)

Saved [107] :

Removed [599] :

Duplicates Found (Similarity = 0.871)

Saved [107] :

Removed [712] :

Duplicates Found (Similarity = 0.938)

Saved [107] :

Removed [853] :

Duplicates Found (Similarity = 0.913)

Saved [107] :

Removed [890] :

Duplicates Found (Similarity = 0.878)

Saved [107] :

Removed [1085] :

Duplicates Found (Similarity = 0.857)

Saved [109] :

Removed [553] :

Duplicates Found (Similarity = 0.853)

Saved [109] :

Removed [793] :

Duplicates Found (Similarity = 0.867)

Saved [110] :

Removed [147] :

Duplicates Found (Similarity = 0.895)

Saved [112] :

Removed [113] :

Duplicates Found (Similarity = 0.996)

Saved [112] :

Removed [253] :

Duplicates Found (Similarity = 0.852)

Saved [112] :

Removed [291] :

Duplicates Found (Similarity = 0.873)

Saved [112] :

Removed [295] :

Duplicates Found (Similarity = 0.879)

Saved [112] :

Removed [887] :

Duplicates Found (Similarity = 0.867)

Saved [112] :

Removed [1088] : ,

Duplicates Found (Similarity = 0.861)

Saved [115] :

Removed [167] :

Duplicates Found (Similarity = 0.879)

Saved [115] :

Removed [1031] :

Duplicates Found (Similarity = 0.886)

Saved [115] :

Removed [1483] :

Duplicates Found (Similarity = 0.991)

Saved [120] :

Removed [271] :

Duplicates Found (Similarity = 0.871)

Saved [121] :

Removed [1452] :

Duplicates Found (Similarity = 0.870)

Saved [127] :

Removed [1479] :

Duplicates Found (Similarity = 0.920)

Saved [130] :

Removed [141] :

Duplicates Found (Similarity = 0.920)

Saved [143] :

Removed [153] : ,

Duplicates Found (Similarity = 0.918)

Saved [143] :

Removed [391] :

Duplicates Found (Similarity = 0.852)

Saved [143] :

Removed [1468] :

Duplicates Found (Similarity = 0.851)

Saved [144] :

Removed [1420] : ,

Duplicates Found (Similarity = 0.868)

Saved [144] :

Removed [1444] :

Duplicates Found (Similarity = 0.853)

Saved [144] :

Removed [1455] :

Duplicates Found (Similarity = 0.955)

Saved [146] :

Removed [1475] :

Duplicates Found (Similarity = 0.881)

Saved [157] :

Removed [960] :

Duplicates Found (Similarity = 0.851)

Saved [157] :

Removed [1244] :

Duplicates Found (Similarity = 0.909)

Saved [168] :

Removed [171] :

Duplicates Found (Similarity = 0.865)

Saved [168] :

Removed [175] :

Duplicates Found (Similarity = 0.883)

Saved [169] : , , ,

Removed [174] :

Duplicates Found (Similarity = 0.854)

Saved [169]: , ,

Removed [807]:

Duplicates Found (Similarity = 0.869)

Saved [172]:

Removed [611]:

Duplicates Found (Similarity = 0.899)

Saved [173]:

Removed [775]:

Duplicates Found (Similarity = 0.894)

Saved [173]:

Removed [806]:

Duplicates Found (Similarity = 0.870)

Saved [178]:

Removed [255]:

Duplicates Found (Similarity = 0.861)

Saved [178]:

Removed [366]:

Duplicates Found (Similarity = 0.863)

Saved [178]:

Removed [506]: ,

Duplicates Found (Similarity = 0.861)

Saved [178]:

Removed [653]:

Duplicates Found (Similarity = 0.857)

Saved [178]:

Removed [660]:

Duplicates Found (Similarity = 0.882)

Saved [178] :

Removed [708] : ,

Duplicates Found (Similarity = 0.861)

Saved [178] :

Removed [1107] : , , ,

Duplicates Found (Similarity = 0.852)

Saved [178] :

Removed [1115] : 30 , , ,

Duplicates Found (Similarity = 0.879)

Saved [179] :

Removed [200] :

Duplicates Found (Similarity = 0.891)

Saved [179] :

Removed [714] :

Duplicates Found (Similarity = 0.877)

Saved [192] :

Removed [602] :

Duplicates Found (Similarity = 0.970)

Saved [194] :

Removed [198] :

Duplicates Found (Similarity = 0.885)

Saved [194] :

Removed [412] :

Duplicates Found (Similarity = 0.879)

Saved [194] :

Removed [572] :

Duplicates Found (Similarity = 0.880)

Saved [194] :
Removed [574] :

Duplicates Found (Similarity = 0.860)
Saved [194] :
Removed [579] :

Duplicates Found (Similarity = 0.869)
Saved [194] :
Removed [595] :

Duplicates Found (Similarity = 0.886)
Saved [194] :
Removed [871] :

Duplicates Found (Similarity = 0.853)
Saved [199] : ,
Removed [239] :

Duplicates Found (Similarity = 0.859)
Saved [201] : ,
Removed [319] :

Duplicates Found (Similarity = 0.881)
Saved [204] :
Removed [534] :

Duplicates Found (Similarity = 0.867)
Saved [204] :
Removed [797] :

Duplicates Found (Similarity = 0.870)
Saved [204] :
Removed [1327] :

Duplicates Found (Similarity = 0.872)
Saved [207] :
Removed [316] :

Duplicates Found (Similarity = 0.874)

Saved [207] :

Removed [598] :

Duplicates Found (Similarity = 0.969)

Saved [209] :

Removed [924] :

Duplicates Found (Similarity = 0.854)

Saved [214] :

Removed [1154] : , , ,
,

Duplicates Found (Similarity = 0.856)

Saved [218] :

Removed [335] :

Duplicates Found (Similarity = 0.875)

Saved [218] :

Removed [809] :

Duplicates Found (Similarity = 0.868)

Saved [221] :

Removed [439] :

Duplicates Found (Similarity = 0.852)

Saved [223] :

Removed [972] :

Duplicates Found (Similarity = 0.917)

Saved [223] :

Removed [1302] :

Duplicates Found (Similarity = 0.890)

Saved [237] : ,

Removed [269] :

Duplicates Found (Similarity = 0.912)

Saved [237] : ,

Removed [270] :

Duplicates Found (Similarity = 0.879)

Saved [237] : ,

Removed [1113] : ,

Duplicates Found (Similarity = 0.874)

Saved [238] :

Removed [939] :

Duplicates Found (Similarity = 0.872)

Saved [241] : ,

Removed [257] : ,

Duplicates Found (Similarity = 0.882)

Saved [242] : ,

Removed [251] : ,

Duplicates Found (Similarity = 0.894)

Saved [243] : ,

Removed [248] : , ,

Duplicates Found (Similarity = 0.917)

Saved [244] :

Removed [941] :

Duplicates Found (Similarity = 0.911)

Saved [247] :

Removed [1040] :

Duplicates Found (Similarity = 0.933)

Saved [247] :

Removed [1041] : ,

Duplicates Found (Similarity = 0.869)

Saved [247] :
Removed [1461] :

Duplicates Found (Similarity = 0.856)
Saved [256] :
Removed [310] :

Duplicates Found (Similarity = 0.886)
Saved [256] :
Removed [337] :

Duplicates Found (Similarity = 0.887)
Saved [256] :
Removed [341] :

Duplicates Found (Similarity = 0.863)
Saved [256] :
Removed [343] :

Duplicates Found (Similarity = 0.905)
Saved [256] :
Removed [344] :

Duplicates Found (Similarity = 0.882)
Saved [256] :
Removed [345] :

Duplicates Found (Similarity = 0.917)
Saved [256] :
Removed [363] :

Duplicates Found (Similarity = 0.905)
Saved [256] :
Removed [364] :

Duplicates Found (Similarity = 0.958)
Saved [256] :
Removed [367] :

Duplicates Found (Similarity = 0.898)

Saved [256] :

Removed [546] :

Duplicates Found (Similarity = 0.912)

Saved [256] :

Removed [548] :

Duplicates Found (Similarity = 0.897)

Saved [256] :

Removed [549] :

Duplicates Found (Similarity = 0.865)

Saved [256] :

Removed [648] : ,

Duplicates Found (Similarity = 0.886)

Saved [256] :

Removed [649] : ,

Duplicates Found (Similarity = 0.873)

Saved [256] :

Removed [651] :

Duplicates Found (Similarity = 0.851)

Saved [256] :

Removed [659] :

Duplicates Found (Similarity = 0.946)

Saved [256] :

Removed [1082] :

Duplicates Found (Similarity = 0.892)

Saved [256] :

Removed [1096] :

Duplicates Found (Similarity = 0.873)

Saved [256] :
Removed [1118] : ,

Duplicates Found (Similarity = 0.907)
Saved [256] :
Removed [1122] : ,

Duplicates Found (Similarity = 0.899)
Saved [256] :
Removed [1125] :

Duplicates Found (Similarity = 0.867)
Saved [256] :
Removed [1139] :

Duplicates Found (Similarity = 0.851)
Saved [259] :
Removed [448] :

Duplicates Found (Similarity = 0.859)
Saved [259] :
Removed [451] :

Duplicates Found (Similarity = 0.856)
Saved [259] :
Removed [862] :

Duplicates Found (Similarity = 0.890)
Saved [259] :
Removed [863] :

Duplicates Found (Similarity = 0.893)
Saved [259] :
Removed [864] :

Duplicates Found (Similarity = 0.869)
Saved [276] :
Removed [789] : ,

Duplicates Found (Similarity = 0.881)

Saved [276] :

Removed [814] :

Duplicates Found (Similarity = 0.858)

Saved [276] :

Removed [1425] :

,

Duplicates Found (Similarity = 0.866)

Saved [276] :

Removed [1426] :

Duplicates Found (Similarity = 0.914)

Saved [276] :

Removed [1462] :

Duplicates Found (Similarity = 0.880)

Saved [278] ,

Removed [1025] :

Duplicates Found (Similarity = 0.850)

Saved [283] :

Removed [1138] :

Duplicates Found (Similarity = 0.914)

Saved [290] :

Removed [670] ,

Duplicates Found (Similarity = 0.854)

Saved [292] :

Removed [785] :

Duplicates Found (Similarity = 0.851)

Saved [292] :

Removed [820] :

Duplicates Found (Similarity = 0.859)

Saved [299] :

Removed [1275] :

Duplicates Found (Similarity = 0.852)

Saved [302] :

Removed [421] :

Duplicates Found (Similarity = 0.909)

Saved [302] :

Removed [488] :

Duplicates Found (Similarity = 0.873)

Saved [302] :

Removed [490] :

Duplicates Found (Similarity = 0.850)

Saved [304] ,

Removed [450] :

Duplicates Found (Similarity = 0.943)

Saved [305] :

Removed [419] :

Duplicates Found (Similarity = 0.906)

Saved [305] :

Removed [427] :

Duplicates Found (Similarity = 0.881)

Saved [305] :

Removed [495] :

Duplicates Found (Similarity = 0.900)

Saved [305] :

Removed [496] :

Duplicates Found (Similarity = 0.908)

Saved [305] :

Removed [504] :

Duplicates Found (Similarity = 0.910)

Saved [305] :

Removed [509] :

Duplicates Found (Similarity = 0.877)

Saved [305] :

Removed [510] :

Duplicates Found (Similarity = 0.851)

Saved [305] :

Removed [647] :

Duplicates Found (Similarity = 0.859)

Saved [311] :

Removed [334] :

Duplicates Found (Similarity = 0.884)

Saved [311] :

Removed [539] :

Duplicates Found (Similarity = 0.855)

Saved [311] :

Removed [550] :

Duplicates Found (Similarity = 0.928)

Saved [318] :

Removed [565] :

Duplicates Found (Similarity = 0.867)

Saved [322] :

Removed [356] :

Duplicates Found (Similarity = 0.876)

Saved [322] :

Removed [568] :

Duplicates Found (Similarity = 0.881)

Saved [322] :

Removed [569] :

Duplicates Found (Similarity = 0.951)

Saved [322] :

Removed [571] : ,

Duplicates Found (Similarity = 0.854)

Saved [322] :

Removed [601] :

Duplicates Found (Similarity = 0.858)

Saved [324] :

Removed [596] :

Duplicates Found (Similarity = 0.851)

Saved [327] :

Removed [577] :

Duplicates Found (Similarity = 0.884)

Saved [328] :

Removed [441] : ,

Duplicates Found (Similarity = 0.873)

Saved [328] :

Removed [482] :

Duplicates Found (Similarity = 0.871)

Saved [328] :

Removed [513] :

Duplicates Found (Similarity = 0.872)

Saved [328] :

Removed [536] :

Duplicates Found (Similarity = 0.872)

Saved [328] :

Removed [588] :

Duplicates Found (Similarity = 0.921)

Saved [328] :

Removed [589] :

Duplicates Found (Similarity = 0.864)

Saved [336] :

Removed [404] :

Duplicates Found (Similarity = 0.877)

Saved [336] :

Removed [410] :

Duplicates Found (Similarity = 0.878)

Saved [336] :

Removed [411] :

Duplicates Found (Similarity = 0.914)

Saved [336] :

Removed [473] :

Duplicates Found (Similarity = 0.851)

Saved [338] :

Removed [449] :

Duplicates Found (Similarity = 0.905)

Saved [338] :

Removed [1092] :

Duplicates Found (Similarity = 0.852)

Saved [338] :

Removed [1140] :

Duplicates Found (Similarity = 0.933)

Saved [339] :

Removed [400] :

Duplicates Found (Similarity = 0.932)

Saved [339] :
Removed [416] :

Duplicates Found (Similarity = 0.861)
Saved [339] :
Removed [469] :

Duplicates Found (Similarity = 0.877)
Saved [339] :
Removed [470] :

Duplicates Found (Similarity = 0.939)
Saved [339] :
Removed [476] :

Duplicates Found (Similarity = 0.869)
Saved [339] :
Removed [485] :

Duplicates Found (Similarity = 0.906)
Saved [348] :
Removed [1132] :

Duplicates Found (Similarity = 0.853)
Saved [352] :
Removed [570] :

Duplicates Found (Similarity = 0.899)
Saved [352] :
Removed [1146] :

Duplicates Found (Similarity = 0.887)
Saved [353] :
Removed [355] :

Duplicates Found (Similarity = 0.876)
Saved [353] :
Removed [503] :

Duplicates Found (Similarity = 0.863)

Saved [354] :

Removed [444] : ,

Duplicates Found (Similarity = 0.860)

Saved [360] :

Removed [1111] :

Duplicates Found (Similarity = 0.870)

Saved [361] :

Removed [362] :

Duplicates Found (Similarity = 0.855)

Saved [365] :

Removed [1062] :

Duplicates Found (Similarity = 0.928)

Saved [368] :

Removed [369] :

Duplicates Found (Similarity = 0.881)

Saved [368] :

Removed [556] :

Duplicates Found (Similarity = 0.875)

Saved [368] :

Removed [609] :

Duplicates Found (Similarity = 0.870)

Saved [368] :

Removed [888] :

Duplicates Found (Similarity = 0.957)

Saved [372] :

Removed [390] :

Duplicates Found (Similarity = 0.865)

Saved [376] : ,

Removed [608] :

Duplicates Found (Similarity = 0.870)

Saved [376] : ,

Removed [802] :

Duplicates Found (Similarity = 0.870)

Saved [376] : ,

Removed [840] :

Duplicates Found (Similarity = 0.877)

Saved [376] : ,

Removed [977] :

Duplicates Found (Similarity = 0.892)

Saved [376] : ,

Removed [980] :

Duplicates Found (Similarity = 0.855)

Saved [376] : ,

Removed [982] :

Duplicates Found (Similarity = 0.917)

Saved [376] : ,

Removed [1043] :

Duplicates Found (Similarity = 0.901)

Saved [376] : ,

Removed [1048] : ,

Duplicates Found (Similarity = 0.965)

Saved [376] : ,

Removed [1051] : ,

Duplicates Found (Similarity = 0.860)

Saved [376] : ,

Removed [1052] :

Duplicates Found (Similarity = 0.962)

Saved [376]: ,

Removed [1053]:

Duplicates Found (Similarity = 0.879)

Saved [376]: ,

Removed [1206]: - ,

Duplicates Found (Similarity = 0.868)

Saved [376]: ,

Removed [1281]: , , ,

Duplicates Found (Similarity = 0.901)

Saved [376]: ,

Removed [1286]: ,

Duplicates Found (Similarity = 0.867)

Saved [380]: , , , ,

Removed [544]:

Duplicates Found (Similarity = 0.902)

Saved [387]:

Removed [396]:

Duplicates Found (Similarity = 0.884)

Saved [395]:

Removed [540]:

Duplicates Found (Similarity = 0.871)

Saved [395]:

Removed [662]:

Duplicates Found (Similarity = 0.864)

Saved [395]:

Removed [673]:

Duplicates Found (Similarity = 0.851)

Saved [395] :
Removed [852] :

Duplicates Found (Similarity = 0.910)
Saved [395] :
Removed [867] :

Duplicates Found (Similarity = 0.975)
Saved [403] :
Removed [443] :

Duplicates Found (Similarity = 0.906)
Saved [403] :
Removed [459] :

Duplicates Found (Similarity = 0.968)
Saved [409] :
Removed [440] :

Duplicates Found (Similarity = 0.904)
Saved [409] :
Removed [484] :

Duplicates Found (Similarity = 0.887)
Saved [409] :
Removed [502] :

Duplicates Found (Similarity = 0.943)
Saved [414] :
Removed [474] :

Duplicates Found (Similarity = 0.870)
Saved [418] : ,
Removed [471] :

Duplicates Found (Similarity = 0.881)
Saved [420] :
Removed [463] :

Duplicates Found (Similarity = 0.850)

Saved [422] :

Removed [446] :

Duplicates Found (Similarity = 0.880)

Saved [422] :

Removed [1149] :

Duplicates Found (Similarity = 0.872)

Saved [423] :

Removed [491] :

Duplicates Found (Similarity = 0.851)

Saved [423] :

Removed [1104] : ,

Duplicates Found (Similarity = 0.853)

Saved [428] :

Removed [1150] : ,

Duplicates Found (Similarity = 0.869)

Saved [430] :

Removed [508] :

Duplicates Found (Similarity = 0.931)

Saved [442] :

Removed [1091] :

Duplicates Found (Similarity = 0.858)

Saved [457] : ,

Removed [891] : ,

Duplicates Found (Similarity = 0.869)

Saved [464] :

Removed [559] :

Duplicates Found (Similarity = 0.874)

Saved [464] :

Removed [587] :

Duplicates Found (Similarity = 0.898)

Saved [465] :

Removed [466] :

Duplicates Found (Similarity = 0.907)

Saved [478] :

Removed [537] :

Duplicates Found (Similarity = 0.922)

Saved [478] :

Removed [586] :

Duplicates Found (Similarity = 0.983)

Saved [479] :

Removed [480] :

Duplicates Found (Similarity = 0.871)

Saved [486] :

Removed [1008] : , , ,

Duplicates Found (Similarity = 0.864)

Saved [493] : , , ,

Removed [632] : , , ,

Duplicates Found (Similarity = 0.871)

Saved [498] :

Removed [1128] :

Duplicates Found (Similarity = 0.863)

Saved [514] :

Removed [1086] : , , ,

Duplicates Found (Similarity = 0.863)

Saved [515] :

Removed [846] :

Duplicates Found (Similarity = 0.851)

Saved [521] :

Removed [640] :

Duplicates Found (Similarity = 0.915)

Saved [521] :

Removed [706] :

Duplicates Found (Similarity = 0.857)

Saved [521] :

Removed [881] :

Duplicates Found (Similarity = 0.876)

Saved [541] :

Removed [576] :

Duplicates Found (Similarity = 0.858)

Saved [551] :

Removed [657] :

Duplicates Found (Similarity = 0.904)

Saved [555] :

Removed [697] :

Duplicates Found (Similarity = 0.865)

Saved [555] :

Removed [902] :

Duplicates Found (Similarity = 0.890)

Saved [555] :

Removed [904] :

Duplicates Found (Similarity = 0.895)

Saved [555] :

Removed [915] : ,

Duplicates Found (Similarity = 0.976)

Saved [561] :

Removed [562] :

Duplicates Found (Similarity = 0.878)

Saved [573] :

Removed [585] :

Duplicates Found (Similarity = 0.913)

Saved [578] :

Removed [612] :

Duplicates Found (Similarity = 0.866)

Saved [630] ,

Removed [821] :

Duplicates Found (Similarity = 0.858)

Saved [630] ,

Removed [1254] ,

Duplicates Found (Similarity = 0.957)

Saved [631] :

Removed [643] :

Duplicates Found (Similarity = 0.896)

Saved [631] :

Removed [873] ,

Duplicates Found (Similarity = 0.885)

Saved [638] :

Removed [1065] :

Duplicates Found (Similarity = 0.864)

Saved [654] :

Removed [687] ,

Duplicates Found (Similarity = 0.853)

Saved [677] , ,

Removed [1050] : ,

Duplicates Found (Similarity = 0.879)
Saved [677] : , ,
Removed [1179] : ,

Duplicates Found (Similarity = 0.893)
Saved [677] : , ,
Removed [1283] : ,

Duplicates Found (Similarity = 0.884)
Saved [682] :
Removed [940] :

Duplicates Found (Similarity = 0.854)
Saved [683] : ,
Removed [912] : ,

Duplicates Found (Similarity = 0.862)
Saved [684] : ,
Removed [1485] :

Duplicates Found (Similarity = 0.858)
Saved [691] :
Removed [833] :

Duplicates Found (Similarity = 0.873)
Saved [699] :
Removed [901] :

Duplicates Found (Similarity = 0.882)
Saved [699] :
Removed [905] :

Duplicates Found (Similarity = 0.883)
Saved [701] : ,
Removed [914] :

Duplicates Found (Similarity = 0.850)

Saved [711] :

Removed [1446] :

Duplicates Found (Similarity = 0.854)

Saved [720] :

Removed [815] :

Duplicates Found (Similarity = 0.887)

Saved [720] :

Removed [818] :

Duplicates Found (Similarity = 0.875)

Saved [720] :

Removed [1335] :

Duplicates Found (Similarity = 0.886)

Saved [728] :

Removed [734] :

Duplicates Found (Similarity = 0.886)

Saved [728] :

Removed [1441] :

Duplicates Found (Similarity = 0.871)

Saved [747] :

Removed [776] :

Duplicates Found (Similarity = 0.979)

Saved [750] :

Removed [751] :

Duplicates Found (Similarity = 0.905)

Saved [750] :

Removed [752] :

Duplicates Found (Similarity = 0.879)

Saved [756] :
Removed [769] :

Duplicates Found (Similarity = 0.872)
Saved [756] :
Removed [1000] :

Duplicates Found (Similarity = 0.874)
Saved [757] :
Removed [1378] :

Duplicates Found (Similarity = 0.905)
Saved [760] :
Removed [761] :

Duplicates Found (Similarity = 0.884)
Saved [763] :
Removed [829] :

Duplicates Found (Similarity = 0.927)
Saved [763] :
Removed [850] :

Duplicates Found (Similarity = 0.855)
Saved [763] :
Removed [918] :

Duplicates Found (Similarity = 0.873)
Saved [763] :
Removed [920] :

Duplicates Found (Similarity = 0.868)
Saved [763] :
Removed [929] :

Duplicates Found (Similarity = 0.886)
Saved [763] :
Removed [947] :

Duplicates Found (Similarity = 0.869)

Saved [763] :

Removed [949] :

Duplicates Found (Similarity = 0.852)

Saved [763] :

Removed [971] :

Duplicates Found (Similarity = 0.852)

Saved [763] :

Removed [1293] :

Duplicates Found (Similarity = 0.886)

Saved [765] :

Removed [930] :

Duplicates Found (Similarity = 0.958)

Saved [767] :

Removed [768] :

Duplicates Found (Similarity = 0.870)

Saved [778] :

Removed [804] :

Duplicates Found (Similarity = 0.889)

Saved [783] :

Removed [824] :

Duplicates Found (Similarity = 0.857)

Saved [783] :

Removed [945] :

Duplicates Found (Similarity = 0.853)

Saved [784] :

Removed [790] :

Duplicates Found (Similarity = 0.947)

Saved [784] :
Removed [791] :

Duplicates Found (Similarity = 0.881)
Saved [784] :
Removed [935] :

Duplicates Found (Similarity = 0.852)
Saved [784] :
Removed [943] :

Duplicates Found (Similarity = 0.886)
Saved [784] :
Removed [994] :

Duplicates Found (Similarity = 0.880)
Saved [798] :
Removed [799] :

Duplicates Found (Similarity = 0.854)
Saved [798] :
Removed [801] :

Duplicates Found (Similarity = 0.929)
Saved [800] :
Removed [883] :

Duplicates Found (Similarity = 0.851)
Saved [805] :
Removed [933] :

Duplicates Found (Similarity = 0.866)
Saved [805] :
Removed [967] :

Duplicates Found (Similarity = 0.934)
Saved [813] :
Removed [1456] :

Duplicates Found (Similarity = 0.877)

Saved [816] :

Removed [1318] :

Duplicates Found (Similarity = 0.851)

Saved [823] :

Removed [825] :

Duplicates Found (Similarity = 0.877)

Saved [826] :

Removed [917] :

Duplicates Found (Similarity = 0.873)

Saved [830] :

Removed [857] :

Duplicates Found (Similarity = 0.861)

Saved [830] :

Removed [961] :

Duplicates Found (Similarity = 0.876)

Saved [832] :

Removed [875] :

Duplicates Found (Similarity = 0.852)

Saved [838] :

Removed [965] :

Duplicates Found (Similarity = 0.917)

Saved [838] :

Removed [988] :

Duplicates Found (Similarity = 0.918)

Saved [839] :

Removed [969] :

Duplicates Found (Similarity = 0.855)

Saved [845] : ,

Removed [865] :

Duplicates Found (Similarity = 0.881)

Saved [851] : ,

Removed [926] : , ,

Duplicates Found (Similarity = 0.897)

Saved [856] :

Removed [953] :

Duplicates Found (Similarity = 0.855)

Saved [856] :

Removed [983] :

Duplicates Found (Similarity = 0.994)

Saved [868] :

Removed [898] :

Duplicates Found (Similarity = 0.918)

Saved [879] :

Removed [938] :

Duplicates Found (Similarity = 0.855)

Saved [879] :

Removed [1454] :

Duplicates Found (Similarity = 0.866)

Saved [896] :

Removed [916] :

Duplicates Found (Similarity = 0.887)

Saved [921] :

Removed [1035] :

Duplicates Found (Similarity = 0.871)

Saved [932] :

Removed [1289] :

Duplicates Found (Similarity = 0.891)

Saved [936] :

Removed [937] :

Duplicates Found (Similarity = 0.910)

Saved [951] :

Removed [958] :

Duplicates Found (Similarity = 0.856)

Saved [952] :

Removed [978] :

Duplicates Found (Similarity = 0.851)

Saved [952] :

Removed [979] :

Duplicates Found (Similarity = 0.869)

Saved [962] ,

Removed [1255] :

Duplicates Found (Similarity = 0.914)

Saved [973] :

Removed [1026] :

Duplicates Found (Similarity = 0.868)

Saved [985] :

Removed [990] :

Duplicates Found (Similarity = 0.853)

Saved [997] :

Removed [998] :

Duplicates Found (Similarity = 0.850)

Saved [1006] :

Removed [1018] :

Duplicates Found (Similarity = 0.903)

Saved [1006] :

Removed [1163] : ,

Duplicates Found (Similarity = 0.876)
Saved [1012] : ,

Removed [1409] :

Duplicates Found (Similarity = 0.867)
Saved [1014] :
Removed [1205] : -

Duplicates Found (Similarity = 0.875)
Saved [1021] :
Removed [1438] :

Duplicates Found (Similarity = 0.921)
Saved [1024] :
Removed [1469] :

Duplicates Found (Similarity = 0.873)
Saved [1024] :
Removed [1470] : -

Duplicates Found (Similarity = 0.874)
Saved [1033] :
Removed [1443] :

Duplicates Found (Similarity = 0.919)
Saved [1049] :
Removed [1165] :

Duplicates Found (Similarity = 0.863)
Saved [1049] :
Removed [1227] :

Duplicates Found (Similarity = 0.870)
Saved [1049] :
Removed [1304] :

Duplicates Found (Similarity = 0.857)

Saved [1099]: , ,

Removed [1119]: ,

Duplicates Found (Similarity = 0.881)

Saved [1120]:

Removed [1121]: , ,

Duplicates Found (Similarity = 0.879)

Saved [1126]:

Removed [1127]:

Duplicates Found (Similarity = 0.865)

Saved [1126]:

Removed [1147]:

Duplicates Found (Similarity = 0.852)

Saved [1187]:

Removed [1188]:

Duplicates Found (Similarity = 0.856)

Saved [1187]:

Removed [1210]:

Duplicates Found (Similarity = 0.862)

Saved [1187]:

Removed [1220]:

Duplicates Found (Similarity = 0.858)

Saved [1187]:

Removed [1276]:

Duplicates Found (Similarity = 0.890)

Saved [1197]:

Removed [1201]: ,

Duplicates Found (Similarity = 0.888)

Saved [1223] :

Removed [1225] : ,

Duplicates Found (Similarity = 0.869)

Saved [1258] :

Removed [1267] :

Duplicates Found (Similarity = 0.862)

Saved [1259] : ,

Removed [1260] : -

Duplicates Found (Similarity = 0.873)

Saved [1331] :

Removed [1332] : ,

Duplicates Found (Similarity = 0.861)

Saved [1365] :

Removed [1400] :

Duplicates Found (Similarity = 0.883)

Saved [1385] :

Removed [1406] :

Duplicates Found (Similarity = 0.897)

Saved [1412] :

Removed [1430] :

Duplicates Found (Similarity = 0.882)

Saved [1412] :

Removed [1457] :

Duplicates Found (Similarity = 0.889)

Saved [1416] : ,

Removed [1423] : ,

```
Total near-duplicates removed: 509
```

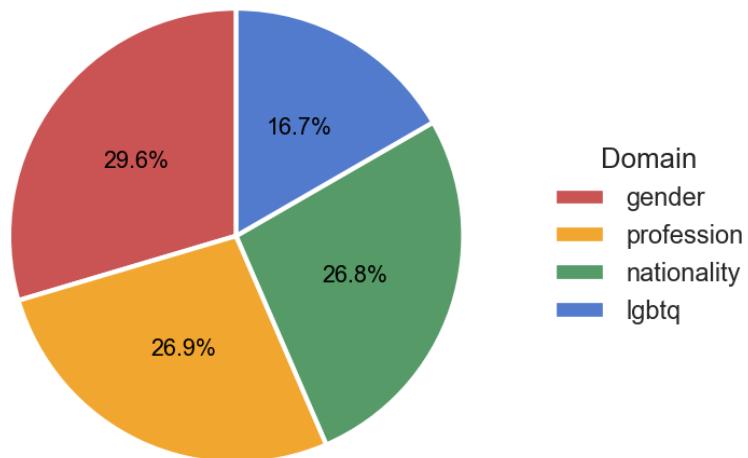
```
RuBiSt Dataset Shape - After: (977, 3)
```

```
[32]: # Display the general information and variable type of the dataset  
rubist_dedup.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 977 entries, 0 to 976  
Data columns (total 3 columns):  
 #   Column           Non-Null Count  Dtype     
---  --     
 0   text            977 non-null    object    
 1   category        977 non-null    object    
 2   stereotype_type 977 non-null    object    
 dtypes: object(3)  
 memory usage: 23.0+ KB
```

```
[33]: # Plot the pie chart  
pie_chart_domain(rubist_dedup, column='stereotype_type', name = "RuBiSt (RuBias  
+ RuSter, 977 samples)")
```

Social Group Distribution: RuBiSt (RuBias + RuSter, 977 samples)



Data Augmentation

```
[34]: from dotenv import load_dotenv  
load_dotenv()  
  
TEAM_ID = os.getenv("BEDROCK_TEAM_ID")
```

```

API_TOKEN = os.getenv("BEDROCK_API_TOKEN")

API_ENDPOINT = "https://ctwa92wg1b.execute-api.us-east-1.amazonaws.com/prod/
    ↪invoke"
MODEL_ID = "us.anthropic.claude-3-5-sonnet-20241022-v2:0"

# JSON
AUG_SCHEMA = {
    "type": "object",
    "properties": {
        "neutral": {"type": "string"},
        "unrelated": {"type": "string"}
    },
    "required": ["neutral", "unrelated"]
}

with open("COMP0173_Prompts/prompt.yaml", "r", encoding="utf-8") as f:
    CONFIG = yaml.safe_load(f)

SYSTEM_PROMPT_RU = CONFIG["instructions"]

```

```

[35]: # Ensure original rows have a label_level column
if "label_level" not in rubist_dedup.columns:
    rubist_dedup["label_level"] = "stereotype"

augmented_rows = []

# Iterate through all selected rows
for _, row in tqdm(rubist_dedup.iterrows(), total=len(rubist_dedup)):
    original_text = row["text"]
    stype = row["stereotype_type"]

    # Store the original stereotype row
    stereo_row = row.copy()
    stereo_row["label_level"] = "stereotype"
    augmented_rows.append(stereo_row)

    # Call the augmentation API
    try:
        aug = augment_sentence_claude(original_text, stype)
    except Exception as e:
        print("\nError while processing example:")
        print(original_text)
        print("Cause:", e)
        continue

    # Neutral version

```

```

neutral_row = row.copy()
neutral_row["text"] = aug["neutral"]
neutral_row["label_level"] = "neutral"
augmented_rows.append(neutral_row)

# Unrelated version
unrelated_row = row.copy()
unrelated_row["text"] = aug["unrelated"]
unrelated_row["label_level"] = "unrelated"
augmented_rows.append(unrelated_row)

# Build final augmented DataFrame
rubist_aug = pd.DataFrame(augmented_rows)

# Save final file
rubist_aug.to_csv("COMP0173_Temp_Data/rubist_aug_second.csv", index=False)

```

0%| 1/977 [00:04<1:17:00, 4.73s/it]

[Quota] LLM=15.656423 | GPU=0.0 | Used=15.656423/50.0 | Remaining=34.343576999999996 | Usage=31.312846%

0%| 2/977 [00:07<59:22, 3.65s/it]

[Quota] LLM=15.660893 | GPU=0.0 | Used=15.660893/50.0 | Remaining=34.339107 | Usage=31.321786000000003%

0%| 3/977 [00:10<51:20, 3.16s/it]

[Quota] LLM=15.665231 | GPU=0.0 | Used=15.665231/50.0 | Remaining=34.334769 | Usage=31.330462000000004%

0%| 4/977 [00:13<50:23, 3.11s/it]

[Quota] LLM=15.669602 | GPU=0.0 | Used=15.669602/50.0 | Remaining=34.330398 | Usage=31.339204%

1%| 5/977 [00:15<46:17, 2.86s/it]

[Quota] LLM=15.673346 | GPU=0.0 | Used=15.673346/50.0 | Remaining=34.326654 | Usage=31.34669199999997%

1%| 6/977 [00:18<45:23, 2.80s/it]

[Quota] LLM=15.67721 | GPU=0.0 | Used=15.67721/50.0 | Remaining=34.32279 | Usage=31.35442%

1%| 7/977 [00:21<47:24, 2.93s/it]

[Quota] LLM=15.681173 | GPU=0.0 | Used=15.681173/50.0 | Remaining=34.318827 | Usage=31.36234599999995%

1%| 8/977 [00:23<44:33, 2.76s/it]

```
[Quota] LLM=15.685154 | GPU=0.0 | Used=15.685154/50.0 | Remaining=34.314846 | Usage=31.370308%
1%|           | 9/977 [00:26<42:45, 2.65s/it]

[Quota] LLM=15.689054 | GPU=0.0 | Used=15.689054/50.0 | Remaining=34.310946 | Usage=31.37810799999997%
1%|           | 10/977 [00:28<39:23, 2.44s/it]

[Quota] LLM=15.692873 | GPU=0.0 | Used=15.692873/50.0 | Remaining=34.307127 | Usage=31.385746000000005%
1%|           | 11/977 [00:30<39:31, 2.46s/it]

[Quota] LLM=15.696632 | GPU=0.0 | Used=15.696632/50.0 | Remaining=34.303368 | Usage=31.393264%
1%|           | 12/977 [00:33<39:57, 2.48s/it]

[Quota] LLM=15.700634 | GPU=0.0 | Used=15.700634/50.0 | Remaining=34.299366 | Usage=31.401268000000005%
1%|           | 13/977 [00:38<53:28, 3.33s/it]

[Quota] LLM=15.705443 | GPU=0.0 | Used=15.705443/50.0 | Remaining=34.294557 | Usage=31.410885999999998%
1%|           | 14/977 [00:41<50:53, 3.17s/it]

[Quota] LLM=15.709424 | GPU=0.0 | Used=15.709424/50.0 | Remaining=34.290576 | Usage=31.418848%
2%|           | 15/977 [00:44<51:22, 3.20s/it]

[Quota] LLM=15.713426 | GPU=0.0 | Used=15.713426/50.0 | Remaining=34.286574 | Usage=31.426852%
2%|           | 16/977 [00:50<1:02:55, 3.93s/it]

[Quota] LLM=15.717377 | GPU=0.0 | Used=15.717377/50.0 | Remaining=34.282623 | Usage=31.434754000000005%
2%|           | 17/977 [00:52<55:47, 3.49s/it]

[Quota] LLM=15.721355 | GPU=0.0 | Used=15.721355/50.0 | Remaining=34.278645 | Usage=31.44271%
2%|           | 18/977 [00:55<52:27, 3.28s/it]

[Quota] LLM=15.725138 | GPU=0.0 | Used=15.725138/50.0 | Remaining=34.274862 | Usage=31.450276%
2%|           | 19/977 [00:58<48:39, 3.05s/it]

[Quota] LLM=15.72917 | GPU=0.0 | Used=15.72917/50.0 | Remaining=34.270830000000004 | Usage=31.45834%
2%|           | 20/977 [01:01<49:11, 3.08s/it]
```

[Quota] LLM=15.733589 | GPU=0.0 | Used=15.733589/50.0 | Remaining=34.266411 | Usage=31.467178%

2%| 21/977 [01:05<56:36, 3.55s/it]

[Quota] LLM=15.738242 | GPU=0.0 | Used=15.738242/50.0 | Remaining=34.261758 | Usage=31.476484%

2%| 22/977 [01:08<53:03, 3.33s/it]

[Quota] LLM=15.74264 | GPU=0.0 | Used=15.74264/50.0 | Remaining=34.25736 | Usage=31.48528%

2%| 23/977 [01:11<50:28, 3.18s/it]

[Quota] LLM=15.747089 | GPU=0.0 | Used=15.747089/50.0 | Remaining=34.252911 | Usage=31.494178%

2%| 24/977 [01:13<46:44, 2.94s/it]

[Quota] LLM=15.751235 | GPU=0.0 | Used=15.751235/50.0 | Remaining=34.248765 | Usage=31.50247%

3%| 25/977 [01:16<45:40, 2.88s/it]

[Quota] LLM=15.755249 | GPU=0.0 | Used=15.755249/50.0 | Remaining=34.244751 | Usage=31.510498%

3%| 26/977 [01:20<49:08, 3.10s/it]

[Quota] LLM=15.759386 | GPU=0.0 | Used=15.759386/50.0 | Remaining=34.240614 | Usage=31.518772000000002%

3%| 27/977 [01:22<46:54, 2.96s/it]

[Quota] LLM=15.763583 | GPU=0.0 | Used=15.763583/50.0 | Remaining=34.236417 | Usage=31.527166%

3%| 28/977 [01:25<46:19, 2.93s/it]

[Quota] LLM=15.767852 | GPU=0.0 | Used=15.767852/50.0 | Remaining=34.232148 | Usage=31.535704%

3%| 29/977 [01:29<51:42, 3.27s/it]

[Quota] LLM=15.772178 | GPU=0.0 | Used=15.772178/50.0 | Remaining=34.227822 | Usage=31.544356%

3%| 30/977 [01:32<49:37, 3.14s/it]

[Quota] LLM=15.776744 | GPU=0.0 | Used=15.776744/50.0 | Remaining=34.223256 | Usage=31.553488%

3%| 31/977 [01:34<45:09, 2.86s/it]

[Quota] LLM=15.780764 | GPU=0.0 | Used=15.780764/50.0 | Remaining=34.219236 | Usage=31.561528%

3%| 32/977 [01:38<48:07, 3.06s/it]

[Quota] LLM=15.785156 | GPU=0.0 | Used=15.785156/50.0 | Remaining=34.214844 | Usage=31.570312%

3% | 33/977 [01:41<47:56, 3.05s/it]

[Quota] LLM=15.789254 | GPU=0.0 | Used=15.789254/50.0 | Remaining=34.210746 | Usage=31.578508%

3% | 34/977 [01:44<47:19, 3.01s/it]

[Quota] LLM=15.793718 | GPU=0.0 | Used=15.793718/50.0 | Remaining=34.206282 | Usage=31.587436000000004%

4% | 35/977 [01:47<46:08, 2.94s/it]

[Quota] LLM=15.797543 | GPU=0.0 | Used=15.797543/50.0 | Remaining=34.202457 | Usage=31.595086%

4% | 36/977 [01:49<43:53, 2.80s/it]

[Quota] LLM=15.80162 | GPU=0.0 | Used=15.80162/50.0 | Remaining=34.19838 | Usage=31.60324%

4% | 37/977 [01:51<41:43, 2.66s/it]

[Quota] LLM=15.805505 | GPU=0.0 | Used=15.805505/50.0 | Remaining=34.194495 | Usage=31.61101%

4% | 38/977 [01:54<39:45, 2.54s/it]

[Quota] LLM=15.809459 | GPU=0.0 | Used=15.809459/50.0 | Remaining=34.190540999999996 | Usage=31.618918%

4% | 39/977 [01:56<39:37, 2.53s/it]

[Quota] LLM=15.813428 | GPU=0.0 | Used=15.813428/50.0 | Remaining=34.186572 | Usage=31.626855999999997%

4% | 40/977 [01:59<38:35, 2.47s/it]

[Quota] LLM=15.817346 | GPU=0.0 | Used=15.817346/50.0 | Remaining=34.182654 | Usage=31.634692000000005%

4% | 41/977 [02:01<38:46, 2.49s/it]

[Quota] LLM=15.821549 | GPU=0.0 | Used=15.821549/50.0 | Remaining=34.178451 | Usage=31.643098%

4% | 42/977 [02:03<37:58, 2.44s/it]

[Quota] LLM=15.825449 | GPU=0.0 | Used=15.825449/50.0 | Remaining=34.174551 | Usage=31.65089799999998%

4% | 43/977 [02:06<37:48, 2.43s/it]

[Quota] LLM=15.829286 | GPU=0.0 | Used=15.829286/50.0 | Remaining=34.170714000000004 | Usage=31.658572000000003%

5% | 44/977 [02:08<37:44, 2.43s/it]

```
[Quota] LLM=15.833312 | GPU=0.0 | Used=15.833312/50.0 | Remaining=34.166688 |  
Usage=31.666624%  
5%|           | 45/977 [02:10<36:55,  2.38s/it]  
  
[Quota] LLM=15.837017 | GPU=0.0 | Used=15.837017/50.0 | Remaining=34.162983 |  
Usage=31.674034000000002%  
5%|           | 46/977 [02:14<40:03,  2.58s/it]  
  
[Quota] LLM=15.840908 | GPU=0.0 | Used=15.840908/50.0 | Remaining=34.159092 |  
Usage=31.681816%  
5%|           | 47/977 [02:16<40:30,  2.61s/it]  
  
[Quota] LLM=15.844859 | GPU=0.0 | Used=15.844859/50.0 | Remaining=34.155141 |  
Usage=31.689718%  
5%|           | 48/977 [02:19<40:44,  2.63s/it]  
  
[Quota] LLM=15.849122 | GPU=0.0 | Used=15.849122/50.0 | Remaining=34.150878 |  
Usage=31.698244%  
5%|           | 49/977 [02:21<40:25,  2.61s/it]  
  
[Quota] LLM=15.853121 | GPU=0.0 | Used=15.853121/50.0 | Remaining=34.146879 |  
Usage=31.706242%  
5%|           | 50/977 [02:25<43:03,  2.79s/it]  
  
[Quota] LLM=15.857318 | GPU=0.0 | Used=15.857318/50.0 | Remaining=34.142682 |  
Usage=31.71463599999995%  
5%|           | 51/977 [02:27<40:15,  2.61s/it]  
  
[Quota] LLM=15.861293 | GPU=0.0 | Used=15.861293/50.0 | Remaining=34.138707 |  
Usage=31.72258599999996%  
5%|           | 52/977 [02:30<40:47,  2.65s/it]  
  
[Quota] LLM=15.865364 | GPU=0.0 | Used=15.865364/50.0 | Remaining=34.134636 |  
Usage=31.73072799999996%  
5%|           | 53/977 [02:33<43:00,  2.79s/it]  
  
[Quota] LLM=15.869381 | GPU=0.0 | Used=15.869381/50.0 |  
Remaining=34.13061899999996 | Usage=31.738762000000005%  
6%|           | 54/977 [02:35<40:18,  2.62s/it]  
  
[Quota] LLM=15.873278 | GPU=0.0 | Used=15.873278/50.0 | Remaining=34.126722 |  
Usage=31.746556000000005%  
6%|           | 55/977 [02:38<40:37,  2.64s/it]  
  
[Quota] LLM=15.877229 | GPU=0.0 | Used=15.877229/50.0 | Remaining=34.122771 |  
Usage=31.754458000000003%  
6%|           | 56/977 [02:41<45:27,  2.96s/it]
```

```
[Quota] LLM=15.881117 | GPU=0.0 | Used=15.881117/50.0 | Remaining=34.118883 | Usage=31.762234%  
6%|           | 57/977 [02:44<42:20,  2.76s/it]  
  
[Quota] LLM=15.884849 | GPU=0.0 | Used=15.884849/50.0 | Remaining=34.115151 | Usage=31.769698%  
6%|           | 58/977 [02:46<41:24,  2.70s/it]  
  
[Quota] LLM=15.888677 | GPU=0.0 | Used=15.888677/50.0 | Remaining=34.111323 | Usage=31.777353999999995%  
6%|           | 59/977 [02:51<50:54,  3.33s/it]  
  
[Quota] LLM=15.892709 | GPU=0.0 | Used=15.892709/50.0 | Remaining=34.107291000000004 | Usage=31.785418%  
6%|           | 60/977 [02:54<48:23,  3.17s/it]  
  
[Quota] LLM=15.896633 | GPU=0.0 | Used=15.896633/50.0 | Remaining=34.103367 | Usage=31.793266%  
6%|           | 61/977 [02:56<44:01,  2.88s/it]  
  
[Quota] LLM=15.900527 | GPU=0.0 | Used=15.900527/50.0 | Remaining=34.099473 | Usage=31.801054%  
6%|           | 62/977 [03:01<54:56,  3.60s/it]  
  
[Quota] LLM=15.904358 | GPU=0.0 | Used=15.904358/50.0 | Remaining=34.095642 | Usage=31.80871599999997%  
6%|           | 63/977 [03:04<51:57,  3.41s/it]  
  
[Quota] LLM=15.908444 | GPU=0.0 | Used=15.908444/50.0 | Remaining=34.091556 | Usage=31.816888%  
7%|           | 64/977 [03:07<49:32,  3.26s/it]  
  
[Quota] LLM=15.912542 | GPU=0.0 | Used=15.912542/50.0 | Remaining=34.087458 | Usage=31.82508399999997%  
7%|           | 65/977 [03:11<49:56,  3.29s/it]  
  
[Quota] LLM=15.916358 | GPU=0.0 | Used=15.916358/50.0 | Remaining=34.083642 | Usage=31.832716%  
7%|           | 66/977 [03:14<49:37,  3.27s/it]  
  
[Quota] LLM=15.920585 | GPU=0.0 | Used=15.920585/50.0 | Remaining=34.079415 | Usage=31.84117%  
7%|           | 67/977 [03:16<46:21,  3.06s/it]  
  
[Quota] LLM=15.924665 | GPU=0.0 | Used=15.924665/50.0 | Remaining=34.075335 | Usage=31.84933%  
7%|           | 68/977 [03:19<44:40,  2.95s/it]
```

```
[Quota] LLM=15.92894 | GPU=0.0 | Used=15.92894/50.0 | Remaining=34.07106 |
Usage=31.857879999999998%
 7%|           | 69/977 [03:22<43:51,  2.90s/it]

[Quota] LLM=15.93323 | GPU=0.0 | Used=15.93323/50.0 | Remaining=34.06677 |
Usage=31.866460000000004%
 7%|           | 70/977 [03:25<44:57,  2.97s/it]

[Quota] LLM=15.937133 | GPU=0.0 | Used=15.937133/50.0 | Remaining=34.062867 |
Usage=31.874266000000002%
 7%|           | 71/977 [03:27<42:24,  2.81s/it]

[Quota] LLM=15.940862 | GPU=0.0 | Used=15.940862/50.0 | Remaining=34.059138000000004 |
Usage=31.881724%
 7%|           | 72/977 [03:32<48:36,  3.22s/it]

[Quota] LLM=15.945677 | GPU=0.0 | Used=15.945677/50.0 | Remaining=34.054323 |
Usage=31.891354%
 7%|           | 73/977 [03:35<51:24,  3.41s/it]

[Quota] LLM=15.950186 | GPU=0.0 | Used=15.950186/50.0 | Remaining=34.049814 |
Usage=31.90037199999997%
 8%|           | 74/977 [03:40<55:05,  3.66s/it]

[Quota] LLM=15.95477 | GPU=0.0 | Used=15.95477/50.0 | Remaining=34.045230000000004 |
Usage=31.909539999999996%
 8%|           | 75/977 [03:42<50:09,  3.34s/it]

[Quota] LLM=15.958637 | GPU=0.0 | Used=15.958637/50.0 | Remaining=34.041363000000004 |
Usage=31.917274000000003%
 8%|           | 76/977 [03:45<46:29,  3.10s/it]

[Quota] LLM=15.962786 | GPU=0.0 | Used=15.962786/50.0 | Remaining=34.037214 |
Usage=31.92557199999995%
 8%|           | 77/977 [03:49<50:26,  3.36s/it]

[Quota] LLM=15.966794 | GPU=0.0 | Used=15.966794/50.0 | Remaining=34.033206 |
Usage=31.933588%
 8%|           | 78/977 [03:52<49:27,  3.30s/it]

[Quota] LLM=15.97097 | GPU=0.0 | Used=15.97097/50.0 | Remaining=34.02903 |
Usage=31.94193999999995%
 8%|           | 79/977 [03:55<46:28,  3.10s/it]

[Quota] LLM=15.97526 | GPU=0.0 | Used=15.97526/50.0 | Remaining=34.02474 |
Usage=31.95051999999997%
 8%|           | 80/977 [03:58<46:01,  3.08s/it]
```

```
[Quota] LLM=15.97937 | GPU=0.0 | Used=15.97937/50.0 | Remaining=34.02063 |  
Usage=31.958739999999995%  
8%|           | 81/977 [04:00<44:27,  2.98s/it]  
  
[Quota] LLM=15.983537 | GPU=0.0 | Used=15.983537/50.0 | Remaining=34.016463 |  
Usage=31.967074%  
8%|           | 82/977 [04:03<41:53,  2.81s/it]  
  
[Quota] LLM=15.987455 | GPU=0.0 | Used=15.987455/50.0 | Remaining=34.012545 |  
Usage=31.97491%  
8%|           | 83/977 [04:05<40:16,  2.70s/it]  
  
[Quota] LLM=15.991919 | GPU=0.0 | Used=15.991919/50.0 |  
Remaining=34.008081000000004 | Usage=31.98383%  
9%|           | 84/977 [04:07<36:38,  2.46s/it]  
  
[Quota] LLM=15.995807 | GPU=0.0 | Used=15.995807/50.0 | Remaining=34.004193 |  
Usage=31.991613999999995%  
9%|           | 85/977 [04:12<47:24,  3.19s/it]  
  
[Quota] LLM=16.00001 | GPU=0.0 | Used=16.00001/50.0 | Remaining=33.99999 |  
Usage=32.00002%  
9%|           | 86/977 [04:14<43:08,  2.91s/it]  
  
[Quota] LLM=16.003781 | GPU=0.0 | Used=16.003781/50.0 |  
Remaining=33.996218999999996 | Usage=32.007562%  
9%|           | 87/977 [04:16<40:23,  2.72s/it]  
  
[Quota] LLM=16.00772 | GPU=0.0 | Used=16.00772/50.0 | Remaining=33.99228 |  
Usage=32.01544%  
9%|           | 88/977 [04:19<40:05,  2.71s/it]  
  
[Quota] LLM=16.011887 | GPU=0.0 | Used=16.011887/50.0 | Remaining=33.988113 |  
Usage=32.023774%  
9%|           | 89/977 [04:22<39:06,  2.64s/it]  
  
[Quota] LLM=16.016054 | GPU=0.0 | Used=16.016054/50.0 | Remaining=33.983946 |  
Usage=32.032108%  
9%|           | 90/977 [04:25<41:08,  2.78s/it]  
  
[Quota] LLM=16.020599 | GPU=0.0 | Used=16.020599/50.0 |  
Remaining=33.979400999999996 | Usage=32.041198%  
9%|           | 91/977 [04:27<39:15,  2.66s/it]  
  
[Quota] LLM=16.024529 | GPU=0.0 | Used=16.024529/50.0 | Remaining=33.975471 |  
Usage=32.049058%  
9%|           | 92/977 [04:30<39:45,  2.70s/it]
```

```
[Quota] LLM=16.028699 | GPU=0.0 | Used=16.028699/50.0 | Remaining=33.971301 |
Usage=32.057398%
10%|          | 93/977 [04:34<45:56,  3.12s/it]

[Quota] LLM=16.032437 | GPU=0.0 | Used=16.032437/50.0 | Remaining=33.967563 |
Usage=32.064874%
10%|          | 94/977 [04:37<46:52,  3.18s/it]

[Quota] LLM=16.03634 | GPU=0.0 | Used=16.03634/50.0 |
Remaining=33.963660000000004 | Usage=32.07268%
10%|          | 95/977 [04:44<1:00:24,  4.11s/it]

[Quota] LLM=16.040387 | GPU=0.0 | Used=16.040387/50.0 |
Remaining=33.959613000000004 | Usage=32.080774%
10%|          | 96/977 [04:48<1:00:35,  4.13s/it]

[Quota] LLM=16.044611 | GPU=0.0 | Used=16.044611/50.0 | Remaining=33.955389 |
Usage=32.089222%
10%|          | 97/977 [04:54<1:08:39,  4.68s/it]

[Quota] LLM=16.048781 | GPU=0.0 | Used=16.048781/50.0 |
Remaining=33.951218999999995 | Usage=32.097562%
10%|          | 98/977 [04:56<58:42,  4.01s/it]

[Quota] LLM=16.052705 | GPU=0.0 | Used=16.052705/50.0 | Remaining=33.947295 |
Usage=32.10541%
10%|          | 99/977 [04:59<53:22,  3.65s/it]

[Quota] LLM=16.05686 | GPU=0.0 | Used=16.05686/50.0 | Remaining=33.94314 |
Usage=32.11372%
10%|          | 100/977 [05:01<48:12,  3.30s/it]

[Quota] LLM=16.060898 | GPU=0.0 | Used=16.060898/50.0 | Remaining=33.939102 |
Usage=32.121796%
10%|          | 101/977 [05:07<56:26,  3.87s/it]

[Quota] LLM=16.066322 | GPU=0.0 | Used=16.066322/50.0 | Remaining=33.933678 |
Usage=32.132644%
10%|          | 102/977 [05:10<53:25,  3.66s/it]

[Quota] LLM=16.070471 | GPU=0.0 | Used=16.070471/50.0 | Remaining=33.929529 |
Usage=32.140942%
11%|          | 103/977 [05:13<51:22,  3.53s/it]

[Quota] LLM=16.074521 | GPU=0.0 | Used=16.074521/50.0 |
Remaining=33.925478999999996 | Usage=32.149042%
11%|          | 104/977 [05:15<44:14,  3.04s/it]
```

```
[Quota] LLM=16.078289 | GPU=0.0 | Used=16.078289/50.0 | Remaining=33.921711 |
Usage=32.156578%
11%|           | 105/977 [05:18<45:17,  3.12s/it]

[Quota] LLM=16.08275 | GPU=0.0 | Used=16.08275/50.0 |
Remaining=33.917249999999996 | Usage=32.1655%
11%|           | 106/977 [05:22<50:03,  3.45s/it]

[Quota] LLM=16.087265 | GPU=0.0 | Used=16.087265/50.0 | Remaining=33.912735 |
Usage=32.17453%
11%|           | 107/977 [05:25<45:58,  3.17s/it]

[Quota] LLM=16.091336 | GPU=0.0 | Used=16.091336/50.0 | Remaining=33.908664 |
Usage=32.182672%
11%|           | 108/977 [05:28<43:12,  2.98s/it]

[Quota] LLM=16.095521 | GPU=0.0 | Used=16.095521/50.0 |
Remaining=33.904478999999995 | Usage=32.191042%
11%|           | 109/977 [05:31<46:38,  3.22s/it]

[Quota] LLM=16.099817 | GPU=0.0 | Used=16.099817/50.0 | Remaining=33.900183 |
Usage=32.199634%
11%|           | 110/977 [05:34<44:26,  3.08s/it]

[Quota] LLM=16.104302 | GPU=0.0 | Used=16.104302/50.0 |
Remaining=33.895697999999996 | Usage=32.208604%
11%|           | 111/977 [05:36<40:42,  2.82s/it]

[Quota] LLM=16.10855 | GPU=0.0 | Used=16.10855/50.0 | Remaining=33.89145 |
Usage=32.2171%
11%|           | 112/977 [05:41<49:42,  3.45s/it]

[Quota] LLM=16.112882 | GPU=0.0 | Used=16.112882/50.0 | Remaining=33.887118 |
Usage=32.225764%
12%|           | 113/977 [05:44<46:16,  3.21s/it]

[Quota] LLM=16.11689 | GPU=0.0 | Used=16.11689/50.0 | Remaining=33.88311 |
Usage=32.23378%
12%|           | 114/977 [05:48<50:14,  3.49s/it]

[Quota] LLM=16.121606 | GPU=0.0 | Used=16.121606/50.0 | Remaining=33.878394 |
Usage=32.243212%
12%|           | 115/977 [05:50<44:13,  3.08s/it]

[Quota] LLM=16.125545 | GPU=0.0 | Used=16.125545/50.0 | Remaining=33.874455 |
Usage=32.25109%
12%|           | 116/977 [05:53<41:30,  2.89s/it]
```

```
[Quota] LLM=16.129538 | GPU=0.0 | Used=16.129538/50.0 | Remaining=33.870462 |
Usage=32.259076%
12%|           | 117/977 [05:56<41:39,  2.91s/it]

[Quota] LLM=16.1339 | GPU=0.0 | Used=16.1339/50.0 | Remaining=33.8661 |
Usage=32.2678%
12%|           | 118/977 [05:58<40:16,  2.81s/it]

[Quota] LLM=16.138061 | GPU=0.0 | Used=16.138061/50.0 | Remaining=33.861939 |
Usage=32.276122%
12%|           | 119/977 [06:01<38:25,  2.69s/it]

[Quota] LLM=16.142177 | GPU=0.0 | Used=16.142177/50.0 | Remaining=33.857822999999996 |
Usage=32.284354%
12%|           | 120/977 [06:03<37:32,  2.63s/it]

[Quota] LLM=16.146407 | GPU=0.0 | Used=16.146407/50.0 | Remaining=33.853593000000004 |
Usage=32.292814%
12%|           | 121/977 [06:06<37:42,  2.64s/it]

[Quota] LLM=16.150709 | GPU=0.0 | Used=16.150709/50.0 | Remaining=33.849291 |
Usage=32.301418%
12%|           | 122/977 [06:08<38:10,  2.68s/it]

[Quota] LLM=16.15463 | GPU=0.0 | Used=16.15463/50.0 | Remaining=33.84537 |
Usage=32.30926%
13%|           | 123/977 [06:12<43:51,  3.08s/it]

[Quota] LLM=16.158983 | GPU=0.0 | Used=16.158983/50.0 | Remaining=33.841017 |
Usage=32.317966%
13%|           | 124/977 [06:17<48:17,  3.40s/it]

[Quota] LLM=16.163474 | GPU=0.0 | Used=16.163474/50.0 | Remaining=33.836526 |
Usage=32.326948%
13%|           | 125/977 [06:19<45:20,  3.19s/it]

[Quota] LLM=16.167698 | GPU=0.0 | Used=16.167698/50.0 | Remaining=33.832302 |
Usage=32.335396%
13%|           | 126/977 [06:22<42:49,  3.02s/it]

[Quota] LLM=16.171634 | GPU=0.0 | Used=16.171634/50.0 | Remaining=33.828366 |
Usage=32.343268%
13%|           | 127/977 [06:25<44:52,  3.17s/it]

[Quota] LLM=16.175585 | GPU=0.0 | Used=16.175585/50.0 | Remaining=33.824415 |
Usage=32.35117%
13%|           | 128/977 [06:28<40:13,  2.84s/it]
```

```
[Quota] LLM=16.179446 | GPU=0.0 | Used=16.179446/50.0 | Remaining=33.820554 |
Usage=32.358892%
13%|           | 129/977 [06:30<38:44,  2.74s/it]

[Quota] LLM=16.183322 | GPU=0.0 | Used=16.183322/50.0 | Remaining=33.816677999999996 |
Usage=32.366644%
13%|           | 130/977 [06:32<37:24,  2.65s/it]

[Quota] LLM=16.18745 | GPU=0.0 | Used=16.18745/50.0 | Remaining=33.81255 |
Usage=32.3749%
13%|           | 131/977 [06:36<41:24,  2.94s/it]

[Quota] LLM=16.192247 | GPU=0.0 | Used=16.192247/50.0 | Remaining=33.807753000000005 |
Usage=32.384494%
14%|           | 132/977 [06:38<38:37,  2.74s/it]

[Quota] LLM=16.196156 | GPU=0.0 | Used=16.196156/50.0 | Remaining=33.803844 |
Usage=32.392312%
14%|           | 133/977 [06:42<43:13,  3.07s/it]

[Quota] LLM=16.200305 | GPU=0.0 | Used=16.200305/50.0 | Remaining=33.799695 |
Usage=32.40061%
14%|           | 134/977 [06:45<41:11,  2.93s/it]

[Quota] LLM=16.204373 | GPU=0.0 | Used=16.204373/50.0 | Remaining=33.795626999999996 |
Usage=32.408746%
14%|           | 135/977 [06:48<41:13,  2.94s/it]

[Quota] LLM=16.208303 | GPU=0.0 | Used=16.208303/50.0 | Remaining=33.791697 |
Usage=32.416606%
14%|           | 136/977 [06:51<41:02,  2.93s/it]

[Quota] LLM=16.212632 | GPU=0.0 | Used=16.212632/50.0 | Remaining=33.787368 |
Usage=32.425264%
14%|           | 137/977 [06:54<43:39,  3.12s/it]

[Quota] LLM=16.216799 | GPU=0.0 | Used=16.216799/50.0 | Remaining=33.783201 |
Usage=32.433598%
14%|           | 138/977 [06:57<43:25,  3.11s/it]

[Quota] LLM=16.220753 | GPU=0.0 | Used=16.220753/50.0 | Remaining=33.779247 |
Usage=32.441506%
14%|           | 139/977 [07:01<45:31,  3.26s/it]

[Quota] LLM=16.224638 | GPU=0.0 | Used=16.224638/50.0 | Remaining=33.775362 |
Usage=32.449276%
14%|           | 140/977 [07:05<47:39,  3.42s/it]
```

[Quota] LLM=16.228604 | GPU=0.0 | Used=16.228604/50.0 | Remaining=33.771395999999996 | Usage=32.457208%
14% | 141/977 [07:08<48:17, 3.47s/it]

[Quota] LLM=16.232501 | GPU=0.0 | Used=16.232501/50.0 | Remaining=33.767499 | Usage=32.465002%
15% | 142/977 [07:12<51:05, 3.67s/it]

[Quota] LLM=16.236338 | GPU=0.0 | Used=16.236338/50.0 | Remaining=33.763662 | Usage=32.472676%
15% | 143/977 [07:16<50:06, 3.60s/it]

[Quota] LLM=16.240436 | GPU=0.0 | Used=16.240436/50.0 | Remaining=33.759564 | Usage=32.480872%
15% | 144/977 [07:19<48:29, 3.49s/it]

[Quota] LLM=16.244417 | GPU=0.0 | Used=16.244417/50.0 | Remaining=33.755583 | Usage=32.488834%
15% | 145/977 [07:23<51:08, 3.69s/it]

[Quota] LLM=16.248368 | GPU=0.0 | Used=16.248368/50.0 | Remaining=33.751632 | Usage=32.496736%
15% | 146/977 [07:26<45:42, 3.30s/it]

[Quota] LLM=16.252094 | GPU=0.0 | Used=16.252094/50.0 | Remaining=33.747906 | Usage=32.504188%
15% | 147/977 [07:29<43:52, 3.17s/it]

[Quota] LLM=16.255847 | GPU=0.0 | Used=16.255847/50.0 | Remaining=33.744153 | Usage=32.511694%
15% | 148/977 [07:32<45:48, 3.32s/it]

[Quota] LLM=16.259897 | GPU=0.0 | Used=16.259897/50.0 | Remaining=33.740103000000005 | Usage=32.519794%
15% | 149/977 [07:36<46:37, 3.38s/it]

[Quota] LLM=16.263977 | GPU=0.0 | Used=16.263977/50.0 | Remaining=33.736023 | Usage=32.527954%
15% | 150/977 [07:38<43:01, 3.12s/it]

[Quota] LLM=16.267745 | GPU=0.0 | Used=16.267745/50.0 | Remaining=33.732254999999995 | Usage=32.53549%
15% | 151/977 [07:41<40:58, 2.98s/it]

[Quota] LLM=16.27175 | GPU=0.0 | Used=16.27175/50.0 | Remaining=33.72825 | Usage=32.5435%
16% | 152/977 [07:44<42:24, 3.08s/it]

[Quota] LLM=16.275704 | GPU=0.0 | Used=16.275704/50.0 | Remaining=33.724295999999995 | Usage=32.551408%
16%| 153/977 [07:47<42:50, 3.12s/it]

[Quota] LLM=16.279868 | GPU=0.0 | Used=16.279868/50.0 | Remaining=33.720132 | Usage=32.559736%
16%| 154/977 [07:50<41:18, 3.01s/it]

[Quota] LLM=16.283621 | GPU=0.0 | Used=16.283621/50.0 | Remaining=33.716379 | Usage=32.567242%
16%| 155/977 [07:54<45:30, 3.32s/it]

[Quota] LLM=16.288166 | GPU=0.0 | Used=16.288166/50.0 | Remaining=33.71183399999996 | Usage=32.576332%
16%| 156/977 [07:58<45:55, 3.36s/it]

[Quota] LLM=16.292141 | GPU=0.0 | Used=16.292141/50.0 | Remaining=33.707859 | Usage=32.584282%
16%| 157/977 [08:02<51:46, 3.79s/it]

[Quota] LLM=16.296554 | GPU=0.0 | Used=16.296554/50.0 | Remaining=33.703446 | Usage=32.593108%
16%| 158/977 [08:06<50:46, 3.72s/it]

[Quota] LLM=16.300424 | GPU=0.0 | Used=16.300424/50.0 | Remaining=33.699576 | Usage=32.600848%
16%| 159/977 [08:09<49:04, 3.60s/it]

[Quota] LLM=16.304387 | GPU=0.0 | Used=16.304387/50.0 | Remaining=33.695613 | Usage=32.608774%
16%| 160/977 [08:11<41:46, 3.07s/it]

[Quota] LLM=16.308176 | GPU=0.0 | Used=16.308176/50.0 | Remaining=33.691824 | Usage=32.616352%
16%| 161/977 [08:14<41:33, 3.06s/it]

[Quota] LLM=16.312187 | GPU=0.0 | Used=16.312187/50.0 | Remaining=33.687813 | Usage=32.624374%
17%| 162/977 [08:18<45:12, 3.33s/it]

[Quota] LLM=16.316531 | GPU=0.0 | Used=16.316531/50.0 | Remaining=33.683469 | Usage=32.633062%
17%| 163/977 [08:21<43:58, 3.24s/it]

[Quota] LLM=16.321073 | GPU=0.0 | Used=16.321073/50.0 | Remaining=33.678927 | Usage=32.642146%
17%| 164/977 [08:25<47:31, 3.51s/it]

```
[Quota] LLM=16.325282 | GPU=0.0 | Used=16.325282/50.0 | Remaining=33.674718 |
Usage=32.650564%
17%|           | 165/977 [08:28<46:04,  3.40s/it]

[Quota] LLM=16.329098 | GPU=0.0 | Used=16.329098/50.0 | Remaining=33.670902 |
Usage=32.658196%
17%|           | 166/977 [08:32<46:50,  3.47s/it]

[Quota] LLM=16.33379 | GPU=0.0 | Used=16.33379/50.0 | Remaining=33.66621 |
Usage=32.66758%
17%|           | 167/977 [08:35<44:33,  3.30s/it]

[Quota] LLM=16.337768 | GPU=0.0 | Used=16.337768/50.0 | Remaining=33.662232 |
Usage=32.675536%
17%|           | 168/977 [08:39<45:35,  3.38s/it]

[Quota] LLM=16.341782 | GPU=0.0 | Used=16.341782/50.0 | Remaining=33.658218000000005 |
Usage=32.683564%
17%|           | 169/977 [08:41<41:22,  3.07s/it]

[Quota] LLM=16.345685 | GPU=0.0 | Used=16.345685/50.0 | Remaining=33.654315 |
Usage=32.69137%
17%|           | 170/977 [08:43<37:29,  2.79s/it]

[Quota] LLM=16.349879 | GPU=0.0 | Used=16.349879/50.0 | Remaining=33.650121 |
Usage=32.699758%
18%|           | 171/977 [08:46<39:15,  2.92s/it]

[Quota] LLM=16.353956 | GPU=0.0 | Used=16.353956/50.0 | Remaining=33.646044 |
Usage=32.707912%
18%|           | 172/977 [08:49<36:36,  2.73s/it]

[Quota] LLM=16.357754 | GPU=0.0 | Used=16.357754/50.0 | Remaining=33.642246 |
Usage=32.715508%
18%|           | 173/977 [08:52<38:40,  2.89s/it]

[Quota] LLM=16.361516 | GPU=0.0 | Used=16.361516/50.0 | Remaining=33.638484 |
Usage=32.723032%
18%|           | 174/977 [08:55<38:12,  2.86s/it]

[Quota] LLM=16.365452 | GPU=0.0 | Used=16.365452/50.0 |
Remaining=33.634547999999995 | Usage=32.730904%
18%|           | 175/977 [08:57<36:52,  2.76s/it]

[Quota] LLM=16.36973 | GPU=0.0 | Used=16.36973/50.0 |
Remaining=33.630269999999996 | Usage=32.73946%
18%|           | 176/977 [09:00<36:10,  2.71s/it]
```

[Quota] LLM=16.373576 | GPU=0.0 | Used=16.373576/50.0 | Remaining=33.626424 | Usage=32.747152%

18%| 177/977 [09:03<38:07, 2.86s/it]

[Quota] LLM=16.377836 | GPU=0.0 | Used=16.377836/50.0 | Remaining=33.622164 | Usage=32.755672%

18%| 178/977 [09:05<35:43, 2.68s/it]

[Quota] LLM=16.381769 | GPU=0.0 | Used=16.381769/50.0 | Remaining=33.618231 | Usage=32.763538%

18%| 179/977 [09:08<35:46, 2.69s/it]

[Quota] LLM=16.385555 | GPU=0.0 | Used=16.385555/50.0 | Remaining=33.614445 | Usage=32.77111%

18%| 180/977 [09:12<39:25, 2.97s/it]

[Quota] LLM=16.390481 | GPU=0.0 | Used=16.390481/50.0 | Remaining=33.609519 | Usage=32.780962%

19%| 181/977 [09:15<40:46, 3.07s/it]

[Quota] LLM=16.394678 | GPU=0.0 | Used=16.394678/50.0 | Remaining=33.605322 | Usage=32.789356%

19%| 182/977 [09:18<41:18, 3.12s/it]

[Quota] LLM=16.398929 | GPU=0.0 | Used=16.398929/50.0 | Remaining=33.601071000000005 | Usage=32.797858%

19%| 183/977 [09:21<38:33, 2.91s/it]

[Quota] LLM=16.402712 | GPU=0.0 | Used=16.402712/50.0 | Remaining=33.597288 | Usage=32.805424%

19%| 184/977 [09:25<43:32, 3.29s/it]

[Quota] LLM=16.406663 | GPU=0.0 | Used=16.406663/50.0 | Remaining=33.593337 | Usage=32.813326%

19%| 185/977 [09:28<43:10, 3.27s/it]

[Quota] LLM=16.411403 | GPU=0.0 | Used=16.411403/50.0 | Remaining=33.588597 | Usage=32.822806%

19%| 186/977 [09:30<38:05, 2.89s/it]

[Quota] LLM=16.415372 | GPU=0.0 | Used=16.415372/50.0 | Remaining=33.584627999999995 | Usage=32.830744%

19%| 187/977 [09:33<36:58, 2.81s/it]

[Quota] LLM=16.419158 | GPU=0.0 | Used=16.419158/50.0 | Remaining=33.580842000000004 | Usage=32.838316%

19%| 188/977 [09:35<34:43, 2.64s/it]

```
[Quota] LLM=16.423244 | GPU=0.0 | Used=16.423244/50.0 | Remaining=33.576756 |
Usage=32.846488%
19%|           | 189/977 [09:37<33:14,  2.53s/it]

[Quota] LLM=16.427282 | GPU=0.0 | Used=16.427282/50.0 | Remaining=33.572717999999995 |
Usage=32.854564%
19%|           | 190/977 [09:40<33:00,  2.52s/it]

[Quota] LLM=16.431296 | GPU=0.0 | Used=16.431296/50.0 | Remaining=33.568704 |
Usage=32.862592%
20%|           | 191/977 [09:42<33:09,  2.53s/it]

[Quota] LLM=16.43522 | GPU=0.0 | Used=16.43522/50.0 | Remaining=33.56478 |
Usage=32.87044%
20%|           | 192/977 [09:45<32:48,  2.51s/it]

[Quota] LLM=16.439048 | GPU=0.0 | Used=16.439048/50.0 | Remaining=33.560952 |
Usage=32.878096%
20%|           | 193/977 [09:47<32:13,  2.47s/it]

[Quota] LLM=16.442969 | GPU=0.0 | Used=16.442969/50.0 | Remaining=33.557030999999995 |
Usage=32.885938%
20%|           | 194/977 [09:50<32:55,  2.52s/it]

[Quota] LLM=16.446887 | GPU=0.0 | Used=16.446887/50.0 | Remaining=33.553112999999996 |
Usage=32.893774%
20%|           | 195/977 [09:52<32:28,  2.49s/it]

[Quota] LLM=16.450862 | GPU=0.0 | Used=16.450862/50.0 | Remaining=33.549138 |
Usage=32.901724%
20%|           | 196/977 [09:55<35:12,  2.71s/it]

[Quota] LLM=16.455173 | GPU=0.0 | Used=16.455173/50.0 | Remaining=33.544827 |
Usage=32.910346%
20%|           | 197/977 [09:58<35:59,  2.77s/it]

[Quota] LLM=16.459475 | GPU=0.0 | Used=16.459475/50.0 | Remaining=33.540525 |
Usage=32.91895%
20%|           | 198/977 [10:01<36:08,  2.78s/it]

[Quota] LLM=16.463582 | GPU=0.0 | Used=16.463582/50.0 | Remaining=33.536418 |
Usage=32.927164%
20%|           | 199/977 [10:06<43:08,  3.33s/it]

[Quota] LLM=16.467665 | GPU=0.0 | Used=16.467665/50.0 | Remaining=33.532335 |
Usage=32.93533%
20%|           | 200/977 [10:08<40:30,  3.13s/it]
```

[Quota] LLM=16.47146 | GPU=0.0 | Used=16.47146/50.0 | Remaining=33.52854 | Usage=32.94292%
21% | 201/977 [10:11<38:18, 2.96s/it]

[Quota] LLM=16.475309 | GPU=0.0 | Used=16.475309/50.0 | Remaining=33.524691000000004 | Usage=32.950618%
21% | 202/977 [10:14<37:25, 2.90s/it]

[Quota] LLM=16.479239 | GPU=0.0 | Used=16.479239/50.0 | Remaining=33.520761 | Usage=32.958478%
21% | 203/977 [10:16<37:26, 2.90s/it]

[Quota] LLM=16.483334 | GPU=0.0 | Used=16.483334/50.0 | Remaining=33.516666 | Usage=32.966668%
21% | 204/977 [10:19<34:50, 2.70s/it]

[Quota] LLM=16.48745 | GPU=0.0 | Used=16.48745/50.0 | Remaining=33.51255000000005 | Usage=32.9749%
21% | 205/977 [10:21<34:50, 2.71s/it]

[Quota] LLM=16.491254 | GPU=0.0 | Used=16.491254/50.0 | Remaining=33.508746 | Usage=32.982508%
21% | 206/977 [10:24<34:22, 2.68s/it]

[Quota] LLM=16.495046 | GPU=0.0 | Used=16.495046/50.0 | Remaining=33.504954 | Usage=32.990092%
21% | 207/977 [10:26<33:10, 2.58s/it]

[Quota] LLM=16.498883 | GPU=0.0 | Used=16.498883/50.0 | Remaining=33.501117 | Usage=32.997766%
21% | 208/977 [10:29<31:37, 2.47s/it]

[Quota] LLM=16.502726 | GPU=0.0 | Used=16.502726/50.0 | Remaining=33.49727400000004 | Usage=33.005452%
21% | 209/977 [10:31<30:50, 2.41s/it]

[Quota] LLM=16.506545 | GPU=0.0 | Used=16.506545/50.0 | Remaining=33.493455 | Usage=33.01309%
21% | 210/977 [10:34<34:54, 2.73s/it]

[Quota] LLM=16.510571 | GPU=0.0 | Used=16.510571/50.0 | Remaining=33.489429 | Usage=33.021142%
22% | 211/977 [10:36<32:20, 2.53s/it]

[Quota] LLM=16.514525 | GPU=0.0 | Used=16.514525/50.0 | Remaining=33.485475 | Usage=33.02905%
22% | 212/977 [10:39<31:49, 2.50s/it]

```
[Quota] LLM=16.518515 | GPU=0.0 | Used=16.518515/50.0 | Remaining=33.481485 | Usage=33.03703%
22%|           | 213/977 [10:41<30:25, 2.39s/it]

[Quota] LLM=16.522568 | GPU=0.0 | Used=16.522568/50.0 | Remaining=33.477432 | Usage=33.045136%
22%|           | 214/977 [10:44<31:57, 2.51s/it]

[Quota] LLM=16.526669 | GPU=0.0 | Used=16.526669/50.0 | Remaining=33.473331 | Usage=33.053338%
22%|           | 215/977 [10:47<34:07, 2.69s/it]

[Quota] LLM=16.5308 | GPU=0.0 | Used=16.5308/50.0 | Remaining=33.4692 | Usage=33.0616%
22%|           | 216/977 [10:49<33:45, 2.66s/it]

[Quota] LLM=16.534646 | GPU=0.0 | Used=16.534646/50.0 | Remaining=33.465354000000005 | Usage=33.069292%
22%|           | 217/977 [10:52<32:26, 2.56s/it]

[Quota] LLM=16.538387 | GPU=0.0 | Used=16.538387/50.0 | Remaining=33.461613 | Usage=33.076774%
22%|           | 218/977 [10:55<35:31, 2.81s/it]

[Quota] LLM=16.542722 | GPU=0.0 | Used=16.542722/50.0 | Remaining=33.457278 | Usage=33.085444%
22%|           | 219/977 [10:58<35:56, 2.84s/it]

[Quota] LLM=16.546646 | GPU=0.0 | Used=16.546646/50.0 | Remaining=33.453354000000004 | Usage=33.093292%
23%|           | 220/977 [11:01<35:43, 2.83s/it]

[Quota] LLM=16.550597 | GPU=0.0 | Used=16.550597/50.0 | Remaining=33.449403000000004 | Usage=33.101194%
23%|           | 221/977 [11:03<33:56, 2.69s/it]

[Quota] LLM=16.554374 | GPU=0.0 | Used=16.554374/50.0 | Remaining=33.445626000000004 | Usage=33.108748%
23%|           | 222/977 [11:06<34:43, 2.76s/it]

[Quota] LLM=16.558574 | GPU=0.0 | Used=16.558574/50.0 | Remaining=33.441426 | Usage=33.117148%
23%|           | 223/977 [11:08<31:51, 2.53s/it]

[Quota] LLM=16.562399 | GPU=0.0 | Used=16.562399/50.0 | Remaining=33.437601 | Usage=33.124798%
23%|           | 224/977 [11:13<41:27, 3.30s/it]
```

```
[Quota] LLM=16.566899 | GPU=0.0 | Used=16.566899/50.0 | Remaining=33.433101 | Usage=33.133798%
23%|           | 225/977 [11:16<39:01,  3.11s/it]

[Quota] LLM=16.571066 | GPU=0.0 | Used=16.571066/50.0 | Remaining=33.428934 | Usage=33.142132%
23%|           | 226/977 [11:18<36:47,  2.94s/it]

[Quota] LLM=16.574873 | GPU=0.0 | Used=16.574873/50.0 | Remaining=33.425127 | Usage=33.149746%
23%|           | 227/977 [11:20<33:12,  2.66s/it]

[Quota] LLM=16.578644 | GPU=0.0 | Used=16.578644/50.0 | Remaining=33.421356 | Usage=33.157288%
23%|           | 228/977 [11:23<34:14,  2.74s/it]

[Quota] LLM=16.582895 | GPU=0.0 | Used=16.582895/50.0 | Remaining=33.417105 | Usage=33.16579%
23%|           | 229/977 [11:26<35:22,  2.84s/it]

[Quota] LLM=16.586828 | GPU=0.0 | Used=16.586828/50.0 | Remaining=33.413172 | Usage=33.173656%
24%|           | 230/977 [11:30<37:06,  2.98s/it]

[Quota] LLM=16.591208 | GPU=0.0 | Used=16.591208/50.0 | Remaining=33.408792 | Usage=33.182416%
24%|           | 231/977 [11:32<34:41,  2.79s/it]

[Quota] LLM=16.595081 | GPU=0.0 | Used=16.595081/50.0 | Remaining=33.404919 | Usage=33.190162%
24%|           | 232/977 [11:35<35:44,  2.88s/it]

[Quota] LLM=16.599314 | GPU=0.0 | Used=16.599314/50.0 | Remaining=33.400686 | Usage=33.198628%
24%|           | 233/977 [11:38<34:33,  2.79s/it]

[Quota] LLM=16.603265 | GPU=0.0 | Used=16.603265/50.0 | Remaining=33.396735 | Usage=33.20653%
24%|           | 234/977 [11:40<32:31,  2.63s/it]

[Quota] LLM=16.607099 | GPU=0.0 | Used=16.607099/50.0 | Remaining=33.392900999999995 | Usage=33.214198%
24%|           | 235/977 [11:43<33:48,  2.73s/it]

[Quota] LLM=16.611533 | GPU=0.0 | Used=16.611533/50.0 | Remaining=33.388467 | Usage=33.223066%
24%|           | 236/977 [11:47<39:42,  3.22s/it]
```

[Quota] LLM=16.615727 | GPU=0.0 | Used=16.615727/50.0 | Remaining=33.384273 | Usage=33.231454%

24%| 237/977 [11:51<40:27, 3.28s/it]

[Quota] LLM=16.62011 | GPU=0.0 | Used=16.62011/50.0 | Remaining=33.37989 | Usage=33.24022%

24%| 238/977 [11:54<38:38, 3.14s/it]

[Quota] LLM=16.624106 | GPU=0.0 | Used=16.624106/50.0 | Remaining=33.375894 | Usage=33.248212%

24%| 239/977 [11:57<37:58, 3.09s/it]

[Quota] LLM=16.627988 | GPU=0.0 | Used=16.627988/50.0 | Remaining=33.372012 | Usage=33.255976%

25%| 240/977 [11:58<33:35, 2.73s/it]

[Quota] LLM=16.631759 | GPU=0.0 | Used=16.631759/50.0 | Remaining=33.368241 | Usage=33.263518%

25%| 241/977 [12:02<35:37, 2.90s/it]

[Quota] LLM=16.635989 | GPU=0.0 | Used=16.635989/50.0 | Remaining=33.364011000000005 | Usage=33.271978%

25%| 242/977 [12:05<34:57, 2.85s/it]

[Quota] LLM=16.640015000000002 | GPU=0.0 | Used=16.640015000000002/50.0 | Remaining=33.359984999999995 | Usage=33.280030000000004%

25%| 243/977 [12:07<34:14, 2.80s/it]

[Quota] LLM=16.643996 | GPU=0.0 | Used=16.643996/50.0 | Remaining=33.356004 | Usage=33.287992%

25%| 244/977 [12:10<34:05, 2.79s/it]

[Quota] LLM=16.648028 | GPU=0.0 | Used=16.648028/50.0 | Remaining=33.351972 | Usage=33.296056%

25%| 245/977 [12:13<35:04, 2.87s/it]

[Quota] LLM=16.652306 | GPU=0.0 | Used=16.652306/50.0 | Remaining=33.347694000000004 | Usage=33.304612%

25%| 246/977 [12:15<32:59, 2.71s/it]

[Quota] LLM=16.656341 | GPU=0.0 | Used=16.656341/50.0 | Remaining=33.343659 | Usage=33.312682%

25%| 247/977 [12:18<33:26, 2.75s/it]

[Quota] LLM=16.66013 | GPU=0.0 | Used=16.66013/50.0 | Remaining=33.339870000000005 | Usage=33.32026%

25%| 248/977 [12:21<33:23, 2.75s/it]

[Quota] LLM=16.664063 | GPU=0.0 | Used=16.664063/50.0 | Remaining=33.335937 | Usage=33.328126%

25% | 249/977 [12:24<32:50, 2.71s/it]

[Quota] LLM=16.668131 | GPU=0.0 | Used=16.668131/50.0 | Remaining=33.331869 | Usage=33.336262%

26% | 250/977 [12:26<33:26, 2.76s/it]

[Quota] LLM=16.672001 | GPU=0.0 | Used=16.672001/50.0 | Remaining=33.327999 | Usage=33.344002%

26% | 251/977 [12:29<32:05, 2.65s/it]

[Quota] LLM=16.676024 | GPU=0.0 | Used=16.676024/50.0 | Remaining=33.323976 | Usage=33.352048%

26% | 252/977 [12:33<38:53, 3.22s/it]

[Quota] LLM=16.679942 | GPU=0.0 | Used=16.679942/50.0 | Remaining=33.320058 | Usage=33.359884%

26% | 253/977 [12:36<38:14, 3.17s/it]

[Quota] LLM=16.684067 | GPU=0.0 | Used=16.684067/50.0 | Remaining=33.315933 | Usage=33.368134%

26% | 254/977 [12:39<37:03, 3.08s/it]

[Quota] LLM=16.688429 | GPU=0.0 | Used=16.688429/50.0 | Remaining=33.311571 | Usage=33.376858%

26% | 255/977 [12:42<35:09, 2.92s/it]

[Quota] LLM=16.692509 | GPU=0.0 | Used=16.692509/50.0 | Remaining=33.307491 | Usage=33.385018%

26% | 256/977 [12:44<33:05, 2.75s/it]

[Quota] LLM=16.69646 | GPU=0.0 | Used=16.69646/50.0 | Remaining=33.30354 | Usage=33.39292%

26% | 257/977 [12:47<31:39, 2.64s/it]

[Quota] LLM=16.700435 | GPU=0.0 | Used=16.700435/50.0 | Remaining=33.299565 | Usage=33.40087%

26% | 258/977 [12:50<33:40, 2.81s/it]

[Quota] LLM=16.704395 | GPU=0.0 | Used=16.704395/50.0 | Remaining=33.295604999999995 | Usage=33.40879%

27% | 259/977 [12:54<39:55, 3.34s/it]

[Quota] LLM=16.708667 | GPU=0.0 | Used=16.708667/50.0 | Remaining=33.291333 | Usage=33.417334%

27% | 260/977 [12:58<39:36, 3.31s/it]

[Quota] LLM=16.71239 | GPU=0.0 | Used=16.71239/50.0 | Remaining=33.28761 | Usage=33.42478%

27% | 261/977 [13:01<40:55, 3.43s/it]

[Quota] LLM=16.716359 | GPU=0.0 | Used=16.716359/50.0 | Remaining=33.283641 | Usage=33.432718%

27% | 262/977 [13:04<37:41, 3.16s/it]

[Quota] LLM=16.720295 | GPU=0.0 | Used=16.720295/50.0 | Remaining=33.279705 | Usage=33.44059%

27% | 263/977 [13:08<40:24, 3.40s/it]

[Quota] LLM=16.724579 | GPU=0.0 | Used=16.724579/50.0 | Remaining=33.275421 | Usage=33.449158%

27% | 264/977 [13:11<39:54, 3.36s/it]

[Quota] LLM=16.728974 | GPU=0.0 | Used=16.728974/50.0 | Remaining=33.271026 | Usage=33.457948%

27% | 265/977 [13:13<35:49, 3.02s/it]

[Quota] LLM=16.732847 | GPU=0.0 | Used=16.732847/50.0 | Remaining=33.267153 | Usage=33.465694%

27% | 266/977 [13:16<33:17, 2.81s/it]

[Quota] LLM=16.73663 | GPU=0.0 | Used=16.73663/50.0 | Remaining=33.263369999999995 | Usage=33.47326%

27% | 267/977 [13:19<36:08, 3.05s/it]

[Quota] LLM=16.740509 | GPU=0.0 | Used=16.740509/50.0 | Remaining=33.259491 | Usage=33.481018%

27% | 268/977 [13:23<36:59, 3.13s/it]

[Quota] LLM=16.744496 | GPU=0.0 | Used=16.744496/50.0 | Remaining=33.255504 | Usage=33.488992%

28% | 269/977 [13:25<34:51, 2.95s/it]

[Quota] LLM=16.748375 | GPU=0.0 | Used=16.748375/50.0 | Remaining=33.251625000000004 | Usage=33.49675%

28% | 270/977 [13:28<33:48, 2.87s/it]

[Quota] LLM=16.752641 | GPU=0.0 | Used=16.752641/50.0 | Remaining=33.247359 | Usage=33.505282%

28% | 271/977 [13:30<33:17, 2.83s/it]

[Quota] LLM=16.756694 | GPU=0.0 | Used=16.756694/50.0 | Remaining=33.243306000000004 | Usage=33.513388%

28% | 272/977 [13:34<34:17, 2.92s/it]

[Quota] LLM=16.760843 | GPU=0.0 | Used=16.760843/50.0 | Remaining=33.239157 | Usage=33.521686%

28%| 273/977 [13:36<32:55, 2.81s/it]

[Quota] LLM=16.764962 | GPU=0.0 | Used=16.764962/50.0 | Remaining=33.235038 | Usage=33.529924%

28%| 274/977 [13:39<32:39, 2.79s/it]

[Quota] LLM=16.768961 | GPU=0.0 | Used=16.768961/50.0 | Remaining=33.231038999999996 | Usage=33.537922%

28%| 275/977 [13:42<34:10, 2.92s/it]

[Quota] LLM=16.773092 | GPU=0.0 | Used=16.773092/50.0 | Remaining=33.226908 | Usage=33.546184%

28%| 276/977 [13:46<36:18, 3.11s/it]

[Quota] LLM=16.77719 | GPU=0.0 | Used=16.77719/50.0 | Remaining=33.222809999999996 | Usage=33.55438%

28%| 277/977 [13:49<36:11, 3.10s/it]

[Quota] LLM=16.781606 | GPU=0.0 | Used=16.781606/50.0 | Remaining=33.218394 | Usage=33.563212%

28%| 278/977 [13:52<38:14, 3.28s/it]

[Quota] LLM=16.785572 | GPU=0.0 | Used=16.785572/50.0 | Remaining=33.214428 | Usage=33.571144%

29%| 279/977 [13:57<41:38, 3.58s/it]

[Quota] LLM=16.789709 | GPU=0.0 | Used=16.789709/50.0 | Remaining=33.210291 | Usage=33.579418%

29%| 280/977 [13:59<37:25, 3.22s/it]

[Quota] LLM=16.793684 | GPU=0.0 | Used=16.793684/50.0 | Remaining=33.206316 | Usage=33.587368%

29%| 281/977 [14:02<36:08, 3.12s/it]

[Quota] LLM=16.798001 | GPU=0.0 | Used=16.798001/50.0 | Remaining=33.201999 | Usage=33.596002%

29%| 282/977 [14:04<32:23, 2.80s/it]

[Quota] LLM=16.80191 | GPU=0.0 | Used=16.80191/50.0 | Remaining=33.19809 | Usage=33.60382%

29%| 283/977 [14:07<31:21, 2.71s/it]

[Quota] LLM=16.805846 | GPU=0.0 | Used=16.805846/50.0 | Remaining=33.194154 | Usage=33.611692%

29%| 284/977 [14:09<31:45, 2.75s/it]

[Quota] LLM=16.810139 | GPU=0.0 | Used=16.810139/50.0 | Remaining=33.189861 | Usage=33.620278%

29%| 285/977 [14:13<33:16, 2.89s/it]

[Quota] LLM=16.814315 | GPU=0.0 | Used=16.814315/50.0 | Remaining=33.185685 | Usage=33.62863%

29%| 286/977 [14:15<30:41, 2.66s/it]

[Quota] LLM=16.818299 | GPU=0.0 | Used=16.818299/50.0 | Remaining=33.181701000000004 | Usage=33.636598%

29%| 287/977 [14:17<29:28, 2.56s/it]

[Quota] LLM=16.822226 | GPU=0.0 | Used=16.822226/50.0 | Remaining=33.177774 | Usage=33.644452%

29%| 288/977 [14:20<29:22, 2.56s/it]

[Quota] LLM=16.826114 | GPU=0.0 | Used=16.826114/50.0 | Remaining=33.173885999999996 | Usage=33.652228%

30%| 289/977 [14:23<31:16, 2.73s/it]

[Quota] LLM=16.830323 | GPU=0.0 | Used=16.830323/50.0 | Remaining=33.169677 | Usage=33.660646%

30%| 290/977 [14:25<30:31, 2.67s/it]

[Quota] LLM=16.834385 | GPU=0.0 | Used=16.834385/50.0 | Remaining=33.165615 | Usage=33.66877%

30%| 291/977 [14:28<30:30, 2.67s/it]

[Quota] LLM=16.838513 | GPU=0.0 | Used=16.838513/50.0 | Remaining=33.161487 | Usage=33.677026%

30%| 292/977 [14:31<32:16, 2.83s/it]

[Quota] LLM=16.842884 | GPU=0.0 | Used=16.842884/50.0 | Remaining=33.157116 | Usage=33.685768%

30%| 293/977 [14:35<34:26, 3.02s/it]

[Quota] LLM=16.846961 | GPU=0.0 | Used=16.846961/50.0 | Remaining=33.153039 | Usage=33.693922%

30%| 294/977 [14:37<31:20, 2.75s/it]

[Quota] LLM=16.850732 | GPU=0.0 | Used=16.850732/50.0 | Remaining=33.149268 | Usage=33.701464%

30%| 295/977 [14:40<33:06, 2.91s/it]

[Quota] LLM=16.855163 | GPU=0.0 | Used=16.855163/50.0 | Remaining=33.144836999999995 | Usage=33.710326%

30%| 296/977 [14:43<33:17, 2.93s/it]

[Quota] LLM=16.85903 | GPU=0.0 | Used=16.85903/50.0 |
Remaining=33.140969999999996 | Usage=33.71806%
30%| 297/977 [14:46<32:14, 2.84s/it]

[Quota] LLM=16.863101 | GPU=0.0 | Used=16.863101/50.0 | Remaining=33.136899 |
Usage=33.726202%
31%| 298/977 [14:48<31:44, 2.81s/it]

[Quota] LLM=16.867088 | GPU=0.0 | Used=16.867088/50.0 |
Remaining=33.132912000000005 | Usage=33.734176%
31%| 299/977 [14:51<30:20, 2.69s/it]

[Quota] LLM=16.871102 | GPU=0.0 | Used=16.871102/50.0 | Remaining=33.128898 |
Usage=33.742204%
31%| 300/977 [14:53<30:00, 2.66s/it]

[Quota] LLM=16.875332 | GPU=0.0 | Used=16.875332/50.0 | Remaining=33.124668 |
Usage=33.750664%
31%| 301/977 [14:56<28:32, 2.53s/it]

[Quota] LLM=16.879169 | GPU=0.0 | Used=16.879169/50.0 |
Remaining=33.120830999999995 | Usage=33.758338%
31%| 302/977 [14:58<28:30, 2.53s/it]

[Quota] LLM=16.883102 | GPU=0.0 | Used=16.883102/50.0 | Remaining=33.116898 |
Usage=33.766204%
31%| 303/977 [15:02<31:38, 2.82s/it]

[Quota] LLM=16.888274 | GPU=0.0 | Used=16.888274/50.0 |
Remaining=33.111726000000004 | Usage=33.776548%
31%| 304/977 [15:04<29:48, 2.66s/it]

[Quota] LLM=16.892273 | GPU=0.0 | Used=16.892273/50.0 | Remaining=33.107727 |
Usage=33.784546%
31%| 305/977 [15:06<28:34, 2.55s/it]

[Quota] LLM=16.896242 | GPU=0.0 | Used=16.896242/50.0 | Remaining=33.103758 |
Usage=33.792484%
31%| 306/977 [15:09<28:09, 2.52s/it]

[Quota] LLM=16.900058 | GPU=0.0 | Used=16.900058/50.0 | Remaining=33.099942 |
Usage=33.800116%
31%| 307/977 [15:11<27:25, 2.46s/it]

[Quota] LLM=16.903991 | GPU=0.0 | Used=16.903991/50.0 |
Remaining=33.096008999999995 | Usage=33.807982%
32%| 308/977 [15:14<30:13, 2.71s/it]

[Quota] LLM=16.908452 | GPU=0.0 | Used=16.908452/50.0 | Remaining=33.091548 | Usage=33.816904%

32% | 309/977 [15:17<29:15, 2.63s/it]

[Quota] LLM=16.912235 | GPU=0.0 | Used=16.912235/50.0 | Remaining=33.087765000000005 | Usage=33.82447%

32% | 310/977 [15:20<30:11, 2.72s/it]

[Quota] LLM=16.916933 | GPU=0.0 | Used=16.916933/50.0 | Remaining=33.083067 | Usage=33.833866%

32% | 311/977 [15:23<31:06, 2.80s/it]

[Quota] LLM=16.920701 | GPU=0.0 | Used=16.920701/50.0 | Remaining=33.079299 | Usage=33.841402%

32% | 312/977 [15:25<28:58, 2.61s/it]

[Quota] LLM=16.924457 | GPU=0.0 | Used=16.924457/50.0 | Remaining=33.075542999999996 | Usage=33.848914%

32% | 313/977 [15:27<27:48, 2.51s/it]

[Quota] LLM=16.928255 | GPU=0.0 | Used=16.928255/50.0 | Remaining=33.071745 | Usage=33.85651%

32% | 314/977 [15:39<1:00:30, 5.48s/it]

[Quota] LLM=16.932371 | GPU=0.0 | Used=16.932371/50.0 | Remaining=33.067629 | Usage=33.864742%

32% | 315/977 [15:43<52:58, 4.80s/it]

[Quota] LLM=16.937033 | GPU=0.0 | Used=16.937033/50.0 | Remaining=33.062967 | Usage=33.874066%

32% | 316/977 [15:45<43:53, 3.98s/it]

[Quota] LLM=16.940888 | GPU=0.0 | Used=16.940888/50.0 | Remaining=33.059112 | Usage=33.881776%

32% | 317/977 [15:47<38:25, 3.49s/it]

[Quota] LLM=16.944614 | GPU=0.0 | Used=16.944614/50.0 | Remaining=33.055386 | Usage=33.889228%

33% | 318/977 [15:50<37:04, 3.38s/it]

[Quota] LLM=16.949195 | GPU=0.0 | Used=16.949195/50.0 | Remaining=33.050805 | Usage=33.89839%

33% | 319/977 [15:53<34:22, 3.14s/it]

[Quota] LLM=16.953065 | GPU=0.0 | Used=16.953065/50.0 | Remaining=33.046935000000005 | Usage=33.90613%

33% | 320/977 [15:56<34:21, 3.14s/it]

```
[Quota] LLM=16.957169 | GPU=0.0 | Used=16.957169/50.0 | Remaining=33.042831 |
Usage=33.914338%
33%|           | 321/977 [15:58<31:49,  2.91s/it]

[Quota] LLM=16.961318 | GPU=0.0 | Used=16.961318/50.0 | Remaining=33.038682 |
Usage=33.922636%
33%|           | 322/977 [16:01<30:11,  2.77s/it]

[Quota] LLM=16.965215 | GPU=0.0 | Used=16.965215/50.0 | Remaining=33.034785 |
Usage=33.93043%
33%|           | 323/977 [16:03<27:53,  2.56s/it]

[Quota] LLM=16.969148 | GPU=0.0 | Used=16.969148/50.0 |
Remaining=33.030851999999996 | Usage=33.938296%
33%|           | 324/977 [16:06<30:58,  2.85s/it]

[Quota] LLM=16.972955 | GPU=0.0 | Used=16.972955/50.0 | Remaining=33.027045 |
Usage=33.94591%
33%|           | 325/977 [16:09<29:01,  2.67s/it]

[Quota] LLM=16.97681 | GPU=0.0 | Used=16.97681/50.0 | Remaining=33.02319 |
Usage=33.95362%
33%|           | 326/977 [16:11<27:42,  2.55s/it]

[Quota] LLM=16.980632 | GPU=0.0 | Used=16.980632/50.0 | Remaining=33.019368 |
Usage=33.961264%
33%|           | 327/977 [16:14<29:15,  2.70s/it]

[Quota] LLM=16.985165 | GPU=0.0 | Used=16.985165/50.0 |
Remaining=33.014835000000005 | Usage=33.97033%
34%|           | 328/977 [16:17<29:31,  2.73s/it]

[Quota] LLM=16.989077 | GPU=0.0 | Used=16.989077/50.0 | Remaining=33.010923 |
Usage=33.978154%
34%|           | 329/977 [16:19<27:43,  2.57s/it]

[Quota] LLM=16.992902 | GPU=0.0 | Used=16.992902/50.0 | Remaining=33.007098 |
Usage=33.985804%
34%|           | 330/977 [16:22<29:37,  2.75s/it]

[Quota] LLM=16.99676 | GPU=0.0 | Used=16.99676/50.0 |
Remaining=33.003240000000005 | Usage=33.99352%
34%|           | 331/977 [16:24<27:45,  2.58s/it]

[Quota] LLM=17.00093 | GPU=0.0 | Used=17.00093/50.0 | Remaining=32.99907 |
Usage=34.00186%
34%|           | 332/977 [16:27<27:24,  2.55s/it]
```

```
[Quota] LLM=17.00534 | GPU=0.0 | Used=17.00534/50.0 |  
Remaining=32.994659999999996 | Usage=34.01068%  
34%|           | 333/977 [16:30<29:00,  2.70s/it]  
  
[Quota] LLM=17.009189 | GPU=0.0 | Used=17.009189/50.0 | Remaining=32.990811 |  
Usage=34.018378%  
34%|           | 334/977 [16:32<27:59,  2.61s/it]  
  
[Quota] LLM=17.013284 | GPU=0.0 | Used=17.013284/50.0 | Remaining=32.986716 |  
Usage=34.026568%  
34%|           | 335/977 [16:34<26:10,  2.45s/it]  
  
[Quota] LLM=17.017118 | GPU=0.0 | Used=17.017118/50.0 |  
Remaining=32.982882000000004 | Usage=34.034236%  
34%|           | 336/977 [16:37<26:17,  2.46s/it]  
  
[Quota] LLM=17.021027 | GPU=0.0 | Used=17.021027/50.0 |  
Remaining=32.978972999999996 | Usage=34.042054%  
34%|           | 337/977 [16:39<26:19,  2.47s/it]  
  
[Quota] LLM=17.025044 | GPU=0.0 | Used=17.025044/50.0 | Remaining=32.974956 |  
Usage=34.050088%  
35%|           | 338/977 [16:42<25:42,  2.41s/it]  
  
[Quota] LLM=17.028914 | GPU=0.0 | Used=17.028914/50.0 | Remaining=32.971086 |  
Usage=34.057828%  
35%|           | 339/977 [16:44<26:07,  2.46s/it]  
  
[Quota] LLM=17.032916 | GPU=0.0 | Used=17.032916/50.0 | Remaining=32.967084 |  
Usage=34.065832%  
35%|           | 340/977 [16:47<26:21,  2.48s/it]  
  
[Quota] LLM=17.036864 | GPU=0.0 | Used=17.036864/50.0 | Remaining=32.963136 |  
Usage=34.073728%  
35%|           | 341/977 [16:49<25:49,  2.44s/it]  
  
[Quota] LLM=17.041046 | GPU=0.0 | Used=17.041046/50.0 | Remaining=32.958954 |  
Usage=34.082092%  
35%|           | 342/977 [16:51<24:51,  2.35s/it]  
  
[Quota] LLM=17.044916 | GPU=0.0 | Used=17.044916/50.0 | Remaining=32.955084 |  
Usage=34.089832%  
35%|           | 343/977 [16:53<24:47,  2.35s/it]  
  
[Quota] LLM=17.048969 | GPU=0.0 | Used=17.048969/50.0 | Remaining=32.951031 |  
Usage=34.097938%  
35%|           | 344/977 [16:56<24:53,  2.36s/it]
```

[Quota] LLM=17.052938 | GPU=0.0 | Used=17.052938/50.0 | Remaining=32.947062 | Usage=34.105876%

35%| 345/977 [16:59<26:27, 2.51s/it]

[Quota] LLM=17.057036 | GPU=0.0 | Used=17.057036/50.0 | Remaining=32.942964 | Usage=34.114072%

35%| 346/977 [17:01<26:53, 2.56s/it]

[Quota] LLM=17.061119 | GPU=0.0 | Used=17.061119/50.0 | Remaining=32.938880999999995 | Usage=34.122238%

36%| 347/977 [17:04<28:25, 2.71s/it]

[Quota] LLM=17.065217 | GPU=0.0 | Used=17.065217/50.0 | Remaining=32.934782999999996 | Usage=34.130434%

36%| 348/977 [17:07<29:02, 2.77s/it]

[Quota] LLM=17.069021 | GPU=0.0 | Used=17.069021/50.0 | Remaining=32.930979 | Usage=34.138042%

36%| 349/977 [17:10<28:32, 2.73s/it]

[Quota] LLM=17.073404 | GPU=0.0 | Used=17.073404/50.0 | Remaining=32.926596 | Usage=34.146808%

36%| 350/977 [17:13<29:35, 2.83s/it]

[Quota] LLM=17.077439000000002 | GPU=0.0 | Used=17.077439000000002/50.0 | Remaining=32.922561 | Usage=34.154878000000004%

36%| 351/977 [17:16<28:40, 2.75s/it]

[Quota] LLM=17.081375 | GPU=0.0 | Used=17.081375/50.0 | Remaining=32.918625 | Usage=34.16275%

36%| 352/977 [17:18<28:29, 2.74s/it]

[Quota] LLM=17.085362 | GPU=0.0 | Used=17.085362/50.0 | Remaining=32.914638 | Usage=34.170724%

36%| 353/977 [17:20<26:28, 2.55s/it]

[Quota] LLM=17.089169 | GPU=0.0 | Used=17.089169/50.0 | Remaining=32.910831 | Usage=34.178338%

36%| 354/977 [17:23<27:46, 2.68s/it]

[Quota] LLM=17.093234 | GPU=0.0 | Used=17.093234/50.0 | Remaining=32.906766000000005 | Usage=34.186468%

36%| 355/977 [17:26<26:19, 2.54s/it]

[Quota] LLM=17.097059 | GPU=0.0 | Used=17.097059/50.0 | Remaining=32.902941 | Usage=34.194118%

36%| 356/977 [17:28<25:47, 2.49s/it]

[Quota] LLM=17.100995 | GPU=0.0 | Used=17.100995/50.0 | Remaining=32.899005 | Usage=34.20199%

37% | 357/977 [17:33<34:58, 3.39s/it]

[Quota] LLM=17.105201 | GPU=0.0 | Used=17.105201/50.0 | Remaining=32.894799 | Usage=34.210402%

37% | 358/977 [17:37<34:26, 3.34s/it]

[Quota] LLM=17.109239 | GPU=0.0 | Used=17.109239/50.0 | Remaining=32.890761 | Usage=34.218478%

37% | 359/977 [17:39<31:14, 3.03s/it]

[Quota] LLM=17.11319 | GPU=0.0 | Used=17.11319/50.0 | Remaining=32.88681 | Usage=34.22638%

37% | 360/977 [17:42<30:45, 2.99s/it]

[Quota] LLM=17.11778 | GPU=0.0 | Used=17.11778/50.0 | Remaining=32.88222000000004 | Usage=34.23556%

37% | 361/977 [17:44<29:27, 2.87s/it]

[Quota] LLM=17.121983 | GPU=0.0 | Used=17.121983/50.0 | Remaining=32.878017 | Usage=34.243966%

37% | 362/977 [17:47<27:45, 2.71s/it]

[Quota] LLM=17.126081 | GPU=0.0 | Used=17.126081/50.0 | Remaining=32.873919 | Usage=34.252162%

37% | 363/977 [17:49<27:35, 2.70s/it]

[Quota] LLM=17.129948 | GPU=0.0 | Used=17.129948/50.0 | Remaining=32.870052 | Usage=34.259896%

37% | 364/977 [17:52<26:56, 2.64s/it]

[Quota] LLM=17.133986 | GPU=0.0 | Used=17.133986/50.0 | Remaining=32.866014 | Usage=34.267972%

37% | 365/977 [17:55<27:59, 2.74s/it]

[Quota] LLM=17.138408 | GPU=0.0 | Used=17.138408/50.0 | Remaining=32.861592 | Usage=34.276816%

37% | 366/977 [17:58<28:37, 2.81s/it]

[Quota] LLM=17.142626 | GPU=0.0 | Used=17.142626/50.0 | Remaining=32.857374 | Usage=34.285252%

38% | 367/977 [18:01<29:44, 2.92s/it]

[Quota] LLM=17.14673 | GPU=0.0 | Used=17.14673/50.0 | Remaining=32.85326999999995 | Usage=34.29346%

38% | 368/977 [18:04<30:28, 3.00s/it]

[Quota] LLM=17.150597 | GPU=0.0 | Used=17.150597/50.0 | Remaining=32.849402999999995 | Usage=34.301194%
38%| 369/977 [18:07<29:33, 2.92s/it]

[Quota] LLM=17.15453 | GPU=0.0 | Used=17.15453/50.0 | Remaining=32.84547 | Usage=34.30906%
38%| 370/977 [18:09<27:55, 2.76s/it]

[Quota] LLM=17.158367 | GPU=0.0 | Used=17.158367/50.0 | Remaining=32.841633 | Usage=34.316734%
38%| 371/977 [18:12<27:36, 2.73s/it]

[Quota] LLM=17.16218 | GPU=0.0 | Used=17.16218/50.0 | Remaining=32.83782 | Usage=34.32436%
38%| 372/977 [18:15<27:31, 2.73s/it]

[Quota] LLM=17.166794 | GPU=0.0 | Used=17.166794/50.0 | Remaining=32.833206000000004 | Usage=34.333588%
38%| 373/977 [18:18<28:43, 2.85s/it]

[Quota] LLM=17.171204 | GPU=0.0 | Used=17.171204/50.0 | Remaining=32.828796 | Usage=34.342408%
38%| 374/977 [18:21<28:09, 2.80s/it]

[Quota] LLM=17.176034 | GPU=0.0 | Used=17.176034/50.0 | Remaining=32.823966 | Usage=34.352068%
38%| 375/977 [18:23<27:08, 2.71s/it]

[Quota] LLM=17.18003 | GPU=0.0 | Used=17.18003/50.0 | Remaining=32.81997 | Usage=34.36006%
38%| 376/977 [18:25<25:51, 2.58s/it]

[Quota] LLM=17.183999 | GPU=0.0 | Used=17.183999/50.0 | Remaining=32.816001 | Usage=34.367998%
39%| 377/977 [18:28<26:25, 2.64s/it]

[Quota] LLM=17.188397 | GPU=0.0 | Used=17.188397/50.0 | Remaining=32.811603000000005 | Usage=34.376794%
39%| 378/977 [18:30<24:51, 2.49s/it]

[Quota] LLM=17.192204 | GPU=0.0 | Used=17.192204/50.0 | Remaining=32.807795999999996 | Usage=34.384408%
39%| 379/977 [18:33<25:58, 2.61s/it]

[Quota] LLM=17.196335 | GPU=0.0 | Used=17.196335/50.0 | Remaining=32.803664999999995 | Usage=34.39267%
39%| 380/977 [18:36<27:42, 2.78s/it]

[Quota] LLM=17.200193 | GPU=0.0 | Used=17.200193/50.0 | Remaining=32.799807 | Usage=34.400386%

39% | 381/977 [18:42<35:56, 3.62s/it]

[Quota] LLM=17.204162 | GPU=0.0 | Used=17.204162/50.0 | Remaining=32.795838 | Usage=34.408324%

39% | 382/977 [18:45<33:20, 3.36s/it]

[Quota] LLM=17.208104 | GPU=0.0 | Used=17.208104/50.0 | Remaining=32.791896 | Usage=34.416208%

39% | 383/977 [18:48<31:49, 3.21s/it]

[Quota] LLM=17.212316 | GPU=0.0 | Used=17.212316/50.0 | Remaining=32.787684 | Usage=34.424632%

39% | 384/977 [18:52<33:47, 3.42s/it]

[Quota] LLM=17.216747 | GPU=0.0 | Used=17.216747/50.0 | Remaining=32.783253 | Usage=34.433494%

39% | 385/977 [18:54<31:18, 3.17s/it]

[Quota] LLM=17.220794 | GPU=0.0 | Used=17.220794/50.0 | Remaining=32.779206 | Usage=34.441588%

40% | 386/977 [18:57<29:25, 2.99s/it]

[Quota] LLM=17.224814 | GPU=0.0 | Used=17.224814/50.0 | Remaining=32.775186000000005 | Usage=34.449628%

40% | 387/977 [19:00<30:11, 3.07s/it]

[Quota] LLM=17.228831 | GPU=0.0 | Used=17.228831/50.0 | Remaining=32.771169 | Usage=34.457662%

40% | 388/977 [19:03<29:43, 3.03s/it]

[Quota] LLM=17.232833 | GPU=0.0 | Used=17.232833/50.0 | Remaining=32.767167 | Usage=34.465666%

40% | 389/977 [19:06<29:18, 2.99s/it]

[Quota] LLM=17.236928 | GPU=0.0 | Used=17.236928/50.0 | Remaining=32.763072 | Usage=34.473856%

40% | 390/977 [19:09<29:43, 3.04s/it]

[Quota] LLM=17.241149 | GPU=0.0 | Used=17.241149/50.0 | Remaining=32.758851 | Usage=34.482298%

40% | 391/977 [19:11<27:54, 2.86s/it]

[Quota] LLM=17.245034 | GPU=0.0 | Used=17.245034/50.0 | Remaining=32.754965999999996 | Usage=34.490068%

40% | 392/977 [19:14<28:30, 2.92s/it]

[Quota] LLM=17.249069 | GPU=0.0 | Used=17.249069/50.0 | Remaining=32.750931 | Usage=34.498138%

40%| 393/977 [19:17<28:08, 2.89s/it]

[Quota] LLM=17.25341 | GPU=0.0 | Used=17.25341/50.0 | Remaining=32.74659 | Usage=34.50682%

40%| 394/977 [19:20<28:42, 2.95s/it]

[Quota] LLM=17.257697 | GPU=0.0 | Used=17.257697/50.0 | Remaining=32.742303 | Usage=34.515394%

40%| 395/977 [19:24<29:25, 3.03s/it]

[Quota] LLM=17.261783 | GPU=0.0 | Used=17.261783/50.0 | Remaining=32.738217 | Usage=34.523566%

41%| 396/977 [19:26<28:10, 2.91s/it]

[Quota] LLM=17.26592 | GPU=0.0 | Used=17.26592/50.0 | Remaining=32.73408 | Usage=34.53184%

41%| 397/977 [19:28<26:21, 2.73s/it]

[Quota] LLM=17.269859 | GPU=0.0 | Used=17.269859/50.0 | Remaining=32.730141 | Usage=34.539718%

41%| 398/977 [19:31<25:01, 2.59s/it]

[Quota] LLM=17.273729 | GPU=0.0 | Used=17.273729/50.0 | Remaining=32.726271 | Usage=34.547458%

41%| 399/977 [19:33<24:52, 2.58s/it]

[Quota] LLM=17.277794 | GPU=0.0 | Used=17.277794/50.0 | Remaining=32.722206 | Usage=34.555588%

41%| 400/977 [19:36<23:45, 2.47s/it]

[Quota] LLM=17.281679 | GPU=0.0 | Used=17.281679/50.0 | Remaining=32.718321 | Usage=34.563358%

41%| 401/977 [19:39<26:46, 2.79s/it]

[Quota] LLM=17.285615 | GPU=0.0 | Used=17.285615/50.0 | Remaining=32.714385 | Usage=34.57123%

41%| 402/977 [19:41<25:03, 2.61s/it]

[Quota] LLM=17.289623 | GPU=0.0 | Used=17.289623/50.0 | Remaining=32.710377 | Usage=34.579246%

41%| 403/977 [19:44<26:15, 2.74s/it]

[Quota] LLM=17.293772 | GPU=0.0 | Used=17.293772/50.0 | Remaining=32.706227999999996 | Usage=34.587544%

41%| 404/977 [19:47<26:14, 2.75s/it]

[Quota] LLM=17.297879000000002 | GPU=0.0 | Used=17.297879000000002/50.0 | Remaining=32.702121 | Usage=34.595758000000004%

41% | 405/977 [19:50<26:15, 2.75s/it]

[Quota] LLM=17.301923 | GPU=0.0 | Used=17.301923/50.0 | Remaining=32.698077 | Usage=34.603846%

42% | 406/977 [19:54<28:54, 3.04s/it]

[Quota] LLM=17.306087 | GPU=0.0 | Used=17.306087/50.0 | Remaining=32.693912999999995 | Usage=34.612174%

42% | 407/977 [19:56<27:32, 2.90s/it]

[Quota] LLM=17.309894 | GPU=0.0 | Used=17.309894/50.0 | Remaining=32.690106 | Usage=34.619788%

42% | 408/977 [19:58<25:22, 2.68s/it]

[Quota] LLM=17.313812 | GPU=0.0 | Used=17.313812/50.0 | Remaining=32.686188 | Usage=34.627624%

42% | 409/977 [20:00<23:50, 2.52s/it]

[Quota] LLM=17.317736 | GPU=0.0 | Used=17.317736/50.0 | Remaining=32.682264 | Usage=34.635472%

42% | 410/977 [20:04<25:30, 2.70s/it]

[Quota] LLM=17.322329 | GPU=0.0 | Used=17.322329/50.0 | Remaining=32.677671000000004 | Usage=34.644658%

42% | 411/977 [20:06<23:22, 2.48s/it]

[Quota] LLM=17.326154 | GPU=0.0 | Used=17.326154/50.0 | Remaining=32.673846 | Usage=34.652308%

42% | 412/977 [20:08<22:12, 2.36s/it]

[Quota] LLM=17.329994 | GPU=0.0 | Used=17.329994/50.0 | Remaining=32.670006 | Usage=34.659988%

42% | 413/977 [20:10<22:18, 2.37s/it]

[Quota] LLM=17.3339 | GPU=0.0 | Used=17.3339/50.0 | Remaining=32.6661 | Usage=34.6678%

42% | 414/977 [20:13<23:52, 2.54s/it]

[Quota] LLM=17.337998 | GPU=0.0 | Used=17.337998/50.0 | Remaining=32.662002 | Usage=34.675996%

42% | 415/977 [20:16<23:54, 2.55s/it]

[Quota] LLM=17.342012 | GPU=0.0 | Used=17.342012/50.0 | Remaining=32.657988 | Usage=34.684024%

43% | 416/977 [20:19<26:10, 2.80s/it]

[Quota] LLM=17.346131 | GPU=0.0 | Used=17.346131/50.0 | Remaining=32.653869 | Usage=34.692262%

43% | 417/977 [20:21<24:46, 2.65s/it]

[Quota] LLM=17.350214 | GPU=0.0 | Used=17.350214/50.0 | Remaining=32.649786 | Usage=34.700428%

43% | 418/977 [20:24<24:46, 2.66s/it]

[Quota] LLM=17.354378 | GPU=0.0 | Used=17.354378/50.0 | Remaining=32.645622 | Usage=34.708756%

43% | 419/977 [20:26<23:56, 2.57s/it]

[Quota] LLM=17.358299 | GPU=0.0 | Used=17.358299/50.0 | Remaining=32.641701 | Usage=34.716598%

43% | 420/977 [20:29<25:20, 2.73s/it]

[Quota] LLM=17.362544 | GPU=0.0 | Used=17.362544/50.0 | Remaining=32.637456 | Usage=34.725088%

43% | 421/977 [20:32<26:29, 2.86s/it]

[Quota] LLM=17.366936 | GPU=0.0 | Used=17.366936/50.0 | Remaining=32.633064000000005 | Usage=34.733872%

43% | 422/977 [20:35<26:35, 2.88s/it]

[Quota] LLM=17.370857 | GPU=0.0 | Used=17.370857/50.0 | Remaining=32.629143 | Usage=34.741714%

43% | 423/977 [20:38<26:17, 2.85s/it]

[Quota] LLM=17.375354 | GPU=0.0 | Used=17.375354/50.0 | Remaining=32.624646 | Usage=34.750708%

43% | 424/977 [20:42<28:04, 3.05s/it]

[Quota] LLM=17.380046 | GPU=0.0 | Used=17.380046/50.0 | Remaining=32.619954 | Usage=34.760092%

44% | 425/977 [20:44<26:32, 2.89s/it]

[Quota] LLM=17.384012 | GPU=0.0 | Used=17.384012/50.0 | Remaining=32.615988 | Usage=34.768024%

44% | 426/977 [20:47<27:09, 2.96s/it]

[Quota] LLM=17.387951 | GPU=0.0 | Used=17.387951/50.0 | Remaining=32.612049 | Usage=34.775902%

44% | 427/977 [20:50<27:29, 3.00s/it]

[Quota] LLM=17.392217 | GPU=0.0 | Used=17.392217/50.0 | Remaining=32.607783 | Usage=34.784434%

44% | 428/977 [20:54<29:13, 3.19s/it]

```
[Quota] LLM=17.396786 | GPU=0.0 | Used=17.396786/50.0 | Remaining=32.603214 | Usage=34.793572%
44%|           | 429/977 [20:57<27:54,  3.06s/it]

[Quota] LLM=17.401004 | GPU=0.0 | Used=17.401004/50.0 | Remaining=32.598996 | Usage=34.802008%
44%|           | 430/977 [21:01<29:57,  3.29s/it]

[Quota] LLM=17.40542 | GPU=0.0 | Used=17.40542/50.0 | Remaining=32.59458 | Usage=34.81084%
44%|           | 431/977 [21:04<30:21,  3.34s/it]

[Quota] LLM=17.409413 | GPU=0.0 | Used=17.409413/50.0 | Remaining=32.590587 | Usage=34.818826%
44%|           | 432/977 [21:07<28:47,  3.17s/it]

[Quota] LLM=17.413553 | GPU=0.0 | Used=17.413553/50.0 | Remaining=32.586447 | Usage=34.827106%
44%|           | 433/977 [21:10<27:51,  3.07s/it]

[Quota] LLM=17.417672 | GPU=0.0 | Used=17.417672/50.0 | Remaining=32.582328000000004 | Usage=34.835344%
44%|           | 434/977 [21:15<33:33,  3.71s/it]

[Quota] LLM=17.421734 | GPU=0.0 | Used=17.421734/50.0 | Remaining=32.578266 | Usage=34.843468%
45%|           | 435/977 [21:19<34:40,  3.84s/it]

[Quota] LLM=17.426348 | GPU=0.0 | Used=17.426348/50.0 | Remaining=32.573651999999996 | Usage=34.852696%
45%|           | 436/977 [21:22<31:23,  3.48s/it]

[Quota] LLM=17.430443 | GPU=0.0 | Used=17.430443/50.0 | Remaining=32.569557 | Usage=34.860886%
45%|           | 437/977 [21:25<31:14,  3.47s/it]

[Quota] LLM=17.434541 | GPU=0.0 | Used=17.434541/50.0 | Remaining=32.565459000000004 | Usage=34.869082%
45%|           | 438/977 [21:28<30:05,  3.35s/it]

[Quota] LLM=17.438864 | GPU=0.0 | Used=17.438864/50.0 | Remaining=32.561136000000005 | Usage=34.877728%
45%|           | 439/977 [21:31<27:32,  3.07s/it]

[Quota] LLM=17.443028 | GPU=0.0 | Used=17.443028/50.0 | Remaining=32.556972 | Usage=34.886056%
45%|           | 440/977 [21:34<27:41,  3.09s/it]
```

[Quota] LLM=17.446898 | GPU=0.0 | Used=17.446898/50.0 | Remaining=32.553101999999996 | Usage=34.893796%
45%| 441/977 [21:38<29:30, 3.30s/it]

[Quota] LLM=17.451077 | GPU=0.0 | Used=17.451077/50.0 | Remaining=32.548923 | Usage=34.902154%
45%| 442/977 [21:41<30:31, 3.42s/it]

[Quota] LLM=17.455046 | GPU=0.0 | Used=17.455046/50.0 | Remaining=32.544954000000004 | Usage=34.910092%
45%| 443/977 [21:44<28:47, 3.23s/it]

[Quota] LLM=17.459402 | GPU=0.0 | Used=17.459402/50.0 | Remaining=32.540598 | Usage=34.918804%
45%| 444/977 [21:46<26:26, 2.98s/it]

[Quota] LLM=17.463221 | GPU=0.0 | Used=17.463221/50.0 | Remaining=32.536778999999996 | Usage=34.926442%
46%| 445/977 [21:50<28:45, 3.24s/it]

[Quota] LLM=17.467451 | GPU=0.0 | Used=17.467451/50.0 | Remaining=32.532549 | Usage=34.934902%
46%| 446/977 [21:53<27:18, 3.08s/it]

[Quota] LLM=17.471348 | GPU=0.0 | Used=17.471348/50.0 | Remaining=32.528652 | Usage=34.942696%
46%| 447/977 [21:57<29:17, 3.32s/it]

[Quota] LLM=17.475311 | GPU=0.0 | Used=17.475311/50.0 | Remaining=32.524688999999995 | Usage=34.950622%
46%| 448/977 [22:00<29:36, 3.36s/it]

[Quota] LLM=17.479709 | GPU=0.0 | Used=17.479709/50.0 | Remaining=32.520291 | Usage=34.959418%
46%| 449/977 [22:04<31:04, 3.53s/it]

[Quota] LLM=17.484149 | GPU=0.0 | Used=17.484149/50.0 | Remaining=32.515851 | Usage=34.968298%
46%| 450/977 [22:09<33:01, 3.76s/it]

[Quota] LLM=17.48861 | GPU=0.0 | Used=17.48861/50.0 | Remaining=32.51139 | Usage=34.97722%
46%| 451/977 [22:11<30:01, 3.42s/it]

[Quota] LLM=17.49284 | GPU=0.0 | Used=17.49284/50.0 | Remaining=32.50716 | Usage=34.98568%
46%| 452/977 [22:14<28:31, 3.26s/it]

```
[Quota] LLM=17.496992 | GPU=0.0 | Used=17.496992/50.0 | Remaining=32.503008 | Usage=34.993984%  
46%|           | 453/977 [22:17<26:31,  3.04s/it]  
  
[Quota] LLM=17.501294 | GPU=0.0 | Used=17.501294/50.0 | Remaining=32.498706 | Usage=35.002588%  
46%|           | 454/977 [22:19<25:52,  2.97s/it]  
  
[Quota] LLM=17.505116 | GPU=0.0 | Used=17.505116/50.0 | Remaining=32.494884 | Usage=35.010232%  
47%|           | 455/977 [22:23<26:12,  3.01s/it]  
  
[Quota] LLM=17.509268 | GPU=0.0 | Used=17.509268/50.0 | Remaining=32.490732 | Usage=35.018536%  
47%|           | 456/977 [22:25<25:55,  2.99s/it]  
  
[Quota] LLM=17.513348 | GPU=0.0 | Used=17.513348/50.0 | Remaining=32.486652 | Usage=35.026696%  
47%|           | 457/977 [22:28<25:04,  2.89s/it]  
  
[Quota] LLM=17.517566 | GPU=0.0 | Used=17.517566/50.0 | Remaining=32.482434 | Usage=35.035132%  
47%|           | 458/977 [22:33<28:56,  3.35s/it]  
  
[Quota] LLM=17.522006 | GPU=0.0 | Used=17.522006/50.0 | Remaining=32.477993999999995 | Usage=35.044012%  
47%|           | 459/977 [22:36<28:06,  3.26s/it]  
  
[Quota] LLM=17.526185 | GPU=0.0 | Used=17.526185/50.0 | Remaining=32.473815 | Usage=35.05237%  
47%|           | 460/977 [22:38<25:43,  2.99s/it]  
  
[Quota] LLM=17.530175 | GPU=0.0 | Used=17.530175/50.0 | Remaining=32.469825 | Usage=35.06035%  
47%|           | 461/977 [22:42<27:29,  3.20s/it]  
  
[Quota] LLM=17.534273 | GPU=0.0 | Used=17.534273/50.0 | Remaining=32.465727 | Usage=35.068546%  
47%|           | 462/977 [22:44<25:36,  2.98s/it]  
  
[Quota] LLM=17.538386 | GPU=0.0 | Used=17.538386/50.0 | Remaining=32.461614 | Usage=35.076772%  
47%|           | 463/977 [22:47<25:33,  2.98s/it]  
  
[Quota] LLM=17.542262 | GPU=0.0 | Used=17.542262/50.0 | Remaining=32.457738 | Usage=35.084524%  
47%|           | 464/977 [22:51<27:58,  3.27s/it]
```

```
[Quota] LLM=17.547128 | GPU=0.0 | Used=17.547128/50.0 | Remaining=32.452872 | Usage=35.094256%
48%|          | 465/977 [22:53<25:22,  2.97s/it]

[Quota] LLM=17.551007 | GPU=0.0 | Used=17.551007/50.0 | Remaining=32.448993 | Usage=35.102014%
48%|          | 466/977 [22:57<26:57,  3.17s/it]

[Quota] LLM=17.555132 | GPU=0.0 | Used=17.555132/50.0 | Remaining=32.444868 | Usage=35.110264%
48%|          | 467/977 [23:00<26:27,  3.11s/it]

[Quota] LLM=17.559167 | GPU=0.0 | Used=17.559167/50.0 | Remaining=32.440833 | Usage=35.118334%
48%|          | 468/977 [23:02<24:25,  2.88s/it]

[Quota] LLM=17.563172 | GPU=0.0 | Used=17.563172/50.0 | Remaining=32.436828 | Usage=35.126344%
48%|          | 469/977 [23:06<25:47,  3.05s/it]

[Quota] LLM=17.567321 | GPU=0.0 | Used=17.567321/50.0 | Remaining=32.432679 | Usage=35.134642%
48%|          | 470/977 [23:10<27:40,  3.28s/it]

[Quota] LLM=17.571593 | GPU=0.0 | Used=17.571593/50.0 | Remaining=32.428407 | Usage=35.143186%
48%|          | 471/977 [23:13<27:06,  3.21s/it]

[Quota] LLM=17.575484 | GPU=0.0 | Used=17.575484/50.0 | Remaining=32.424516 | Usage=35.150968%
48%|          | 472/977 [23:16<26:51,  3.19s/it]

[Quota] LLM=17.579585 | GPU=0.0 | Used=17.579585/50.0 | Remaining=32.420415 | Usage=35.15917%
48%|          | 473/977 [23:18<24:59,  2.97s/it]

[Quota] LLM=17.583638 | GPU=0.0 | Used=17.583638/50.0 | Remaining=32.416362 | Usage=35.167276%
49%|          | 474/977 [23:22<25:55,  3.09s/it]

[Quota] LLM=17.588165 | GPU=0.0 | Used=17.588165/50.0 | Remaining=32.41183499999996 | Usage=35.17633%
49%|          | 475/977 [23:24<24:18,  2.91s/it]

[Quota] LLM=17.592101 | GPU=0.0 | Used=17.592101/50.0 | Remaining=32.407899 | Usage=35.184202%
49%|          | 476/977 [23:26<22:45,  2.73s/it]
```

[Quota] LLM=17.59619 | GPU=0.0 | Used=17.59619/50.0 | Remaining=32.40381 | Usage=35.19238%

49% | 477/977 [23:31<26:53, 3.23s/it]

[Quota] LLM=17.60024 | GPU=0.0 | Used=17.60024/50.0 | Remaining=32.39976 | Usage=35.20048%

49% | 478/977 [23:33<23:48, 2.86s/it]

[Quota] LLM=17.604041 | GPU=0.0 | Used=17.604041/50.0 | Remaining=32.395959000000005 | Usage=35.208082%

49% | 479/977 [23:35<23:15, 2.80s/it]

[Quota] LLM=17.608076 | GPU=0.0 | Used=17.608076/50.0 | Remaining=32.391924 | Usage=35.216152%

49% | 480/977 [23:38<22:12, 2.68s/it]

[Quota] LLM=17.611907 | GPU=0.0 | Used=17.611907/50.0 | Remaining=32.388093 | Usage=35.223814%

49% | 481/977 [23:41<23:15, 2.81s/it]

[Quota] LLM=17.616221 | GPU=0.0 | Used=17.616221/50.0 | Remaining=32.383779000000004 | Usage=35.232442%

49% | 482/977 [23:43<22:16, 2.70s/it]

[Quota] LLM=17.620067 | GPU=0.0 | Used=17.620067/50.0 | Remaining=32.379933 | Usage=35.240134%

49% | 483/977 [23:47<24:45, 3.01s/it]

[Quota] LLM=17.624222 | GPU=0.0 | Used=17.624222/50.0 | Remaining=32.375778 | Usage=35.248444%

50% | 484/977 [23:51<26:54, 3.27s/it]

[Quota] LLM=17.628056 | GPU=0.0 | Used=17.628056/50.0 | Remaining=32.371944 | Usage=35.256112%

50% | 485/977 [23:53<24:11, 2.95s/it]

[Quota] LLM=17.631791 | GPU=0.0 | Used=17.631791/50.0 | Remaining=32.368209 | Usage=35.263582%

50% | 486/977 [23:55<22:01, 2.69s/it]

[Quota] LLM=17.635745 | GPU=0.0 | Used=17.635745/50.0 | Remaining=32.364255 | Usage=35.27149%

50% | 487/977 [23:58<22:46, 2.79s/it]

[Quota] LLM=17.640212000000002 | GPU=0.0 | Used=17.640212000000002/50.0 | Remaining=32.359787999999995 | Usage=35.280424000000004%

50% | 488/977 [24:01<21:58, 2.70s/it]

```
[Quota] LLM=17.644196 | GPU=0.0 | Used=17.644196/50.0 | Remaining=32.355804 |
Usage=35.288392%
50%|      | 489/977 [24:04<22:52,  2.81s/it]

[Quota] LLM=17.648357 | GPU=0.0 | Used=17.648357/50.0 | Remaining=32.351642999999996 |
Usage=35.296714%
50%|      | 490/977 [24:06<22:20,  2.75s/it]

[Quota] LLM=17.652401 | GPU=0.0 | Used=17.652401/50.0 | Remaining=32.347599 |
Usage=35.304802%
50%|      | 491/977 [24:09<20:58,  2.59s/it]

[Quota] LLM=17.656172 | GPU=0.0 | Used=17.656172/50.0 | Remaining=32.343828 |
Usage=35.312344%
50%|      | 492/977 [24:12<21:52,  2.71s/it]

[Quota] LLM=17.660177 | GPU=0.0 | Used=17.660177/50.0 | Remaining=32.339822999999996 |
Usage=35.320354%
50%|      | 493/977 [24:14<21:45,  2.70s/it]

[Quota] LLM=17.664257 | GPU=0.0 | Used=17.664257/50.0 | Remaining=32.335743 |
Usage=35.328514%
51%|      | 494/977 [24:17<22:01,  2.74s/it]

[Quota] LLM=17.668232 | GPU=0.0 | Used=17.668232/50.0 | Remaining=32.331768 |
Usage=35.336464%
51%|      | 495/977 [24:20<23:23,  2.91s/it]

[Quota] LLM=17.672507 | GPU=0.0 | Used=17.672507/50.0 | Remaining=32.327493000000004 |
Usage=35.345014%
51%|      | 496/977 [24:23<22:39,  2.83s/it]

[Quota] LLM=17.676545 | GPU=0.0 | Used=17.676545/50.0 | Remaining=32.323454999999996 |
Usage=35.35309%
51%|      | 497/977 [24:25<21:23,  2.67s/it]

[Quota] LLM=17.680646 | GPU=0.0 | Used=17.680646/50.0 | Remaining=32.319354000000004 |
Usage=35.361292%
51%|      | 498/977 [24:28<20:30,  2.57s/it]

[Quota] LLM=17.684519 | GPU=0.0 | Used=17.684519/50.0 | Remaining=32.315481 |
Usage=35.369038%
51%|      | 499/977 [24:31<23:13,  2.92s/it]

[Quota] LLM=17.68898 | GPU=0.0 | Used=17.68898/50.0 | Remaining=32.31102 |
Usage=35.37796%
51%|      | 500/977 [24:34<22:15,  2.80s/it]
```

```
[Quota] LLM=17.692952000000002 | GPU=0.0 | Used=17.692952000000002/50.0 |  
Remaining=32.307047999999995 | Usage=35.385904000000004%  
51%|      | 501/977 [24:38<25:52,  3.26s/it]  
  
[Quota] LLM=17.697149 | GPU=0.0 | Used=17.697149/50.0 |  
Remaining=32.302851000000004 | Usage=35.394298%  
51%|      | 502/977 [24:41<24:10,  3.05s/it]  
  
[Quota] LLM=17.701265 | GPU=0.0 | Used=17.701265/50.0 | Remaining=32.298735 |  
Usage=35.40253%  
51%|      | 503/977 [24:43<22:54,  2.90s/it]  
  
[Quota] LLM=17.705327 | GPU=0.0 | Used=17.705327/50.0 | Remaining=32.294673 |  
Usage=35.410654%  
52%|      | 504/977 [24:46<21:00,  2.67s/it]  
  
[Quota] LLM=17.709278 | GPU=0.0 | Used=17.709278/50.0 | Remaining=32.290722 |  
Usage=35.418556%  
52%|      | 505/977 [24:47<19:03,  2.42s/it]  
  
[Quota] LLM=17.713157 | GPU=0.0 | Used=17.713157/50.0 |  
Remaining=32.286843000000005 | Usage=35.426314%  
52%|      | 506/977 [24:50<20:07,  2.56s/it]  
  
[Quota] LLM=17.717192 | GPU=0.0 | Used=17.717192/50.0 | Remaining=32.282808 |  
Usage=35.434384%  
52%|      | 507/977 [24:53<20:48,  2.66s/it]  
  
[Quota] LLM=17.721113 | GPU=0.0 | Used=17.721113/50.0 | Remaining=32.278887 |  
Usage=35.442226%  
52%|      | 508/977 [24:57<22:44,  2.91s/it]  
  
[Quota] LLM=17.724881 | GPU=0.0 | Used=17.724881/50.0 |  
Remaining=32.275119000000004 | Usage=35.449762%  
52%|      | 509/977 [25:00<22:53,  2.94s/it]  
  
[Quota] LLM=17.729234 | GPU=0.0 | Used=17.729234/50.0 |  
Remaining=32.270765999999995 | Usage=35.458468%  
52%|      | 510/977 [25:03<22:52,  2.94s/it]  
  
[Quota] LLM=17.733911 | GPU=0.0 | Used=17.733911/50.0 | Remaining=32.266089 |  
Usage=35.467822%  
52%|      | 511/977 [25:05<21:22,  2.75s/it]  
  
[Quota] LLM=17.737979 | GPU=0.0 | Used=17.737979/50.0 |  
Remaining=32.262021000000004 | Usage=35.475958%  
52%|      | 512/977 [25:07<20:19,  2.62s/it]
```

[Quota] LLM=17.741921 | GPU=0.0 | Used=17.741921/50.0 | Remaining=32.258078999999995 | Usage=35.483842%
53%| 513/977 [25:09<19:15, 2.49s/it]

[Quota] LLM=17.745911 | GPU=0.0 | Used=17.745911/50.0 | Remaining=32.254089 | Usage=35.491822%
53%| 514/977 [25:13<22:03, 2.86s/it]

[Quota] LLM=17.749877 | GPU=0.0 | Used=17.749877/50.0 | Remaining=32.250123 | Usage=35.499754%
53%| 515/977 [25:16<21:41, 2.82s/it]

[Quota] LLM=17.753765 | GPU=0.0 | Used=17.753765/50.0 | Remaining=32.246235 | Usage=35.50753%
53%| 516/977 [25:18<20:38, 2.69s/it]

[Quota] LLM=17.757725 | GPU=0.0 | Used=17.757725/50.0 | Remaining=32.242275 | Usage=35.51545%
53%| 517/977 [25:21<20:20, 2.65s/it]

[Quota] LLM=17.761988 | GPU=0.0 | Used=17.761988/50.0 | Remaining=32.238012 | Usage=35.523976%
53%| 518/977 [25:23<19:27, 2.54s/it]

[Quota] LLM=17.765885 | GPU=0.0 | Used=17.765885/50.0 | Remaining=32.234115 | Usage=35.53177%
53%| 519/977 [25:25<18:44, 2.46s/it]

[Quota] LLM=17.769893 | GPU=0.0 | Used=17.769893/50.0 | Remaining=32.230107000000004 | Usage=35.539786%
53%| 520/977 [25:28<18:14, 2.39s/it]

[Quota] LLM=17.77376 | GPU=0.0 | Used=17.77376/50.0 | Remaining=32.226240000000004 | Usage=35.54752%
53%| 521/977 [25:30<18:31, 2.44s/it]

[Quota] LLM=17.777828 | GPU=0.0 | Used=17.777828/50.0 | Remaining=32.222172 | Usage=35.555656%
53%| 522/977 [25:33<20:02, 2.64s/it]

[Quota] LLM=17.781875 | GPU=0.0 | Used=17.781875/50.0 | Remaining=32.218125 | Usage=35.56375%
54%| 523/977 [25:36<19:10, 2.53s/it]

[Quota] LLM=17.785745 | GPU=0.0 | Used=17.785745/50.0 | Remaining=32.214255 | Usage=35.57149%
54%| 524/977 [25:38<18:56, 2.51s/it]

```
[Quota] LLM=17.789582 | GPU=0.0 | Used=17.789582/50.0 |  
Remaining=32.210418000000004 | Usage=35.579164%  
54%|      | 525/977 [25:40<17:21,  2.30s/it]  
  
[Quota] LLM=17.793485 | GPU=0.0 | Used=17.793485/50.0 |  
Remaining=32.206514999999996 | Usage=35.58697%  
54%|      | 526/977 [25:43<18:39,  2.48s/it]  
  
[Quota] LLM=17.797508 | GPU=0.0 | Used=17.797508/50.0 | Remaining=32.202492 |  
Usage=35.595016%  
54%|      | 527/977 [25:45<18:22,  2.45s/it]  
  
[Quota] LLM=17.801486 | GPU=0.0 | Used=17.801486/50.0 | Remaining=32.198514 |  
Usage=35.602972%  
54%|      | 528/977 [25:49<20:34,  2.75s/it]  
  
[Quota] LLM=17.80568 | GPU=0.0 | Used=17.80568/50.0 |  
Remaining=32.19432000000005 | Usage=35.61136%  
54%|      | 529/977 [25:51<20:50,  2.79s/it]  
  
[Quota] LLM=17.809505 | GPU=0.0 | Used=17.809505/50.0 | Remaining=32.190495 |  
Usage=35.61901%  
54%|      | 530/977 [25:54<20:16,  2.72s/it]  
  
[Quota] LLM=17.813423 | GPU=0.0 | Used=17.813423/50.0 | Remaining=32.186577 |  
Usage=35.626846%  
54%|      | 531/977 [26:01<29:42,  4.00s/it]  
  
[Quota] LLM=17.817338 | GPU=0.0 | Used=17.817338/50.0 | Remaining=32.182662 |  
Usage=35.634676%  
54%|      | 532/977 [26:04<26:43,  3.60s/it]  
  
[Quota] LLM=17.821238 | GPU=0.0 | Used=17.821238/50.0 | Remaining=32.178762 |  
Usage=35.642476%  
55%|      | 533/977 [26:06<24:02,  3.25s/it]  
  
[Quota] LLM=17.825072 | GPU=0.0 | Used=17.825072/50.0 | Remaining=32.174928 |  
Usage=35.650144%  
55%|      | 534/977 [26:09<22:46,  3.08s/it]  
  
[Quota] LLM=17.828984 | GPU=0.0 | Used=17.828984/50.0 | Remaining=32.171016 |  
Usage=35.657968%  
55%|      | 535/977 [26:11<21:28,  2.91s/it]  
  
[Quota] LLM=17.833043 | GPU=0.0 | Used=17.833043/50.0 | Remaining=32.166957 |  
Usage=35.666086%  
55%|      | 536/977 [26:14<21:11,  2.88s/it]
```

```
[Quota] LLM=17.836874 | GPU=0.0 | Used=17.836874/50.0 | Remaining=32.163126 | Usage=35.673748%
55%|      | 537/977 [26:17<20:51, 2.85s/it]

[Quota] LLM=17.840708 | GPU=0.0 | Used=17.840708/50.0 | Remaining=32.159292 | Usage=35.681416%
55%|      | 538/977 [26:21<23:57, 3.27s/it]

[Quota] LLM=17.844578 | GPU=0.0 | Used=17.844578/50.0 | Remaining=32.155422 | Usage=35.689156%
55%|      | 539/977 [26:50<1:19:06, 10.84s/it]

[Quota] LLM=17.848793 | GPU=0.0 | Used=17.848793/50.0 | Remaining=32.151207 | Usage=35.697586%
55%|      | 540/977 [26:52<1:01:03, 8.38s/it]

[Quota] LLM=17.852612 | GPU=0.0 | Used=17.852612/50.0 | Remaining=32.147388 | Usage=35.705224%
55%|      | 541/977 [26:56<50:41, 6.98s/it]

[Quota] LLM=17.856503 | GPU=0.0 | Used=17.856503/50.0 | Remaining=32.143496999999996 | Usage=35.713006%
55%|      | 542/977 [26:59<42:36, 5.88s/it]

[Quota] LLM=17.860472 | GPU=0.0 | Used=17.860472/50.0 | Remaining=32.139528 | Usage=35.720944%
56%|      | 543/977 [27:02<36:02, 4.98s/it]

[Quota] LLM=17.864282 | GPU=0.0 | Used=17.864282/50.0 | Remaining=32.135718 | Usage=35.728564%
56%|      | 544/977 [27:05<30:26, 4.22s/it]

[Quota] LLM=17.868083 | GPU=0.0 | Used=17.868083/50.0 | Remaining=32.131917 | Usage=35.736166%
56%|      | 545/977 [27:07<27:14, 3.78s/it]

[Quota] LLM=17.872049 | GPU=0.0 | Used=17.872049/50.0 | Remaining=32.127950999999996 | Usage=35.744098%
56%|      | 546/977 [27:12<28:01, 3.90s/it]

[Quota] LLM=17.876102 | GPU=0.0 | Used=17.876102/50.0 | Remaining=32.123898 | Usage=35.752204%
56%|      | 547/977 [27:14<25:30, 3.56s/it]

[Quota] LLM=17.879954 | GPU=0.0 | Used=17.879954/50.0 | Remaining=32.120046 | Usage=35.759908%
56%|      | 548/977 [27:17<24:06, 3.37s/it]
```

[Quota] LLM=17.88422 | GPU=0.0 | Used=17.88422/50.0 | Remaining=32.11578 | Usage=35.76844%

56%| 549/977 [27:20<22:31, 3.16s/it]

[Quota] LLM=17.888156 | GPU=0.0 | Used=17.888156/50.0 | Remaining=32.111844000000005 | Usage=35.776312%

56%| 550/977 [27:23<22:36, 3.18s/it]

[Quota] LLM=17.892743 | GPU=0.0 | Used=17.892743/50.0 | Remaining=32.107257000000004 | Usage=35.785486%

56%| 551/977 [27:25<20:24, 2.87s/it]

[Quota] LLM=17.896511 | GPU=0.0 | Used=17.896511/50.0 | Remaining=32.103488999999996 | Usage=35.793022%

56%| 552/977 [27:28<20:05, 2.84s/it]

[Quota] LLM=17.900447 | GPU=0.0 | Used=17.900447/50.0 | Remaining=32.099553 | Usage=35.800894%

57%| 553/977 [27:31<19:17, 2.73s/it]

[Quota] LLM=17.904374 | GPU=0.0 | Used=17.904374/50.0 | Remaining=32.095625999999996 | Usage=35.808748%

57%| 554/977 [27:33<17:56, 2.54s/it]

[Quota] LLM=17.908166 | GPU=0.0 | Used=17.908166/50.0 | Remaining=32.091834 | Usage=35.816332%

57%| 555/977 [27:35<16:59, 2.42s/it]

[Quota] LLM=17.911964 | GPU=0.0 | Used=17.911964/50.0 | Remaining=32.088036 | Usage=35.823928%

57%| 556/977 [27:37<16:47, 2.39s/it]

[Quota] LLM=17.91572 | GPU=0.0 | Used=17.91572/50.0 | Remaining=32.08428 | Usage=35.83144%

57%| 557/977 [27:40<17:05, 2.44s/it]

[Quota] LLM=17.919806 | GPU=0.0 | Used=17.919806/50.0 | Remaining=32.080194 | Usage=35.839612%

57%| 558/977 [27:42<16:06, 2.31s/it]

[Quota] LLM=17.923646 | GPU=0.0 | Used=17.923646/50.0 | Remaining=32.076353999999995 | Usage=35.847292%

57%| 559/977 [27:44<15:34, 2.24s/it]

[Quota] LLM=17.927369 | GPU=0.0 | Used=17.927369/50.0 | Remaining=32.072631 | Usage=35.854738%

57%| 560/977 [27:46<15:10, 2.18s/it]

[Quota] LLM=17.93114 | GPU=0.0 | Used=17.93114/50.0 | Remaining=32.06886 | Usage=35.86228%
57%| 561/977 [27:48<15:53, 2.29s/it]

[Quota] LLM=17.934905 | GPU=0.0 | Used=17.934905/50.0 | Remaining=32.065095 | Usage=35.86981%
58%| 562/977 [27:51<15:53, 2.30s/it]

[Quota] LLM=17.938883 | GPU=0.0 | Used=17.938883/50.0 | Remaining=32.061116999999996 | Usage=35.877766%
58%| 563/977 [27:54<18:22, 2.66s/it]

[Quota] LLM=17.942666 | GPU=0.0 | Used=17.942666/50.0 | Remaining=32.057334 | Usage=35.885332%
58%| 564/977 [27:57<18:04, 2.63s/it]

[Quota] LLM=17.946722 | GPU=0.0 | Used=17.946722/50.0 | Remaining=32.053278 | Usage=35.893444%
58%| 565/977 [28:00<18:51, 2.75s/it]

[Quota] LLM=17.950586 | GPU=0.0 | Used=17.950586/50.0 | Remaining=32.049414 | Usage=35.901172%
58%| 566/977 [28:02<17:48, 2.60s/it]

[Quota] LLM=17.954438 | GPU=0.0 | Used=17.954438/50.0 | Remaining=32.045562000000004 | Usage=35.908876%
58%| 567/977 [28:04<17:33, 2.57s/it]

[Quota] LLM=17.958275 | GPU=0.0 | Used=17.958275/50.0 | Remaining=32.041725 | Usage=35.91655%
58%| 568/977 [28:07<18:20, 2.69s/it]

[Quota] LLM=17.962163 | GPU=0.0 | Used=17.962163/50.0 | Remaining=32.037836999999996 | Usage=35.924326%
58%| 569/977 [28:10<18:19, 2.70s/it]

[Quota] LLM=17.966246 | GPU=0.0 | Used=17.966246/50.0 | Remaining=32.033754 | Usage=35.932492%
58%| 570/977 [28:13<19:07, 2.82s/it]

[Quota] LLM=17.970032 | GPU=0.0 | Used=17.970032/50.0 | Remaining=32.029968 | Usage=35.940064%
58%| 571/977 [28:16<18:21, 2.71s/it]

[Quota] LLM=17.973896 | GPU=0.0 | Used=17.973896/50.0 | Remaining=32.026104000000004 | Usage=35.947792%
59%| 572/977 [28:19<18:45, 2.78s/it]

```
[Quota] LLM=17.977985 | GPU=0.0 | Used=17.977985/50.0 |  
Remaining=32.022014999999996 | Usage=35.95597%  
59%|      | 573/977 [28:21<18:30,  2.75s/it]  
  
[Quota] LLM=17.981816 | GPU=0.0 | Used=17.981816/50.0 |  
Remaining=32.018184000000005 | Usage=35.963632%  
59%|      | 574/977 [28:24<18:50,  2.80s/it]  
  
[Quota] LLM=17.98604 | GPU=0.0 | Used=17.98604/50.0 | Remaining=32.01396 |  
Usage=35.97208%  
59%|      | 575/977 [28:27<18:35,  2.78s/it]  
  
[Quota] LLM=17.99009 | GPU=0.0 | Used=17.99009/50.0 |  
Remaining=32.009910000000005 | Usage=35.98018%  
59%|      | 576/977 [28:30<18:49,  2.82s/it]  
  
[Quota] LLM=17.994107 | GPU=0.0 | Used=17.994107/50.0 | Remaining=32.005893 |  
Usage=35.988214%  
59%|      | 577/977 [28:32<18:08,  2.72s/it]  
  
[Quota] LLM=17.99786 | GPU=0.0 | Used=17.99786/50.0 | Remaining=32.00214 |  
Usage=35.99572%  
59%|      | 578/977 [28:35<17:30,  2.63s/it]  
  
[Quota] LLM=18.001829 | GPU=0.0 | Used=18.001829/50.0 | Remaining=31.998171 |  
Usage=36.003658%  
59%|      | 579/977 [28:37<16:40,  2.51s/it]  
  
[Quota] LLM=18.005822 | GPU=0.0 | Used=18.005822/50.0 | Remaining=31.994178 |  
Usage=36.011644%  
59%|      | 580/977 [28:40<16:38,  2.52s/it]  
  
[Quota] LLM=18.009902 | GPU=0.0 | Used=18.009902/50.0 | Remaining=31.990098 |  
Usage=36.019804%  
59%|      | 581/977 [28:42<16:50,  2.55s/it]  
  
[Quota] LLM=18.014099 | GPU=0.0 | Used=18.014099/50.0 | Remaining=31.985901 |  
Usage=36.028198%  
60%|      | 582/977 [28:45<17:08,  2.60s/it]  
  
[Quota] LLM=18.018173 | GPU=0.0 | Used=18.018173/50.0 | Remaining=31.981827 |  
Usage=36.036346%  
60%|      | 583/977 [28:47<16:32,  2.52s/it]  
  
[Quota] LLM=18.022145 | GPU=0.0 | Used=18.022145/50.0 | Remaining=31.977855 |  
Usage=36.04429%  
60%|      | 584/977 [28:50<17:16,  2.64s/it]
```

[Quota] LLM=18.026162 | GPU=0.0 | Used=18.026162/50.0 | Remaining=31.973838 | Usage=36.052324%

60%| 585/977 [28:53<17:57, 2.75s/it]

[Quota] LLM=18.030368 | GPU=0.0 | Used=18.030368/50.0 | Remaining=31.969632 | Usage=36.060736%

60%| 586/977 [28:56<17:31, 2.69s/it]

[Quota] LLM=18.034301 | GPU=0.0 | Used=18.034301/50.0 | Remaining=31.965699 | Usage=36.068602%

60%| 587/977 [28:59<18:36, 2.86s/it]

[Quota] LLM=18.038288 | GPU=0.0 | Used=18.038288/50.0 | Remaining=31.961712 | Usage=36.076576%

60%| 588/977 [29:02<18:30, 2.86s/it]

[Quota] LLM=18.042437 | GPU=0.0 | Used=18.042437/50.0 | Remaining=31.957563 | Usage=36.084874%

60%| 589/977 [29:06<20:03, 3.10s/it]

[Quota] LLM=18.047051 | GPU=0.0 | Used=18.047051/50.0 | Remaining=31.952949 | Usage=36.094102%

60%| 590/977 [29:09<20:59, 3.26s/it]

[Quota] LLM=18.051218 | GPU=0.0 | Used=18.051218/50.0 | Remaining=31.948782 | Usage=36.102436%

60%| 591/977 [29:12<19:48, 3.08s/it]

[Quota] LLM=18.055256 | GPU=0.0 | Used=18.055256/50.0 | Remaining=31.944744 | Usage=36.110512%

61%| 592/977 [29:14<18:34, 2.89s/it]

[Quota] LLM=18.059513 | GPU=0.0 | Used=18.059513/50.0 | Remaining=31.940487 | Usage=36.119026%

61%| 593/977 [29:17<18:49, 2.94s/it]

[Quota] LLM=18.063476 | GPU=0.0 | Used=18.063476/50.0 | Remaining=31.936524 | Usage=36.126952%

61%| 594/977 [29:20<17:45, 2.78s/it]

[Quota] LLM=18.067283 | GPU=0.0 | Used=18.067283/50.0 | Remaining=31.932717 | Usage=36.134566%

61%| 595/977 [29:23<17:42, 2.78s/it]

[Quota] LLM=18.07121 | GPU=0.0 | Used=18.07121/50.0 | Remaining=31.92879 | Usage=36.14242%

61%| 596/977 [29:25<17:34, 2.77s/it]

[Quota] LLM=18.075191 | GPU=0.0 | Used=18.075191/50.0 | Remaining=31.924809 | Usage=36.150382%

61% | 597/977 [29:28<17:41, 2.79s/it]

[Quota] LLM=18.078992 | GPU=0.0 | Used=18.078992/50.0 | Remaining=31.921008 | Usage=36.157984%

61% | 598/977 [29:31<17:23, 2.75s/it]

[Quota] LLM=18.082778 | GPU=0.0 | Used=18.082778/50.0 | Remaining=31.917222 | Usage=36.165556%

61% | 599/977 [29:33<16:53, 2.68s/it]

[Quota] LLM=18.086558 | GPU=0.0 | Used=18.086558/50.0 | Remaining=31.913442 | Usage=36.173116%

61% | 600/977 [29:36<16:45, 2.67s/it]

[Quota] LLM=18.090464 | GPU=0.0 | Used=18.090464/50.0 | Remaining=31.909536 | Usage=36.180928%

62% | 601/977 [29:39<17:37, 2.81s/it]

[Quota] LLM=18.094475 | GPU=0.0 | Used=18.094475/50.0 | Remaining=31.905525 | Usage=36.18895%

62% | 602/977 [29:42<18:39, 2.99s/it]

[Quota] LLM=18.098888 | GPU=0.0 | Used=18.098888/50.0 | Remaining=31.901112 | Usage=36.197776%

62% | 603/977 [29:45<18:07, 2.91s/it]

[Quota] LLM=18.102761 | GPU=0.0 | Used=18.102761/50.0 | Remaining=31.897239 | Usage=36.205522%

62% | 604/977 [29:47<16:35, 2.67s/it]

[Quota] LLM=18.106661 | GPU=0.0 | Used=18.106661/50.0 | Remaining=31.893339 | Usage=36.213322%

62% | 605/977 [29:51<18:30, 2.98s/it]

[Quota] LLM=18.111257 | GPU=0.0 | Used=18.111257/50.0 | Remaining=31.888743 | Usage=36.222514%

62% | 606/977 [29:55<20:10, 3.26s/it]

[Quota] LLM=18.115379 | GPU=0.0 | Used=18.115379/50.0 | Remaining=31.884621 | Usage=36.230758%

62% | 607/977 [29:58<19:43, 3.20s/it]

[Quota] LLM=18.119366 | GPU=0.0 | Used=18.119366/50.0 | Remaining=31.880634 | Usage=36.238732%

62% | 608/977 [30:02<21:07, 3.44s/it]

[Quota] LLM=18.123368 | GPU=0.0 | Used=18.123368/50.0 | Remaining=31.876632 | Usage=36.246736%

62%| 609/977 [30:05<19:56, 3.25s/it]

[Quota] LLM=18.127439 | GPU=0.0 | Used=18.127439/50.0 | Remaining=31.872561 | Usage=36.254878%

62%| 610/977 [30:08<19:47, 3.24s/it]

[Quota] LLM=18.131453 | GPU=0.0 | Used=18.131453/50.0 | Remaining=31.868547 | Usage=36.262906%

63%| 611/977 [30:10<18:07, 2.97s/it]

[Quota] LLM=18.135386 | GPU=0.0 | Used=18.135386/50.0 | Remaining=31.864614 | Usage=36.270772%

63%| 612/977 [30:13<17:00, 2.80s/it]

[Quota] LLM=18.139235 | GPU=0.0 | Used=18.139235/50.0 | Remaining=31.860765 | Usage=36.27847%

63%| 613/977 [30:15<16:28, 2.72s/it]

[Quota] LLM=18.143087 | GPU=0.0 | Used=18.143087/50.0 | Remaining=31.856913 | Usage=36.286174%

63%| 614/977 [30:18<16:37, 2.75s/it]

[Quota] LLM=18.147062000000002 | GPU=0.0 | Used=18.147062000000002/50.0 | Remaining=31.852937999999998 | Usage=36.294124000000004%

63%| 615/977 [30:21<16:50, 2.79s/it]

[Quota] LLM=18.151001 | GPU=0.0 | Used=18.151001/50.0 | Remaining=31.848999 | Usage=36.302002%

63%| 616/977 [30:23<16:08, 2.68s/it]

[Quota] LLM=18.155003 | GPU=0.0 | Used=18.155003/50.0 | Remaining=31.844997 | Usage=36.310006%

63%| 617/977 [30:26<15:49, 2.64s/it]

[Quota] LLM=18.158873 | GPU=0.0 | Used=18.158873/50.0 | Remaining=31.841127 | Usage=36.317746%

63%| 618/977 [30:29<16:49, 2.81s/it]

[Quota] LLM=18.162626 | GPU=0.0 | Used=18.162626/50.0 | Remaining=31.837374 | Usage=36.325252%

63%| 619/977 [30:33<17:45, 2.98s/it]

[Quota] LLM=18.16649 | GPU=0.0 | Used=18.16649/50.0 | Remaining=31.83351 | Usage=36.33298%

63%| 620/977 [30:34<15:44, 2.64s/it]

```
[Quota] LLM=18.170222 | GPU=0.0 | Used=18.170222/50.0 | Remaining=31.829778 |
Usage=36.340444%  
64%|       | 621/977 [30:37<15:42,  2.65s/it]  
  
[Quota] LLM=18.17429 | GPU=0.0 | Used=18.17429/50.0 | Remaining=31.82571 |
Usage=36.34858%  
64%|       | 622/977 [30:41<18:23,  3.11s/it]  
  
[Quota] LLM=18.178643 | GPU=0.0 | Used=18.178643/50.0 | Remaining=31.821357 |
Usage=36.357286%  
64%|       | 623/977 [30:44<17:13,  2.92s/it]  
  
[Quota] LLM=18.182612 | GPU=0.0 | Used=18.182612/50.0 | Remaining=31.817388 |
Usage=36.365224%  
64%|       | 624/977 [30:46<16:36,  2.82s/it]  
  
[Quota] LLM=18.186572 | GPU=0.0 | Used=18.186572/50.0 | Remaining=31.813428 |
Usage=36.373144%  
64%|       | 625/977 [30:49<15:29,  2.64s/it]  
  
[Quota] LLM=18.19046 | GPU=0.0 | Used=18.19046/50.0 | Remaining=31.80954 |
Usage=36.38092%  
64%|       | 626/977 [30:51<14:44,  2.52s/it]  
  
[Quota] LLM=18.194171 | GPU=0.0 | Used=18.194171/50.0 | Remaining=31.805829 |
Usage=36.388342%  
64%|       | 627/977 [30:53<14:50,  2.55s/it]  
  
[Quota] LLM=18.198029000000002 | GPU=0.0 | Used=18.198029000000002/50.0 |
Remaining=31.801970999999998 | Usage=36.396058000000004%  
64%|       | 628/977 [30:56<14:39,  2.52s/it]  
  
[Quota] LLM=18.201743 | GPU=0.0 | Used=18.201743/50.0 | Remaining=31.798257 |
Usage=36.403486%  
64%|       | 629/977 [30:58<14:38,  2.52s/it]  
  
[Quota] LLM=18.205601 | GPU=0.0 | Used=18.205601/50.0 | Remaining=31.794399 |
Usage=36.411202%  
64%|       | 630/977 [31:01<14:48,  2.56s/it]  
  
[Quota] LLM=18.209705 | GPU=0.0 | Used=18.209705/50.0 | Remaining=31.790295 |
Usage=36.41941%  
65%|       | 631/977 [31:03<14:14,  2.47s/it]  
  
[Quota] LLM=18.21356 | GPU=0.0 | Used=18.21356/50.0 | Remaining=31.78644 |
Usage=36.42712%  
65%|       | 632/977 [31:05<13:43,  2.39s/it]
```

```
[Quota] LLM=18.217496 | GPU=0.0 | Used=18.217496/50.0 | Remaining=31.782504 | Usage=36.434992%
65%|       | 633/977 [31:07<13:05, 2.28s/it]

[Quota] LLM=18.221369 | GPU=0.0 | Used=18.221369/50.0 | Remaining=31.778631 | Usage=36.442738%
65%|       | 634/977 [31:10<13:56, 2.44s/it]

[Quota] LLM=18.225293 | GPU=0.0 | Used=18.225293/50.0 | Remaining=31.774707 | Usage=36.450586%
65%|       | 635/977 [31:13<13:50, 2.43s/it]

[Quota] LLM=18.229211 | GPU=0.0 | Used=18.229211/50.0 | Remaining=31.770789 | Usage=36.458422%
65%|       | 636/977 [31:15<13:37, 2.40s/it]

[Quota] LLM=18.232979 | GPU=0.0 | Used=18.232979/50.0 | Remaining=31.767021 | Usage=36.465958%
65%|       | 637/977 [31:17<12:33, 2.22s/it]

[Quota] LLM=18.236858 | GPU=0.0 | Used=18.236858/50.0 | Remaining=31.763142 | Usage=36.473716%
65%|       | 638/977 [31:19<13:08, 2.33s/it]

[Quota] LLM=18.24092 | GPU=0.0 | Used=18.24092/50.0 | Remaining=31.75908 | Usage=36.48184%
65%|       | 639/977 [31:22<13:59, 2.48s/it]

[Quota] LLM=18.244952 | GPU=0.0 | Used=18.244952/50.0 | Remaining=31.755048 | Usage=36.489904%
66%|       | 640/977 [31:25<13:59, 2.49s/it]

[Quota] LLM=18.248927 | GPU=0.0 | Used=18.248927/50.0 | Remaining=31.751073 | Usage=36.497854%
66%|       | 641/977 [31:27<13:55, 2.49s/it]

[Quota] LLM=18.253037 | GPU=0.0 | Used=18.253037/50.0 | Remaining=31.746963 | Usage=36.506074%
66%|       | 642/977 [31:30<14:06, 2.53s/it]

[Quota] LLM=18.256919 | GPU=0.0 | Used=18.256919/50.0 | Remaining=31.743081 | Usage=36.513838%
66%|       | 643/977 [31:32<13:51, 2.49s/it]

[Quota] LLM=18.26078 | GPU=0.0 | Used=18.26078/50.0 | Remaining=31.73922 | Usage=36.52156%
66%|       | 644/977 [31:35<13:35, 2.45s/it]
```

[Quota] LLM=18.264635 | GPU=0.0 | Used=18.264635/50.0 | Remaining=31.735365 | Usage=36.52927%
66%| 645/977 [31:37<14:02, 2.54s/it]

[Quota] LLM=18.268442 | GPU=0.0 | Used=18.268442/50.0 | Remaining=31.731558 | Usage=36.536884%
66%| 646/977 [31:41<15:05, 2.74s/it]

[Quota] LLM=18.272612 | GPU=0.0 | Used=18.272612/50.0 | Remaining=31.727388 | Usage=36.545224%
66%| 647/977 [31:43<14:53, 2.71s/it]

[Quota] LLM=18.276476 | GPU=0.0 | Used=18.276476/50.0 | Remaining=31.723524 | Usage=36.552952%
66%| 648/977 [31:46<14:45, 2.69s/it]

[Quota] LLM=18.280634 | GPU=0.0 | Used=18.280634/50.0 | Remaining=31.719366 | Usage=36.561268%
66%| 649/977 [31:49<15:15, 2.79s/it]

[Quota] LLM=18.284687 | GPU=0.0 | Used=18.284687/50.0 | Remaining=31.715313 | Usage=36.569374%
67%| 650/977 [31:51<14:52, 2.73s/it]

[Quota] LLM=18.288653 | GPU=0.0 | Used=18.288653/50.0 | Remaining=31.711347 | Usage=36.577306%
67%| 651/977 [31:54<14:30, 2.67s/it]

[Quota] LLM=18.292748 | GPU=0.0 | Used=18.292748/50.0 | Remaining=31.707252 | Usage=36.585496%
67%| 652/977 [31:56<13:47, 2.55s/it]

[Quota] LLM=18.296579 | GPU=0.0 | Used=18.296579/50.0 | Remaining=31.703421 | Usage=36.593158%
67%| 653/977 [31:59<14:48, 2.74s/it]

[Quota] LLM=18.30047 | GPU=0.0 | Used=18.30047/50.0 | Remaining=31.69953 | Usage=36.60094%
67%| 654/977 [32:02<14:37, 2.72s/it]

[Quota] LLM=18.30425 | GPU=0.0 | Used=18.30425/50.0 | Remaining=31.69575 | Usage=36.6085%
67%| 655/977 [32:04<13:52, 2.59s/it]

[Quota] LLM=18.308102 | GPU=0.0 | Used=18.308102/50.0 | Remaining=31.691898 | Usage=36.616204%
67%| 656/977 [32:07<13:26, 2.51s/it]

[Quota] LLM=18.31193 | GPU=0.0 | Used=18.31193/50.0 | Remaining=31.68807 | Usage=36.62386%

67%| 657/977 [32:09<13:41, 2.57s/it]

[Quota] LLM=18.315722 | GPU=0.0 | Used=18.315722/50.0 | Remaining=31.684278 | Usage=36.631444%

67%| 658/977 [32:11<12:51, 2.42s/it]

[Quota] LLM=18.319505 | GPU=0.0 | Used=18.319505/50.0 | Remaining=31.680495 | Usage=36.63901%

67%| 659/977 [32:14<13:13, 2.50s/it]

[Quota] LLM=18.323258 | GPU=0.0 | Used=18.323258/50.0 | Remaining=31.676742 | Usage=36.646516%

68%| 660/977 [32:17<13:21, 2.53s/it]

[Quota] LLM=18.327143 | GPU=0.0 | Used=18.327143/50.0 | Remaining=31.672857 | Usage=36.654286%

68%| 661/977 [32:19<12:50, 2.44s/it]

[Quota] LLM=18.33116 | GPU=0.0 | Used=18.33116/50.0 | Remaining=31.66884 | Usage=36.66232%

68%| 662/977 [32:22<12:59, 2.47s/it]

[Quota] LLM=18.335165 | GPU=0.0 | Used=18.335165/50.0 | Remaining=31.664835 | Usage=36.67033%

68%| 663/977 [32:24<12:51, 2.46s/it]

[Quota] LLM=18.339185 | GPU=0.0 | Used=18.339185/50.0 | Remaining=31.660815 | Usage=36.67837%

68%| 664/977 [32:26<12:49, 2.46s/it]

[Quota] LLM=18.343121 | GPU=0.0 | Used=18.343121/50.0 | Remaining=31.656879 | Usage=36.686242%

68%| 665/977 [32:29<13:41, 2.63s/it]

[Quota] LLM=18.347114 | GPU=0.0 | Used=18.347114/50.0 | Remaining=31.652886 | Usage=36.694228%

68%| 666/977 [32:32<14:07, 2.73s/it]

[Quota] LLM=18.351011 | GPU=0.0 | Used=18.351011/50.0 | Remaining=31.648989 | Usage=36.702022%

68%| 667/977 [32:36<14:58, 2.90s/it]

[Quota] LLM=18.355637 | GPU=0.0 | Used=18.355637/50.0 | Remaining=31.644363 | Usage=36.711274%

68%| 668/977 [32:38<14:16, 2.77s/it]

```
[Quota] LLM=18.359999 | GPU=0.0 | Used=18.359999/50.0 | Remaining=31.640001 | Usage=36.719998%
68%|       | 669/977 [32:41<15:01, 2.93s/it]

[Quota] LLM=18.364037 | GPU=0.0 | Used=18.364037/50.0 | Remaining=31.635963 | Usage=36.728074%
69%|       | 670/977 [32:46<16:45, 3.28s/it]

[Quota] LLM=18.368657 | GPU=0.0 | Used=18.368657/50.0 | Remaining=31.631343 | Usage=36.737314%
69%|       | 671/977 [32:48<15:48, 3.10s/it]

[Quota] LLM=18.372629 | GPU=0.0 | Used=18.372629/50.0 | Remaining=31.627371 | Usage=36.745258%
69%|       | 672/977 [32:51<14:38, 2.88s/it]

[Quota] LLM=18.376562 | GPU=0.0 | Used=18.376562/50.0 | Remaining=31.623438 | Usage=36.753124%
69%|       | 673/977 [32:54<15:19, 3.03s/it]

[Quota] LLM=18.380912 | GPU=0.0 | Used=18.380912/50.0 | Remaining=31.619088 | Usage=36.761824%
69%|       | 674/977 [32:57<15:11, 3.01s/it]

[Quota] LLM=18.384746 | GPU=0.0 | Used=18.384746/50.0 | Remaining=31.615254 | Usage=36.769492%
69%|       | 675/977 [33:00<14:41, 2.92s/it]

[Quota] LLM=18.388571 | GPU=0.0 | Used=18.388571/50.0 | Remaining=31.611429 | Usage=36.777142%
69%|       | 676/977 [33:03<15:02, 3.00s/it]

[Quota] LLM=18.392939 | GPU=0.0 | Used=18.392939/50.0 | Remaining=31.607061 | Usage=36.785878%
69%|       | 677/977 [33:06<14:52, 2.98s/it]

[Quota] LLM=18.396974 | GPU=0.0 | Used=18.396974/50.0 | Remaining=31.603026 | Usage=36.793948%
69%|       | 678/977 [33:08<14:16, 2.86s/it]

[Quota] LLM=18.400931 | GPU=0.0 | Used=18.400931/50.0 | Remaining=31.599069 | Usage=36.801862%
69%|       | 679/977 [33:11<14:10, 2.85s/it]

[Quota] LLM=18.405443 | GPU=0.0 | Used=18.405443/50.0 | Remaining=31.594557 | Usage=36.810886%
70%|       | 680/977 [33:14<14:20, 2.90s/it]
```

[Quota] LLM=18.409817 | GPU=0.0 | Used=18.409817/50.0 | Remaining=31.590183 | Usage=36.819634%

70%| 681/977 [33:16<13:15, 2.69s/it]

[Quota] LLM=18.413747 | GPU=0.0 | Used=18.413747/50.0 | Remaining=31.586253 | Usage=36.827494%

70%| 682/977 [33:20<13:54, 2.83s/it]

[Quota] LLM=18.417986 | GPU=0.0 | Used=18.417986/50.0 | Remaining=31.582014 | Usage=36.835972%

70%| 683/977 [33:22<13:36, 2.78s/it]

[Quota] LLM=18.421817 | GPU=0.0 | Used=18.421817/50.0 | Remaining=31.578183 | Usage=36.843634%

70%| 684/977 [33:25<13:10, 2.70s/it]

[Quota] LLM=18.425756 | GPU=0.0 | Used=18.425756/50.0 | Remaining=31.574244 | Usage=36.851512%

70%| 685/977 [33:27<12:43, 2.62s/it]

[Quota] LLM=18.429782 | GPU=0.0 | Used=18.429782/50.0 | Remaining=31.570218 | Usage=36.859564%

70%| 686/977 [33:30<13:23, 2.76s/it]

[Quota] LLM=18.433811 | GPU=0.0 | Used=18.433811/50.0 | Remaining=31.566189 | Usage=36.867622%

70%| 687/977 [33:34<15:03, 3.11s/it]

[Quota] LLM=18.437681 | GPU=0.0 | Used=18.437681/50.0 | Remaining=31.562319 | Usage=36.875362%

70%| 688/977 [33:46<26:52, 5.58s/it]

[Quota] LLM=18.441578 | GPU=0.0 | Used=18.441578/50.0 | Remaining=31.558422 | Usage=36.883156%

71%| 689/977 [33:48<22:34, 4.70s/it]

[Quota] LLM=18.445724 | GPU=0.0 | Used=18.445724/50.0 | Remaining=31.554276 | Usage=36.891448%

71%| 690/977 [33:51<19:29, 4.07s/it]

[Quota] LLM=18.449858 | GPU=0.0 | Used=18.449858/50.0 | Remaining=31.550142 | Usage=36.899716%

71%| 691/977 [33:53<17:21, 3.64s/it]

[Quota] LLM=18.453959 | GPU=0.0 | Used=18.453959/50.0 | Remaining=31.546041 | Usage=36.907918%

71%| 692/977 [33:56<15:52, 3.34s/it]

```
[Quota] LLM=18.458207 | GPU=0.0 | Used=18.458207/50.0 | Remaining=31.541793 | Usage=36.916414%
71%|      | 693/977 [33:59<15:47,  3.34s/it]

[Quota] LLM=18.462542 | GPU=0.0 | Used=18.462542/50.0 | Remaining=31.537458 | Usage=36.925084%
71%|      | 694/977 [34:02<14:32,  3.08s/it]

[Quota] LLM=18.466475 | GPU=0.0 | Used=18.466475/50.0 | Remaining=31.533525 | Usage=36.93295%
71%|      | 695/977 [34:05<14:03,  2.99s/it]

[Quota] LLM=18.470663 | GPU=0.0 | Used=18.470663/50.0 | Remaining=31.529337 | Usage=36.941326%
71%|      | 696/977 [34:07<13:20,  2.85s/it]

[Quota] LLM=18.474468 | GPU=0.0 | Used=18.474468/50.0 | Remaining=31.52532 | Usage=36.94936%
71%|      | 697/977 [34:10<13:46,  2.95s/it]

[Quota] LLM=18.478796 | GPU=0.0 | Used=18.478796/50.0 | Remaining=31.521204 | Usage=36.957592%
71%|      | 698/977 [34:14<14:20,  3.08s/it]

[Quota] LLM=18.483158 | GPU=0.0 | Used=18.483158/50.0 | Remaining=31.516842 | Usage=36.966316%
72%|      | 699/977 [34:16<13:36,  2.94s/it]

[Quota] LLM=18.487502 | GPU=0.0 | Used=18.487502/50.0 | Remaining=31.512498 | Usage=36.975004%
72%|      | 700/977 [34:20<14:41,  3.18s/it]

[Quota] LLM=18.491699 | GPU=0.0 | Used=18.491699/50.0 | Remaining=31.508301 | Usage=36.983398%
72%|      | 701/977 [34:23<14:43,  3.20s/it]

[Quota] LLM=18.495674 | GPU=0.0 | Used=18.495674/50.0 | Remaining=31.504326 | Usage=36.991348%
72%|      | 702/977 [34:27<15:16,  3.33s/it]

[Quota] LLM=18.500075 | GPU=0.0 | Used=18.500075/50.0 | Remaining=31.499925 | Usage=37.00015%
72%|      | 703/977 [34:30<14:27,  3.16s/it]

[Quota] LLM=18.50414 | GPU=0.0 | Used=18.50414/50.0 | Remaining=31.49586 | Usage=37.00828%
72%|      | 704/977 [34:32<13:19,  2.93s/it]
```

[Quota] LLM=18.50819 | GPU=0.0 | Used=18.50819/50.0 | Remaining=31.49181 | Usage=37.01638%

72%| 705/977 [34:35<12:58, 2.86s/it]

[Quota] LLM=18.512303 | GPU=0.0 | Used=18.512303/50.0 | Remaining=31.487697 | Usage=37.024606%

72%| 706/977 [34:38<12:44, 2.82s/it]

[Quota] LLM=18.51662 | GPU=0.0 | Used=18.51662/50.0 | Remaining=31.48338 | Usage=37.03324%

72%| 707/977 [34:40<12:10, 2.71s/it]

[Quota] LLM=18.520508 | GPU=0.0 | Used=18.520508/50.0 | Remaining=31.479492 | Usage=37.041016%

72%| 708/977 [34:43<13:09, 2.94s/it]

[Quota] LLM=18.524423 | GPU=0.0 | Used=18.524423/50.0 | Remaining=31.475577 | Usage=37.048846%

73%| 709/977 [34:46<12:29, 2.80s/it]

[Quota] LLM=18.528626 | GPU=0.0 | Used=18.528626/50.0 | Remaining=31.471374 | Usage=37.057252%

73%| 710/977 [34:50<13:48, 3.10s/it]

[Quota] LLM=18.533435 | GPU=0.0 | Used=18.533435/50.0 | Remaining=31.466565 | Usage=37.06687%

73%| 711/977 [34:53<13:52, 3.13s/it]

[Quota] LLM=18.537971 | GPU=0.0 | Used=18.537971/50.0 | Remaining=31.462029 | Usage=37.075942%

73%| 712/977 [34:56<14:18, 3.24s/it]

[Quota] LLM=18.54257 | GPU=0.0 | Used=18.54257/50.0 | Remaining=31.45743 | Usage=37.08514%

73%| 713/977 [35:00<14:06, 3.21s/it]

[Quota] LLM=18.546638 | GPU=0.0 | Used=18.546638/50.0 | Remaining=31.453362 | Usage=37.093276%

73%| 714/977 [35:03<13:47, 3.14s/it]

[Quota] LLM=18.550997 | GPU=0.0 | Used=18.550997/50.0 | Remaining=31.449003 | Usage=37.101994%

73%| 715/977 [35:06<13:46, 3.15s/it]

[Quota] LLM=18.555098 | GPU=0.0 | Used=18.555098/50.0 | Remaining=31.444902 | Usage=37.110196%

73%| 716/977 [35:09<13:41, 3.15s/it]

[Quota] LLM=18.559094 | GPU=0.0 | Used=18.559094/50.0 | Remaining=31.440906 | Usage=37.118188%

73%| 717/977 [35:12<13:54, 3.21s/it]

[Quota] LLM=18.563771 | GPU=0.0 | Used=18.563771/50.0 | Remaining=31.436229 | Usage=37.127542%

73%| 718/977 [35:17<15:20, 3.55s/it]

[Quota] LLM=18.568478 | GPU=0.0 | Used=18.568478/50.0 | Remaining=31.431522 | Usage=37.136956%

74%| 719/977 [35:20<15:23, 3.58s/it]

[Quota] LLM=18.573269 | GPU=0.0 | Used=18.573269/50.0 | Remaining=31.426731 | Usage=37.146538%

74%| 720/977 [35:25<16:36, 3.88s/it]

[Quota] LLM=18.578093 | GPU=0.0 | Used=18.578093/50.0 | Remaining=31.421907 | Usage=37.156186%

74%| 721/977 [35:27<14:38, 3.43s/it]

[Quota] LLM=18.582095 | GPU=0.0 | Used=18.582095/50.0 | Remaining=31.417905 | Usage=37.16419%

74%| 722/977 [35:30<13:45, 3.24s/it]

[Quota] LLM=18.58628 | GPU=0.0 | Used=18.58628/50.0 | Remaining=31.41372 | Usage=37.17256%

74%| 723/977 [35:34<14:13, 3.36s/it]

[Quota] LLM=18.590327 | GPU=0.0 | Used=18.590327/50.0 | Remaining=31.409673 | Usage=37.180654%

74%| 724/977 [35:37<14:44, 3.49s/it]

[Quota] LLM=18.594563 | GPU=0.0 | Used=18.594563/50.0 | Remaining=31.405437 | Usage=37.189126%

74%| 725/977 [35:41<14:27, 3.44s/it]

[Quota] LLM=18.598679 | GPU=0.0 | Used=18.598679/50.0 | Remaining=31.401321 | Usage=37.197358%

74%| 726/977 [35:44<13:53, 3.32s/it]

[Quota] LLM=18.602972 | GPU=0.0 | Used=18.602972/50.0 | Remaining=31.397028 | Usage=37.205944%

74%| 727/977 [35:46<12:47, 3.07s/it]

[Quota] LLM=18.606938 | GPU=0.0 | Used=18.606938/50.0 | Remaining=31.393062 | Usage=37.213876%

75%| 728/977 [35:51<14:09, 3.41s/it]

[Quota] LLM=18.611549 | GPU=0.0 | Used=18.611549/50.0 | Remaining=31.388451 | Usage=37.223098%

75%| 729/977 [35:54<14:17, 3.46s/it]

[Quota] LLM=18.615848 | GPU=0.0 | Used=18.615848/50.0 | Remaining=31.384152 | Usage=37.231696%

75%| 730/977 [35:57<13:19, 3.24s/it]

[Quota] LLM=18.619799 | GPU=0.0 | Used=18.619799/50.0 | Remaining=31.380201 | Usage=37.239598%

75%| 731/977 [36:00<12:37, 3.08s/it]

[Quota] LLM=18.624008 | GPU=0.0 | Used=18.624008/50.0 | Remaining=31.375992 | Usage=37.248016%

75%| 732/977 [36:02<12:24, 3.04s/it]

[Quota] LLM=18.628208 | GPU=0.0 | Used=18.628208/50.0 | Remaining=31.371792 | Usage=37.256416%

75%| 733/977 [36:05<12:19, 3.03s/it]

[Quota] LLM=18.632486 | GPU=0.0 | Used=18.632486/50.0 | Remaining=31.367514 | Usage=37.264972%

75%| 734/977 [36:08<12:15, 3.03s/it]

[Quota] LLM=18.636701 | GPU=0.0 | Used=18.636701/50.0 | Remaining=31.363299 | Usage=37.273402%

75%| 735/977 [36:12<12:14, 3.04s/it]

[Quota] LLM=18.641111 | GPU=0.0 | Used=18.641111/50.0 | Remaining=31.358889 | Usage=37.282222%

75%| 736/977 [36:16<13:28, 3.36s/it]

[Quota] LLM=18.645308 | GPU=0.0 | Used=18.645308/50.0 | Remaining=31.354692 | Usage=37.290616%

75%| 737/977 [36:19<13:29, 3.37s/it]

[Quota] LLM=18.649355 | GPU=0.0 | Used=18.649355/50.0 | Remaining=31.350645 | Usage=37.29871%

76%| 738/977 [36:22<13:08, 3.30s/it]

[Quota] LLM=18.653285 | GPU=0.0 | Used=18.653285/50.0 | Remaining=31.346715 | Usage=37.30657%

76%| 739/977 [36:25<12:32, 3.16s/it]

[Quota] LLM=18.657185 | GPU=0.0 | Used=18.657185/50.0 | Remaining=31.342815 | Usage=37.31437%

76%| 740/977 [36:28<11:53, 3.01s/it]

[Quota] LLM=18.661283 | GPU=0.0 | Used=18.661283/50.0 | Remaining=31.338717 | Usage=37.322566%

76%| 741/977 [36:30<11:09, 2.84s/it]

[Quota] LLM=18.665267 | GPU=0.0 | Used=18.665267/50.0 | Remaining=31.334733 | Usage=37.330534%

76%| 742/977 [36:33<11:12, 2.86s/it]

[Quota] LLM=18.669419 | GPU=0.0 | Used=18.669419/50.0 | Remaining=31.330581 | Usage=37.338838%

76%| 743/977 [36:37<12:29, 3.20s/it]

[Quota] LLM=18.673976 | GPU=0.0 | Used=18.673976/50.0 | Remaining=31.326024 | Usage=37.347952%

76%| 744/977 [36:41<12:55, 3.33s/it]

[Quota] LLM=18.67877 | GPU=0.0 | Used=18.67877/50.0 | Remaining=31.32123 | Usage=37.35754%

76%| 745/977 [36:44<12:35, 3.26s/it]

[Quota] LLM=18.683396 | GPU=0.0 | Used=18.683396/50.0 | Remaining=31.316604 | Usage=37.366792%

76%| 746/977 [36:47<12:11, 3.17s/it]

[Quota] LLM=18.687884 | GPU=0.0 | Used=18.687884/50.0 | Remaining=31.312116 | Usage=37.375768%

76%| 747/977 [36:49<11:30, 3.00s/it]

[Quota] LLM=18.692129 | GPU=0.0 | Used=18.692129/50.0 | Remaining=31.307871 | Usage=37.384258%

77%| 748/977 [36:52<11:29, 3.01s/it]

[Quota] LLM=18.69602 | GPU=0.0 | Used=18.69602/50.0 | Remaining=31.30398 | Usage=37.39204%

77%| 749/977 [36:56<12:26, 3.27s/it]

[Quota] LLM=18.700694 | GPU=0.0 | Used=18.700694/50.0 | Remaining=31.299306 | Usage=37.401388%

77%| 750/977 [37:00<12:40, 3.35s/it]

[Quota] LLM=18.704726 | GPU=0.0 | Used=18.704726/50.0 | Remaining=31.295274 | Usage=37.409452%

77%| 751/977 [37:02<11:46, 3.12s/it]

[Quota] LLM=18.708725 | GPU=0.0 | Used=18.708725/50.0 | Remaining=31.291275 | Usage=37.41745%

77%| 752/977 [37:06<11:53, 3.17s/it]

[Quota] LLM=18.712886 | GPU=0.0 | Used=18.712886/50.0 | Remaining=31.287114 | Usage=37.425772%

77%| 753/977 [37:08<11:22, 3.05s/it]

[Quota] LLM=18.716759 | GPU=0.0 | Used=18.716759/50.0 | Remaining=31.283241 | Usage=37.433518%

77%| 754/977 [37:11<10:33, 2.84s/it]

[Quota] LLM=18.720581 | GPU=0.0 | Used=18.720581/50.0 | Remaining=31.279419 | Usage=37.441162%

77%| 755/977 [37:14<10:34, 2.86s/it]

[Quota] LLM=18.724325 | GPU=0.0 | Used=18.724325/50.0 | Remaining=31.275675 | Usage=37.44865%

77%| 756/977 [37:16<09:59, 2.71s/it]

[Quota] LLM=18.728297 | GPU=0.0 | Used=18.728297/50.0 | Remaining=31.271703 | Usage=37.456594%

77%| 757/977 [37:18<09:37, 2.63s/it]

[Quota] LLM=18.732134 | GPU=0.0 | Used=18.732134/50.0 | Remaining=31.267866 | Usage=37.464268%

78%| 758/977 [37:21<09:57, 2.73s/it]

[Quota] LLM=18.736622 | GPU=0.0 | Used=18.736622/50.0 | Remaining=31.263378 | Usage=37.473244%

78%| 759/977 [37:24<10:11, 2.80s/it]

[Quota] LLM=18.740651 | GPU=0.0 | Used=18.740651/50.0 | Remaining=31.259349 | Usage=37.481302%

78%| 760/977 [37:27<09:58, 2.76s/it]

[Quota] LLM=18.744683 | GPU=0.0 | Used=18.744683/50.0 | Remaining=31.255317 | Usage=37.489366%

78%| 761/977 [37:31<11:01, 3.06s/it]

[Quota] LLM=18.74966 | GPU=0.0 | Used=18.74966/50.0 | Remaining=31.25034 | Usage=37.49932%

78%| 762/977 [37:36<13:23, 3.74s/it]

[Quota] LLM=18.753737 | GPU=0.0 | Used=18.753737/50.0 | Remaining=31.246263 | Usage=37.507474%

78%| 763/977 [37:39<12:29, 3.50s/it]

[Quota] LLM=18.758408 | GPU=0.0 | Used=18.758408/50.0 | Remaining=31.241592 | Usage=37.516816%

78%| 764/977 [37:42<11:34, 3.26s/it]

[Quota] LLM=18.763142 | GPU=0.0 | Used=18.763142/50.0 | Remaining=31.236858 | Usage=37.526284%

78%| 765/977 [37:45<11:42, 3.31s/it]

[Quota] LLM=18.76823 | GPU=0.0 | Used=18.76823/50.0 | Remaining=31.23177 | Usage=37.53646%

78%| 766/977 [37:48<10:51, 3.09s/it]

[Quota] LLM=18.77246 | GPU=0.0 | Used=18.77246/50.0 | Remaining=31.22754 | Usage=37.54492%

79%| 767/977 [37:51<10:24, 2.98s/it]

[Quota] LLM=18.776648 | GPU=0.0 | Used=18.776648/50.0 | Remaining=31.223352 | Usage=37.553296%

79%| 768/977 [37:53<09:39, 2.77s/it]

[Quota] LLM=18.780692 | GPU=0.0 | Used=18.780692/50.0 | Remaining=31.219308 | Usage=37.561384%

79%| 769/977 [37:55<09:29, 2.74s/it]

[Quota] LLM=18.784547 | GPU=0.0 | Used=18.784547/50.0 | Remaining=31.215453 | Usage=37.569094%

79%| 770/977 [37:58<09:15, 2.69s/it]

[Quota] LLM=18.788708 | GPU=0.0 | Used=18.788708/50.0 | Remaining=31.211292 | Usage=37.577416%

79%| 771/977 [38:01<09:10, 2.67s/it]

[Quota] LLM=18.793244 | GPU=0.0 | Used=18.793244/50.0 | Remaining=31.206756 | Usage=37.586488%

79%| 772/977 [38:05<10:41, 3.13s/it]

[Quota] LLM=18.798062 | GPU=0.0 | Used=18.798062/50.0 | Remaining=31.201938 | Usage=37.596124%

79%| 773/977 [38:07<10:03, 2.96s/it]

[Quota] LLM=18.802151 | GPU=0.0 | Used=18.802151/50.0 | Remaining=31.197849 | Usage=37.604302%

79%| 774/977 [38:10<09:49, 2.91s/it]

[Quota] LLM=18.806357 | GPU=0.0 | Used=18.806357/50.0 | Remaining=31.193643 | Usage=37.612714%

79%| 775/977 [38:12<09:05, 2.70s/it]

[Quota] LLM=18.810389 | GPU=0.0 | Used=18.810389/50.0 | Remaining=31.189611 | Usage=37.620778%

79%| 776/977 [38:15<08:50, 2.64s/it]

[Quota] LLM=18.814277 | GPU=0.0 | Used=18.814277/50.0 | Remaining=31.185723 | Usage=37.628554%

80%| 777/977 [38:18<09:01, 2.71s/it]

[Quota] LLM=18.818561 | GPU=0.0 | Used=18.818561/50.0 | Remaining=31.181439 | Usage=37.637122%

80%| 778/977 [38:21<09:09, 2.76s/it]

[Quota] LLM=18.822692 | GPU=0.0 | Used=18.822692/50.0 | Remaining=31.177308 | Usage=37.645384%

80%| 779/977 [38:23<08:41, 2.63s/it]

[Quota] LLM=18.826742 | GPU=0.0 | Used=18.826742/50.0 | Remaining=31.173258 | Usage=37.653484%

80%| 780/977 [38:26<08:33, 2.61s/it]

[Quota] LLM=18.830942 | GPU=0.0 | Used=18.830942/50.0 | Remaining=31.169058 | Usage=37.661884%

80%| 781/977 [38:29<09:34, 2.93s/it]

[Quota] LLM=18.834923 | GPU=0.0 | Used=18.834923/50.0 | Remaining=31.165077 | Usage=37.669846%

80%| 782/977 [38:32<09:14, 2.84s/it]

[Quota] LLM=18.839123 | GPU=0.0 | Used=18.839123/50.0 | Remaining=31.160877 | Usage=37.678246%

80%| 783/977 [38:35<09:32, 2.95s/it]

[Quota] LLM=18.843206 | GPU=0.0 | Used=18.843206/50.0 | Remaining=31.156794 | Usage=37.686412%

80%| 784/977 [38:38<09:33, 2.97s/it]

[Quota] LLM=18.847214 | GPU=0.0 | Used=18.847214/50.0 | Remaining=31.152786 | Usage=37.694428%

80%| 785/977 [38:41<09:43, 3.04s/it]

[Quota] LLM=18.851651 | GPU=0.0 | Used=18.851651/50.0 | Remaining=31.148349 | Usage=37.703302%

80%| 786/977 [38:44<09:40, 3.04s/it]

[Quota] LLM=18.855923 | GPU=0.0 | Used=18.855923/50.0 | Remaining=31.144077 | Usage=37.711846%

81%| 787/977 [38:47<09:02, 2.85s/it]

[Quota] LLM=18.860288 | GPU=0.0 | Used=18.860288/50.0 | Remaining=31.139712 | Usage=37.720576%

81%| 788/977 [38:54<13:19, 4.23s/it]

[Quota] LLM=18.864437 | GPU=0.0 | Used=18.864437/50.0 | Remaining=31.135563 | Usage=37.728874%

81%| 789/977 [38:57<11:54, 3.80s/it]

[Quota] LLM=18.868874 | GPU=0.0 | Used=18.868874/50.0 | Remaining=31.131126 | Usage=37.737748%

81%| 790/977 [39:00<11:00, 3.53s/it]

[Quota] LLM=18.873092 | GPU=0.0 | Used=18.873092/50.0 | Remaining=31.126908 | Usage=37.746184%

81%| 791/977 [39:04<11:02, 3.56s/it]

[Quota] LLM=18.877244 | GPU=0.0 | Used=18.877244/50.0 | Remaining=31.122756 | Usage=37.754488%

81%| 792/977 [39:06<09:53, 3.21s/it]

[Quota] LLM=18.881642 | GPU=0.0 | Used=18.881642/50.0 | Remaining=31.118358 | Usage=37.763284%

81%| 793/977 [39:08<09:00, 2.94s/it]

[Quota] LLM=18.885755 | GPU=0.0 | Used=18.885755/50.0 | Remaining=31.114245 | Usage=37.77151%

81%| 794/977 [39:11<08:34, 2.81s/it]

[Quota] LLM=18.88985 | GPU=0.0 | Used=18.88985/50.0 | Remaining=31.11015 | Usage=37.7797%

81%| 795/977 [39:14<08:57, 2.95s/it]

[Quota] LLM=18.894113 | GPU=0.0 | Used=18.894113/50.0 | Remaining=31.105887 | Usage=37.788226%

81%| 796/977 [39:16<08:28, 2.81s/it]

[Quota] LLM=18.898247 | GPU=0.0 | Used=18.898247/50.0 | Remaining=31.101753 | Usage=37.796494%

82%| 797/977 [39:19<08:36, 2.87s/it]

[Quota] LLM=18.902867 | GPU=0.0 | Used=18.902867/50.0 | Remaining=31.097133 | Usage=37.805734%

82%| 798/977 [39:23<08:59, 3.01s/it]

[Quota] LLM=18.90743 | GPU=0.0 | Used=18.90743/50.0 | Remaining=31.09257 | Usage=37.81486%

82%| 799/977 [39:26<08:36, 2.90s/it]

[Quota] LLM=18.911363 | GPU=0.0 | Used=18.911363/50.0 | Remaining=31.088637 | Usage=37.822726%

82%| 800/977 [39:27<07:45, 2.63s/it]

```
[Quota] LLM=18.915311 | GPU=0.0 | Used=18.915311/50.0 | Remaining=31.084689 | Usage=37.830622%
82%|     | 801/977 [39:30<07:53, 2.69s/it]

[Quota] LLM=18.919556 | GPU=0.0 | Used=18.919556/50.0 | Remaining=31.080444 | Usage=37.839112%
82%|     | 802/977 [39:33<07:48, 2.68s/it]

[Quota] LLM=18.923606 | GPU=0.0 | Used=18.923606/50.0 | Remaining=31.076394 | Usage=37.847212%
82%|     | 803/977 [39:36<07:45, 2.68s/it]

[Quota] LLM=18.927677 | GPU=0.0 | Used=18.927677/50.0 | Remaining=31.072323 | Usage=37.855354%
82%|     | 804/977 [39:38<07:28, 2.60s/it]

[Quota] LLM=18.931547 | GPU=0.0 | Used=18.931547/50.0 | Remaining=31.068453 | Usage=37.863094%
82%|     | 805/977 [39:40<07:04, 2.47s/it]

[Quota] LLM=18.935435000000002 | GPU=0.0 | Used=18.935435000000002/50.0 | Remaining=31.064564999999998 | Usage=37.870870000000004%
82%|     | 806/977 [39:43<06:59, 2.45s/it]

[Quota] LLM=18.939482 | GPU=0.0 | Used=18.939482/50.0 | Remaining=31.060518 | Usage=37.878964%
83%|     | 807/977 [39:48<09:04, 3.20s/it]

[Quota] LLM=18.943796 | GPU=0.0 | Used=18.943796/50.0 | Remaining=31.056204 | Usage=37.887592%
83%|     | 808/977 [39:50<08:36, 3.06s/it]

[Quota] LLM=18.948197 | GPU=0.0 | Used=18.948197/50.0 | Remaining=31.051803 | Usage=37.896394%
83%|     | 809/977 [39:53<08:15, 2.95s/it]

[Quota] LLM=18.952727 | GPU=0.0 | Used=18.952727/50.0 | Remaining=31.047273 | Usage=37.905454%
83%|     | 810/977 [39:56<08:20, 3.00s/it]

[Quota] LLM=18.957191 | GPU=0.0 | Used=18.957191/50.0 | Remaining=31.042809 | Usage=37.914382%
83%|     | 811/977 [39:59<08:06, 2.93s/it]

[Quota] LLM=18.961436 | GPU=0.0 | Used=18.961436/50.0 | Remaining=31.038564 | Usage=37.922872%
83%|     | 812/977 [40:02<08:18, 3.02s/it]
```

[Quota] LLM=18.966176 | GPU=0.0 | Used=18.966176/50.0 | Remaining=31.033824 | Usage=37.932352%

83%| 813/977 [40:05<07:55, 2.90s/it]

[Quota] LLM=18.970028 | GPU=0.0 | Used=18.970028/50.0 | Remaining=31.029972 | Usage=37.940056%

83%| 814/977 [40:07<07:41, 2.83s/it]

[Quota] LLM=18.974174 | GPU=0.0 | Used=18.974174/50.0 | Remaining=31.025826 | Usage=37.948348%

83%| 815/977 [40:10<07:27, 2.76s/it]

[Quota] LLM=18.978536 | GPU=0.0 | Used=18.978536/50.0 | Remaining=31.021464 | Usage=37.957072%

84%| 816/977 [40:13<07:19, 2.73s/it]

[Quota] LLM=18.982982 | GPU=0.0 | Used=18.982982/50.0 | Remaining=31.017018 | Usage=37.965964%

84%| 817/977 [40:15<07:16, 2.73s/it]

[Quota] LLM=18.987359 | GPU=0.0 | Used=18.987359/50.0 | Remaining=31.012641 | Usage=37.974718%

84%| 818/977 [40:18<07:21, 2.78s/it]

[Quota] LLM=18.991397 | GPU=0.0 | Used=18.991397/50.0 | Remaining=31.008603 | Usage=37.982794%

84%| 819/977 [40:21<07:16, 2.76s/it]

[Quota] LLM=18.995837 | GPU=0.0 | Used=18.995837/50.0 | Remaining=31.004163 | Usage=37.991674%

84%| 820/977 [40:24<07:12, 2.76s/it]

[Quota] LLM=18.999734 | GPU=0.0 | Used=18.999734/50.0 | Remaining=31.000266 | Usage=37.999468%

84%| 821/977 [40:27<07:19, 2.82s/it]

[Quota] LLM=19.003772 | GPU=0.0 | Used=19.003772/50.0 | Remaining=30.996228 | Usage=38.007544%

84%| 822/977 [40:29<07:03, 2.73s/it]

[Quota] LLM=19.007576 | GPU=0.0 | Used=19.007576/50.0 | Remaining=30.992424 | Usage=38.015152%

84%| 823/977 [40:32<07:02, 2.74s/it]

[Quota] LLM=19.011593 | GPU=0.0 | Used=19.011593/50.0 | Remaining=30.988407 | Usage=38.023186%

84%| 824/977 [40:35<07:22, 2.89s/it]

[Quota] LLM=19.015991 | GPU=0.0 | Used=19.015991/50.0 | Remaining=30.984009 | Usage=38.031982%
84%| 825/977 [40:38<07:09, 2.83s/it]

[Quota] LLM=19.019993 | GPU=0.0 | Used=19.019993/50.0 | Remaining=30.980007 | Usage=38.039986%
85%| 826/977 [40:41<07:05, 2.81s/it]

[Quota] LLM=19.02389 | GPU=0.0 | Used=19.02389/50.0 | Remaining=30.97611 | Usage=38.04778%
85%| 827/977 [40:43<06:32, 2.61s/it]

[Quota] LLM=19.027844 | GPU=0.0 | Used=19.027844/50.0 | Remaining=30.972156 | Usage=38.055688%
85%| 828/977 [40:46<07:07, 2.87s/it]

[Quota] LLM=19.031696 | GPU=0.0 | Used=19.031696/50.0 | Remaining=30.968304 | Usage=38.063392%
85%| 829/977 [40:50<07:27, 3.03s/it]

[Quota] LLM=19.035593 | GPU=0.0 | Used=19.035593/50.0 | Remaining=30.964407 | Usage=38.071186%
85%| 830/977 [40:54<08:27, 3.45s/it]

[Quota] LLM=19.040267 | GPU=0.0 | Used=19.040267/50.0 | Remaining=30.959733 | Usage=38.080534%
85%| 831/977 [40:57<07:57, 3.27s/it]

[Quota] LLM=19.044497 | GPU=0.0 | Used=19.044497/50.0 | Remaining=30.955503 | Usage=38.088994%
85%| 832/977 [41:00<07:41, 3.18s/it]

[Quota] LLM=19.048619 | GPU=0.0 | Used=19.048619/50.0 | Remaining=30.951381 | Usage=38.097238%
85%| 833/977 [41:03<07:24, 3.08s/it]

[Quota] LLM=19.052825 | GPU=0.0 | Used=19.052825/50.0 | Remaining=30.947175 | Usage=38.10565%
85%| 834/977 [41:06<07:28, 3.14s/it]

[Quota] LLM=19.057082 | GPU=0.0 | Used=19.057082/50.0 | Remaining=30.942918 | Usage=38.114164%
85%| 835/977 [41:09<07:29, 3.17s/it]

[Quota] LLM=19.061429 | GPU=0.0 | Used=19.061429/50.0 | Remaining=30.938571 | Usage=38.122858%
86%| 836/977 [41:13<07:27, 3.17s/it]

[Quota] LLM=19.065557 | GPU=0.0 | Used=19.065557/50.0 | Remaining=30.934443 | Usage=38.131114%

86%| 837/977 [41:16<07:30, 3.22s/it]

[Quota] LLM=19.069616 | GPU=0.0 | Used=19.069616/50.0 | Remaining=30.930384 | Usage=38.139232%

86%| 838/977 [41:18<07:01, 3.03s/it]

[Quota] LLM=19.073786 | GPU=0.0 | Used=19.073786/50.0 | Remaining=30.926214 | Usage=38.147572%

86%| 839/977 [41:21<06:58, 3.03s/it]

[Quota] LLM=19.078226 | GPU=0.0 | Used=19.078226/50.0 | Remaining=30.921774 | Usage=38.156452%

86%| 840/977 [41:24<06:19, 2.77s/it]

[Quota] LLM=19.082156 | GPU=0.0 | Used=19.082156/50.0 | Remaining=30.917844 | Usage=38.164312%

86%| 841/977 [41:26<06:04, 2.68s/it]

[Quota] LLM=19.085996 | GPU=0.0 | Used=19.085996/50.0 | Remaining=30.914004 | Usage=38.171992%

86%| 842/977 [41:29<06:29, 2.89s/it]

[Quota] LLM=19.090256 | GPU=0.0 | Used=19.090256/50.0 | Remaining=30.909744 | Usage=38.180512%

86%| 843/977 [41:33<06:37, 2.96s/it]

[Quota] LLM=19.094321 | GPU=0.0 | Used=19.094321/50.0 | Remaining=30.905679 | Usage=38.188642%

86%| 844/977 [41:37<07:10, 3.24s/it]

[Quota] LLM=19.098398 | GPU=0.0 | Used=19.098398/50.0 | Remaining=30.901602 | Usage=38.196796%

86%| 845/977 [41:39<06:43, 3.06s/it]

[Quota] LLM=19.102346 | GPU=0.0 | Used=19.102346/50.0 | Remaining=30.897654 | Usage=38.204692%

87%| 846/977 [41:43<07:16, 3.33s/it]

[Quota] LLM=19.106468 | GPU=0.0 | Used=19.106468/50.0 | Remaining=30.893532 | Usage=38.212936%

87%| 847/977 [41:46<07:10, 3.31s/it]

[Quota] LLM=19.110476 | GPU=0.0 | Used=19.110476/50.0 | Remaining=30.889524 | Usage=38.220952%

87%| 848/977 [41:49<06:29, 3.02s/it]

[Quota] LLM=19.114364 | GPU=0.0 | Used=19.114364/50.0 | Remaining=30.885636 | Usage=38.228728%
87%| 849/977 [41:52<06:44, 3.16s/it]

[Quota] LLM=19.118513 | GPU=0.0 | Used=19.118513/50.0 | Remaining=30.881487 | Usage=38.237026%
87%| 850/977 [41:55<06:09, 2.91s/it]

[Quota] LLM=19.122434 | GPU=0.0 | Used=19.122434/50.0 | Remaining=30.877566 | Usage=38.244868%
87%| 851/977 [41:57<06:07, 2.92s/it]

[Quota] LLM=19.126499 | GPU=0.0 | Used=19.126499/50.0 | Remaining=30.873501 | Usage=38.252998%
87%| 852/977 [42:01<06:23, 3.07s/it]

[Quota] LLM=19.130471 | GPU=0.0 | Used=19.130471/50.0 | Remaining=30.869529 | Usage=38.260942%
87%| 853/977 [42:03<05:43, 2.77s/it]

[Quota] LLM=19.134422 | GPU=0.0 | Used=19.134422/50.0 | Remaining=30.865578 | Usage=38.268844%
87%| 854/977 [42:05<05:20, 2.61s/it]

[Quota] LLM=19.138325000000002 | GPU=0.0 | Used=19.138325000000002/50.0 | Remaining=30.861674999999998 | Usage=38.276650000000004%
88%| 855/977 [42:08<05:35, 2.75s/it]

[Quota] LLM=19.142264 | GPU=0.0 | Used=19.142264/50.0 | Remaining=30.857736 | Usage=38.284528%
88%| 856/977 [42:11<05:29, 2.73s/it]

[Quota] LLM=19.146038 | GPU=0.0 | Used=19.146038/50.0 | Remaining=30.853962 | Usage=38.292076%
88%| 857/977 [42:13<05:09, 2.58s/it]

[Quota] LLM=19.149839 | GPU=0.0 | Used=19.149839/50.0 | Remaining=30.850161 | Usage=38.299678%
88%| 858/977 [42:16<05:27, 2.76s/it]

[Quota] LLM=19.153676 | GPU=0.0 | Used=19.153676/50.0 | Remaining=30.846324 | Usage=38.307352%
88%| 859/977 [42:20<05:58, 3.04s/it]

[Quota] LLM=19.158317 | GPU=0.0 | Used=19.158317/50.0 | Remaining=30.841683 | Usage=38.316634%
88%| 860/977 [42:24<06:16, 3.21s/it]

[Quota] LLM=19.162268 | GPU=0.0 | Used=19.162268/50.0 | Remaining=30.837732 | Usage=38.324536%

88%| 861/977 [42:27<06:03, 3.13s/it]

[Quota] LLM=19.166150000000002 | GPU=0.0 | Used=19.166150000000002/50.0 | Remaining=30.83384999999998 | Usage=38.332300000000004%

88%| 862/977 [42:30<05:55, 3.09s/it]

[Quota] LLM=19.170497 | GPU=0.0 | Used=19.170497/50.0 | Remaining=30.829503 | Usage=38.340994%

88%| 863/977 [42:33<05:52, 3.10s/it]

[Quota] LLM=19.174403 | GPU=0.0 | Used=19.174403/50.0 | Remaining=30.825597 | Usage=38.348806%

88%| 864/977 [42:36<05:59, 3.18s/it]

[Quota] LLM=19.178435 | GPU=0.0 | Used=19.178435/50.0 | Remaining=30.821565 | Usage=38.35687%

89%| 865/977 [42:40<06:13, 3.34s/it]

[Quota] LLM=19.182698 | GPU=0.0 | Used=19.182698/50.0 | Remaining=30.817302 | Usage=38.365396%

89%| 866/977 [42:43<05:50, 3.16s/it]

[Quota] LLM=19.186541 | GPU=0.0 | Used=19.186541/50.0 | Remaining=30.813459 | Usage=38.373082%

89%| 867/977 [42:46<05:53, 3.21s/it]

[Quota] LLM=19.19042 | GPU=0.0 | Used=19.19042/50.0 | Remaining=30.80958 | Usage=38.38084%

89%| 868/977 [42:49<05:51, 3.22s/it]

[Quota] LLM=19.194356 | GPU=0.0 | Used=19.194356/50.0 | Remaining=30.805644 | Usage=38.388712%

89%| 869/977 [42:52<05:30, 3.06s/it]

[Quota] LLM=19.19819 | GPU=0.0 | Used=19.19819/50.0 | Remaining=30.80181 | Usage=38.39638%

89%| 870/977 [42:55<05:25, 3.04s/it]

[Quota] LLM=19.202291 | GPU=0.0 | Used=19.202291/50.0 | Remaining=30.797709 | Usage=38.404582%

89%| 871/977 [42:58<05:20, 3.03s/it]

[Quota] LLM=19.206302 | GPU=0.0 | Used=19.206302/50.0 | Remaining=30.793698 | Usage=38.412604%

89%| 872/977 [43:00<04:58, 2.85s/it]

[Quota] LLM=19.210157 | GPU=0.0 | Used=19.210157/50.0 | Remaining=30.789843 | Usage=38.420314%

89%| 873/977 [43:03<05:02, 2.90s/it]

[Quota] LLM=19.214099 | GPU=0.0 | Used=19.214099/50.0 | Remaining=30.785901 | Usage=38.428198%

89%| 874/977 [43:07<05:36, 3.26s/it]

[Quota] LLM=19.218083 | GPU=0.0 | Used=19.218083/50.0 | Remaining=30.781917 | Usage=38.436166%

90%| 875/977 [43:10<05:27, 3.21s/it]

[Quota] LLM=19.221965 | GPU=0.0 | Used=19.221965/50.0 | Remaining=30.778035 | Usage=38.44393%

90%| 876/977 [43:13<04:53, 2.91s/it]

[Quota] LLM=19.225784 | GPU=0.0 | Used=19.225784/50.0 | Remaining=30.774216 | Usage=38.451568%

90%| 877/977 [43:16<04:50, 2.91s/it]

[Quota] LLM=19.229726 | GPU=0.0 | Used=19.229726/50.0 | Remaining=30.770274 | Usage=38.459452%

90%| 878/977 [43:19<05:16, 3.19s/it]

[Quota] LLM=19.233575 | GPU=0.0 | Used=19.233575/50.0 | Remaining=30.766425 | Usage=38.46715%

90%| 879/977 [43:22<04:55, 3.02s/it]

[Quota] LLM=19.237556 | GPU=0.0 | Used=19.237556/50.0 | Remaining=30.762444 | Usage=38.475112%

90%| 880/977 [43:26<05:17, 3.28s/it]

[Quota] LLM=19.242047 | GPU=0.0 | Used=19.242047/50.0 | Remaining=30.757953 | Usage=38.484094%

90%| 881/977 [43:30<05:35, 3.50s/it]

[Quota] LLM=19.246433 | GPU=0.0 | Used=19.246433/50.0 | Remaining=30.753567 | Usage=38.492866%

90%| 882/977 [43:35<06:17, 3.97s/it]

[Quota] LLM=19.251041 | GPU=0.0 | Used=19.251041/50.0 | Remaining=30.748959 | Usage=38.502082%

90%| 883/977 [43:39<06:09, 3.93s/it]

[Quota] LLM=19.255055 | GPU=0.0 | Used=19.255055/50.0 | Remaining=30.744945 | Usage=38.51011%

90%| 884/977 [43:43<06:02, 3.89s/it]

[Quota] LLM=19.259156 | GPU=0.0 | Used=19.259156/50.0 | Remaining=30.740844 | Usage=38.518312%
91% | 885/977 [43:46<05:54, 3.86s/it]

[Quota] LLM=19.263122 | GPU=0.0 | Used=19.263122/50.0 | Remaining=30.736878 | Usage=38.526244%
91% | 886/977 [43:49<05:17, 3.49s/it]

[Quota] LLM=19.267028 | GPU=0.0 | Used=19.267028/50.0 | Remaining=30.732972 | Usage=38.534056%
91% | 887/977 [43:52<04:48, 3.21s/it]

[Quota] LLM=19.270988 | GPU=0.0 | Used=19.270988/50.0 | Remaining=30.729012 | Usage=38.541976%
91% | 888/977 [43:55<04:55, 3.32s/it]

[Quota] LLM=19.275317 | GPU=0.0 | Used=19.275317/50.0 | Remaining=30.724683 | Usage=38.550634%
91% | 889/977 [43:59<04:56, 3.37s/it]

[Quota] LLM=19.279754 | GPU=0.0 | Used=19.279754/50.0 | Remaining=30.720246 | Usage=38.559508%
91% | 890/977 [44:02<04:57, 3.42s/it]

[Quota] LLM=19.283543 | GPU=0.0 | Used=19.283543/50.0 | Remaining=30.716457 | Usage=38.567086%
91% | 891/977 [44:10<06:39, 4.64s/it]

[Quota] LLM=19.287377 | GPU=0.0 | Used=19.287377/50.0 | Remaining=30.712623 | Usage=38.574754%
91% | 892/977 [44:13<06:08, 4.34s/it]

[Quota] LLM=19.291649 | GPU=0.0 | Used=19.291649/50.0 | Remaining=30.708351 | Usage=38.583298%
91% | 893/977 [44:16<05:22, 3.83s/it]

[Quota] LLM=19.295456 | GPU=0.0 | Used=19.295456/50.0 | Remaining=30.704544 | Usage=38.590912%
92% | 894/977 [44:19<05:02, 3.64s/it]

[Quota] LLM=19.299632 | GPU=0.0 | Used=19.299632/50.0 | Remaining=30.700368 | Usage=38.599264%
92% | 895/977 [44:22<04:36, 3.38s/it]

[Quota] LLM=19.303484 | GPU=0.0 | Used=19.303484/50.0 | Remaining=30.696516 | Usage=38.606968%
92% | 896/977 [44:25<04:36, 3.41s/it]

[Quota] LLM=19.307594 | GPU=0.0 | Used=19.307594/50.0 | Remaining=30.692406 | Usage=38.615188%
92%| 897/977 [44:29<04:43, 3.55s/it]

[Quota] LLM=19.311737 | GPU=0.0 | Used=19.311737/50.0 | Remaining=30.688263 | Usage=38.623474%
92%| 898/977 [44:32<04:23, 3.33s/it]

[Quota] LLM=19.315754 | GPU=0.0 | Used=19.315754/50.0 | Remaining=30.684246 | Usage=38.631508%
92%| 899/977 [44:36<04:26, 3.42s/it]

[Quota] LLM=19.319801000000002 | GPU=0.0 | Used=19.319801000000002/50.0 | Remaining=30.68019899999998 | Usage=38.639602000000004%
92%| 900/977 [44:39<04:12, 3.28s/it]

[Quota] LLM=19.323689 | GPU=0.0 | Used=19.323689/50.0 | Remaining=30.676311 | Usage=38.647378%
92%| 901/977 [44:41<03:58, 3.13s/it]

[Quota] LLM=19.327661 | GPU=0.0 | Used=19.327661/50.0 | Remaining=30.672339 | Usage=38.655322%
92%| 902/977 [44:44<03:49, 3.07s/it]

[Quota] LLM=19.331606 | GPU=0.0 | Used=19.331606/50.0 | Remaining=30.668394 | Usage=38.663212%
92%| 903/977 [44:48<03:51, 3.13s/it]

[Quota] LLM=19.335755 | GPU=0.0 | Used=19.335755/50.0 | Remaining=30.664245 | Usage=38.67151%
93%| 904/977 [44:52<04:08, 3.40s/it]

[Quota] LLM=19.339745 | GPU=0.0 | Used=19.339745/50.0 | Remaining=30.660255 | Usage=38.67949%
93%| 905/977 [44:55<04:02, 3.37s/it]

[Quota] LLM=19.343702 | GPU=0.0 | Used=19.343702/50.0 | Remaining=30.656298 | Usage=38.687404%
93%| 906/977 [44:58<04:01, 3.40s/it]

[Quota] LLM=19.347911 | GPU=0.0 | Used=19.347911/50.0 | Remaining=30.652089 | Usage=38.695822%
93%| 907/977 [45:01<03:42, 3.18s/it]

[Quota] LLM=19.351703 | GPU=0.0 | Used=19.351703/50.0 | Remaining=30.648297 | Usage=38.703406%
93%| 908/977 [45:04<03:24, 2.97s/it]

[Quota] LLM=19.355636 | GPU=0.0 | Used=19.355636/50.0 | Remaining=30.644364 | Usage=38.711272%
93%| 909/977 [45:06<03:12, 2.83s/it]

[Quota] LLM=19.35956 | GPU=0.0 | Used=19.35956/50.0 | Remaining=30.64044 | Usage=38.71912%
93%| 910/977 [45:09<03:15, 2.92s/it]

[Quota] LLM=19.363403 | GPU=0.0 | Used=19.363403/50.0 | Remaining=30.636597 | Usage=38.726806%
93%| 911/977 [45:12<03:03, 2.79s/it]

[Quota] LLM=19.367288 | GPU=0.0 | Used=19.367288/50.0 | Remaining=30.632712 | Usage=38.734576%
93%| 912/977 [45:14<02:58, 2.74s/it]

[Quota] LLM=19.371089 | GPU=0.0 | Used=19.371089/50.0 | Remaining=30.628911 | Usage=38.742178%
93%| 913/977 [45:18<03:08, 2.94s/it]

[Quota] LLM=19.375187 | GPU=0.0 | Used=19.375187/50.0 | Remaining=30.624813 | Usage=38.750374%
94%| 914/977 [45:20<02:48, 2.68s/it]

[Quota] LLM=19.378994 | GPU=0.0 | Used=19.378994/50.0 | Remaining=30.621006 | Usage=38.757988%
94%| 915/977 [45:23<02:50, 2.75s/it]

[Quota] LLM=19.382795 | GPU=0.0 | Used=19.382795/50.0 | Remaining=30.617205 | Usage=38.76559%
94%| 916/977 [45:25<02:42, 2.67s/it]

[Quota] LLM=19.386782 | GPU=0.0 | Used=19.386782/50.0 | Remaining=30.613218 | Usage=38.773564%
94%| 917/977 [45:28<02:42, 2.70s/it]

[Quota] LLM=19.390994 | GPU=0.0 | Used=19.390994/50.0 | Remaining=30.609006 | Usage=38.781988%
94%| 918/977 [45:30<02:33, 2.61s/it]

[Quota] LLM=19.394816 | GPU=0.0 | Used=19.394816/50.0 | Remaining=30.605184 | Usage=38.789632%
94%| 919/977 [45:33<02:27, 2.55s/it]

[Quota] LLM=19.398896 | GPU=0.0 | Used=19.398896/50.0 | Remaining=30.601104 | Usage=38.797792%
94%| 920/977 [45:37<02:45, 2.91s/it]

[Quota] LLM=19.40288 | GPU=0.0 | Used=19.40288/50.0 | Remaining=30.59712 | Usage=38.80576%
94%| 921/977 [45:39<02:30, 2.68s/it]

[Quota] LLM=19.40666 | GPU=0.0 | Used=19.40666/50.0 | Remaining=30.59334 | Usage=38.81332%
94%| 922/977 [45:42<02:33, 2.79s/it]

[Quota] LLM=19.411253 | GPU=0.0 | Used=19.411253/50.0 | Remaining=30.588747 | Usage=38.822506%
94%| 923/977 [45:46<02:45, 3.07s/it]

[Quota] LLM=19.415282 | GPU=0.0 | Used=19.415282/50.0 | Remaining=30.584718 | Usage=38.830564%
95%| 924/977 [45:48<02:37, 2.97s/it]

[Quota] LLM=19.419344 | GPU=0.0 | Used=19.419344/50.0 | Remaining=30.580656 | Usage=38.838688%
95%| 925/977 [45:51<02:28, 2.85s/it]

[Quota] LLM=19.423478 | GPU=0.0 | Used=19.423478/50.0 | Remaining=30.576522 | Usage=38.846956%
95%| 926/977 [45:53<02:20, 2.76s/it]

[Quota] LLM=19.427597 | GPU=0.0 | Used=19.427597/50.0 | Remaining=30.572403 | Usage=38.855194%
95%| 927/977 [45:56<02:21, 2.82s/it]

[Quota] LLM=19.431851 | GPU=0.0 | Used=19.431851/50.0 | Remaining=30.568149 | Usage=38.863702%
95%| 928/977 [45:59<02:21, 2.88s/it]

[Quota] LLM=19.436054 | GPU=0.0 | Used=19.436054/50.0 | Remaining=30.563946 | Usage=38.872108%
95%| 929/977 [46:02<02:15, 2.83s/it]

[Quota] LLM=19.440023 | GPU=0.0 | Used=19.440023/50.0 | Remaining=30.559977 | Usage=38.880046%
95%| 930/977 [46:06<02:22, 3.03s/it]

[Quota] LLM=19.444235 | GPU=0.0 | Used=19.444235/50.0 | Remaining=30.555765 | Usage=38.88847%
95%| 931/977 [46:08<02:16, 2.97s/it]

[Quota] LLM=19.448168 | GPU=0.0 | Used=19.448168/50.0 | Remaining=30.551832 | Usage=38.896336%
95%| 932/977 [46:11<02:09, 2.89s/it]

[Quota] LLM=19.45193 | GPU=0.0 | Used=19.45193/50.0 | Remaining=30.54807 | Usage=38.90386%

95%| 933/977 [46:14<02:01, 2.76s/it]

[Quota] LLM=19.455878 | GPU=0.0 | Used=19.455878/50.0 | Remaining=30.544122 | Usage=38.911756%

96%| 934/977 [46:16<01:56, 2.71s/it]

[Quota] LLM=19.459697 | GPU=0.0 | Used=19.459697/50.0 | Remaining=30.540303 | Usage=38.919394%

96%| 935/977 [46:19<02:00, 2.87s/it]

[Quota] LLM=19.463912 | GPU=0.0 | Used=19.463912/50.0 | Remaining=30.536088 | Usage=38.927824%

96%| 936/977 [46:22<01:52, 2.73s/it]

[Quota] LLM=19.467782 | GPU=0.0 | Used=19.467782/50.0 | Remaining=30.532218 | Usage=38.935564%

96%| 937/977 [46:25<01:55, 2.89s/it]

[Quota] LLM=19.472225 | GPU=0.0 | Used=19.472225/50.0 | Remaining=30.527775 | Usage=38.94445%

96%| 938/977 [46:28<01:48, 2.79s/it]

[Quota] LLM=19.476128 | GPU=0.0 | Used=19.476128/50.0 | Remaining=30.523872 | Usage=38.952256%

96%| 939/977 [46:30<01:46, 2.80s/it]

[Quota] LLM=19.48052 | GPU=0.0 | Used=19.48052/50.0 | Remaining=30.51948 | Usage=38.96104%

96%| 940/977 [46:34<01:51, 3.00s/it]

[Quota] LLM=19.485218 | GPU=0.0 | Used=19.485218/50.0 | Remaining=30.514782 | Usage=38.970436%

96%| 941/977 [46:39<02:08, 3.56s/it]

[Quota] LLM=19.490051 | GPU=0.0 | Used=19.490051/50.0 | Remaining=30.509949 | Usage=38.980102%

96%| 942/977 [46:44<02:20, 4.00s/it]

[Quota] LLM=19.495124 | GPU=0.0 | Used=19.495124/50.0 | Remaining=30.504876 | Usage=38.990248%

97%| 943/977 [46:53<03:10, 5.59s/it]

[Quota] LLM=19.499891 | GPU=0.0 | Used=19.499891/50.0 | Remaining=30.500109 | Usage=38.999782%

97%| 944/977 [46:56<02:41, 4.89s/it]

[Quota] LLM=19.504595 | GPU=0.0 | Used=19.504595/50.0 | Remaining=30.495405 | Usage=39.00919%

97%| 945/977 [46:59<02:17, 4.29s/it]

[Quota] LLM=19.508864 | GPU=0.0 | Used=19.508864/50.0 | Remaining=30.491136 | Usage=39.017728%

97%| 946/977 [47:03<02:05, 4.04s/it]

[Quota] LLM=19.513253 | GPU=0.0 | Used=19.513253/50.0 | Remaining=30.486747 | Usage=39.026506%

97%| 947/977 [47:06<01:53, 3.78s/it]

[Quota] LLM=19.517384 | GPU=0.0 | Used=19.517384/50.0 | Remaining=30.482616 | Usage=39.034768%

97%| 948/977 [47:09<01:46, 3.69s/it]

[Quota] LLM=19.521971 | GPU=0.0 | Used=19.521971/50.0 | Remaining=30.478029 | Usage=39.043942%

97%| 949/977 [47:12<01:30, 3.25s/it]

[Quota] LLM=19.52588 | GPU=0.0 | Used=19.52588/50.0 | Remaining=30.47412 | Usage=39.05176%

97%| 950/977 [47:15<01:27, 3.24s/it]

[Quota] LLM=19.529903 | GPU=0.0 | Used=19.529903/50.0 | Remaining=30.470097 | Usage=39.059806%

97%| 951/977 [47:18<01:22, 3.15s/it]

[Quota] LLM=19.534088 | GPU=0.0 | Used=19.534088/50.0 | Remaining=30.465912 | Usage=39.068176%

97%| 952/977 [47:20<01:15, 3.02s/it]

[Quota] LLM=19.537862 | GPU=0.0 | Used=19.537862/50.0 | Remaining=30.462138 | Usage=39.075724%

98%| 953/977 [47:23<01:10, 2.95s/it]

[Quota] LLM=19.541903 | GPU=0.0 | Used=19.541903/50.0 | Remaining=30.458097 | Usage=39.083806%

98%| 954/977 [47:27<01:10, 3.08s/it]

[Quota] LLM=19.546025 | GPU=0.0 | Used=19.546025/50.0 | Remaining=30.453975 | Usage=39.09205%

98%| 955/977 [47:31<01:14, 3.40s/it]

[Quota] LLM=19.549898 | GPU=0.0 | Used=19.549898/50.0 | Remaining=30.450102 | Usage=39.099796%

98%| 956/977 [47:35<01:19, 3.79s/it]

```
[Quota] LLM=19.553939 | GPU=0.0 | Used=19.553939/50.0 | Remaining=30.446061 |
Usage=39.107878%
98%|      | 957/977 [47:39<01:16,  3.81s/it]

[Quota] LLM=19.558478 | GPU=0.0 | Used=19.558478/50.0 | Remaining=30.441522 |
Usage=39.116956%
98%|      | 958/977 [47:42<01:04,  3.40s/it]

[Quota] LLM=19.5623 | GPU=0.0 | Used=19.5623/50.0 | Remaining=30.4377 |
Usage=39.1246%
98%|      | 959/977 [47:45<00:58,  3.24s/it]

[Quota] LLM=19.566242 | GPU=0.0 | Used=19.566242/50.0 | Remaining=30.433758 |
Usage=39.132484%
98%|      | 960/977 [47:48<00:53,  3.16s/it]

[Quota] LLM=19.570229 | GPU=0.0 | Used=19.570229/50.0 | Remaining=30.429771 |
Usage=39.140458%
98%|      | 961/977 [47:51<00:50,  3.16s/it]

[Quota] LLM=19.574249000000002 | GPU=0.0 | Used=19.574249000000002/50.0 |
Remaining=30.425750999999998 | Usage=39.148498000000004%
98%|      | 962/977 [47:56<00:56,  3.77s/it]

[Quota] LLM=19.578626 | GPU=0.0 | Used=19.578626/50.0 | Remaining=30.421374 |
Usage=39.157252%
99%|      | 963/977 [48:01<00:57,  4.11s/it]

[Quota] LLM=19.583114 | GPU=0.0 | Used=19.583114/50.0 | Remaining=30.416886 |
Usage=39.166228%
99%|      | 964/977 [48:05<00:52,  4.04s/it]

[Quota] LLM=19.58729 | GPU=0.0 | Used=19.58729/50.0 | Remaining=30.41271 |
Usage=39.17458%
99%|      | 965/977 [48:09<00:48,  4.02s/it]

[Quota] LLM=19.591679 | GPU=0.0 | Used=19.591679/50.0 | Remaining=30.408321 |
Usage=39.183358%
99%|      | 966/977 [48:12<00:42,  3.86s/it]

[Quota] LLM=19.596413 | GPU=0.0 | Used=19.596413/50.0 | Remaining=30.403587 |
Usage=39.192826%
99%|      | 967/977 [48:16<00:37,  3.77s/it]

[Quota] LLM=19.601507 | GPU=0.0 | Used=19.601507/50.0 | Remaining=30.398493 |
Usage=39.203014%
99%|      | 968/977 [48:20<00:34,  3.81s/it]
```

[Quota] LLM=19.605971 | GPU=0.0 | Used=19.605971/50.0 | Remaining=30.394029 | Usage=39.211942%

99%| 969/977 [48:24<00:31, 3.89s/it]

[Quota] LLM=19.610336 | GPU=0.0 | Used=19.610336/50.0 | Remaining=30.389664 | Usage=39.220672%

99%| 970/977 [48:28<00:27, 3.88s/it]

[Quota] LLM=19.615079 | GPU=0.0 | Used=19.615079/50.0 | Remaining=30.384921 | Usage=39.230158%

99%| 971/977 [48:31<00:23, 3.86s/it]

[Quota] LLM=19.61912 | GPU=0.0 | Used=19.61912/50.0 | Remaining=30.38088 | Usage=39.23824%

99%| 972/977 [48:36<00:19, 3.94s/it]

[Quota] LLM=19.623671 | GPU=0.0 | Used=19.623671/50.0 | Remaining=30.376329 | Usage=39.247342%

100%| 973/977 [48:41<00:17, 4.31s/it]

[Quota] LLM=19.628402 | GPU=0.0 | Used=19.628402/50.0 | Remaining=30.371598 | Usage=39.256804%

100%| 974/977 [48:44<00:11, 3.92s/it]

[Quota] LLM=19.632686 | GPU=0.0 | Used=19.632686/50.0 | Remaining=30.367314 | Usage=39.265372%

100%| 975/977 [48:47<00:07, 3.70s/it]

[Quota] LLM=19.636859 | GPU=0.0 | Used=19.636859/50.0 | Remaining=30.363141 | Usage=39.273718%

100%| 976/977 [48:49<00:03, 3.33s/it]

[Quota] LLM=19.640681 | GPU=0.0 | Used=19.640681/50.0 | Remaining=30.359319 | Usage=39.281362%

100%| 977/977 [48:52<00:00, 3.00s/it]

[Quota] LLM=19.644689 | GPU=0.0 | Used=19.644689/50.0 | Remaining=30.355311 | Usage=39.289378%

3.3 References

- [1] Theo King, Zekun Wu, Adriano Koshiyama, Emre Kazim, and Philip Treleaven. 2024. HEARTS: A holistic framework for explainable, sustainable and robust text stereotype detection. arXiv preprint arXiv:2409.11579. Available at: <https://arxiv.org/abs/2409.11579> (Accessed: 4 December 2025). <https://doi.org/10.48550/arXiv.2409.11579>

- [2] Theo King, Zekun Wu, Adriano Koshiyama, Emre Kazim, and Philip Treleaven. 2024. HEARTS-Text-Stereotype-Detection (GitHub Repository). Available at: <https://github.com/holistic-ai/HEARTS-Text-Stereotype-Detection> (Accessed: 4 December 2025).
- [3] Theo King, Zekun Wu, Adriano Koshiyama, Emre Kazim, and Philip Treleaven. Holistic AI. 2024. EMGSD: Expanded Multi-Group Stereotype Dataset (HuggingFace Dataset). Available at: <https://huggingface.co/datasets/holistic-ai/EMGSD> (Accessed: 4 December 2025).
- [4] University College London Technical Support Group (TSG). 2025. GPU Access and Usage Documentation. Available at: <https://tsg.cs.ucl.ac.uk/gpus/> (Accessed: 6 December 2025).
- [5] United Nations. 2025. The 2030 Agenda for Sustainable Development. Available at: <https://sdgs.un.org/2030agenda> (Accessed: 6 December 2025).