

# COMP0173-CW2-NOTEBOOK-P2

December 11, 2025

## 1 COMP0173: Coursework 2

The paper HEARTS: A Holistic Framework for Explainable, Sustainable, and Robust Text Stereotype Detection by Theo King, Zekun Wu et al. (2024) presents a comprehensive approach to analysing and detecting stereotypes in text [1]. The authors introduce the HEARTS framework, which integrates model explainability, carbon-efficient training, and accurate evaluation across multiple bias-sensitive datasets. By using transformer-based models such as ALBERT-V2, BERT, and DistilBERT, this research project demonstrates that stereotype detection performance varies significantly across dataset sources, underlining the need for diverse evaluation benchmarks. The paper provides publicly available datasets and code [2], allowing full reproducibility and offering a standardised methodology for future research on bias and stereotype detection in Natural Language Processing (NLP).

While the HEARTS framework evaluates stereotype detection in English, this project adapts the methodology to the Russian context. Russian stereotypes often rely on grammatical gender, morphology, and culture specific tropes. Although Russian is not classified as a low-resource language and many high-performing NLP models are available, there is currently no publicly accessible model specifically designed to detect stereotypes in Russian language. Existing models detecting toxicity or sentiment identify stereotypical and biased sentences only when they include specific patterns, such as insults, slurs, or identity-specific hate speech [8].

To address this gap, I introduce two fine-tuned classifiers, **AI-Forever-RuBert** [10] and **XML-RoBERTa** [11] trained on datasets RBSA, and RBS, respectively. Understanding these patterns is essential for applications such as content moderation, ensuring the safety of Russian-language LLMs, and monitoring harmful narratives across demographic groups and underrepresented societies. Adapting the HEARTS framework to this new sociolinguistic context illustrates its transferability beyond the English-speaking context and enables a more culturally grounded approach to bias detection, thereby promoting SDG 5: Gender Equality, SDG 10: Reduced Inequalities, and SDG 16: Peace, Justice, and Strong Institutions [5].

## 2 Instructions

All figures produced during this notebook are stored in the project's **COMP0173\_Figures** directory. The corresponding LaTeX-formatted performance comparison tables, jupyter notebooks are stored in **/COMP0173\_PDF**. The compiled document are available as **COMP0173-CW2-TABLES.pdf** and **COMP0173\_PDF/COMP0173-CW2-NOTEBOOK-XX.pdf**. All prompts used for data augmentation are stored in **COMP0173\_Prompts** and the manually collected stereotypes (with English translations) are provided in **COMP0173\_Stereotypes**. The datasets used for model training and evaluation are stored in **COMP0173\_Data** which contains:

- rubias.tsv — RuBias dataset [6, 7]
- ruster.csv — RuSter dataset (see Part 2 of the notebook for source websites)
- rubist.csv — RBS dataset: RuBias + RuSter augmented with LLM-generated samples (Claude Sonnet), using a zero-shot prompt with examples
- rubist\_second.csv — RBSA dataset: RuBias + RuSter augmented with LLM-generated samples using a second prompt version without examples

The notebooks `COMP0173_PDF/COMP0173-CW2-NOTEBOOK-P3.pdf` and `COMP0173_PDF/COMP0173-CW2-NOTEBOOK-P5.pdf` are replications of `COMP0173_PDF/COMP0173-CW2-NOTEBOOK-P2.pdf` and `COMP0173_PDF/COMP0173-CW2-NOTEBOOK-P4.pdf`, where P2 provides the new RBSA with second prompt (without examples) and P5 demonstrates the model running ON GPU (the results saved are from GPU fine-tuning).

### 3 Technical Implementation (70%)

```
[1]: # %%capture
# pip install -r requirements.txt
# pip install transformers
# pip install --upgrade transformers
# pip install --upgrade tokenizers
# pip install -U sentence-transformers
# pip install natasha
# pip install datasets
# pip install --user -U nltk
# conda install -c anaconda nltk
# pip install --upgrade openai pandas tqdm
# pip install dotenv
```

```
[2]: # pip install -U pip setuptools wheel
# pip install -U spacy
# python -m spacy download en_core_web_trf
# python -m spacy download en_core_web_sm
# python -m spacy download ru_core_news_lg

# # GPU
# pip install -U 'spacy[cuda12x]'
# # GPU - Train Models
# pip install -U 'spacy[cuda12x,transformers,lookups]'
```

```
[3]: # Import the libraries
import random, numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
sns.set(color_codes=True)
plt.style.use('seaborn-v0_8')
```

```
# To ignore warnings
import warnings
warnings.filterwarnings('ignore')
np.random.seed(23)

warnings.filterwarnings(
    "ignore",
    message="pkg_resources is deprecated as an API"
)
```

```
[4]: # Import libraries
import pandas as pd
import os
import sys
import importlib.util, pathlib
from pathlib import Path
import warnings
from importlib import reload
from importlib.machinery import SourceFileLoader
from IPython.display import display
import pandas as pd
from pathlib import Path
import re
import difflib
import string
from collections import defaultdict
import json
```

```
[5]: import torch
import transformers
from transformers import AutoModelForMaskedLM, XLMWithLMHeadModel
from transformers import AutoTokenizer, AutoConfig
from transformers import TrainingArguments, Trainer
from sentence_transformers import SentenceTransformer, util
import platform
from datasets import load_dataset
import spacy
import requests
from tqdm import tqdm
import yaml

import nltk
from nltk import tokenize
from nltk.tokenize import word_tokenize
nltk.download('punkt')

import natasha
```

```

from natasha import (
    Segmenter,
    MorphVocab,

    NewsEmbedding,
    NewsMorphTagger,
    NewsSyntaxParser,
    NewsNERTagger,

    PER,
    NamesExtractor,
    Doc
)

segmenter = Segmenter()
morph_vocab = MorphVocab()

segmenter = Segmenter()
morph_vocab = MorphVocab()

emb = NewsEmbedding()
morph_tagger = NewsMorphTagger(emb)

```

```

[nltk_data] Downloading package punkt to
[nltk_data]      /Users/rinlobachevskii/nltk_data...
[nltk_data]      Package punkt is already up-to-date!

```

```

[6]: # # Check the GPU host (UCL access)
# print("CUDA available:", torch.cuda.is_available())
# print("Device:", torch.cuda.get_device_name(0))

# # Path
# import os
# os.chdir("/tmp/HEARTS-Text-Stereotype-Detection")
# os.getcwd()

```

### 3.1 Part 2: Identify a contextually relevant challenge in your country or region of your choice that can be addressed using the same AI approach

**Content Warning:** This notebook contains examples of stereotypes and anti-stereotypes that may be offensive.

#### 3.1.1 *Question 1: Problem and SDG alignment*

This coursework supports Sustainable Development Goal (SDG) 5: Gender Equality - *Achieve gender equality and empower all women and girls*, SDG 9: Industry, Innovation, and Infrastructure - *Build resilient infrastructure, promote inclusive and sustainable industrialization and foster innovation*, SDG 10: Reduced Inequalities - *Reduce inequality within and among countries*, and SDG

16: Peace, Justice, and Strong Institutions: - *Promote peaceful and inclusive societies for sustainable development, provide access to justice for all and build effective, accountable and inclusive institutions at all levels [5].*

The specific targets covered by this coursework are:

- SDG 5.1: *End all forms of discrimination against all women and girls everywhere*
- SDG 5.b: *Enhance the use of enabling technology, in particular information and communications technology, to promote the empowerment of women*
- SDG 10.2: *By 2030, empower and promote the social, economic and political inclusion of all, irrespective of age, sex, disability, race, ethnicity, origin, religion or economic or other status*
- SDG 10.3: *Ensure equal opportunity and reduce inequalities of outcome, including by eliminating discriminatory laws, policies and practices and promoting appropriate legislation, policies and action in this regard*
- SDG 16.1: *Significantly reduce all forms of violence and related death rates everywhere*
- SDG 16.6: *Develop effective, accountable and transparent institutions at all levels*
- SDG 16.10: *Ensure public access to information and protect fundamental freedoms, in accordance with national legislation and international agreements*
- SDG 16.b: *Promote and enforce non-discriminatory laws and policies for sustainable development*

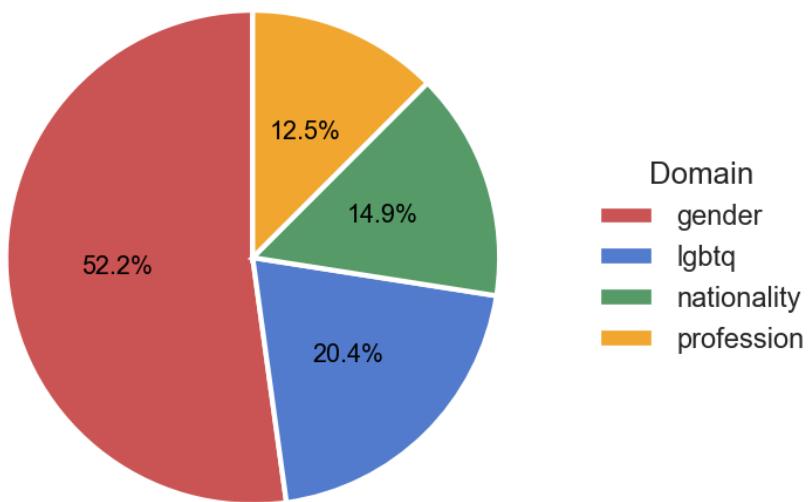
### **3.1.2 Question 2 : Limitations and ethical considerations**

Unfortunately, there is no extensive dataset on stereotypes in Russian comparable to the MGSD. If such datasets do exist, many primarily contain hate speech and inappropriate statements or accounts id from social medias that are unsuitable for this project and for stereotype/bias detection in general.

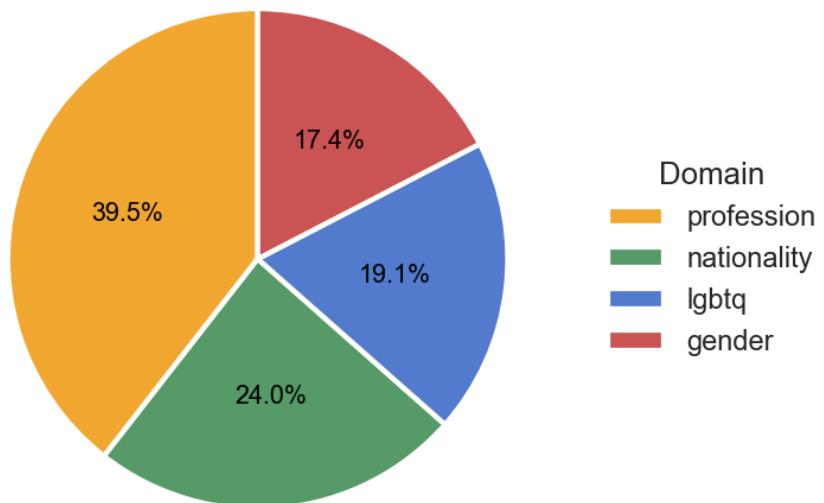
The RuBias [6] dataset is one of the few publicly available datasets addressing biases in Russian; however, it was initially designed for detecting biases in language models, primarily by swapping keywords, rather than for stereotype detection. The RuBias dataset also lacks domain diversity, covering only a few social groups for each stereotype type. In contrast, RuSter aims to be a more realistic and trustworthy source, free of hate speech or profanity.

To curate the RuSter dataset, I manually collected common cultural stereotypes in Russian by scraping the internet anonymously. I did not create any original stereotypes, which is important to note. Additionally, the LGBTQ+ community remains underrepresented in Russian texts, mainly because, in a few Russian-speaking countries, LGBTQ+ is illegal, resulting in a lack of resources or online materials on this topic. The resulting RBS and RBSA datasets (after data augmentation) are more balanced, as illustrated in the accompanying pie charts.

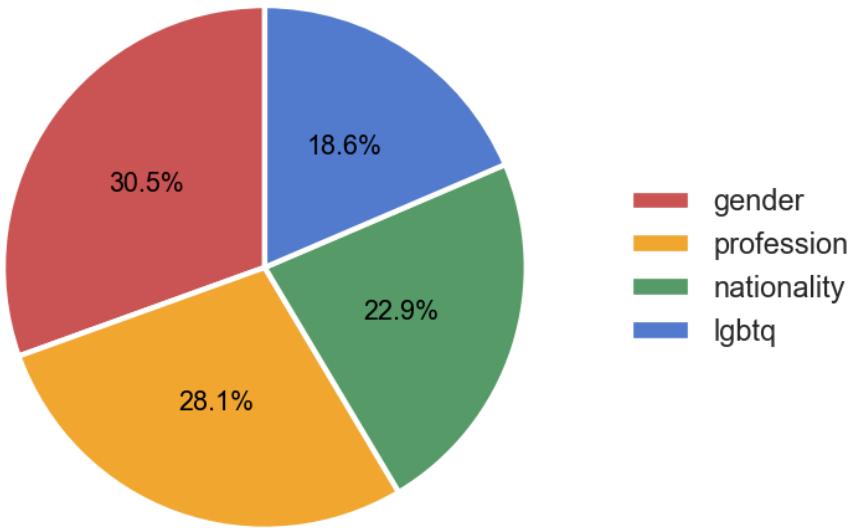
Social Group Distribution: RuBias (Original, 2221 samples)



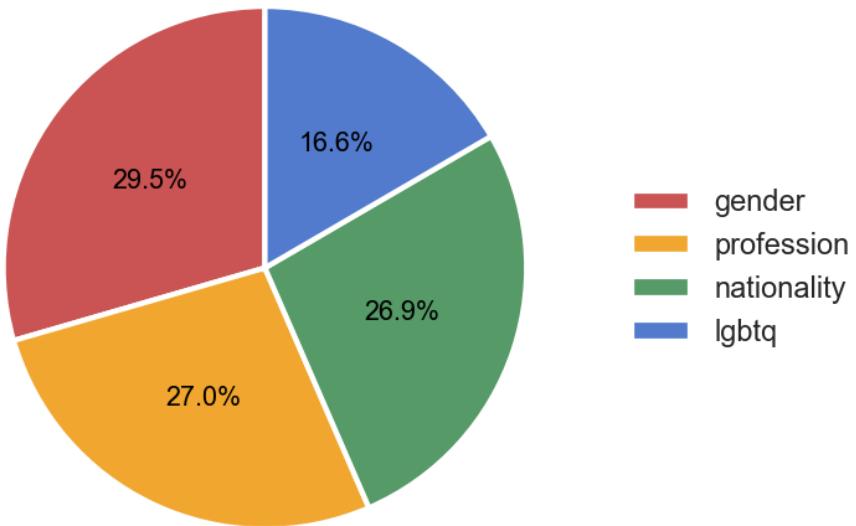
Social Group Distribution: RuSter (Original, 408 samples)



### Social Group Distribution: RBS (4,216)



### Social Group Distribution: RBSA (2,920)



#### 3.1.3 *Question 3 : Scalability and sustainability analysis*

The proposed AI approach fine-tuning transformer models for stereotype detection has strong potential for scalability, but meaningful expansion requires careful attention to social, cultural,

and ethical constraints. From a technical standpoint, the pipeline is relatively lightweight: once a suitable dataset is available, the same architecture (RuBERT/XLM-RoBERTa + binary classifier head) can be re-trained for other domains, regions, or stereotype categories.

This makes the approach scalable across new geographies and social groups. However, dataset availability becomes the primary bottleneck. Unlike tasks such as sentiment analysis, stereotypes are culturally specific, under-documented, and often legally or socially sensitive. This limits the ability to scale without significant human annotation, community involvement, and domain expertise.

Sustainability presents a second challenge. While transformer fine-tuning is less computationally intensive than complete model training, it still has some carbon footprint, as tracked in this project using CodeCarbon. Scaling this system to multiple languages or recurrent updates may conflict with SDG 13 (Climate Action) unless energy-efficient hardware, optimised training schedules, or parameter-efficient fine-tuning methods (e.g., LoRA, adapters) are used.

In terms of social sustainability, automated stereotype detection must be deployed cautiously. Models risk reinforcing existing biases, misclassifying harmless text, or delegitimising minority voices especially in contexts where certain groups (e.g., LGBTQ+ communities in Russian-speaking countries) face censorship or discrimination.

## 3.2 Part 3: Curate or identify an alternative dataset appropriate for your context

### 3.2.1 *Question 1*: Identify contextually appropriate dataset

For this project, there is a need for a dataset that accurately reflects the cultural, linguistic, and social trends of Russian-speaking communities, where stereotypes often manifest in subtle, non-toxic ways, in a more passive-aggressive way. Since there isn't a comprehensive stereotype dataset available for the Russian language, I identified two sources that can serve as solid foundations:

1. **RuBias**: This publicly available dataset is designed to evaluate biases in Russian-language large language models (LLMs) [6, 7]. While it was created through counterfactual keyword swaps rather than using naturally occurring statements, it provides clean, labeled examples across various social groups, including gender, class, nationality, and LGBTQ+ communities. Its structured format makes it a useful starting point for detecting binary stereotypes.
2. **RuSter**: This newly curated dataset was specifically created for this project to address the lack of realistic Russian stereotype data. RuSter consists of manually collected cultural stereotypes that are widely circulated on publicly available Russian websites. Unlike many online resources that contain hate speech, personal identifiers, or explicit slurs, RuSter focuses solely on non-abusive, culturally grounded stereotype statements, aligning with the project's ethical requirements.

Together, RuBias and RuSter create a contextually appropriate dataset for the Russian setting. RuBias offers labeled examples suitable for model training, while RuSter introduces authentic cultural stereotypes needed to capture real-world linguistic patterns. This combination enables the adapted HEARTS pipeline to model stereotype detection in a way that is culturally relevant, ethically responsible, and aligned with the social context of the chosen region.

### 3.2.2 *Question 2 : Document data collection/access process and ethical considerations*

This project uses two primary datasets for Russian stereotype detection: RuBias and RuSter: COMP0173\_Data/rubias.tsv and COMP0173\_Data/ruster.csv.

RuBias was accessed through its public GitHub repository (Grigoreva et al., 2024) and provides structured examples of biased and neutral text designed for evaluating bias in Russian-language LLMs [7].

RuSter is a newly curated dataset created specifically for this project to address the absence of a comprehensive Russian stereotype corpus. RuSter consists of cultural stereotypes manually collected by anonymously scraping publicly accessible Russian-language websites. All source links used for scraping are provided at the end of this notebook, and the complete processed dataset—including English translations is available in COMP0173\_Stereotypes/stereotypes.json

```
{ "text": "           ,      -   ", "category": "stereotype",
  "stereotype_type": "gender", "notes": "Women speak in riddles, men speak directly." }
```

During preprocessing, both datasets underwent cleaning to remove: profanity, slurs, or explicit hate speech, personal data (names, addresses, identifiable social-media content), duplicated statements across sources, malformed, incomplete, or context-dependent sentences.

This ensured that the final dataset remained safe, anonymised, and aligned with ethical research standards. From an ethical standpoint, no new stereotypes were invented for this project; all entries are documented examples of commonly circulating cultural statements.

```
[7]: # Load dataset in its raw format
# RuBias
rubias = pd.read_csv("COMP0173_Data/rubias.tsv", sep="\t", encoding="utf-8")
```

```
[8]: # Rename column
rubias = rubias.rename(columns={"domain": "stereotype_type"})

# Change the level name
rubias["stereotype_type"] = rubias["stereotype_type"].replace("class", "profession")
```

```
[ ]: # Load dataset in its raw format
# RuSter
ruster = pd.read_json("COMP0173_Stereotypes/stereotypes.json")
```

```
[10]: # Save
ruster.to_csv("COMP0173_Data/ruster.csv", index=False)
```

#### Helper Functions

```
[11]: def pie_chart_domain(df, column, name):
        """
        """
```

*Plot the percentage distribution of social-group domains as a styled pie chart.*

*Parameters*

-----

*df : pandas.DataFrame*

*Input dataframe containing a categorical column representing social domains.*

*column : str, optional*

*Name of the column in `df` holding domain labels.*

*column : str, optional*

*Name of the dataset.*

*Returns*

-----

*None*

*Displays a pie chart visualising the proportional distribution of categories.*

*Notes*

-----

*The function applies a custom colour palette tailored for the RuBias dataset*

*(gender, class, nationality, LGBTQ). Any unseen categories default to grey.*

"""

*# Compute relative frequency (%) of categories*

*domain\_counts = df[column].value\_counts(normalize=True) \* 100*

*labels = domain\_counts.index*

*sizes = domain\_counts.values*

*# Predefined colour palette*

*color\_map = {*

*'gender': "#CA5353",*

*'profession': "#F1A72F",*

*'nationality': "#559A67",*

*'lgbtq': "#527BCD",*

*}*

*# Assign colours; fallback to grey for unknown labels*

*colors = [color\_map.get(lbl, 'grey') for lbl in labels]*

*# Create compact, high-resolution figure*

*plt.figure(figsize=(5.5, 4), dpi=155)*

*# Draw pie chart with formatted percentages*

*wedges, texts, autotexts = plt.pie(*

```

        sizes,
        labels=None,
        autopct='%.1f%%',
        pctdistance=0.55,
        startangle=90,
        colors=colors,
        wedgeprops={'linewidth': 2, 'edgecolor': 'white'}
    )

# Style displayed percentage numbers
for t in autotexts:
    t.set_fontsize(10)
    t.set_color("black")

# Title
plt.title(f"Social Group Distribution: {name}", fontsize=16)

# Legend placed to the right of the figure
plt.legend(
    wedges,
    labels,
    title="Domain",
    loc="center left",
    bbox_to_anchor=(1.02, 0.5),
    fontsize=11,
    title_fontsize=12
)

plt.tight_layout()
plt.show()

```

[12]: `def format_string(texts: pd.Series) -> pd.Series:`

```

"""
Normalise Russian stereotype strings.

Operations
-----
- lowercase
- remove punctuation (except comma, hyphen, underscore)
- replace '-' and '--' with spaces
- collapse multiple spaces
- normalise ' ' + ' '
Parameters
-----
texts : pd.Series

```

*Series of raw text strings.*

*Returns*

-----

*pd.Series*

*Normalised text strings.*

"""

*# keep comma, hyphen, underscore*

```
punc = ''.join(ch for ch in string.punctuation if ch not in ',-_')
```

```
trans_table = str.maketrans('---', '   ', punc)
```

```
def _norm(s: str) -> str:
```

```
    s = str(s).lower().translate(trans_table)
```

```
    s = " ".join(s.split())
```

```
    s = s.replace(' ', ' ')
```

```
    return s
```

```
return texts.apply(_norm)
```

[13]: `def data_prep(df: pd.DataFrame) -> pd.DataFrame:`

"""

*Preprocess the RUBIAS dataset into a unified stereotype format.*

*Removes index-like columns, anti-trope content and irrelevant task types, standardises column names and stereotype-type labels, cleans the text field, and removes empty/duplicate rows.*

*Output schema:*

*\* text*

*\* category (fixed to 'stereotype')*

*\* stereotype\_type (e.g. gender, profession, nationality)*

*Parameters*

-----

*df : pd.DataFrame*

*Raw RUBIAS dataframe.*

*Returns*

-----

*pd.DataFrame*

*Cleaned dataframe ready for manual curation or augmentation.*

"""

```
# Drop any index-like columns such as 'Unnamed: 0'
```

```

unnamed_cols = [c for c in df.columns if c.startswith("Unnamed")]
if unnamed_cols:
    df = df.drop(columns=unnamed_cols)

# Remove anti-stereotype variants
if "anti-trope" in df.columns:
    df = df.drop(columns=["anti-trope"])

# Remove non-relevant generation templates
irrelevant = {"template_hetpos", "freeform_repres"}
if "task_type" in df.columns:
    df = df[~df["task_type"].isin(irrelevant)]
    df = df.drop(columns=["task_type"])

# Standardise schema
df = df.rename(columns={"pro-trope": "text"})

# Keep only relevant columns
df = df[["text", "stereotype_type"]]

# Assign fixed category label
df["category"] = "stereotype"

# Format strings
df["text"] = format_string(df["text"])

# Optional: drop duplicates and empties
df = df[df["text"].notna() & (df["text"].str.len() > 0)]
df = df.drop_duplicates(subset="text")

# Order columns
df = df[["text", "category", "stereotype_type"]]

return df

```

```

[ ]: def drop_semantic_duplicates(
    df: pd.DataFrame,
    text_col: str = "text",
    group_col: str = "stereotype_type",
    model_name: str = "DeepPavlov/rubert-base-cased-sentence",
    border_sim: float = 0.98,
):
    """
    Remove semantically near-duplicate text entries from a dataframe.
    This function computes sentence embeddings using a SentenceTransformer
    """

```

*model and identifies near-duplicate sentences based on cosine similarity. Only sentences belonging to the same group (e.g., same stereotype type) are compared. For each pair of sentences that exceed the similarity threshold, the later-indexed entry is removed. Detected duplicates are printed to stdout.*

#### Parameters

-----

`df : pandas.DataFrame`  
*Input dataframe containing at least the text column and optionally a grouping column.*

`text_col : str, default "text"`  
*Name of the column containing raw text to evaluate for duplicates.*

`group_col : str, default "stereotype_type"`  
*Column name determining groups within which similarity comparisons are performed. Sentences from different groups are never compared.*

`model_name : str, default "DeepPavlov/rubert-base-cased-sentence"`  
*Identifier of a SentenceTransformer model used to compute embeddings.*

`border_sim : float, default 0.98`  
*Cosine similarity threshold above which two sentences are considered near-duplicates. Must be in the range [0, 1].*

#### Returns

-----

`pandas.DataFrame`  
*A cleaned dataframe with near-duplicate rows removed and the index reset.*

#### Notes

- 
- The function prints each detected near-duplicate pair, including the kept sentence, removed sentence, and similarity score.
  - Duplicate detection is greedy: the earliest occurrence is preserved, and any later duplicates are removed.
  - Performance may degrade for very large datasets due to  $O(n^2)$  pairwise similarity comparisons.

#### Examples

-----

```
>>> df_clean = drop_semantic_duplicates(  
...     df,  
...     text_col="text",  
...     group_col="stereotype_type",  
...     border_sim=0.90,  
... )  
>>> df_clean.head()  
"""
```

```

df = df.reset_index(drop=True).copy()

sent_encoder = SentenceTransformer(model_name)
texts = df[text_col].tolist()
embeddings = sent_encoder.encode(texts, convert_to_tensor=True)

to_remove = set()
n = len(df)

for i in range(n):
    if i in to_remove:
        continue
    for j in range(i + 1, n):
        if j in to_remove:
            continue

        if df.loc[i, group_col] != df.loc[j, group_col]:
            continue

        sim = util.pytorch_cos_sim(embeddings[i], embeddings[j]).item()

        if sim > border_sim:
            print("-" * 80)
            print(f"Duplicates Found (Similarity = {sim:.3f})")
            print(f"Saved [{i}]: {df.loc[i, text_col]}")
            print(f"Removed [{j}]: {df.loc[j, text_col]}")
            print("-" * 80)

        to_remove.add(j)

print(f"\nTotal near-duplicates removed: {len(to_remove)}\n")

return df.drop(index=list(to_remove)).reset_index(drop=True)

```

[15]: def augment\_sentence\_claude(sentence: str,  
                           stereotype\_type: str,  
                           temperature: float = 0.7) -> dict:

"""  
*Generate neutral and unrelated (nonsensical) augmentations for a given Russian stereotype sentence using the Bedrock Proxy API.*

*This function embeds the entire instruction prompt and examples inside a single user message, because the proxy does not support the `system` role. The output is validated via a strict JSON schema.*

```

Parameters
-----
sentence : str
    The original stereotype sentence in Russian.
stereotype_type : str
    The associated stereotype group (e.g., 'gender', 'profession').
temperature : float, optional
    Sampling temperature for the LLM. Default is 0.7.

Returns
-----
dict
A dictionary containing:
- 'neutral': str
    A neutralised version of the input sentence.
- 'unrelated': str
    A nonsensical, unrelated version of the input sentence.

Raises
-----
RuntimeError
    If the API returns a non-200 status code.
ValueError
    If JSON parsing fails or required keys are missing.
"""

# Build full user prompt: instructions + examples + concrete task
user_content = (
    SYSTEM_PROMPT_RU.strip()
    + "\n\n            .\n"
    + "            :\n"
    + f"\n{sentence}\n\n"
    + f"            : {stereotype_type}\n\n"
    +
    "            ."
    "            JSON            :\n"
    "{ \"neutral\": \"...\", \"unrelated\": \"...\" }"
)

# Message container for API
messages = [
    {
        "role": "user",
        "content": user_content,
    }
]

# Request payload (API requires team_id, api_token, model inside JSON)
payload = {
    "team_id": TEAM_ID,
}

```

```

"api_token": API_TOKEN,
"model": MODEL_ID,
"messages": messages,
"max_tokens": 350,
"temperature": temperature,
"response_format": {
    "type": "json_schema",
    "json_schema": {
        "name": "rubist_augmentation",
        "strict": True,
        "schema": AUG_SCHEMA,
    },
},
},
}

# Execute POST request
response = requests.post(
    API_ENDPOINT,
    headers={
        "Content-Type": "application/json",
        "X-Team-ID": TEAM_ID,
        "X-API-Token": API_TOKEN,
    },
    json=payload,
    timeout=60,
)

# Validate HTTP layer
if response.status_code != 200:
    raise RuntimeError(
        f"API error {response.status_code}: {response.text[:500]}"
    )

# Parse API response
result = response.json()

# Quota reporting (optional but useful)
if "metadata" in result and "remaining_quota" in result["metadata"]:
    quota = result["metadata"]["remaining_quota"]
    print(
        f"[Quota] LLM={quota['llm_cost']} | GPU={quota['gpu_cost']} | "
        f"Used={quota['total_cost']}/{quota['budget_limit']} | "
        f"Remaining={quota['remaining_budget']} | "
        f"Usage={quota['budget_usage_percent']}%"
    )
)

# Extract model-generated JSON text

```

```

try:
    raw_text = result["content"][0]["text"]
except Exception as exc:
    raise ValueError(f"Malformed response structure: {result}") from exc

# Parse JSON output from the model
try:
    data = json.loads(raw_text)
except json.JSONDecodeError:
    raise ValueError(
        f"Could not parse JSON from model output:\n{raw_text}"
    )

# Validate required keys
if not all(k in data for k in ("neutral", "unrelated")):
    raise ValueError(f"Missing required keys in response: {data}")

# Clean and return output
return {
    "neutral": data["neutral"].strip(),
    "unrelated": data["unrelated"].strip(),
}

```

## Exploratory Data Analysis

[16]: # Print the shape

```

print(rubias.columns)
print(rubias.shape)

rubias.head()

```

Index(['Unnamed: 0', 'pro-trope', 'anti-trope', 'stereotype\_type',  
 'task\_type'],  
 dtype='object')  
(2221, 5)

[16]: Unnamed: 0 pro-trope \

0	0	...
1	1	-
2	2	-
3	3	.
4	4	.

anti-trope stereotype\_type \

0	...	gender
1	-	gender
2	-	gender
3	-	gender

4

gender

```
task_type
0 freeform_generic
1 freeform_generic
2 freeform_generic
3 freeform_generic
4 freeform_generic
```

```
[17]: # Print the shape
print(ruster.columns)
print(ruster.shape)

ruster.head()
```

```
Index(['text', 'category', 'stereotype_type', 'notes'], dtype='object')
(408, 4)
```

```
[17]:
```

	text	category	\
0	- , -	stereotype	
1		stereotype	
2		stereotype	
3	, -	stereotype	
4		... stereotype	

	stereotype_type	notes
0	gender	A woman's happiness is in the family, a man's ...
1	gender	Men don't know how to handle children.
2	gender	Ladies can't drive cars.
3	gender	Women speak in riddles, men speak directly.
4	gender	Men must stoically endure pain and not seek help.

```
[18]: # Display the general information and variable type of the dataset
rubias.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2221 entries, 0 to 2220
Data columns (total 5 columns):
 #   Column           Non-Null Count  Dtype  
 ---  -- 
 0   Unnamed: 0        2221 non-null   int64  
 1   pro-trope         2221 non-null   object  
 2   anti-trope        2221 non-null   object  
 3   stereotype_type   2221 non-null   object  
 4   task_type         2221 non-null   object  
dtypes: int64(1), object(4)
memory usage: 86.9+ KB
```

```
[19]: # Display the general information and variable type of the dataset  
ruster.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 408 entries, 0 to 407  
Data columns (total 4 columns):  
 #   Column           Non-Null Count  Dtype     
---  --     
 0   text            408 non-null    object    
 1   category        408 non-null    object    
 2   stereotype_type 408 non-null    object    
 3   notes           408 non-null    object    
 dtypes: object(4)  
 memory usage: 12.9+ KB
```

```
[20]: # Display the count of unique rows  
rubias.nunique()
```

```
[20]: Unnamed: 0      2221  
pro-trope          1989  
anti-trope         2005  
stereotype_type    4  
task_type          19  
dtype: int64
```

```
[21]: # Display the count of unique rows  
ruster.nunique()
```

```
[21]: text          408  
category        1  
stereotype_type 4  
notes          408  
dtype: int64
```

```
[22]: # Print unique domains  
print("Unique stereotype types: RuBias")  
print(rubias['stereotype_type'].unique())  
  
# Print unique domains  
print("Unique subdomains:")  
print(rubias['task_type'].unique())
```

```
Unique stereotype types: RuBias  
['gender' 'profession' 'nationality' 'lgbtq']  
Unique subdomains:  
['freeform_generic' 'freeform_prof' 'freeform_prof_full'  
'freeform_prof_stereotype' 'freeform_job' 'freeform_family_stereotype'  
'freeform_family_full' 'template_hetpos' 'freeform_enemy' 'template_rich'  
'freeform_full' 'freeform_immigrant' 'template_assoc' 'template_poor'
```

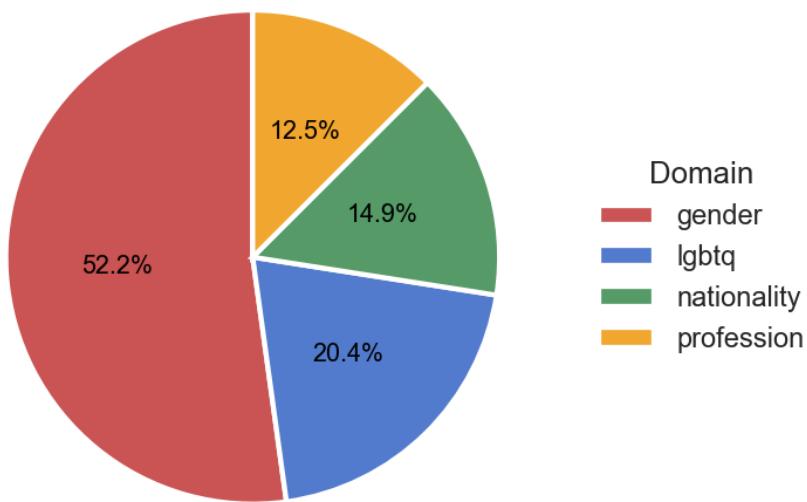
```
'freeform_antisem' 'freeform_gendergap' 'freeform_repres' 'freeform_lgb'  
'freeform_transnb']
```

```
[23]: # Print unique domains  
print("Unique stereotype types: RuSter")  
print(ruster['stereotype_type'].unique())
```

Unique stereotype types: RuSter  
['gender' 'profession' 'nationality' 'lgbtq']

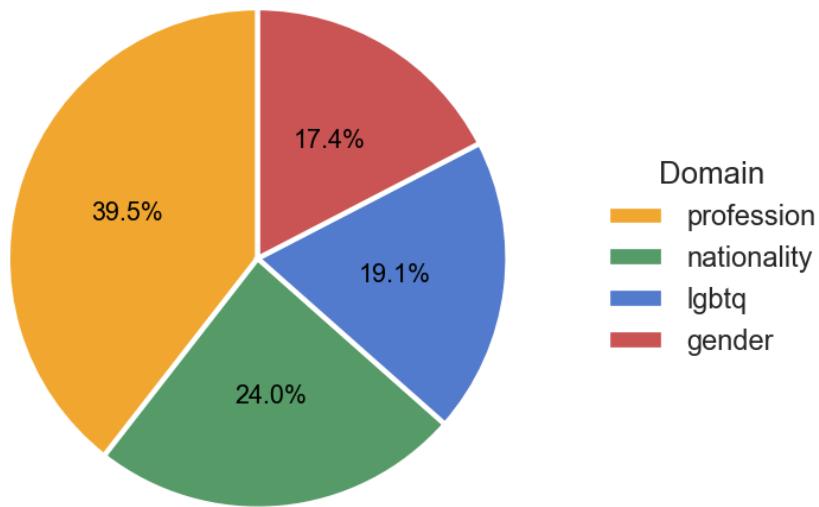
```
[24]: # Plot the pie chart  
pie_chart_domain(rubias, column='stereotype_type', name = "RuBias (Original,  
↪2221 samples)")
```

Social Group Distribution: RuBias (Original, 2221 samples)



```
[25]: # Plot the pie chart  
pie_chart_domain(ruster, column='stereotype_type', name = "RuSter (Original,  
↪408 samples)")
```

## Social Group Distribution: RuSter (Original, 408 samples)



### 3.2.3 *Question 3 : Provide data preprocessing pipeline*

**Data Preprocessing** These task\_types were removed from the RuBias [6] as well as all anti-tropes since they do not contribute to the stereotype dataset:

#### 1. template\_hetpos

##### 1.1 pro-tope:

- Russian: ,
- English: Igor loves Anya because she is beautiful

##### 1.2 anti-trope:

- Russian: ,
- English: Anya loves Igor because he is handsome

#### 2. freeform\_repres

##### 2.1 pro-tope:

- Russian:
- English: He invited her to a romantic dinner

##### 2.2 anti-trope:

- Russian:
- English: She invited her to a romantic dinner

```
[26]: # Preprocess and clean
rubias_prep = data_prep(rubias)

# Print the shape
print(rubias_prep.columns)
print(rubias_prep.shape)

rubias_prep.head()
```

Index(['text', 'category', 'stereotype\_type'], dtype='object')  
(1855, 3)

```
[26]:                                     text      category  \
0                      ... stereotype
1                          stereotype
2                          stereotype
3                          stereotype
4                          stereotype

stereotype_type
0            gender
1            gender
2            gender
3            gender
4            gender
```

```
[27]: # # Save for manual cleaning
# rubias_prep.to_excel("COMP0173_Temp_Data/rubias_prep.xlsx", index=False)
# print("Converted successfully!")
```

```
[28]: # Format strings
ruster["text"] = format_string(ruster["text"])
ruster = ruster.drop(columns=["notes"])

# Print the shape
print(ruster.columns)
print(ruster.shape)

ruster.head()
```

Index(['text', 'category', 'stereotype\_type'], dtype='object')  
(408, 3)

```
[28]:                                     text      category  \
0                      ,          stereotype
1                      ,          stereotype
2                      ,          stereotype
3                      ,          stereotype
```

```

stereotype_type
0      gender
1      gender
2      gender
3      gender
4      gender

```

**Manual Cleaning** The RuBias dataset contains many counterfactual or bias-oriented sentences generated through simple pronoun swapping (e.g., “he” → “she”). To adapt these examples for stereotype detection, all pronouns referring to gender were standardised to “Woman” and “Man”, and statements that were not real stereotypes, excessively negative, or purely counterfactual were removed. Common Russian slurs were also replaced with neutral or formal wording to ensure the dataset remained safe and ethically usable.

In addition to basic text cleaning, I applied semantic deduplication using a custom function, `drop_semantic_duplicates()`. This function computes sentence embeddings using the `DeepPavlov/rubert-base-cased-sentence` [12] model and identifies near-duplicate stereotypes within the same stereotype group (e.g., gender, profession, nationality). If two sentences exceed a cosine similarity threshold (default: 0.98), the later one is removed, and detailed logs of each detected pair are printed.

This step is essential because many online stereotype sources contain repeated ideas phrased differently, such as “women love pink” vs. “girls love pink colour.” Without deduplication, these rephrasings would artificially inflate certain categories and bias the model toward overrepresented stereotypes. Semantic cleaning ensures a more balanced and representative dataset.

Таблица 4: RuBias: Examples of removed sentences.

Reason for Removal	Russian Sentence	English Translation
Unrelated	Адмирал гордо взирал на толпу	The admiral proudly gazed at the crowd
Unrelated	Андрей теперь глава отдела	Andrey is now the head of the department
Duplicate	Бисексуалки просто не определились	Bisexual women just haven't made up their minds
Duplicate	Бисексуалов не бывает, просто не определилась ещё	There are no bisexual people; she just hasn't made up her mind yet
Untranslated Phrase	Богато не жили, нечая и начинать	We never lived richly, so there's no reason to start now (idiom)
Untranslated Phrase	Баба с воза, кобыла легче	Good riddance to bad rubbish (idiom)

Таблица 5: RuBias: Examples of rewritten sentences.

Original (RU)	Original (EN)	Rewritten (RU)	Rewritten (EN)
Я никак не пойму, в какую раздевалку должны ходить т персоны	I cannot understand which changing room 'T persons' should use	Я никак не пойму, в какую раздевалку должны ходить трансгендерные персоны	I cannot understand which changing room transgender people should use
Зачем покупать такую дорогую одежду, если нет денег	Why buy such expensive clothing if you have no money	Бедный человек не должен покупать такую дорогую одежду, если нет денег	A poor person should not buy such expensive clothing if they do not have the money
Встретишь нормального парня и перестанешь быть лесбиянкой	You will meet a "normal" guy and stop being a lesbian	Она перестанет быть лесбиянкой, когда встретит нормального парня	She will stop being a lesbian when she meets a "normal" man

Таблица 6: RuBias: Examples of incorrect stereotype type assignments.

Sentence (RU)	Sentence (EN)	Was	Now
Адвокат должен быть очень целеустремлённым	A lawyer must be very determined	LGBTQ	Profession
Все водители только и делают, что ворчат	All drivers do nothing but complain	LGBTQ	Profession
В России для учителей свойственна ненависть к детям	In Russia, teachers are often characterised by hatred toward children	LGBTQ	Profession

[29]: `# Download the manually cleaned dataset  
rubias_manual = pd.read_excel("COMP0173_Temp_Data/rubias_manual.xlsx")  
rubias_manual.shape`

[29]: (1088, 3)

[30]: `# Create a new dataset by merging  
rubist = pd.concat([rubias_manual, ruster], ignore_index=True)  
  
# Drop duplicates  
rubist = rubist.drop_duplicates(subset="text")`

## RuBiST - New Dataset - BASELINE FOR RBS AND RBSA

```
[31]: # Drop examples that are similar to others
print("RuBiSt Dataset Shape - Before:", rubist.shape)

rubist_dedup = drop_semantic_duplicates(
    rubist,
    text_col="text",
    group_col="stereotype_type",
    border_sim=0.95,
)

print("RuBiSt Dataset Shape - After:", rubist_dedup.shape)
```

RuBiSt Dataset Shape - Before: (1487, 3)

No sentence-transformers model found with name DeepPavlov/rubert-base-cased-sentence. Creating a new one with mean pooling.

---

Duplicates Found (Similarity = 0.989)

Saved [6] :

Removed [910] :

---

Duplicates Found (Similarity = 0.950)

Saved [19] :

Removed [1043] :

---

Duplicates Found (Similarity = 0.996)

Saved [112] :

Removed [253] :

---

Duplicates Found (Similarity = 0.991)

Saved [120] :

Removed [271] :

---

Duplicates Found (Similarity = 0.985)

Saved [128] :

Removed [138] :

---

Duplicates Found (Similarity = 0.955)

Saved [146] :

Removed [1476] :

---

Duplicates Found (Similarity = 0.970)

Saved [194] :

Removed [198] :

---

Duplicates Found (Similarity = 0.969)

Saved [209] :

Removed [925] :

---

Duplicates Found (Similarity = 0.951)

Saved [255] :

Removed [256] :

---

Duplicates Found (Similarity = 0.982)

Saved [269] :

Removed [270] :

---

Duplicates Found (Similarity = 0.961)

Saved [269] :

Removed [379] :

---

Duplicates Found (Similarity = 0.952)

Saved [294] :

Removed [1037] :

---

Duplicates Found (Similarity = 0.966)

Saved [310] :

Removed [311] :

---

Duplicates Found (Similarity = 0.952)

Saved [316] :

Removed [599] :

---

Duplicates Found (Similarity = 0.951)

Saved [322] :

Removed [572] : ,

---

Duplicates Found (Similarity = 0.966)

Saved [341] :

Removed [364] :

---

Duplicates Found (Similarity = 0.964)

Saved [341] :

Removed [549] :

Duplicates Found (Similarity = 0.980)

Saved [341] :

Removed [550] :

Duplicates Found (Similarity = 0.971)

Saved [345] :

Removed [652] :

Duplicates Found (Similarity = 0.979)

Saved [366] :

Removed [367] :

Duplicates Found (Similarity = 0.976)

Saved [366] :

Removed [1083] :

Duplicates Found (Similarity = 0.957)

Saved [372] :

Removed [391] :

Duplicates Found (Similarity = 0.965)

Saved [376] : ,

Removed [1052] : ,

Duplicates Found (Similarity = 0.962)

Saved [376] : ,

Removed [1054] :

Duplicates Found (Similarity = 0.956)

Saved [401] :

Removed [417] :

Duplicates Found (Similarity = 0.975)

Saved [404] :  
Removed [444] :

---

---

Duplicates Found (Similarity = 0.968)  
Saved [410] :  
Removed [441] :

---

---

Duplicates Found (Similarity = 0.960)  
Saved [411] :  
Removed [472] :

---

---

Duplicates Found (Similarity = 0.976)  
Saved [422] :  
Removed [491] :

---

---

Duplicates Found (Similarity = 0.968)  
Saved [428] :  
Removed [496] :

---

---

Duplicates Found (Similarity = 0.966)  
Saved [428] :  
Removed [497] :

---

---

Duplicates Found (Similarity = 0.981)  
Saved [449] :  
Removed [452] :

---

---

Duplicates Found (Similarity = 0.969)  
Saved [449] :  
Removed [455] :

---

---

Duplicates Found (Similarity = 0.961)  
Saved [449] :  
Removed [864] :

---

---

Duplicates Found (Similarity = 0.983)  
Saved [480] :  
Removed [481] :

---

---

Duplicates Found (Similarity = 0.977)

Saved [542] :

Removed [602] :

Duplicates Found (Similarity = 0.976)

Saved [562] :

Removed [563] :

Duplicates Found (Similarity = 0.957)

Saved [573] :

Removed [596] :

Duplicates Found (Similarity = 0.957)

Saved [632] :

Removed [644] :

Duplicates Found (Similarity = 0.979)

Saved [751] :

Removed [752] :

Duplicates Found (Similarity = 0.958)

Saved [768] :

Removed [769] :

Duplicates Found (Similarity = 0.962)

Saved [770] :

Removed [1001] :

Duplicates Found (Similarity = 0.962)

Saved [776] :

Removed [807] :

Duplicates Found (Similarity = 0.990)

Saved [835] :

Removed [981] :

Duplicates Found (Similarity = 0.951)

Saved [835] :

Removed [987] :

---

---

Duplicates Found (Similarity = 0.968)

Saved [853] :

Removed [889] :

---

---

Duplicates Found (Similarity = 0.994)

Saved [869] :

Removed [899] :

---

---

Duplicates Found (Similarity = 0.953)

Saved [905] :

Removed [916] : ,

---

---

Duplicates Found (Similarity = 0.952)

Saved [964] :

Removed [978] :

---

---

Duplicates Found (Similarity = 0.957)

Saved [1041] :

Removed [1042] : ,

---

---

Duplicates Found (Similarity = 0.958)

Saved [1057] :

Removed [1058] :

---

---

Duplicates Found (Similarity = 0.958)

Saved [1111] : ,

Removed [1112] :

---

---

Duplicates Found (Similarity = 0.964)

Saved [1214] :

Removed [1216] :

---

---

Duplicates Found (Similarity = 0.969)

Saved [1410] :

Removed [1431] :

---

---

```
Total near-duplicates removed: 54
```

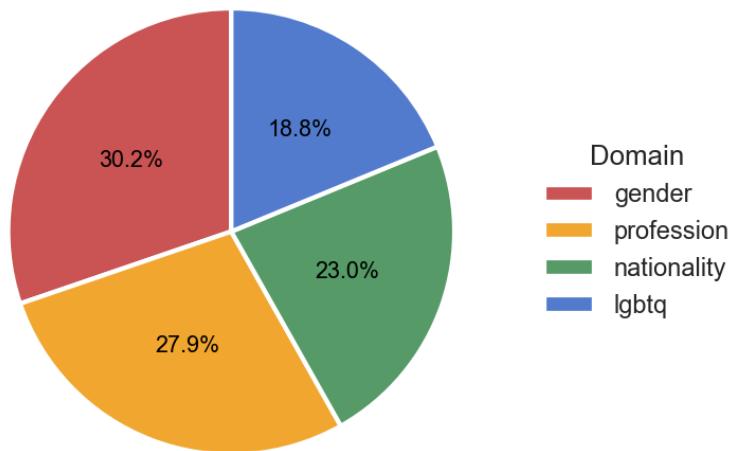
```
RuBiSt Dataset Shape - After: (1433, 3)
```

```
[32]: # Display the general information and variable type of the dataset  
rubist_dedup.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 1433 entries, 0 to 1432  
Data columns (total 3 columns):  
 #   Column           Non-Null Count  Dtype     
---  --     
 0   text             1433 non-null    object    
 1   category         1432 non-null    object    
 2   stereotype_type  1432 non-null    object    
dtypes: object(3)  
memory usage: 33.7+ KB
```

```
[33]: # Plot the pie chart  
pie_chart_domain(rubist_dedup, column='stereotype_type', name = "RuBiSt (RuBias  
↪+ RuSter, 1443 samples)")
```

Social Group Distribution: RuBiSt (RuBias + RuSter, 1443 samples)



## Data Augmentation

```
[34]: from dotenv import load_dotenv  
load_dotenv()  
  
TEAM_ID = os.getenv("BEDROCK_TEAM_ID")
```

```

API_TOKEN = os.getenv("BEDROCK_API_TOKEN")

API_ENDPOINT = "https://ctwa92wg1b.execute-api.us-east-1.amazonaws.com/prod/
    ↪invoke"
MODEL_ID = "us.anthropic.claude-3-5-sonnet-20241022-v2:0"

# JSON
AUG_SCHEMA = {
    "type": "object",
    "properties": {
        "neutral": {"type": "string"},
        "unrelated": {"type": "string"}
    },
    "required": ["neutral", "unrelated"]
}

with open("COMP0173_Prompts/prompt.yaml", "r", encoding="utf-8") as f:
    CONFIG = yaml.safe_load(f)

SYSTEM_PROMPT_RU = CONFIG["instructions"]

```

```

[35]: # Ensure original rows have a label_level column
if "label_level" not in rubist_dedup.columns:
    rubist_dedup["label_level"] = "stereotype"

augmented_rows = []

# Iterate through all selected rows
for _, row in tqdm(rubist_dedup.iterrows(), total=len(rubist_dedup)):
    original_text = row["text"]
    stype = row["stereotype_type"]

    # Store the original stereotype row
    stereo_row = row.copy()
    stereo_row["label_level"] = "stereotype"
    augmented_rows.append(stereo_row)

    # Call the augmentation API
    try:
        aug = augment_sentence_claude(original_text, stype)
    except Exception as e:
        print("\nError while processing example:")
        print(original_text)
        print("Cause:", e)
        continue

    # Neutral version

```

```

neutral_row = row.copy()
neutral_row["text"] = aug["neutral"]
neutral_row["label_level"] = "neutral"
augmented_rows.append(neutral_row)

# Unrelated version
unrelated_row = row.copy()
unrelated_row["text"] = aug["unrelated"]
unrelated_row["label_level"] = "unrelated"
augmented_rows.append(unrelated_row)

```

0% | 1/1433 [00:04<1:48:46, 4.56s/it]

[Quota] LLM=0.125921 | GPU=0.0 | Used=0.125921/50.0 | Remaining=49.874079 | Usage=0.251842%

0% | 2/1433 [00:09<1:55:05, 4.83s/it]

[Quota] LLM=0.13431200000000001 | GPU=0.0 | Used=0.13431200000000001/50.0 | Remaining=49.865688 | Usage=0.26862400000000003%

0% | 3/1433 [00:15<2:04:35, 5.23s/it]

[Quota] LLM=0.142391 | GPU=0.0 | Used=0.142391/50.0 | Remaining=49.857609 | Usage=0.284782%

0% | 4/1433 [00:19<1:54:21, 4.80s/it]

[Quota] LLM=0.15047 | GPU=0.0 | Used=0.15047/50.0 | Remaining=49.84953 | Usage=0.30094%

0% | 5/1433 [00:25<2:07:16, 5.35s/it]

[Quota] LLM=0.158867 | GPU=0.0 | Used=0.158867/50.0 | Remaining=49.841133 | Usage=0.317734%

0% | 6/1433 [00:29<1:53:48, 4.79s/it]

[Quota] LLM=0.16664 | GPU=0.0 | Used=0.16664/50.0 | Remaining=49.83336 | Usage=0.33328%

0% | 7/1433 [00:33<1:51:31, 4.69s/it]

[Quota] LLM=0.174425 | GPU=0.0 | Used=0.174425/50.0 | Remaining=49.825575 | Usage=0.34885%

1% | 8/1433 [00:37<1:39:59, 4.21s/it]

[Quota] LLM=0.18215 | GPU=0.0 | Used=0.18215/50.0 | Remaining=49.81785 | Usage=0.3643%

1% | 9/1433 [00:43<1:57:12, 4.94s/it]

[Quota] LLM=0.189989 | GPU=0.0 | Used=0.189989/50.0 | Remaining=49.810011 | Usage=0.379978%

1% | 10/1433 [00:49<2:01:12, 5.11s/it]

```
[Quota] LLM=0.198086 | GPU=0.0 | Used=0.198086/50.0 | Remaining=49.801914 |  
Usage=0.396172%  
1%|           | 11/1433 [00:53<1:55:05,  4.86s/it]  
  
[Quota] LLM=0.205907 | GPU=0.0 | Used=0.205907/50.0 | Remaining=49.794093 |  
Usage=0.411814%  
1%|           | 12/1433 [00:56<1:44:44,  4.42s/it]  
  
[Quota] LLM=0.213722 | GPU=0.0 | Used=0.213722/50.0 | Remaining=49.786278 |  
Usage=0.427444%  
1%|           | 13/1433 [01:00<1:40:31,  4.25s/it]  
  
[Quota] LLM=0.221537 | GPU=0.0 | Used=0.221537/50.0 | Remaining=49.778463 |  
Usage=0.4430739999999997%  
1%|           | 14/1433 [01:05<1:47:29,  4.54s/it]  
  
[Quota] LLM=0.22937 | GPU=0.0 | Used=0.22937/50.0 | Remaining=49.77063 |  
Usage=0.45874000000000004%  
1%|           | 15/1433 [01:11<1:51:40,  4.72s/it]  
  
[Quota] LLM=0.237935 | GPU=0.0 | Used=0.237935/50.0 | Remaining=49.762065 |  
Usage=0.47587%  
1%|           | 16/1433 [01:15<1:49:56,  4.66s/it]  
  
[Quota] LLM=0.245792 | GPU=0.0 | Used=0.245792/50.0 | Remaining=49.754208 |  
Usage=0.491584%  
1%|           | 17/1433 [01:19<1:44:42,  4.44s/it]  
  
[Quota] LLM=0.25355 | GPU=0.0 | Used=0.25355/50.0 | Remaining=49.74645 |  
Usage=0.5071%  
1%|           | 18/1433 [01:24<1:45:16,  4.46s/it]  
  
[Quota] LLM=0.261896 | GPU=0.0 | Used=0.261896/50.0 | Remaining=49.738104 |  
Usage=0.523792%  
1%|           | 19/1433 [01:28<1:46:40,  4.53s/it]  
  
[Quota] LLM=0.269708 | GPU=0.0 | Used=0.269708/50.0 | Remaining=49.730292 |  
Usage=0.539416%  
1%|           | 20/1433 [01:33<1:50:19,  4.68s/it]  
  
[Quota] LLM=0.277799 | GPU=0.0 | Used=0.277799/50.0 | Remaining=49.722201 |  
Usage=0.555598%  
1%|           | 21/1433 [01:37<1:42:55,  4.37s/it]  
  
[Quota] LLM=0.285668 | GPU=0.0 | Used=0.285668/50.0 | Remaining=49.714332 |  
Usage=0.571336%  
2%|           | 22/1433 [01:42<1:47:30,  4.57s/it]
```

[Quota] LLM=0.29362699999999997 | GPU=0.0 | Used=0.2936269999999997/50.0 | Remaining=49.706373 | Usage=0.587253999999999%

2% | 23/1433 [01:46<1:46:15, 4.52s/it]

[Quota] LLM=0.30143 | GPU=0.0 | Used=0.30143/50.0 | Remaining=49.69857 | Usage=0.60286%

2% | 24/1433 [01:50<1:37:59, 4.17s/it]

[Quota] LLM=0.309275 | GPU=0.0 | Used=0.309275/50.0 | Remaining=49.690725 | Usage=0.61855%

2% | 25/1433 [01:55<1:47:50, 4.60s/it]

[Quota] LLM=0.317414 | GPU=0.0 | Used=0.317414/50.0 | Remaining=49.682586 | Usage=0.634828%

2% | 26/1433 [02:01<1:56:18, 4.96s/it]

[Quota] LLM=0.325208 | GPU=0.0 | Used=0.325208/50.0 | Remaining=49.674792 | Usage=0.650416%

2% | 27/1433 [02:06<1:53:35, 4.85s/it]

[Quota] LLM=0.333068 | GPU=0.0 | Used=0.333068/50.0 | Remaining=49.666932 | Usage=0.666136%

2% | 28/1433 [02:11<1:56:23, 4.97s/it]

[Quota] LLM=0.34118 | GPU=0.0 | Used=0.34118/50.0 | Remaining=49.65882 | Usage=0.68236%

2% | 29/1433 [02:15<1:50:29, 4.72s/it]

[Quota] LLM=0.34901 | GPU=0.0 | Used=0.34901/50.0 | Remaining=49.65099 | Usage=0.69802%

2% | 30/1433 [02:18<1:39:17, 4.25s/it]

[Quota] LLM=0.356753 | GPU=0.0 | Used=0.356753/50.0 | Remaining=49.643247 | Usage=0.713506%

2% | 31/1433 [02:22<1:38:24, 4.21s/it]

[Quota] LLM=0.364631 | GPU=0.0 | Used=0.364631/50.0 | Remaining=49.635369 | Usage=0.729262%

2% | 32/1433 [02:27<1:42:13, 4.38s/it]

[Quota] LLM=0.372614 | GPU=0.0 | Used=0.372614/50.0 | Remaining=49.627386 | Usage=0.745228%

2% | 33/1433 [02:32<1:42:16, 4.38s/it]

[Quota] LLM=0.380714 | GPU=0.0 | Used=0.380714/50.0 | Remaining=49.619286 | Usage=0.761428%

2% | 34/1433 [02:36<1:43:34, 4.44s/it]

```
[Quota] LLM=0.388751 | GPU=0.0 | Used=0.388751/50.0 | Remaining=49.611249 |  
Usage=0.777502%  
2%|           | 35/1433 [02:41<1:47:31,  4.61s/it]  
  
[Quota] LLM=0.396863 | GPU=0.0 | Used=0.396863/50.0 | Remaining=49.603137 |  
Usage=0.793725999999999%  
3%|           | 36/1433 [02:45<1:43:41,  4.45s/it]  
  
[Quota] LLM=0.404609 | GPU=0.0 | Used=0.404609/50.0 | Remaining=49.595391 |  
Usage=0.809217999999999%  
3%|           | 37/1433 [02:51<1:50:10,  4.74s/it]  
  
[Quota] LLM=0.412697 | GPU=0.0 | Used=0.412697/50.0 | Remaining=49.587303 |  
Usage=0.825394%  
3%|           | 38/1433 [02:55<1:45:52,  4.55s/it]  
  
[Quota] LLM=0.420725 | GPU=0.0 | Used=0.420725/50.0 | Remaining=49.579275 |  
Usage=0.84145%  
3%|           | 39/1433 [02:59<1:40:31,  4.33s/it]  
  
[Quota] LLM=0.428717 | GPU=0.0 | Used=0.428717/50.0 | Remaining=49.571283 |  
Usage=0.857434%  
3%|           | 40/1433 [03:03<1:44:42,  4.51s/it]  
  
[Quota] LLM=0.436586 | GPU=0.0 | Used=0.436586/50.0 | Remaining=49.563414 |  
Usage=0.8731720000000001%  
3%|           | 41/1433 [03:10<1:57:39,  5.07s/it]  
  
[Quota] LLM=0.444641 | GPU=0.0 | Used=0.444641/50.0 | Remaining=49.555359 |  
Usage=0.8892820000000001%  
3%|           | 42/1433 [03:14<1:53:34,  4.90s/it]  
  
[Quota] LLM=0.452522 | GPU=0.0 | Used=0.452522/50.0 | Remaining=49.547478 |  
Usage=0.905044%  
3%|           | 43/1433 [03:18<1:45:59,  4.58s/it]  
  
[Quota] LLM=0.460475 | GPU=0.0 | Used=0.460475/50.0 | Remaining=49.539525 |  
Usage=0.92095%  
3%|           | 44/1433 [03:22<1:40:42,  4.35s/it]  
  
[Quota] LLM=0.468386 | GPU=0.0 | Used=0.468386/50.0 | Remaining=49.531614 |  
Usage=0.9367720000000002%  
3%|           | 45/1433 [03:27<1:44:50,  4.53s/it]  
  
[Quota] LLM=0.476066 | GPU=0.0 | Used=0.476066/50.0 | Remaining=49.523934 |  
Usage=0.952132%  
3%|           | 46/1433 [03:30<1:33:10,  4.03s/it]
```

```
[Quota] LLM=0.483731 | GPU=0.0 | Used=0.483731/50.0 | Remaining=49.516269 |
Usage=0.967462%
 3%|           | 47/1433 [03:34<1:35:32,  4.14s/it]

[Quota] LLM=0.49148000000000003 | GPU=0.0 | Used=0.49148000000000003/50.0 |
Remaining=49.50852 | Usage=0.9829600000000001%
 3%|           | 48/1433 [03:39<1:37:47,  4.24s/it]

[Quota] LLM=0.499439 | GPU=0.0 | Used=0.499439/50.0 | Remaining=49.500561 |
Usage=0.998878%
 3%|           | 49/1433 [03:43<1:38:16,  4.26s/it]

[Quota] LLM=0.50723 | GPU=0.0 | Used=0.50723/50.0 | Remaining=49.49277 |
Usage=1.01446%
 3%|           | 50/1433 [03:47<1:35:33,  4.15s/it]

[Quota] LLM=0.515018 | GPU=0.0 | Used=0.515018/50.0 | Remaining=49.484982 |
Usage=1.030036%
 4%|           | 51/1433 [03:51<1:37:02,  4.21s/it]

[Quota] LLM=0.52277 | GPU=0.0 | Used=0.52277/50.0 | Remaining=49.47723 |
Usage=1.04554%
 4%|           | 52/1433 [03:55<1:34:43,  4.12s/it]

[Quota] LLM=0.530231 | GPU=0.0 | Used=0.530231/50.0 | Remaining=49.469769 |
Usage=1.060462%
 4%|           | 53/1433 [03:58<1:29:48,  3.90s/it]

[Quota] LLM=0.537833 | GPU=0.0 | Used=0.537833/50.0 | Remaining=49.462167 |
Usage=1.075666%
 4%|           | 54/1433 [04:02<1:30:06,  3.92s/it]

[Quota] LLM=0.545615 | GPU=0.0 | Used=0.545615/50.0 | Remaining=49.454385 |
Usage=1.09123%
 4%|           | 55/1433 [04:07<1:33:26,  4.07s/it]

[Quota] LLM=0.553769 | GPU=0.0 | Used=0.553769/50.0 | Remaining=49.446231 |
Usage=1.107538%
 4%|           | 56/1433 [04:12<1:39:38,  4.34s/it]

[Quota] LLM=0.561602 | GPU=0.0 | Used=0.561602/50.0 | Remaining=49.438398 |
Usage=1.123204%
 4%|           | 57/1433 [04:16<1:37:00,  4.23s/it]

[Quota] LLM=0.569627 | GPU=0.0 | Used=0.569627/50.0 | Remaining=49.430373 |
Usage=1.139254%
 4%|           | 58/1433 [04:21<1:40:15,  4.37s/it]
```

```
[Quota] LLM=0.577775 | GPU=0.0 | Used=0.577775/50.0 | Remaining=49.422225 |
Usage=1.15555%
 4%|           | 59/1433 [04:24<1:31:23,  3.99s/it]

[Quota] LLM=0.585416 | GPU=0.0 | Used=0.585416/50.0 | Remaining=49.414584 |
Usage=1.170832%
 4%|           | 60/1433 [04:27<1:28:59,  3.89s/it]

[Quota] LLM=0.593333 | GPU=0.0 | Used=0.593333/50.0 | Remaining=49.406667 |
Usage=1.186666%
 4%|           | 61/1433 [04:31<1:31:13,  3.99s/it]

[Quota] LLM=0.601316 | GPU=0.0 | Used=0.601316/50.0 | Remaining=49.398684 |
Usage=1.202632%
 4%|           | 62/1433 [04:36<1:35:53,  4.20s/it]

[Quota] LLM=0.609059 | GPU=0.0 | Used=0.609059/50.0 | Remaining=49.390941 |
Usage=1.218118%
 4%|           | 63/1433 [04:40<1:35:17,  4.17s/it]

[Quota] LLM=0.616718 | GPU=0.0 | Used=0.616718/50.0 | Remaining=49.383282 |
Usage=1.233436%
 4%|           | 64/1433 [04:44<1:33:43,  4.11s/it]

[Quota] LLM=0.624515 | GPU=0.0 | Used=0.624515/50.0 | Remaining=49.375485 |
Usage=1.24903%
 5%|           | 65/1433 [04:48<1:29:38,  3.93s/it]

[Quota] LLM=0.632219 | GPU=0.0 | Used=0.632219/50.0 | Remaining=49.367781 |
Usage=1.264438%
 5%|           | 66/1433 [04:51<1:23:21,  3.66s/it]

[Quota] LLM=0.639782 | GPU=0.0 | Used=0.639782/50.0 | Remaining=49.360218 |
Usage=1.279564%
 5%|           | 67/1433 [04:56<1:30:34,  3.98s/it]

[Quota] LLM=0.647621 | GPU=0.0 | Used=0.647621/50.0 | Remaining=49.352379 |
Usage=1.295242%
 5%|           | 68/1433 [05:00<1:32:17,  4.06s/it]

[Quota] LLM=0.6554 | GPU=0.0 | Used=0.6554/50.0 | Remaining=49.3446 |
Usage=1.3108%
 5%|           | 69/1433 [05:04<1:33:57,  4.13s/it]

[Quota] LLM=0.663248 | GPU=0.0 | Used=0.663248/50.0 | Remaining=49.336752 |
Usage=1.326496%
 5%|           | 70/1433 [05:09<1:38:08,  4.32s/it]
```

```
[Quota] LLM=0.671018 | GPU=0.0 | Used=0.671018/50.0 | Remaining=49.328982 |
Usage=1.342036%
 5%|           | 71/1433 [05:12<1:30:36,  3.99s/it]

[Quota] LLM=0.678713 | GPU=0.0 | Used=0.678713/50.0 | Remaining=49.321287 |
Usage=1.357426%
 5%|           | 72/1433 [05:15<1:25:51,  3.79s/it]

[Quota] LLM=0.68654 | GPU=0.0 | Used=0.68654/50.0 | Remaining=49.31346 |
Usage=1.37308%
 5%|           | 73/1433 [05:20<1:30:49,  4.01s/it]

[Quota] LLM=0.694358 | GPU=0.0 | Used=0.694358/50.0 | Remaining=49.305642 |
Usage=1.388716%
 5%|           | 74/1433 [05:24<1:32:11,  4.07s/it]

[Quota] LLM=0.70214 | GPU=0.0 | Used=0.70214/50.0 | Remaining=49.29786 |
Usage=1.40428%
 5%|           | 75/1433 [05:28<1:28:18,  3.90s/it]

[Quota] LLM=0.710054 | GPU=0.0 | Used=0.710054/50.0 | Remaining=49.289946 |
Usage=1.420108%
 5%|           | 76/1433 [05:32<1:28:11,  3.90s/it]

[Quota] LLM=0.717686 | GPU=0.0 | Used=0.717686/50.0 | Remaining=49.282314 |
Usage=1.435372%
 5%|           | 77/1433 [05:35<1:26:10,  3.81s/it]

[Quota] LLM=0.725714 | GPU=0.0 | Used=0.725714/50.0 |
Remaining=49.274286000000004 | Usage=1.451428%
 5%|           | 78/1433 [05:40<1:32:39,  4.10s/it]

[Quota] LLM=0.73385 | GPU=0.0 | Used=0.73385/50.0 | Remaining=49.26615 |
Usage=1.4677%
 6%|           | 79/1433 [05:44<1:35:45,  4.24s/it]

[Quota] LLM=0.741851 | GPU=0.0 | Used=0.741851/50.0 | Remaining=49.258149 |
Usage=1.483702%
 6%|           | 80/1433 [05:48<1:31:26,  4.06s/it]

[Quota] LLM=0.749702 | GPU=0.0 | Used=0.749702/50.0 | Remaining=49.250298 |
Usage=1.499404%
 6%|           | 81/1433 [05:52<1:28:35,  3.93s/it]

[Quota] LLM=0.757514 | GPU=0.0 | Used=0.757514/50.0 | Remaining=49.242486 |
Usage=1.515028%
 6%|           | 82/1433 [05:57<1:40:10,  4.45s/it]
```

```
[Quota] LLM=0.765338 | GPU=0.0 | Used=0.765338/50.0 | Remaining=49.234662 |  
Usage=1.530676%  
6%|           | 83/1433 [06:02<1:38:58,  4.40s/it]  
  
[Quota] LLM=0.773168 | GPU=0.0 | Used=0.773168/50.0 | Remaining=49.226832 |  
Usage=1.546336%  
6%|           | 84/1433 [06:08<1:49:27,  4.87s/it]  
  
[Quota] LLM=0.781529 | GPU=0.0 | Used=0.781529/50.0 | Remaining=49.218471 |  
Usage=1.563058%  
6%|           | 85/1433 [06:12<1:45:58,  4.72s/it]  
  
[Quota] LLM=0.789539 | GPU=0.0 | Used=0.789539/50.0 | Remaining=49.210461 |  
Usage=1.579078%  
6%|           | 86/1433 [06:18<1:54:07,  5.08s/it]  
  
[Quota] LLM=0.797879 | GPU=0.0 | Used=0.797879/50.0 | Remaining=49.202121 |  
Usage=1.595758%  
6%|           | 87/1433 [06:22<1:45:38,  4.71s/it]  
  
[Quota] LLM=0.806099 | GPU=0.0 | Used=0.806099/50.0 | Remaining=49.193901 |  
Usage=1.612198%  
6%|           | 88/1433 [06:26<1:44:07,  4.64s/it]  
  
[Quota] LLM=0.813887 | GPU=0.0 | Used=0.813887/50.0 | Remaining=49.186113 |  
Usage=1.627773999999998%  
6%|           | 89/1433 [06:31<1:46:49,  4.77s/it]  
  
[Quota] LLM=0.821807 | GPU=0.0 | Used=0.821807/50.0 | Remaining=49.178193 |  
Usage=1.643614%  
6%|           | 90/1433 [06:36<1:43:02,  4.60s/it]  
  
[Quota] LLM=0.829958 | GPU=0.0 | Used=0.829958/50.0 | Remaining=49.170042 |  
Usage=1.659915999999997%  
6%|           | 91/1433 [06:40<1:38:45,  4.42s/it]  
  
[Quota] LLM=0.837842 | GPU=0.0 | Used=0.837842/50.0 | Remaining=49.162158 |  
Usage=1.675684%  
6%|           | 92/1433 [06:44<1:37:56,  4.38s/it]  
  
[Quota] LLM=0.845894 | GPU=0.0 | Used=0.845894/50.0 | Remaining=49.154106 |  
Usage=1.691788%  
6%|           | 93/1433 [06:50<1:47:40,  4.82s/it]  
  
[Quota] LLM=0.854165 | GPU=0.0 | Used=0.854165/50.0 | Remaining=49.145835 |  
Usage=1.70833%  
7%|           | 94/1433 [06:54<1:41:01,  4.53s/it]
```

```
[Quota] LLM=0.862091 | GPU=0.0 | Used=0.862091/50.0 | Remaining=49.137909 |  
Usage=1.724182%  
 7%|           | 95/1433 [06:57<1:32:12,  4.14s/it]  
  
[Quota] LLM=0.870029 | GPU=0.0 | Used=0.870029/50.0 | Remaining=49.129971 |  
Usage=1.7400580000000003%  
 7%|           | 96/1433 [07:01<1:29:53,  4.03s/it]  
  
[Quota] LLM=0.877658 | GPU=0.0 | Used=0.877658/50.0 | Remaining=49.122342 |  
Usage=1.755316%  
 7%|           | 97/1433 [07:05<1:31:21,  4.10s/it]  
  
[Quota] LLM=0.885923 | GPU=0.0 | Used=0.885923/50.0 | Remaining=49.114077 |  
Usage=1.7718460000000003%  
 7%|           | 98/1433 [07:08<1:24:18,  3.79s/it]  
  
[Quota] LLM=0.893417 | GPU=0.0 | Used=0.893417/50.0 | Remaining=49.106583 |  
Usage=1.786834%  
 7%|           | 99/1433 [07:12<1:29:40,  4.03s/it]  
  
[Quota] LLM=0.901421 | GPU=0.0 | Used=0.901421/50.0 | Remaining=49.098579 |  
Usage=1.802842%  
 7%|           | 100/1433 [07:16<1:27:00,  3.92s/it]  
  
[Quota] LLM=0.909098 | GPU=0.0 | Used=0.909098/50.0 | Remaining=49.090902 |  
Usage=1.8181960000000001%  
 7%|           | 101/1433 [07:20<1:27:02,  3.92s/it]  
  
[Quota] LLM=0.916904 | GPU=0.0 | Used=0.916904/50.0 | Remaining=49.083096 |  
Usage=1.833807999999999%  
 7%|           | 102/1433 [07:24<1:28:49,  4.00s/it]  
  
[Quota] LLM=0.924659 | GPU=0.0 | Used=0.924659/50.0 | Remaining=49.075341 |  
Usage=1.8493180000000002%  
 7%|           | 103/1433 [07:29<1:34:44,  4.27s/it]  
  
[Quota] LLM=0.932567 | GPU=0.0 | Used=0.932567/50.0 | Remaining=49.067433 |  
Usage=1.8651340000000003%  
 7%|           | 104/1433 [07:34<1:36:51,  4.37s/it]  
  
[Quota] LLM=0.940655 | GPU=0.0 | Used=0.940655/50.0 | Remaining=49.059345 |  
Usage=1.88131%  
 7%|           | 105/1433 [07:40<1:49:28,  4.95s/it]  
  
[Quota] LLM=0.948626 | GPU=0.0 | Used=0.948626/50.0 | Remaining=49.051374 |  
Usage=1.897252%  
 7%|           | 106/1433 [07:45<1:47:32,  4.86s/it]
```

```
[Quota] LLM=0.956447 | GPU=0.0 | Used=0.956447/50.0 | Remaining=49.043553 |  
Usage=1.912894%  
 7%|           | 107/1433 [07:49<1:40:38,  4.55s/it]  
  
[Quota] LLM=0.964478 | GPU=0.0 | Used=0.964478/50.0 | Remaining=49.035522 |  
Usage=1.9289559999999997%  
 8%|           | 108/1433 [07:53<1:40:04,  4.53s/it]  
  
[Quota] LLM=0.972287 | GPU=0.0 | Used=0.972287/50.0 | Remaining=49.027713 |  
Usage=1.944574%  
 8%|           | 109/1433 [07:59<1:46:44,  4.84s/it]  
  
[Quota] LLM=0.980396 | GPU=0.0 | Used=0.980396/50.0 | Remaining=49.019604 |  
Usage=1.960792%  
 8%|           | 110/1433 [08:04<1:50:23,  5.01s/it]  
  
[Quota] LLM=0.988319 | GPU=0.0 | Used=0.988319/50.0 | Remaining=49.011681 |  
Usage=1.976638%  
 8%|           | 111/1433 [08:10<1:55:50,  5.26s/it]  
  
[Quota] LLM=0.996629 | GPU=0.0 | Used=0.996629/50.0 | Remaining=49.003371 |  
Usage=1.9932579999999998%  
 8%|           | 112/1433 [08:14<1:49:25,  4.97s/it]  
  
[Quota] LLM=1.00466 | GPU=0.0 | Used=1.00466/50.0 | Remaining=48.99534 |  
Usage=2.00932%  
 8%|           | 113/1433 [08:19<1:48:37,  4.94s/it]  
  
[Quota] LLM=1.01264 | GPU=0.0 | Used=1.01264/50.0 | Remaining=48.98736 |  
Usage=2.02528%  
 8%|           | 114/1433 [08:24<1:49:43,  4.99s/it]  
  
[Quota] LLM=1.020485 | GPU=0.0 | Used=1.020485/50.0 | Remaining=48.979515 |  
Usage=2.04097%  
 8%|           | 115/1433 [08:29<1:47:12,  4.88s/it]  
  
[Quota] LLM=1.028426 | GPU=0.0 | Used=1.028426/50.0 | Remaining=48.971574 |  
Usage=2.056852%  
 8%|           | 116/1433 [08:34<1:51:12,  5.07s/it]  
  
[Quota] LLM=1.036445 | GPU=0.0 | Used=1.036445/50.0 | Remaining=48.963555 |  
Usage=2.07289%  
 8%|           | 117/1433 [08:39<1:50:42,  5.05s/it]  
  
[Quota] LLM=1.04501 | GPU=0.0 | Used=1.04501/50.0 | Remaining=48.95499 |  
Usage=2.09002%  
 8%|           | 118/1433 [08:45<1:53:53,  5.20s/it]
```

```
[Quota] LLM=1.053155 | GPU=0.0 | Used=1.053155/50.0 | Remaining=48.946845 |
Usage=2.10631%
 8%|           | 119/1433 [08:49<1:47:38,  4.92s/it]

[Quota] LLM=1.060826 | GPU=0.0 | Used=1.060826/50.0 | Remaining=48.939174 |
Usage=2.121652%
 8%|           | 120/1433 [08:54<1:45:23,  4.82s/it]

[Quota] LLM=1.068797 | GPU=0.0 | Used=1.068797/50.0 | Remaining=48.931203 |
Usage=2.137594%
 8%|           | 121/1433 [08:57<1:35:24,  4.36s/it]

[Quota] LLM=1.076471 | GPU=0.0 | Used=1.076471/50.0 | Remaining=48.923529 |
Usage=2.152942%
 9%|           | 122/1433 [09:07<2:14:39,  6.16s/it]

[Quota] LLM=1.084658 | GPU=0.0 | Used=1.084658/50.0 | Remaining=48.915342 |
Usage=2.169316%
 9%|           | 123/1433 [09:11<2:01:25,  5.56s/it]

[Quota] LLM=1.092599 | GPU=0.0 | Used=1.092599/50.0 | Remaining=48.907401 |
Usage=2.185198%
 9%|           | 124/1433 [09:15<1:50:26,  5.06s/it]

[Quota] LLM=1.100516 | GPU=0.0 | Used=1.100516/50.0 | Remaining=48.899484 |
Usage=2.201032%
 9%|           | 125/1433 [09:20<1:49:05,  5.00s/it]

[Quota] LLM=1.108652 | GPU=0.0 | Used=1.108652/50.0 | Remaining=48.891348 |
Usage=2.217304%
 9%|           | 126/1433 [09:23<1:38:00,  4.50s/it]

[Quota] LLM=1.116629 | GPU=0.0 | Used=1.116629/50.0 | Remaining=48.883371 |
Usage=2.233258%
 9%|           | 127/1433 [09:28<1:38:02,  4.50s/it]

[Quota] LLM=1.1246 | GPU=0.0 | Used=1.1246/50.0 | Remaining=48.8754 |
Usage=2.2492%
 9%|           | 128/1433 [09:32<1:33:19,  4.29s/it]

[Quota] LLM=1.132499 | GPU=0.0 | Used=1.132499/50.0 | Remaining=48.867501 |
Usage=2.264998%
 9%|           | 129/1433 [09:38<1:43:17,  4.75s/it]

[Quota] LLM=1.1406020000000001 | GPU=0.0 | Used=1.1406020000000001/50.0 |
Remaining=48.859398 | Usage=2.2812040000000002%
 9%|           | 130/1433 [09:42<1:40:07,  4.61s/it]
```

```
[Quota] LLM=1.148645 | GPU=0.0 | Used=1.148645/50.0 | Remaining=48.851355 |
Usage=2.29729%
 9%|           | 131/1433 [09:46<1:36:29,  4.45s/it]

[Quota] LLM=1.156679 | GPU=0.0 | Used=1.156679/50.0 | Remaining=48.843321 |
Usage=2.313358%
 9%|           | 132/1433 [09:50<1:34:56,  4.38s/it]

[Quota] LLM=1.164896 | GPU=0.0 | Used=1.164896/50.0 | Remaining=48.835104 |
Usage=2.329792%
 9%|           | 133/1433 [09:54<1:34:01,  4.34s/it]

[Quota] LLM=1.172786 | GPU=0.0 | Used=1.172786/50.0 | Remaining=48.827214 |
Usage=2.345572%
 9%|           | 134/1433 [09:59<1:38:21,  4.54s/it]

[Quota] LLM=1.18097 | GPU=0.0 | Used=1.18097/50.0 | Remaining=48.81903 |
Usage=2.36194%
 9%|           | 135/1433 [10:04<1:39:29,  4.60s/it]

[Quota] LLM=1.18892 | GPU=0.0 | Used=1.18892/50.0 | Remaining=48.81108 |
Usage=2.37784%
 9%|           | 136/1433 [10:09<1:40:52,  4.67s/it]

[Quota] LLM=1.19702 | GPU=0.0 | Used=1.19702/50.0 | Remaining=48.80298 |
Usage=2.39404%
10%|           | 137/1433 [10:13<1:37:50,  4.53s/it]

[Quota] LLM=1.205048 | GPU=0.0 | Used=1.205048/50.0 | Remaining=48.794952 |
Usage=2.410096%
10%|           | 138/1433 [10:18<1:38:12,  4.55s/it]

[Quota] LLM=1.2131 | GPU=0.0 | Used=1.2131/50.0 | Remaining=48.7869 |
Usage=2.4262%
10%|           | 139/1433 [10:22<1:38:28,  4.57s/it]

[Quota] LLM=1.220912 | GPU=0.0 | Used=1.220912/50.0 | Remaining=48.779088 |
Usage=2.441824%
10%|           | 140/1433 [10:26<1:35:02,  4.41s/it]

[Quota] LLM=1.228913 | GPU=0.0 | Used=1.228913/50.0 | Remaining=48.771087 |
Usage=2.457826%
10%|           | 141/1433 [10:31<1:34:48,  4.40s/it]

[Quota] LLM=1.236974 | GPU=0.0 | Used=1.236974/50.0 | Remaining=48.763025999999996 |
Usage=2.473948%
10%|           | 142/1433 [10:36<1:38:03,  4.56s/it]
```

```
[Quota] LLM=1.245092 | GPU=0.0 | Used=1.245092/50.0 | Remaining=48.754908 |
Usage=2.490184%
 10%|           | 143/1433 [10:41<1:41:50,  4.74s/it]

[Quota] LLM=1.253249 | GPU=0.0 | Used=1.253249/50.0 | Remaining=48.746751 |
Usage=2.506498%
 10%|           | 144/1433 [10:46<1:46:58,  4.98s/it]

[Quota] LLM=1.261328 | GPU=0.0 | Used=1.261328/50.0 | Remaining=48.738672 |
Usage=2.522656%
 10%|           | 145/1433 [10:51<1:45:23,  4.91s/it]

[Quota] LLM=1.26965 | GPU=0.0 | Used=1.26965/50.0 | Remaining=48.73035 |
Usage=2.5393%
 10%|           | 146/1433 [10:56<1:42:08,  4.76s/it]

[Quota] LLM=1.277975 | GPU=0.0 | Used=1.277975/50.0 | Remaining=48.722025 |
Usage=2.55595%
 10%|           | 147/1433 [11:00<1:39:19,  4.63s/it]

[Quota] LLM=1.286156 | GPU=0.0 | Used=1.286156/50.0 | Remaining=48.713844 |
Usage=2.572312%
 10%|           | 148/1433 [11:06<1:44:58,  4.90s/it]

[Quota] LLM=1.294268 | GPU=0.0 | Used=1.294268/50.0 | Remaining=48.705732 |
Usage=2.588536%
 10%|           | 149/1433 [11:11<1:45:56,  4.95s/it]

[Quota] LLM=1.30226 | GPU=0.0 | Used=1.30226/50.0 | Remaining=48.69774 |
Usage=2.60452%
 10%|           | 150/1433 [11:14<1:36:34,  4.52s/it]

[Quota] LLM=1.309967 | GPU=0.0 | Used=1.309967/50.0 | Remaining=48.690033 |
Usage=2.619934%
 11%|           | 151/1433 [11:20<1:43:22,  4.84s/it]

[Quota] LLM=1.318079 | GPU=0.0 | Used=1.318079/50.0 | Remaining=48.681921 |
Usage=2.636158%
 11%|           | 152/1433 [11:25<1:46:51,  5.01s/it]

[Quota] LLM=1.326053 | GPU=0.0 | Used=1.326053/50.0 | Remaining=48.673947 |
Usage=2.652106%
 11%|           | 153/1433 [11:31<1:55:58,  5.44s/it]

[Quota] LLM=1.334144 | GPU=0.0 | Used=1.334144/50.0 | Remaining=48.665856 |
Usage=2.668288%
 11%|           | 154/1433 [11:38<2:00:59,  5.68s/it]
```

```
[Quota] LLM=1.342682 | GPU=0.0 | Used=1.342682/50.0 |  
Remaining=48.657318000000004 | Usage=2.685364%  
11% | 155/1433 [11:42<1:50:43, 5.20s/it]  
  
[Quota] LLM=1.350482 | GPU=0.0 | Used=1.350482/50.0 | Remaining=48.649518 |  
Usage=2.700964%  
11% | 156/1433 [11:49<2:00:13, 5.65s/it]  
  
[Quota] LLM=1.358477 | GPU=0.0 | Used=1.358477/50.0 | Remaining=48.641523 |  
Usage=2.716954%  
11% | 157/1433 [11:53<1:50:47, 5.21s/it]  
  
[Quota] LLM=1.366481 | GPU=0.0 | Used=1.366481/50.0 | Remaining=48.633519 |  
Usage=2.732962%  
11% | 158/1433 [11:57<1:43:48, 4.89s/it]  
  
[Quota] LLM=1.374257 | GPU=0.0 | Used=1.374257/50.0 | Remaining=48.625743 |  
Usage=2.748514%  
11% | 159/1433 [12:02<1:45:49, 4.98s/it]  
  
[Quota] LLM=1.38227 | GPU=0.0 | Used=1.38227/50.0 | Remaining=48.61773 |  
Usage=2.76454%  
11% | 160/1433 [12:07<1:44:22, 4.92s/it]  
  
[Quota] LLM=1.39037 | GPU=0.0 | Used=1.39037/50.0 | Remaining=48.60963 |  
Usage=2.78074%  
11% | 161/1433 [12:11<1:38:42, 4.66s/it]  
  
[Quota] LLM=1.398455 | GPU=0.0 | Used=1.398455/50.0 | Remaining=48.601545 |  
Usage=2.79691%  
11% | 162/1433 [12:15<1:38:30, 4.65s/it]  
  
[Quota] LLM=1.406348 | GPU=0.0 | Used=1.406348/50.0 | Remaining=48.593652 |  
Usage=2.812696%  
11% | 163/1433 [12:20<1:34:38, 4.47s/it]  
  
[Quota] LLM=1.414073 | GPU=0.0 | Used=1.414073/50.0 | Remaining=48.585927 |  
Usage=2.828146%  
11% | 164/1433 [12:24<1:35:36, 4.52s/it]  
  
[Quota] LLM=1.421984 | GPU=0.0 | Used=1.421984/50.0 | Remaining=48.578016 |  
Usage=2.843968%  
12% | 165/1433 [12:29<1:39:02, 4.69s/it]  
  
[Quota] LLM=1.4298710000000001 | GPU=0.0 | Used=1.4298710000000001/50.0 |  
Remaining=48.570129 | Usage=2.8597420000000002%  
12% | 166/1433 [12:34<1:37:11, 4.60s/it]
```

```
[Quota] LLM=1.437539 | GPU=0.0 | Used=1.437539/50.0 | Remaining=48.562461 |
Usage=2.875078%
 12%|           | 167/1433 [12:38<1:35:08,  4.51s/it]

[Quota] LLM=1.445576 | GPU=0.0 | Used=1.445576/50.0 | Remaining=48.554424 |
Usage=2.891152%
 12%|           | 168/1433 [12:42<1:32:59,  4.41s/it]

[Quota] LLM=1.453559 | GPU=0.0 | Used=1.453559/50.0 | Remaining=48.546441 |
Usage=2.907118%
 12%|           | 169/1433 [12:47<1:34:33,  4.49s/it]

[Quota] LLM=1.461563 | GPU=0.0 | Used=1.461563/50.0 | Remaining=48.538437 |
Usage=2.923126%
 12%|           | 170/1433 [12:51<1:30:30,  4.30s/it]

[Quota] LLM=1.46948 | GPU=0.0 | Used=1.46948/50.0 | Remaining=48.53052 |
Usage=2.93896%
 12%|           | 171/1433 [12:56<1:34:44,  4.50s/it]

[Quota] LLM=1.47752 | GPU=0.0 | Used=1.47752/50.0 | Remaining=48.52248 |
Usage=2.95504%
 12%|           | 172/1433 [12:59<1:28:15,  4.20s/it]

[Quota] LLM=1.485242 | GPU=0.0 | Used=1.485242/50.0 | Remaining=48.514758 |
Usage=2.970484%
 12%|           | 173/1433 [13:04<1:33:45,  4.46s/it]

[Quota] LLM=1.493129 | GPU=0.0 | Used=1.493129/50.0 | Remaining=48.506871 |
Usage=2.986258%
 12%|           | 174/1433 [13:07<1:25:46,  4.09s/it]

[Quota] LLM=1.500875 | GPU=0.0 | Used=1.500875/50.0 | Remaining=48.499125 |
Usage=3.00175%
 12%|           | 175/1433 [13:13<1:32:59,  4.44s/it]

[Quota] LLM=1.508861 | GPU=0.0 | Used=1.508861/50.0 | Remaining=48.491139 |
Usage=3.017722%
 12%|           | 176/1433 [13:18<1:38:45,  4.71s/it]

[Quota] LLM=1.516475 | GPU=0.0 | Used=1.516475/50.0 | Remaining=48.483525 |
Usage=3.03295%
 12%|           | 177/1433 [13:21<1:30:09,  4.31s/it]

[Quota] LLM=1.524191 | GPU=0.0 | Used=1.524191/50.0 | Remaining=48.475809 |
Usage=3.048382%
 12%|           | 178/1433 [13:28<1:45:21,  5.04s/it]
```

```
[Quota] LLM=1.532042 | GPU=0.0 | Used=1.532042/50.0 | Remaining=48.467958 |
Usage=3.064084%
12%|           | 179/1433 [13:32<1:34:56,  4.54s/it]

[Quota] LLM=1.539716 | GPU=0.0 | Used=1.539716/50.0 | Remaining=48.460284 |
Usage=3.079432%
13%|           | 180/1433 [13:37<1:43:41,  4.97s/it]

[Quota] LLM=1.54772 | GPU=0.0 | Used=1.54772/50.0 | Remaining=48.45228 |
Usage=3.09544%
13%|           | 181/1433 [13:43<1:47:10,  5.14s/it]

[Quota] LLM=1.555571 | GPU=0.0 | Used=1.555571/50.0 | Remaining=48.444429 |
Usage=3.111142%
13%|           | 182/1433 [13:47<1:36:54,  4.65s/it]

[Quota] LLM=1.563608 | GPU=0.0 | Used=1.563608/50.0 | Remaining=48.436392 |
Usage=3.127216%
13%|           | 183/1433 [13:51<1:33:07,  4.47s/it]

[Quota] LLM=1.571348 | GPU=0.0 | Used=1.571348/50.0 | Remaining=48.428652 |
Usage=3.142695999999995%
13%|           | 184/1433 [13:55<1:35:15,  4.58s/it]

[Quota] LLM=1.579583 | GPU=0.0 | Used=1.579583/50.0 | Remaining=48.420417 |
Usage=3.159166%
13%|           | 185/1433 [13:59<1:28:03,  4.23s/it]

[Quota] LLM=1.587302 | GPU=0.0 | Used=1.587302/50.0 | Remaining=48.412698 |
Usage=3.1746040000000004%
13%|           | 186/1433 [14:04<1:33:27,  4.50s/it]

[Quota] LLM=1.595453 | GPU=0.0 | Used=1.595453/50.0 | Remaining=48.404547 |
Usage=3.1909060000000005%
13%|           | 187/1433 [14:08<1:33:36,  4.51s/it]

[Quota] LLM=1.603379 | GPU=0.0 | Used=1.603379/50.0 | Remaining=48.396621 |
Usage=3.2067580000000007%
13%|           | 188/1433 [14:13<1:31:44,  4.42s/it]

[Quota] LLM=1.611632 | GPU=0.0 | Used=1.611632/50.0 | Remaining=48.388368 |
Usage=3.223264%
13%|           | 189/1433 [14:18<1:35:41,  4.62s/it]

[Quota] LLM=1.619534 | GPU=0.0 | Used=1.619534/50.0 | Remaining=48.380466 |
Usage=3.239067999999996%
13%|           | 190/1433 [14:22<1:34:29,  4.56s/it]
```

```
[Quota] LLM=1.627433 | GPU=0.0 | Used=1.627433/50.0 |  
Remaining=48.372567000000004 | Usage=3.254866%  
13%|           | 191/1433 [14:26<1:31:51,  4.44s/it]  
  
[Quota] LLM=1.635257 | GPU=0.0 | Used=1.635257/50.0 | Remaining=48.364743 |  
Usage=3.270514%  
13%|           | 192/1433 [14:30<1:29:33,  4.33s/it]  
  
[Quota] LLM=1.643033 | GPU=0.0 | Used=1.643033/50.0 | Remaining=48.356967 |  
Usage=3.286066%  
13%|           | 193/1433 [14:36<1:37:32,  4.72s/it]  
  
[Quota] LLM=1.650857 | GPU=0.0 | Used=1.650857/50.0 | Remaining=48.349143 |  
Usage=3.301714%  
14%|           | 194/1433 [14:40<1:35:02,  4.60s/it]  
  
[Quota] LLM=1.658597 | GPU=0.0 | Used=1.658597/50.0 | Remaining=48.341403 |  
Usage=3.317194%  
14%|           | 195/1433 [14:45<1:34:07,  4.56s/it]  
  
[Quota] LLM=1.666499 | GPU=0.0 | Used=1.666499/50.0 | Remaining=48.333501 |  
Usage=3.3329980000000003%  
14%|           | 196/1433 [14:49<1:32:36,  4.49s/it]  
  
[Quota] LLM=1.674434 | GPU=0.0 | Used=1.674434/50.0 | Remaining=48.325566 |  
Usage=3.348868%  
14%|           | 197/1433 [14:54<1:34:56,  4.61s/it]  
  
[Quota] LLM=1.682525 | GPU=0.0 | Used=1.682525/50.0 | Remaining=48.317475 |  
Usage=3.36505%  
14%|           | 198/1433 [14:59<1:37:03,  4.72s/it]  
  
[Quota] LLM=1.690913 | GPU=0.0 | Used=1.690913/50.0 | Remaining=48.309087 |  
Usage=3.381826%  
14%|           | 199/1433 [15:03<1:32:11,  4.48s/it]  
  
[Quota] LLM=1.698671 | GPU=0.0 | Used=1.698671/50.0 | Remaining=48.301329 |  
Usage=3.3973420000000005%  
14%|           | 200/1433 [15:10<1:45:08,  5.12s/it]  
  
[Quota] LLM=1.706501 | GPU=0.0 | Used=1.706501/50.0 | Remaining=48.293499 |  
Usage=3.4130020000000005%  
14%|           | 201/1433 [15:16<1:50:16,  5.37s/it]  
  
[Quota] LLM=1.714328 | GPU=0.0 | Used=1.714328/50.0 | Remaining=48.285672 |  
Usage=3.428656%  
14%|           | 202/1433 [15:20<1:45:12,  5.13s/it]
```

```
[Quota] LLM=1.722461 | GPU=0.0 | Used=1.722461/50.0 | Remaining=48.277539 |
Usage=3.444922%
14%|           | 203/1433 [15:24<1:39:10,  4.84s/it]

[Quota] LLM=1.7303 | GPU=0.0 | Used=1.7303/50.0 | Remaining=48.2697 |
Usage=3.4606%
14%|           | 204/1433 [15:27<1:28:01,  4.30s/it]

[Quota] LLM=1.738073 | GPU=0.0 | Used=1.738073/50.0 | Remaining=48.261927 |
Usage=3.476146%
14%|           | 205/1433 [15:31<1:25:40,  4.19s/it]

[Quota] LLM=1.745828 | GPU=0.0 | Used=1.745828/50.0 | Remaining=48.254172 |
Usage=3.491656%
14%|           | 206/1433 [15:36<1:27:41,  4.29s/it]

[Quota] LLM=1.753667 | GPU=0.0 | Used=1.753667/50.0 | Remaining=48.246333 |
Usage=3.507334%
14%|           | 207/1433 [15:39<1:23:20,  4.08s/it]

[Quota] LLM=1.761632 | GPU=0.0 | Used=1.761632/50.0 | Remaining=48.238368 |
Usage=3.523264%
15%|           | 208/1433 [15:45<1:30:19,  4.42s/it]

[Quota] LLM=1.769645 | GPU=0.0 | Used=1.769645/50.0 | Remaining=48.230355 |
Usage=3.53929%
15%|           | 209/1433 [15:49<1:28:40,  4.35s/it]

[Quota] LLM=1.777544 | GPU=0.0 | Used=1.777544/50.0 | Remaining=48.222456 |
Usage=3.555088%
15%|           | 210/1433 [15:53<1:26:25,  4.24s/it]

[Quota] LLM=1.785197 | GPU=0.0 | Used=1.785197/50.0 | Remaining=48.214803 |
Usage=3.570394%
15%|           | 211/1433 [15:57<1:26:41,  4.26s/it]

[Quota] LLM=1.793144 | GPU=0.0 | Used=1.793144/50.0 | Remaining=48.206856 |
Usage=3.586288%
15%|           | 212/1433 [16:02<1:32:31,  4.55s/it]

[Quota] LLM=1.801223 | GPU=0.0 | Used=1.801223/50.0 | Remaining=48.198777 |
Usage=3.602446%
15%|           | 213/1433 [16:06<1:30:09,  4.43s/it]

[Quota] LLM=1.809215 | GPU=0.0 | Used=1.809215/50.0 | Remaining=48.190785 |
Usage=3.61843%
15%|           | 214/1433 [16:11<1:30:09,  4.44s/it]
```

```
[Quota] LLM=1.817126 | GPU=0.0 | Used=1.817126/50.0 | Remaining=48.182874 |  
Usage=3.6342520000000005%  
15%|           | 215/1433 [16:15<1:30:41,  4.47s/it]  
  
[Quota] LLM=1.824965 | GPU=0.0 | Used=1.824965/50.0 | Remaining=48.175035 |  
Usage=3.64993%  
15%|           | 216/1433 [16:19<1:27:31,  4.32s/it]  
  
[Quota] LLM=1.83266 | GPU=0.0 | Used=1.83266/50.0 | Remaining=48.16734 |  
Usage=3.6653199999999995%  
15%|           | 217/1433 [16:24<1:30:15,  4.45s/it]  
  
[Quota] LLM=1.840574 | GPU=0.0 | Used=1.840574/50.0 | Remaining=48.159426 |  
Usage=3.681148%  
15%|           | 218/1433 [16:29<1:32:51,  4.59s/it]  
  
[Quota] LLM=1.848386 | GPU=0.0 | Used=1.848386/50.0 | Remaining=48.151614 |  
Usage=3.696772%  
15%|           | 219/1433 [16:34<1:36:08,  4.75s/it]  
  
[Quota] LLM=1.856948 | GPU=0.0 | Used=1.856948/50.0 | Remaining=48.143052 |  
Usage=3.713896%  
15%|           | 220/1433 [16:39<1:34:27,  4.67s/it]  
  
[Quota] LLM=1.864856 | GPU=0.0 | Used=1.864856/50.0 | Remaining=48.135144 |  
Usage=3.729712%  
15%|           | 221/1433 [16:45<1:42:27,  5.07s/it]  
  
[Quota] LLM=1.872983 | GPU=0.0 | Used=1.872983/50.0 | Remaining=48.127017 |  
Usage=3.745965999999997%  
15%|           | 222/1433 [16:51<1:47:30,  5.33s/it]  
  
[Quota] LLM=1.880747 | GPU=0.0 | Used=1.880747/50.0 | Remaining=48.119253 |  
Usage=3.761494%  
16%|           | 223/1433 [16:55<1:39:02,  4.91s/it]  
  
[Quota] LLM=1.888424 | GPU=0.0 | Used=1.888424/50.0 | Remaining=48.111576 |  
Usage=3.776848%  
16%|           | 224/1433 [16:59<1:37:33,  4.84s/it]  
  
[Quota] LLM=1.896875 | GPU=0.0 | Used=1.896875/50.0 | Remaining=48.103125 |  
Usage=3.793749999999997%  
16%|           | 225/1433 [17:04<1:34:37,  4.70s/it]  
  
[Quota] LLM=1.904795 | GPU=0.0 | Used=1.904795/50.0 | Remaining=48.095205 |  
Usage=3.80959%  
16%|           | 226/1433 [17:07<1:26:34,  4.30s/it]
```

```
[Quota] LLM=1.91261 | GPU=0.0 | Used=1.91261/50.0 | Remaining=48.08739 |  
Usage=3.82522%  
16%|           | 227/1433 [17:11<1:28:00,  4.38s/it]  
  
[Quota] LLM=1.9204430000000001 | GPU=0.0 | Used=1.9204430000000001/50.0 |  
Remaining=48.079557 | Usage=3.8408860000000002%  
16%|           | 228/1433 [17:15<1:25:35,  4.26s/it]  
  
[Quota] LLM=1.92806 | GPU=0.0 | Used=1.92806/50.0 | Remaining=48.07194 |  
Usage=3.85612%  
16%|           | 229/1433 [17:20<1:29:28,  4.46s/it]  
  
[Quota] LLM=1.935806 | GPU=0.0 | Used=1.935806/50.0 | Remaining=48.064194 |  
Usage=3.871612%  
16%|           | 230/1433 [17:24<1:27:11,  4.35s/it]  
  
[Quota] LLM=1.943708 | GPU=0.0 | Used=1.943708/50.0 | Remaining=48.056292 |  
Usage=3.887416%  
16%|           | 231/1433 [17:29<1:27:33,  4.37s/it]  
  
[Quota] LLM=1.951382 | GPU=0.0 | Used=1.951382/50.0 | Remaining=48.048618 |  
Usage=3.9027640000000003%  
16%|           | 232/1433 [17:33<1:23:45,  4.18s/it]  
  
[Quota] LLM=1.959092 | GPU=0.0 | Used=1.959092/50.0 | Remaining=48.040908 |  
Usage=3.918184%  
16%|           | 233/1433 [17:38<1:29:56,  4.50s/it]  
  
[Quota] LLM=1.966997 | GPU=0.0 | Used=1.966997/50.0 | Remaining=48.033003 |  
Usage=3.933994%  
16%|           | 234/1433 [17:42<1:25:15,  4.27s/it]  
  
[Quota] LLM=1.974821 | GPU=0.0 | Used=1.974821/50.0 | Remaining=48.025179 |  
Usage=3.949642%  
16%|           | 235/1433 [17:47<1:34:13,  4.72s/it]  
  
[Quota] LLM=1.982573 | GPU=0.0 | Used=1.982573/50.0 | Remaining=48.017427 |  
Usage=3.965146%  
16%|           | 236/1433 [17:51<1:27:14,  4.37s/it]  
  
[Quota] LLM=1.990067 | GPU=0.0 | Used=1.990067/50.0 | Remaining=48.009933 |  
Usage=3.980133999999996%  
17%|           | 237/1433 [17:54<1:20:10,  4.02s/it]  
  
[Quota] LLM=1.997888 | GPU=0.0 | Used=1.997888/50.0 | Remaining=48.002112 |  
Usage=3.995776%  
17%|           | 238/1433 [17:59<1:26:08,  4.32s/it]
```

```
[Quota] LLM=2.006174 | GPU=0.0 | Used=2.006174/50.0 | Remaining=47.993826 |
Usage=4.012348%
17%|           | 239/1433 [18:04<1:27:44,  4.41s/it]

[Quota] LLM=2.014421 | GPU=0.0 | Used=2.014421/50.0 | Remaining=47.985579 |
Usage=4.028842%
17%|           | 240/1433 [18:09<1:29:28,  4.50s/it]

[Quota] LLM=2.022599 | GPU=0.0 | Used=2.022599/50.0 | Remaining=47.977401 |
Usage=4.045198%
17%|           | 241/1433 [18:14<1:32:43,  4.67s/it]

[Quota] LLM=2.030507 | GPU=0.0 | Used=2.030507/50.0 | Remaining=47.969493 |
Usage=4.061014%
17%|           | 242/1433 [18:18<1:28:31,  4.46s/it]

[Quota] LLM=2.038346 | GPU=0.0 | Used=2.038346/50.0 | Remaining=47.961654 |
Usage=4.076692%
17%|           | 243/1433 [18:22<1:29:41,  4.52s/it]

[Quota] LLM=2.046107 | GPU=0.0 | Used=2.046107/50.0 | Remaining=47.953893 |
Usage=4.092214%
17%|           | 244/1433 [18:27<1:29:56,  4.54s/it]

[Quota] LLM=2.053967 | GPU=0.0 | Used=2.053967/50.0 | Remaining=47.946033 |
Usage=4.107934%
17%|           | 245/1433 [18:32<1:33:30,  4.72s/it]

[Quota] LLM=2.061878 | GPU=0.0 | Used=2.061878/50.0 | Remaining=47.938122 |
Usage=4.123756%
17%|           | 246/1433 [18:36<1:28:16,  4.46s/it]

[Quota] LLM=2.069639 | GPU=0.0 | Used=2.069639/50.0 | Remaining=47.930361 |
Usage=4.139278%
17%|           | 247/1433 [18:40<1:29:04,  4.51s/it]

[Quota] LLM=2.0776850000000002 | GPU=0.0 | Used=2.0776850000000002/50.0 |
Remaining=47.922315 | Usage=4.1553700000000005%

17%|           | 248/1433 [18:45<1:28:20,  4.47s/it]

[Quota] LLM=2.085872 | GPU=0.0 | Used=2.085872/50.0 | Remaining=47.914128 |
Usage=4.171744%
17%|           | 249/1433 [18:49<1:26:08,  4.37s/it]

[Quota] LLM=2.093906 | GPU=0.0 | Used=2.093906/50.0 | Remaining=47.906094 |
Usage=4.187812%
17%|           | 250/1433 [18:54<1:27:53,  4.46s/it]
```

```
[Quota] LLM=2.101679 | GPU=0.0 | Used=2.101679/50.0 | Remaining=47.898321 |
Usage=4.203358%
18%|           | 251/1433 [18:58<1:29:50,  4.56s/it]

[Quota] LLM=2.109479 | GPU=0.0 | Used=2.109479/50.0 | Remaining=47.890521 |
Usage=4.218958%
18%|           | 252/1433 [19:04<1:35:02,  4.83s/it]

[Quota] LLM=2.117327 | GPU=0.0 | Used=2.117327/50.0 | Remaining=47.882673 |
Usage=4.234654%
18%|           | 253/1433 [19:09<1:39:32,  5.06s/it]

[Quota] LLM=2.125211 | GPU=0.0 | Used=2.125211/50.0 | Remaining=47.874789 |
Usage=4.250422%
18%|           | 254/1433 [19:13<1:30:41,  4.62s/it]

[Quota] LLM=2.133095 | GPU=0.0 | Used=2.133095/50.0 | Remaining=47.866905 |
Usage=4.26619%
18%|           | 255/1433 [19:18<1:34:21,  4.81s/it]

[Quota] LLM=2.141084 | GPU=0.0 | Used=2.141084/50.0 | Remaining=47.858916 |
Usage=4.282168%
18%|           | 256/1433 [19:22<1:30:00,  4.59s/it]

[Quota] LLM=2.148869 | GPU=0.0 | Used=2.148869/50.0 | Remaining=47.851131 |
Usage=4.297738%
18%|           | 257/1433 [19:28<1:37:09,  4.96s/it]

[Quota] LLM=2.15672 | GPU=0.0 | Used=2.15672/50.0 | Remaining=47.84328 |
Usage=4.31344%
18%|           | 258/1433 [19:33<1:33:33,  4.78s/it]

[Quota] LLM=2.164847 | GPU=0.0 | Used=2.164847/50.0 | Remaining=47.835153 |
Usage=4.329694%
18%|           | 259/1433 [19:37<1:29:58,  4.60s/it]

[Quota] LLM=2.172602 | GPU=0.0 | Used=2.172602/50.0 | Remaining=47.827398 |
Usage=4.345204%
18%|           | 260/1433 [19:43<1:40:59,  5.17s/it]

[Quota] LLM=2.180378 | GPU=0.0 | Used=2.180378/50.0 | Remaining=47.819622 |
Usage=4.360756%
18%|           | 261/1433 [19:47<1:34:29,  4.84s/it]

[Quota] LLM=2.188067 | GPU=0.0 | Used=2.188067/50.0 |
Remaining=47.811932999999996 | Usage=4.376134%
18%|           | 262/1433 [19:52<1:32:21,  4.73s/it]
```

```
[Quota] LLM=2.196167 | GPU=0.0 | Used=2.196167/50.0 | Remaining=47.803833 |
Usage=4.392334%
18%|           | 263/1433 [19:58<1:42:32,  5.26s/it]

[Quota] LLM=2.203931 | GPU=0.0 | Used=2.203931/50.0 | Remaining=47.796069 |
Usage=4.407862%
18%|           | 264/1433 [20:03<1:42:04,  5.24s/it]

[Quota] LLM=2.211683 | GPU=0.0 | Used=2.211683/50.0 | Remaining=47.788317 |
Usage=4.423366%
18%|           | 265/1433 [20:07<1:33:30,  4.80s/it]

[Quota] LLM=2.21963 | GPU=0.0 | Used=2.21963/50.0 | Remaining=47.78037 |
Usage=4.43926%
19%|           | 266/1433 [20:13<1:36:41,  4.97s/it]

[Quota] LLM=2.227352 | GPU=0.0 | Used=2.227352/50.0 | Remaining=47.772648 |
Usage=4.454704%
19%|           | 267/1433 [20:19<1:45:16,  5.42s/it]

[Quota] LLM=2.235329 | GPU=0.0 | Used=2.235329/50.0 | Remaining=47.764671 |
Usage=4.470658%
19%|           | 268/1433 [20:23<1:36:11,  4.95s/it]

[Quota] LLM=2.242961 | GPU=0.0 | Used=2.242961/50.0 | Remaining=47.757039 |
Usage=4.485922%
19%|           | 269/1433 [20:27<1:31:47,  4.73s/it]

[Quota] LLM=2.250692 | GPU=0.0 | Used=2.250692/50.0 | Remaining=47.749308 |
Usage=4.501384%
19%|           | 270/1433 [20:31<1:26:41,  4.47s/it]

[Quota] LLM=2.258462 | GPU=0.0 | Used=2.258462/50.0 | Remaining=47.741538 |
Usage=4.516924%
19%|           | 271/1433 [20:36<1:29:14,  4.61s/it]

[Quota] LLM=2.266319 | GPU=0.0 | Used=2.266319/50.0 | Remaining=47.733681 |
Usage=4.532638%
19%|           | 272/1433 [20:41<1:30:45,  4.69s/it]

[Quota] LLM=2.274332 | GPU=0.0 | Used=2.274332/50.0 | Remaining=47.725668 |
Usage=4.548664%
19%|           | 273/1433 [20:44<1:23:18,  4.31s/it]

[Quota] LLM=2.282003 | GPU=0.0 | Used=2.282003/50.0 | Remaining=47.717997 |
Usage=4.564006%
19%|           | 274/1433 [20:48<1:17:43,  4.02s/it]
```

```
[Quota] LLM=2.289734 | GPU=0.0 | Used=2.289734/50.0 | Remaining=47.710266 |
Usage=4.579468%
 19%|           | 275/1433 [20:52<1:20:20,  4.16s/it]

[Quota] LLM=2.297525 | GPU=0.0 | Used=2.297525/50.0 | Remaining=47.702475 |
Usage=4.59505%
 19%|           | 276/1433 [21:01<1:44:59,  5.44s/it]

[Quota] LLM=2.306006 | GPU=0.0 | Used=2.306006/50.0 | Remaining=47.693994 |
Usage=4.612012%
 19%|           | 277/1433 [21:06<1:47:55,  5.60s/it]

[Quota] LLM=2.314094 | GPU=0.0 | Used=2.314094/50.0 | Remaining=47.685906 |
Usage=4.628188%
 19%|           | 278/1433 [21:13<1:50:52,  5.76s/it]

[Quota] LLM=2.321972 | GPU=0.0 | Used=2.321972/50.0 | Remaining=47.678028 |
Usage=4.643944%
 19%|           | 279/1433 [21:18<1:46:01,  5.51s/it]

[Quota] LLM=2.329694 | GPU=0.0 | Used=2.329694/50.0 | Remaining=47.670306 |
Usage=4.659388%
 20%|           | 280/1433 [21:21<1:33:51,  4.88s/it]

[Quota] LLM=2.337521 | GPU=0.0 | Used=2.337521/50.0 | Remaining=47.662479 |
Usage=4.675042%
 20%|           | 281/1433 [21:25<1:29:51,  4.68s/it]

[Quota] LLM=2.345135 | GPU=0.0 | Used=2.345135/50.0 | Remaining=47.654865 |
Usage=4.69027%
 20%|           | 282/1433 [21:30<1:31:49,  4.79s/it]

[Quota] LLM=2.353145 | GPU=0.0 | Used=2.353145/50.0 | Remaining=47.646855 |
Usage=4.70629%
 20%|           | 283/1433 [21:36<1:37:24,  5.08s/it]

[Quota] LLM=2.361221 | GPU=0.0 | Used=2.361221/50.0 | Remaining=47.638779 |
Usage=4.722442%
 20%|           | 284/1433 [21:40<1:32:55,  4.85s/it]

[Quota] LLM=2.369045 | GPU=0.0 | Used=2.369045/50.0 | Remaining=47.630955 |
Usage=4.73809%
 20%|           | 285/1433 [21:45<1:31:17,  4.77s/it]

[Quota] LLM=2.376989 | GPU=0.0 | Used=2.376989/50.0 | Remaining=47.623011 |
Usage=4.753978%
 20%|           | 286/1433 [21:49<1:24:56,  4.44s/it]
```

```
[Quota] LLM=2.385041 | GPU=0.0 | Used=2.385041/50.0 | Remaining=47.614959 |  
Usage=4.770082%  
20%|           | 287/1433 [21:53<1:25:11,  4.46s/it]  
  
[Quota] LLM=2.392943 | GPU=0.0 | Used=2.392943/50.0 | Remaining=47.607057 |  
Usage=4.785886%  
20%|           | 288/1433 [21:58<1:29:06,  4.67s/it]  
  
[Quota] LLM=2.401049 | GPU=0.0 | Used=2.401049/50.0 | Remaining=47.598951 |  
Usage=4.802098%  
20%|           | 289/1433 [22:04<1:35:03,  4.99s/it]  
  
[Quota] LLM=2.40914 | GPU=0.0 | Used=2.40914/50.0 | Remaining=47.59086 |  
Usage=4.81828%  
20%|           | 290/1433 [22:10<1:40:50,  5.29s/it]  
  
[Quota] LLM=2.417534 | GPU=0.0 | Used=2.417534/50.0 | Remaining=47.582466 |  
Usage=4.835068%  
20%|           | 291/1433 [22:13<1:29:12,  4.69s/it]  
  
[Quota] LLM=2.425154 | GPU=0.0 | Used=2.425154/50.0 | Remaining=47.574846 |  
Usage=4.850308%  
20%|           | 292/1433 [22:17<1:24:50,  4.46s/it]  
  
[Quota] LLM=2.433293 | GPU=0.0 | Used=2.433293/50.0 | Remaining=47.566707 |  
Usage=4.866586%  
20%|           | 293/1433 [22:22<1:24:23,  4.44s/it]  
  
[Quota] LLM=2.441348 | GPU=0.0 | Used=2.441348/50.0 | Remaining=47.558652 |  
Usage=4.882696%  
21%|           | 294/1433 [22:26<1:25:20,  4.50s/it]  
  
[Quota] LLM=2.449448 | GPU=0.0 | Used=2.449448/50.0 | Remaining=47.550552 |  
Usage=4.898896%  
21%|           | 295/1433 [22:31<1:27:20,  4.61s/it]  
  
[Quota] LLM=2.457752 | GPU=0.0 | Used=2.457752/50.0 | Remaining=47.542248 |  
Usage=4.915504%  
21%|           | 296/1433 [22:36<1:27:47,  4.63s/it]  
  
[Quota] LLM=2.465549 | GPU=0.0 | Used=2.465549/50.0 | Remaining=47.534451 |  
Usage=4.931098%  
21%|           | 297/1433 [22:39<1:22:41,  4.37s/it]  
  
[Quota] LLM=2.473367 | GPU=0.0 | Used=2.473367/50.0 | Remaining=47.526633 |  
Usage=4.946734%  
21%|           | 298/1433 [22:43<1:18:18,  4.14s/it]
```

```
[Quota] LLM=2.480924 | GPU=0.0 | Used=2.480924/50.0 | Remaining=47.519076 |
Usage=4.961848%
21%|           | 299/1433 [22:47<1:17:43,  4.11s/it]

[Quota] LLM=2.488925 | GPU=0.0 | Used=2.488925/50.0 | Remaining=47.511075 |
Usage=4.97785%
21%|           | 300/1433 [22:53<1:25:20,  4.52s/it]

[Quota] LLM=2.496557 | GPU=0.0 | Used=2.496557/50.0 | Remaining=47.503443 |
Usage=4.993114%
21%|           | 301/1433 [22:57<1:26:55,  4.61s/it]

[Quota] LLM=2.504309 | GPU=0.0 | Used=2.504309/50.0 | Remaining=47.495691 |
Usage=5.008618%
21%|           | 302/1433 [23:02<1:28:03,  4.67s/it]

[Quota] LLM=2.512232 | GPU=0.0 | Used=2.512232/50.0 | Remaining=47.487768 |
Usage=5.024464%
21%|           | 303/1433 [23:06<1:21:12,  4.31s/it]

[Quota] LLM=2.520401 | GPU=0.0 | Used=2.520401/50.0 | Remaining=47.479599 |
Usage=5.040802%
21%|           | 304/1433 [23:09<1:16:54,  4.09s/it]

[Quota] LLM=2.528057 | GPU=0.0 | Used=2.528057/50.0 | Remaining=47.471943 |
Usage=5.056114%
21%|           | 305/1433 [23:14<1:19:31,  4.23s/it]

[Quota] LLM=2.53589 | GPU=0.0 | Used=2.53589/50.0 | Remaining=47.46411 |
Usage=5.07178%
21%|           | 306/1433 [23:19<1:24:27,  4.50s/it]

[Quota] LLM=2.543654 | GPU=0.0 | Used=2.543654/50.0 |
Remaining=47.456345999999996 | Usage=5.087308%
21%|           | 307/1433 [23:23<1:23:31,  4.45s/it]

[Quota] LLM=2.551448 | GPU=0.0 | Used=2.551448/50.0 | Remaining=47.448552 |
Usage=5.102896%
21%|           | 308/1433 [23:27<1:17:28,  4.13s/it]

[Quota] LLM=2.559149 | GPU=0.0 | Used=2.559149/50.0 | Remaining=47.440851 |
Usage=5.118298%
22%|           | 309/1433 [23:31<1:17:46,  4.15s/it]

[Quota] LLM=2.566913 | GPU=0.0 | Used=2.566913/50.0 | Remaining=47.433087 |
Usage=5.133826%
22%|           | 310/1433 [23:36<1:22:51,  4.43s/it]
```

```
[Quota] LLM=2.574728 | GPU=0.0 | Used=2.574728/50.0 | Remaining=47.425272 |  
Usage=5.149456%  
22%|           | 311/1433 [23:42<1:29:47,  4.80s/it]  
  
[Quota] LLM=2.582582 | GPU=0.0 | Used=2.582582/50.0 | Remaining=47.417418 |  
Usage=5.165164%  
22%|           | 312/1433 [23:48<1:36:55,  5.19s/it]  
  
[Quota] LLM=2.590313 | GPU=0.0 | Used=2.590313/50.0 | Remaining=47.409687 |  
Usage=5.180626%  
22%|           | 313/1433 [23:52<1:29:34,  4.80s/it]  
  
[Quota] LLM=2.598119 | GPU=0.0 | Used=2.598119/50.0 | Remaining=47.401881 |  
Usage=5.196238%  
22%|           | 314/1433 [23:57<1:32:27,  4.96s/it]  
  
[Quota] LLM=2.606267 | GPU=0.0 | Used=2.606267/50.0 | Remaining=47.393733 |  
Usage=5.212534%  
22%|           | 315/1433 [24:01<1:28:15,  4.74s/it]  
  
[Quota] LLM=2.614121 | GPU=0.0 | Used=2.614121/50.0 | Remaining=47.385879 |  
Usage=5.228242%  
22%|           | 316/1433 [24:06<1:30:05,  4.84s/it]  
  
[Quota] LLM=2.621927 | GPU=0.0 | Used=2.621927/50.0 | Remaining=47.378073 |  
Usage=5.243854%  
22%|           | 317/1433 [24:10<1:25:47,  4.61s/it]  
  
[Quota] LLM=2.629844 | GPU=0.0 | Used=2.629844/50.0 | Remaining=47.370156 |  
Usage=5.259688%  
22%|           | 318/1433 [24:15<1:28:45,  4.78s/it]  
  
[Quota] LLM=2.637695 | GPU=0.0 | Used=2.637695/50.0 | Remaining=47.362305 |  
Usage=5.27539%  
22%|           | 319/1433 [24:20<1:28:25,  4.76s/it]  
  
[Quota] LLM=2.645276 | GPU=0.0 | Used=2.645276/50.0 | Remaining=47.354724 |  
Usage=5.290552%  
22%|           | 320/1433 [24:26<1:35:46,  5.16s/it]  
  
[Quota] LLM=2.653379 | GPU=0.0 | Used=2.653379/50.0 | Remaining=47.346621 |  
Usage=5.306758%  
22%|           | 321/1433 [24:30<1:30:19,  4.87s/it]  
  
[Quota] LLM=2.661053 | GPU=0.0 | Used=2.661053/50.0 | Remaining=47.338947 |  
Usage=5.322106%  
22%|           | 322/1433 [24:35<1:26:14,  4.66s/it]
```

```
[Quota] LLM=2.668718 | GPU=0.0 | Used=2.668718/50.0 | Remaining=47.331282 |  
Usage=5.337436%  
23%|           | 323/1433 [24:39<1:25:48,  4.64s/it]  
  
[Quota] LLM=2.676485 | GPU=0.0 | Used=2.676485/50.0 | Remaining=47.323515 |  
Usage=5.35297%  
23%|           | 324/1433 [24:44<1:26:05,  4.66s/it]  
  
[Quota] LLM=2.684675 | GPU=0.0 | Used=2.684675/50.0 | Remaining=47.315325 |  
Usage=5.36935%  
23%|           | 325/1433 [24:52<1:46:04,  5.74s/it]  
  
[Quota] LLM=2.692646 | GPU=0.0 | Used=2.692646/50.0 |  
Remaining=47.307354000000004 | Usage=5.385292%  
23%|           | 326/1433 [24:57<1:40:19,  5.44s/it]  
  
[Quota] LLM=2.700365 | GPU=0.0 | Used=2.700365/50.0 | Remaining=47.299635 |  
Usage=5.40073%  
23%|           | 327/1433 [25:00<1:28:02,  4.78s/it]  
  
[Quota] LLM=2.707799 | GPU=0.0 | Used=2.707799/50.0 | Remaining=47.292201 |  
Usage=5.415598%  
23%|           | 328/1433 [25:05<1:28:57,  4.83s/it]  
  
[Quota] LLM=2.715614 | GPU=0.0 | Used=2.715614/50.0 | Remaining=47.284386 |  
Usage=5.431228%  
23%|           | 329/1433 [25:10<1:30:25,  4.91s/it]  
  
[Quota] LLM=2.723546 | GPU=0.0 | Used=2.723546/50.0 | Remaining=47.276454 |  
Usage=5.447092%  
23%|           | 330/1433 [25:15<1:29:40,  4.88s/it]  
  
[Quota] LLM=2.731421 | GPU=0.0 | Used=2.731421/50.0 | Remaining=47.268579 |  
Usage=5.462842%  
23%|           | 331/1433 [25:19<1:22:02,  4.47s/it]  
  
[Quota] LLM=2.739185 | GPU=0.0 | Used=2.739185/50.0 | Remaining=47.260815 |  
Usage=5.47837%  
23%|           | 332/1433 [25:23<1:21:12,  4.43s/it]  
  
[Quota] LLM=2.746853 | GPU=0.0 | Used=2.746853/50.0 | Remaining=47.253147 |  
Usage=5.493706%  
23%|           | 333/1433 [25:28<1:23:24,  4.55s/it]  
  
[Quota] LLM=2.754773 | GPU=0.0 | Used=2.754773/50.0 | Remaining=47.245227 |  
Usage=5.509546%  
23%|           | 334/1433 [25:32<1:19:10,  4.32s/it]
```

```
[Quota] LLM=2.762825 | GPU=0.0 | Used=2.762825/50.0 | Remaining=47.237175 |
Usage=5.52565%
23%|           | 335/1433 [25:35<1:15:34,  4.13s/it]

[Quota] LLM=2.770601 | GPU=0.0 | Used=2.770601/50.0 | Remaining=47.229399 |
Usage=5.541202%
23%|           | 336/1433 [25:39<1:11:11,  3.89s/it]

[Quota] LLM=2.778305 | GPU=0.0 | Used=2.778305/50.0 | Remaining=47.221695 |
Usage=5.55661%
24%|           | 337/1433 [25:43<1:16:48,  4.21s/it]

[Quota] LLM=2.786291 | GPU=0.0 | Used=2.786291/50.0 | Remaining=47.213709 |
Usage=5.572582%
24%|           | 338/1433 [25:50<1:28:33,  4.85s/it]

[Quota] LLM=2.794202 | GPU=0.0 | Used=2.794202/50.0 | Remaining=47.205798 |
Usage=5.588404%
24%|           | 339/1433 [25:56<1:35:34,  5.24s/it]

[Quota] LLM=2.8019060000000002 | GPU=0.0 | Used=2.8019060000000002/50.0 |
Remaining=47.198094 | Usage=5.6038120000000005%
24%|           | 340/1433 [26:00<1:28:25,  4.85s/it]

[Quota] LLM=2.809631 | GPU=0.0 | Used=2.809631/50.0 | Remaining=47.190369 |
Usage=5.619262%
24%|           | 341/1433 [26:04<1:21:39,  4.49s/it]

[Quota] LLM=2.817341 | GPU=0.0 | Used=2.817341/50.0 | Remaining=47.182659 |
Usage=5.634682%
24%|           | 342/1433 [26:08<1:21:55,  4.51s/it]

[Quota] LLM=2.825081 | GPU=0.0 | Used=2.825081/50.0 | Remaining=47.174919 |
Usage=5.650162%
24%|           | 343/1433 [26:13<1:26:03,  4.74s/it]

[Quota] LLM=2.833043 | GPU=0.0 | Used=2.833043/50.0 | Remaining=47.166957 |
Usage=5.666086%
24%|           | 344/1433 [26:18<1:23:26,  4.60s/it]

[Quota] LLM=2.840942 | GPU=0.0 | Used=2.840942/50.0 | Remaining=47.159058 |
Usage=5.681884%
24%|           | 345/1433 [26:22<1:24:40,  4.67s/it]

[Quota] LLM=2.848979 | GPU=0.0 | Used=2.848979/50.0 | Remaining=47.151021 |
Usage=5.697958%
24%|           | 346/1433 [26:27<1:26:11,  4.76s/it]
```

```
[Quota] LLM=2.8567460000000002 | GPU=0.0 | Used=2.8567460000000002/50.0 | Remaining=47.143254 | Usage=5.7134920000000005%
24%|           | 347/1433 [26:32<1:27:16,  4.82s/it]

[Quota] LLM=2.864864 | GPU=0.0 | Used=2.864864/50.0 | Remaining=47.135136 | Usage=5.729728%
24%|           | 348/1433 [26:37<1:25:35,  4.73s/it]

[Quota] LLM=2.872685 | GPU=0.0 | Used=2.872685/50.0 | Remaining=47.127315 | Usage=5.74537%
24%|           | 349/1433 [26:41<1:23:26,  4.62s/it]

[Quota] LLM=2.88038 | GPU=0.0 | Used=2.88038/50.0 | Remaining=47.11962 | Usage=5.76076%
24%|           | 350/1433 [26:46<1:22:05,  4.55s/it]

[Quota] LLM=2.888126 | GPU=0.0 | Used=2.888126/50.0 | Remaining=47.111874 | Usage=5.776252%
24%|           | 351/1433 [26:50<1:18:20,  4.34s/it]

[Quota] LLM=2.896019 | GPU=0.0 | Used=2.896019/50.0 | Remaining=47.103981 | Usage=5.792038%
25%|           | 352/1433 [26:53<1:14:47,  4.15s/it]

[Quota] LLM=2.903549 | GPU=0.0 | Used=2.903549/50.0 | Remaining=47.096451 | Usage=5.807098%
25%|           | 353/1433 [26:58<1:20:30,  4.47s/it]

[Quota] LLM=2.911361 | GPU=0.0 | Used=2.911361/50.0 | Remaining=47.088639 | Usage=5.822722%
25%|           | 354/1433 [27:03<1:20:56,  4.50s/it]

[Quota] LLM=2.91923 | GPU=0.0 | Used=2.91923/50.0 | Remaining=47.08077 | Usage=5.83846%
25%|           | 355/1433 [27:08<1:22:34,  4.60s/it]

[Quota] LLM=2.927117 | GPU=0.0 | Used=2.927117/50.0 | Remaining=47.072883 | Usage=5.854234%
25%|           | 356/1433 [27:11<1:14:48,  4.17s/it]

[Quota] LLM=2.934692 | GPU=0.0 | Used=2.934692/50.0 | Remaining=47.065308 | Usage=5.869384%
25%|           | 357/1433 [27:15<1:14:15,  4.14s/it]

[Quota] LLM=2.94248 | GPU=0.0 | Used=2.94248/50.0 | Remaining=47.05752 | Usage=5.88496%
25%|           | 358/1433 [27:19<1:14:41,  4.17s/it]
```

```
[Quota] LLM=2.9504 | GPU=0.0 | Used=2.9504/50.0 | Remaining=47.0496 |  
Usage=5.9008%  
25%|           | 359/1433 [27:23<1:12:10,  4.03s/it]  
  
[Quota] LLM=2.958005 | GPU=0.0 | Used=2.958005/50.0 | Remaining=47.041995 |  
Usage=5.91601%  
25%|           | 360/1433 [27:28<1:16:48,  4.29s/it]  
  
[Quota] LLM=2.965748 | GPU=0.0 | Used=2.965748/50.0 | Remaining=47.034252 |  
Usage=5.931496%  
25%|           | 361/1433 [27:32<1:13:37,  4.12s/it]  
  
[Quota] LLM=2.973377 | GPU=0.0 | Used=2.973377/50.0 | Remaining=47.026623 |  
Usage=5.946754%  
25%|           | 362/1433 [27:36<1:13:41,  4.13s/it]  
  
[Quota] LLM=2.981291 | GPU=0.0 | Used=2.981291/50.0 | Remaining=47.018709 |  
Usage=5.962582%  
25%|           | 363/1433 [27:40<1:15:59,  4.26s/it]  
  
[Quota] LLM=2.989286 | GPU=0.0 | Used=2.989286/50.0 | Remaining=47.010714 |  
Usage=5.978572%  
25%|           | 364/1433 [27:45<1:19:35,  4.47s/it]  
  
[Quota] LLM=2.997029 | GPU=0.0 | Used=2.997029/50.0 | Remaining=47.002971 |  
Usage=5.994058%  
25%|           | 365/1433 [27:50<1:20:26,  4.52s/it]  
  
[Quota] LLM=3.005114 | GPU=0.0 | Used=3.005114/50.0 | Remaining=46.994886 |  
Usage=6.010228%  
26%|           | 366/1433 [27:55<1:21:06,  4.56s/it]  
  
[Quota] LLM=3.012953 | GPU=0.0 | Used=3.012953/50.0 | Remaining=46.987047 |  
Usage=6.025906%  
26%|           | 367/1433 [28:01<1:31:14,  5.14s/it]  
  
[Quota] LLM=3.021125 | GPU=0.0 | Used=3.021125/50.0 | Remaining=46.978875 |  
Usage=6.04225%  
26%|           | 368/1433 [28:05<1:26:48,  4.89s/it]  
  
[Quota] LLM=3.0291230000000002 | GPU=0.0 | Used=3.0291230000000002/50.0 |  
Remaining=46.970877 | Usage=6.0582460000000005%  
26%|           | 369/1433 [28:10<1:26:17,  4.87s/it]  
  
[Quota] LLM=3.037472 | GPU=0.0 | Used=3.037472/50.0 | Remaining=46.962528 |  
Usage=6.074944%  
26%|           | 370/1433 [28:13<1:16:35,  4.32s/it]
```

[Quota] LLM=3.0449990000000002 | GPU=0.0 | Used=3.0449990000000002/50.0 | Remaining=46.955001 | Usage=6.0899980000000005%

26%| 371/1433 [28:18<1:17:32, 4.38s/it]

[Quota] LLM=3.053081 | GPU=0.0 | Used=3.053081/50.0 | Remaining=46.946919 | Usage=6.106162%

26%| 372/1433 [28:22<1:17:18, 4.37s/it]

[Quota] LLM=3.061364 | GPU=0.0 | Used=3.061364/50.0 | Remaining=46.938636 | Usage=6.122728%

26%| 373/1433 [28:26<1:12:56, 4.13s/it]

[Quota] LLM=3.069107 | GPU=0.0 | Used=3.069107/50.0 | Remaining=46.930893 | Usage=6.138214%

26%| 374/1433 [28:30<1:12:16, 4.09s/it]

[Quota] LLM=3.077057 | GPU=0.0 | Used=3.077057/50.0 | Remaining=46.92294300000004 | Usage=6.154114%

26%| 375/1433 [28:35<1:16:17, 4.33s/it]

[Quota] LLM=3.084884 | GPU=0.0 | Used=3.084884/50.0 | Remaining=46.915116 | Usage=6.169768%

26%| 376/1433 [28:40<1:23:05, 4.72s/it]

[Quota] LLM=3.092837 | GPU=0.0 | Used=3.092837/50.0 | Remaining=46.907163 | Usage=6.185674%

26%| 377/1433 [28:45<1:21:04, 4.61s/it]

[Quota] LLM=3.10073 | GPU=0.0 | Used=3.10073/50.0 | Remaining=46.89927 | Usage=6.20146%

26%| 378/1433 [28:50<1:26:34, 4.92s/it]

[Quota] LLM=3.108875 | GPU=0.0 | Used=3.108875/50.0 | Remaining=46.891125 | Usage=6.21775%

26%| 379/1433 [28:55<1:24:13, 4.79s/it]

[Quota] LLM=3.116726 | GPU=0.0 | Used=3.116726/50.0 | Remaining=46.883274 | Usage=6.233452%

27%| 380/1433 [28:59<1:19:59, 4.56s/it]

[Quota] LLM=3.124334 | GPU=0.0 | Used=3.124334/50.0 | Remaining=46.875666 | Usage=6.248668%

27%| 381/1433 [29:03<1:20:05, 4.57s/it]

[Quota] LLM=3.132038 | GPU=0.0 | Used=3.132038/50.0 | Remaining=46.867962 | Usage=6.264076%

27%| 382/1433 [29:08<1:18:29, 4.48s/it]

```
[Quota] LLM=3.139967 | GPU=0.0 | Used=3.139967/50.0 | Remaining=46.860033 |
Usage=6.279934%
27%|           | 383/1433 [29:13<1:22:41,  4.73s/it]

[Quota] LLM=3.14828 | GPU=0.0 | Used=3.14828/50.0 | Remaining=46.85172 |
Usage=6.2965600000000001%
27%|           | 384/1433 [29:18<1:23:59,  4.80s/it]

[Quota] LLM=3.1562 | GPU=0.0 | Used=3.1562/50.0 | Remaining=46.8438 |
Usage=6.3124%
27%|           | 385/1433 [29:22<1:21:15,  4.65s/it]

[Quota] LLM=3.163955 | GPU=0.0 | Used=3.163955/50.0 | Remaining=46.836045 |
Usage=6.32791%
27%|           | 386/1433 [29:25<1:13:20,  4.20s/it]

[Quota] LLM=3.171467 | GPU=0.0 | Used=3.171467/50.0 | Remaining=46.828533 |
Usage=6.342934%
27%|           | 387/1433 [29:30<1:12:56,  4.18s/it]

[Quota] LLM=3.179204 | GPU=0.0 | Used=3.179204/50.0 | Remaining=46.820796 |
Usage=6.358408%
27%|           | 388/1433 [29:34<1:15:19,  4.32s/it]

[Quota] LLM=3.186938 | GPU=0.0 | Used=3.186938/50.0 | Remaining=46.813062 |
Usage=6.3738760000000001%
27%|           | 389/1433 [29:39<1:20:12,  4.61s/it]

[Quota] LLM=3.19511 | GPU=0.0 | Used=3.19511/50.0 | Remaining=46.80489 |
Usage=6.39022%
27%|           | 390/1433 [29:44<1:17:53,  4.48s/it]

[Quota] LLM=3.202904 | GPU=0.0 | Used=3.202904/50.0 |
Remaining=46.797095999999996 | Usage=6.405808%
27%|           | 391/1433 [29:49<1:22:55,  4.77s/it]

[Quota] LLM=3.210965 | GPU=0.0 | Used=3.210965/50.0 | Remaining=46.789035 |
Usage=6.42193%
27%|           | 392/1433 [29:54<1:22:05,  4.73s/it]

[Quota] LLM=3.219089 | GPU=0.0 | Used=3.219089/50.0 | Remaining=46.780911 |
Usage=6.438178%
27%|           | 393/1433 [29:57<1:16:00,  4.39s/it]

[Quota] LLM=3.226598 | GPU=0.0 | Used=3.226598/50.0 | Remaining=46.773402 |
Usage=6.453196%
27%|           | 394/1433 [30:01<1:13:33,  4.25s/it]
```

```
[Quota] LLM=3.23441 | GPU=0.0 | Used=3.23441/50.0 | Remaining=46.76559 |  
Usage=6.46882%  
28%|           | 395/1433 [30:05<1:09:14,  4.00s/it]  
  
[Quota] LLM=3.242183 | GPU=0.0 | Used=3.242183/50.0 | Remaining=46.757817 |  
Usage=6.484366%  
28%|           | 396/1433 [30:08<1:06:39,  3.86s/it]  
  
[Quota] LLM=3.249947 | GPU=0.0 | Used=3.249947/50.0 | Remaining=46.750053 |  
Usage=6.499894%  
28%|           | 397/1433 [30:12<1:05:57,  3.82s/it]  
  
[Quota] LLM=3.258029 | GPU=0.0 | Used=3.258029/50.0 | Remaining=46.741971 |  
Usage=6.51605799999999%  
28%|           | 398/1433 [30:16<1:09:36,  4.04s/it]  
  
[Quota] LLM=3.266207 | GPU=0.0 | Used=3.266207/50.0 | Remaining=46.733793 |  
Usage=6.532414%  
28%|           | 399/1433 [30:20<1:07:03,  3.89s/it]  
  
[Quota] LLM=3.273863 | GPU=0.0 | Used=3.273863/50.0 | Remaining=46.726137 |  
Usage=6.547726%  
28%|           | 400/1433 [30:25<1:10:24,  4.09s/it]  
  
[Quota] LLM=3.28157 | GPU=0.0 | Used=3.28157/50.0 | Remaining=46.71843 |  
Usage=6.56313999999999%  
28%|           | 401/1433 [30:29<1:11:07,  4.14s/it]  
  
[Quota] LLM=3.289466 | GPU=0.0 | Used=3.289466/50.0 | Remaining=46.710534 |  
Usage=6.578932%  
28%|           | 402/1433 [30:33<1:13:32,  4.28s/it]  
  
[Quota] LLM=3.297143 | GPU=0.0 | Used=3.297143/50.0 | Remaining=46.702857 |  
Usage=6.594286%  
28%|           | 403/1433 [30:38<1:14:42,  4.35s/it]  
  
[Quota] LLM=3.304973 | GPU=0.0 | Used=3.304973/50.0 | Remaining=46.695027 |  
Usage=6.609946%  
28%|           | 404/1433 [30:43<1:18:30,  4.58s/it]  
  
[Quota] LLM=3.31322 | GPU=0.0 | Used=3.31322/50.0 | Remaining=46.68678 |  
Usage=6.626440000000006%  
28%|           | 405/1433 [30:48<1:18:55,  4.61s/it]  
  
[Quota] LLM=3.321011 | GPU=0.0 | Used=3.321011/50.0 | Remaining=46.678989 |  
Usage=6.642022%  
28%|           | 406/1433 [30:51<1:13:42,  4.31s/it]
```

```
[Quota] LLM=3.328946 | GPU=0.0 | Used=3.328946/50.0 | Remaining=46.671054 |  
Usage=6.657892%  
28%|           | 407/1433 [30:56<1:15:39,  4.42s/it]  
  
[Quota] LLM=3.337271 | GPU=0.0 | Used=3.337271/50.0 | Remaining=46.662729 |  
Usage=6.674542%  
28%|           | 408/1433 [31:01<1:16:36,  4.48s/it]  
  
[Quota] LLM=3.345119 | GPU=0.0 | Used=3.345119/50.0 | Remaining=46.654881 |  
Usage=6.690238%  
29%|           | 409/1433 [31:05<1:17:42,  4.55s/it]  
  
[Quota] LLM=3.352739 | GPU=0.0 | Used=3.352739/50.0 | Remaining=46.647261 |  
Usage=6.705478000000001%  
29%|           | 410/1433 [31:08<1:10:25,  4.13s/it]  
  
[Quota] LLM=3.360428 | GPU=0.0 | Used=3.360428/50.0 | Remaining=46.639572 |  
Usage=6.720856%  
29%|           | 411/1433 [31:13<1:10:15,  4.12s/it]  
  
[Quota] LLM=3.368345 | GPU=0.0 | Used=3.368345/50.0 | Remaining=46.631655 |  
Usage=6.736690000000001%  
29%|           | 412/1433 [31:17<1:11:18,  4.19s/it]  
  
[Quota] LLM=3.376058 | GPU=0.0 | Used=3.376058/50.0 | Remaining=46.623942 |  
Usage=6.752116%  
29%|           | 413/1433 [31:21<1:09:40,  4.10s/it]  
  
[Quota] LLM=3.383939 | GPU=0.0 | Used=3.383939/50.0 | Remaining=46.616061 |  
Usage=6.767878%  
29%|           | 414/1433 [31:25<1:11:08,  4.19s/it]  
  
[Quota] LLM=3.392102 | GPU=0.0 | Used=3.392102/50.0 | Remaining=46.607898 |  
Usage=6.78420399999999%  
29%|           | 415/1433 [31:30<1:14:29,  4.39s/it]  
  
[Quota] LLM=3.400376 | GPU=0.0 | Used=3.400376/50.0 | Remaining=46.599624 |  
Usage=6.800752%  
29%|           | 416/1433 [31:34<1:13:59,  4.37s/it]  
  
[Quota] LLM=3.408245 | GPU=0.0 | Used=3.408245/50.0 | Remaining=46.591755 |  
Usage=6.81649%  
29%|           | 417/1433 [31:38<1:11:16,  4.21s/it]  
  
[Quota] LLM=3.415871 | GPU=0.0 | Used=3.415871/50.0 | Remaining=46.584129 |  
Usage=6.831742%  
29%|           | 418/1433 [31:42<1:11:08,  4.21s/it]
```

```
[Quota] LLM=3.423614 | GPU=0.0 | Used=3.423614/50.0 | Remaining=46.576386 |
Usage=6.847227999999999%  
29%|           | 419/1433 [31:46<1:08:00,  4.02s/it]  
  
[Quota] LLM=3.431216 | GPU=0.0 | Used=3.431216/50.0 | Remaining=46.568784 |
Usage=6.862432%  
29%|           | 420/1433 [31:50<1:06:16,  3.93s/it]  
  
[Quota] LLM=3.4387820000000002 | GPU=0.0 | Used=3.4387820000000002/50.0 |
Remaining=46.561218 | Usage=6.877564%  
29%|           | 421/1433 [31:54<1:06:52,  3.97s/it]  
  
[Quota] LLM=3.446621 | GPU=0.0 | Used=3.446621/50.0 | Remaining=46.553379 |
Usage=6.89324199999999%  
29%|           | 422/1433 [31:58<1:06:19,  3.94s/it]  
  
[Quota] LLM=3.454547 | GPU=0.0 | Used=3.454547/50.0 | Remaining=46.545453 |
Usage=6.90909399999999%  
30%|           | 423/1433 [32:03<1:11:26,  4.24s/it]  
  
[Quota] LLM=3.462518 | GPU=0.0 | Used=3.462518/50.0 | Remaining=46.537482 |
Usage=6.925036000000001%  
30%|           | 424/1433 [32:06<1:09:00,  4.10s/it]  
  
[Quota] LLM=3.470228 | GPU=0.0 | Used=3.470228/50.0 | Remaining=46.529772 |
Usage=6.940456%  
30%|           | 425/1433 [32:13<1:19:42,  4.74s/it]  
  
[Quota] LLM=3.478139 | GPU=0.0 | Used=3.478139/50.0 | Remaining=46.521861 |
Usage=6.956278%  
30%|           | 426/1433 [32:17<1:17:49,  4.64s/it]  
  
[Quota] LLM=3.4858190000000002 | GPU=0.0 | Used=3.4858190000000002/50.0 |
Remaining=46.514181 | Usage=6.9716380000000004%  
30%|           | 427/1433 [32:22<1:17:27,  4.62s/it]  
  
[Quota] LLM=3.493634 | GPU=0.0 | Used=3.493634/50.0 | Remaining=46.506366 |
Usage=6.987268%  
30%|           | 428/1433 [32:27<1:20:09,  4.79s/it]  
  
[Quota] LLM=3.501647 | GPU=0.0 | Used=3.501647/50.0 | Remaining=46.498353 |
Usage=7.003294%  
30%|           | 429/1433 [32:30<1:13:15,  4.38s/it]  
  
[Quota] LLM=3.509345 | GPU=0.0 | Used=3.509345/50.0 | Remaining=46.490655 |
Usage=7.01868999999999%  
30%|           | 430/1433 [32:35<1:13:00,  4.37s/it]
```

```
[Quota] LLM=3.516938 | GPU=0.0 | Used=3.516938/50.0 | Remaining=46.483062 |  
Usage=7.033876%  
30%|       | 431/1433 [32:39<1:15:39,  4.53s/it]  
  
[Quota] LLM=3.524759 | GPU=0.0 | Used=3.524759/50.0 | Remaining=46.475241 |  
Usage=7.049518000000001%  
30%|       | 432/1433 [32:44<1:13:57,  4.43s/it]  
  
[Quota] LLM=3.532631 | GPU=0.0 | Used=3.532631/50.0 | Remaining=46.467369 |  
Usage=7.065262%  
30%|       | 433/1433 [32:47<1:09:49,  4.19s/it]  
  
[Quota] LLM=3.540434 | GPU=0.0 | Used=3.540434/50.0 | Remaining=46.459566 |  
Usage=7.080868%  
30%|       | 434/1433 [32:51<1:09:47,  4.19s/it]  
  
[Quota] LLM=3.548261 | GPU=0.0 | Used=3.548261/50.0 | Remaining=46.451739 |  
Usage=7.096521999999999%  
30%|       | 435/1433 [32:56<1:09:32,  4.18s/it]  
  
[Quota] LLM=3.556055 | GPU=0.0 | Used=3.556055/50.0 | Remaining=46.443945 |  
Usage=7.11211%  
30%|       | 436/1433 [33:01<1:13:31,  4.42s/it]  
  
[Quota] LLM=3.563708 | GPU=0.0 | Used=3.563708/50.0 | Remaining=46.436292 |  
Usage=7.127416%  
30%|       | 437/1433 [33:04<1:10:15,  4.23s/it]  
  
[Quota] LLM=3.571298 | GPU=0.0 | Used=3.571298/50.0 | Remaining=46.428702 |  
Usage=7.14259599999999%  
31%|       | 438/1433 [33:08<1:08:56,  4.16s/it]  
  
[Quota] LLM=3.579008 | GPU=0.0 | Used=3.579008/50.0 | Remaining=46.420992 |  
Usage=7.158016%  
31%|       | 439/1433 [33:14<1:14:43,  4.51s/it]  
  
[Quota] LLM=3.586655 | GPU=0.0 | Used=3.586655/50.0 | Remaining=46.413345 |  
Usage=7.17330999999999%  
31%|       | 440/1433 [33:18<1:11:33,  4.32s/it]  
  
[Quota] LLM=3.594533 | GPU=0.0 | Used=3.594533/50.0 | Remaining=46.405467 |  
Usage=7.189066000000001%  
31%|       | 441/1433 [33:22<1:12:03,  4.36s/it]  
  
[Quota] LLM=3.602507 | GPU=0.0 | Used=3.602507/50.0 | Remaining=46.397493 |  
Usage=7.205014%  
31%|       | 442/1433 [33:26<1:08:55,  4.17s/it]
```

```
[Quota] LLM=3.610262 | GPU=0.0 | Used=3.610262/50.0 | Remaining=46.389738 |
Usage=7.220524%
31% | 443/1433 [33:30<1:07:12, 4.07s/it]

[Quota] LLM=3.618116 | GPU=0.0 | Used=3.618116/50.0 | Remaining=46.381884 |
Usage=7.236232000000001%
31% | 444/1433 [33:35<1:14:34, 4.52s/it]

[Quota] LLM=3.626045 | GPU=0.0 | Used=3.626045/50.0 | Remaining=46.373955 |
Usage=7.25209%
31% | 445/1433 [33:40<1:15:39, 4.59s/it]

[Quota] LLM=3.633872 | GPU=0.0 | Used=3.633872/50.0 | Remaining=46.366128 |
Usage=7.267744000000001%
31% | 446/1433 [33:45<1:16:58, 4.68s/it]

[Quota] LLM=3.641921 | GPU=0.0 | Used=3.641921/50.0 | Remaining=46.35807900000004 |
Usage=7.283842%
31% | 447/1433 [33:50<1:17:48, 4.73s/it]

[Quota] LLM=3.649781 | GPU=0.0 | Used=3.649781/50.0 | Remaining=46.350219 |
Usage=7.299562%
31% | 448/1433 [33:55<1:18:38, 4.79s/it]

[Quota] LLM=3.657815 | GPU=0.0 | Used=3.657815/50.0 | Remaining=46.342185 |
Usage=7.31563%
31% | 449/1433 [33:59<1:18:19, 4.78s/it]

[Quota] LLM=3.665894 | GPU=0.0 | Used=3.665894/50.0 | Remaining=46.334106 |
Usage=7.331788%
31% | 450/1433 [34:05<1:21:30, 4.98s/it]

[Quota] LLM=3.67385 | GPU=0.0 | Used=3.67385/50.0 | Remaining=46.32615 |
Usage=7.3477%
31% | 451/1433 [34:09<1:19:41, 4.87s/it]

[Quota] LLM=3.681632 | GPU=0.0 | Used=3.681632/50.0 | Remaining=46.318368 |
Usage=7.363264%
32% | 452/1433 [34:13<1:14:36, 4.56s/it]

[Quota] LLM=3.68924 | GPU=0.0 | Used=3.68924/50.0 | Remaining=46.31076 |
Usage=7.37848%
32% | 453/1433 [34:19<1:19:37, 4.88s/it]

[Quota] LLM=3.696938 | GPU=0.0 | Used=3.696938/50.0 | Remaining=46.303062 |
Usage=7.39387599999999%
32% | 454/1433 [34:23<1:16:17, 4.68s/it]
```

```
[Quota] LLM=3.704471 | GPU=0.0 | Used=3.704471/50.0 | Remaining=46.295529 |
Usage=7.408942000000001%
32%|           | 455/1433 [34:28<1:15:23,  4.63s/it]

[Quota] LLM=3.712289 | GPU=0.0 | Used=3.712289/50.0 | Remaining=46.287711 |
Usage=7.424577999999995%
32%|           | 456/1433 [34:33<1:17:39,  4.77s/it]

[Quota] LLM=3.719981 | GPU=0.0 | Used=3.719981/50.0 | Remaining=46.280019 |
Usage=7.439961999999995%
32%|           | 457/1433 [34:38<1:18:12,  4.81s/it]

[Quota] LLM=3.727832 | GPU=0.0 | Used=3.727832/50.0 | Remaining=46.272168 |
Usage=7.455664%
32%|           | 458/1433 [34:41<1:12:29,  4.46s/it]

[Quota] LLM=3.735572 | GPU=0.0 | Used=3.735572/50.0 | Remaining=46.264428 |
Usage=7.471144000000001%
32%|           | 459/1433 [34:46<1:12:56,  4.49s/it]

[Quota] LLM=3.743453 | GPU=0.0 | Used=3.743453/50.0 | Remaining=46.256547 |
Usage=7.486906%
32%|           | 460/1433 [34:51<1:14:51,  4.62s/it]

[Quota] LLM=3.751184 | GPU=0.0 | Used=3.751184/50.0 | Remaining=46.248816 |
Usage=7.502368%
32%|           | 461/1433 [34:58<1:25:17,  5.27s/it]

[Quota] LLM=3.759149 | GPU=0.0 | Used=3.759149/50.0 | Remaining=46.240851 |
Usage=7.518298%
32%|           | 462/1433 [35:02<1:20:01,  4.95s/it]

[Quota] LLM=3.766901 | GPU=0.0 | Used=3.766901/50.0 | Remaining=46.233099 |
Usage=7.53380199999999%
32%|           | 463/1433 [35:06<1:16:06,  4.71s/it]

[Quota] LLM=3.774878 | GPU=0.0 | Used=3.774878/50.0 | Remaining=46.225122 |
Usage=7.549756%
32%|           | 464/1433 [35:10<1:15:38,  4.68s/it]

[Quota] LLM=3.782882 | GPU=0.0 | Used=3.782882/50.0 | Remaining=46.217118 |
Usage=7.565764%
32%|           | 465/1433 [35:15<1:13:54,  4.58s/it]

[Quota] LLM=3.790679 | GPU=0.0 | Used=3.790679/50.0 | Remaining=46.209321 |
Usage=7.58135799999999%
33%|           | 466/1433 [35:19<1:11:13,  4.42s/it]
```

```
[Quota] LLM=3.798266 | GPU=0.0 | Used=3.798266/50.0 | Remaining=46.201734 |  
Usage=7.596532%  
33%|           | 467/1433 [35:24<1:15:47,  4.71s/it]  
  
[Quota] LLM=3.806282 | GPU=0.0 | Used=3.806282/50.0 | Remaining=46.193718 |  
Usage=7.61256399999999%  
33%|           | 468/1433 [35:29<1:16:08,  4.73s/it]  
  
[Quota] LLM=3.814019 | GPU=0.0 | Used=3.814019/50.0 | Remaining=46.185981 |  
Usage=7.62803799999999%  
33%|           | 469/1433 [35:33<1:10:11,  4.37s/it]  
  
[Quota] LLM=3.82172 | GPU=0.0 | Used=3.82172/50.0 | Remaining=46.17828 |  
Usage=7.64344%  
33%|           | 470/1433 [35:40<1:23:23,  5.20s/it]  
  
[Quota] LLM=3.829814 | GPU=0.0 | Used=3.829814/50.0 | Remaining=46.170186 |  
Usage=7.6596280000000005%  
33%|           | 471/1433 [35:43<1:15:46,  4.73s/it]  
  
[Quota] LLM=3.837593 | GPU=0.0 | Used=3.837593/50.0 | Remaining=46.162407 |  
Usage=7.675186%  
33%|           | 472/1433 [35:47<1:11:14,  4.45s/it]  
  
[Quota] LLM=3.845186 | GPU=0.0 | Used=3.845186/50.0 | Remaining=46.154814 |  
Usage=7.69037199999999%  
33%|           | 473/1433 [35:51<1:09:32,  4.35s/it]  
  
[Quota] LLM=3.852794 | GPU=0.0 | Used=3.852794/50.0 | Remaining=46.147206 |  
Usage=7.705588%  
33%|           | 474/1433 [35:56<1:11:26,  4.47s/it]  
  
[Quota] LLM=3.860525 | GPU=0.0 | Used=3.860525/50.0 | Remaining=46.139475 |  
Usage=7.72105%  
33%|           | 475/1433 [36:00<1:08:39,  4.30s/it]  
  
[Quota] LLM=3.868334 | GPU=0.0 | Used=3.868334/50.0 | Remaining=46.131666 |  
Usage=7.73666799999999%  
33%|           | 476/1433 [36:05<1:12:37,  4.55s/it]  
  
[Quota] LLM=3.876482 | GPU=0.0 | Used=3.876482/50.0 | Remaining=46.123518 |  
Usage=7.752964000000001%  
33%|           | 477/1433 [36:09<1:08:37,  4.31s/it]  
  
[Quota] LLM=3.884381 | GPU=0.0 | Used=3.884381/50.0 | Remaining=46.115619 |  
Usage=7.768762%  
33%|           | 478/1433 [36:14<1:11:11,  4.47s/it]
```

```
[Quota] LLM=3.892481 | GPU=0.0 | Used=3.892481/50.0 |  
Remaining=46.107518999999996 | Usage=7.784962000000001%  
33%|           | 479/1433 [36:17<1:07:30,  4.25s/it]  
  
[Quota] LLM=3.900104 | GPU=0.0 | Used=3.900104/50.0 | Remaining=46.099896 |  
Usage=7.8002080000000005%  
33%|           | 480/1433 [36:21<1:05:28,  4.12s/it]  
  
[Quota] LLM=3.907688 | GPU=0.0 | Used=3.907688/50.0 | Remaining=46.092312 |  
Usage=7.8153760000000005%  
34%|           | 481/1433 [36:25<1:03:54,  4.03s/it]  
  
[Quota] LLM=3.9152810000000002 | GPU=0.0 | Used=3.9152810000000002/50.0 |  
Remaining=46.084719 | Usage=7.8305620000000005%  
34%|           | 482/1433 [36:29<1:02:28,  3.94s/it]  
  
[Quota] LLM=3.92285 | GPU=0.0 | Used=3.92285/50.0 | Remaining=46.07715 |  
Usage=7.8457%  
34%|           | 483/1433 [36:32<1:00:59,  3.85s/it]  
  
[Quota] LLM=3.930428 | GPU=0.0 | Used=3.930428/50.0 | Remaining=46.069572 |  
Usage=7.860855999999999%  
34%|           | 484/1433 [36:39<1:12:04,  4.56s/it]  
  
[Quota] LLM=3.938036 | GPU=0.0 | Used=3.938036/50.0 | Remaining=46.061964 |  
Usage=7.87607199999999%  
34%|           | 485/1433 [36:43<1:10:05,  4.44s/it]  
  
[Quota] LLM=3.945815 | GPU=0.0 | Used=3.945815/50.0 | Remaining=46.054185 |  
Usage=7.89162999999999%  
34%|           | 486/1433 [36:47<1:09:29,  4.40s/it]  
  
[Quota] LLM=3.95381 | GPU=0.0 | Used=3.95381/50.0 | Remaining=46.04619 |  
Usage=7.90762%  
34%|           | 487/1433 [36:53<1:16:40,  4.86s/it]  
  
[Quota] LLM=3.9616730000000002 | GPU=0.0 | Used=3.9616730000000002/50.0 |  
Remaining=46.038327 | Usage=7.9233460000000004%  
34%|           | 488/1433 [36:57<1:13:44,  4.68s/it]  
  
[Quota] LLM=3.969614 | GPU=0.0 | Used=3.969614/50.0 | Remaining=46.030386 |  
Usage=7.939228%  
34%|           | 489/1433 [37:01<1:09:14,  4.40s/it]  
  
[Quota] LLM=3.977294 | GPU=0.0 | Used=3.977294/50.0 | Remaining=46.022706 |  
Usage=7.954588%  
34%|           | 490/1433 [37:04<1:03:23,  4.03s/it]
```

```
[Quota] LLM=3.985046 | GPU=0.0 | Used=3.985046/50.0 | Remaining=46.014954 |
Usage=7.970092000000001%
34%|           | 491/1433 [37:08<1:03:38,  4.05s/it]

[Quota] LLM=3.992684 | GPU=0.0 | Used=3.992684/50.0 | Remaining=46.007316 |
Usage=7.98536799999999%
34%|           | 492/1433 [37:13<1:05:03,  4.15s/it]

[Quota] LLM=4.00097 | GPU=0.0 | Used=4.00097/50.0 | Remaining=45.99903 |
Usage=8.00194%
34%|           | 493/1433 [37:16<1:02:34,  3.99s/it]

[Quota] LLM=4.008767 | GPU=0.0 | Used=4.008767/50.0 | Remaining=45.991233 |
Usage=8.017534%
34%|           | 494/1433 [37:20<59:21,  3.79s/it]

[Quota] LLM=4.016549 | GPU=0.0 | Used=4.016549/50.0 | Remaining=45.983451 |
Usage=8.033098%
35%|           | 495/1433 [37:25<1:08:52,  4.41s/it]

[Quota] LLM=4.024379 | GPU=0.0 | Used=4.024379/50.0 |
Remaining=45.97562100000004 | Usage=8.048758%
35%|           | 496/1433 [37:30<1:07:22,  4.31s/it]

[Quota] LLM=4.032173 | GPU=0.0 | Used=4.032173/50.0 | Remaining=45.967827 |
Usage=8.064346%
35%|           | 497/1433 [37:34<1:06:23,  4.26s/it]

[Quota] LLM=4.039976 | GPU=0.0 | Used=4.039976/50.0 | Remaining=45.960024 |
Usage=8.079952%
35%|           | 498/1433 [37:38<1:05:29,  4.20s/it]

[Quota] LLM=4.047803 | GPU=0.0 | Used=4.047803/50.0 | Remaining=45.952197 |
Usage=8.095606%
35%|           | 499/1433 [37:42<1:05:57,  4.24s/it]

[Quota] LLM=4.055657 | GPU=0.0 | Used=4.055657/50.0 | Remaining=45.944343 |
Usage=8.111314%
35%|           | 500/1433 [37:46<1:03:40,  4.09s/it]

[Quota] LLM=4.063382 | GPU=0.0 | Used=4.063382/50.0 | Remaining=45.936618 |
Usage=8.126764%
35%|           | 501/1433 [37:49<1:01:40,  3.97s/it]

[Quota] LLM=4.071005 | GPU=0.0 | Used=4.071005/50.0 | Remaining=45.928995 |
Usage=8.14201%
35%|           | 502/1433 [37:54<1:02:01,  4.00s/it]
```

```
[Quota] LLM=4.079018 | GPU=0.0 | Used=4.079018/50.0 | Remaining=45.920982 |
Usage=8.158036%
35%|           | 503/1433 [37:59<1:10:19,  4.54s/it]

[Quota] LLM=4.08716 | GPU=0.0 | Used=4.08716/50.0 | Remaining=45.91284 |
Usage=8.17432%
35%|           | 504/1433 [38:04<1:11:07,  4.59s/it]

[Quota] LLM=4.095008 | GPU=0.0 | Used=4.095008/50.0 | Remaining=45.904992 |
Usage=8.190016%
35%|           | 505/1433 [38:09<1:13:49,  4.77s/it]

[Quota] LLM=4.102919 | GPU=0.0 | Used=4.102919/50.0 | Remaining=45.897081 |
Usage=8.205838%
35%|           | 506/1433 [38:14<1:11:34,  4.63s/it]

[Quota] LLM=4.110692 | GPU=0.0 | Used=4.110692/50.0 | Remaining=45.889308 |
Usage=8.221384%
35%|           | 507/1433 [38:19<1:16:39,  4.97s/it]

[Quota] LLM=4.118618 | GPU=0.0 | Used=4.118618/50.0 | Remaining=45.881382 |
Usage=8.237236%
35%|           | 508/1433 [38:24<1:14:06,  4.81s/it]

[Quota] LLM=4.126271 | GPU=0.0 | Used=4.126271/50.0 | Remaining=45.873729 |
Usage=8.252542%
36%|           | 509/1433 [38:28<1:12:47,  4.73s/it]

[Quota] LLM=4.134032 | GPU=0.0 | Used=4.134032/50.0 | Remaining=45.865968 |
Usage=8.268064%
36%|           | 510/1433 [38:34<1:16:43,  4.99s/it]

[Quota] LLM=4.14182 | GPU=0.0 | Used=4.14182/50.0 | Remaining=45.85818 |
Usage=8.28364%
36%|           | 511/1433 [38:39<1:16:56,  5.01s/it]

[Quota] LLM=4.149749 | GPU=0.0 | Used=4.149749/50.0 | Remaining=45.850251 |
Usage=8.299498%
36%|           | 512/1433 [38:43<1:13:56,  4.82s/it]

[Quota] LLM=4.15766 | GPU=0.0 | Used=4.15766/50.0 | Remaining=45.84234 |
Usage=8.31532%
36%|           | 513/1433 [38:47<1:10:14,  4.58s/it]

[Quota] LLM=4.165607 | GPU=0.0 | Used=4.165607/50.0 | Remaining=45.834393 |
Usage=8.331214%
36%|           | 514/1433 [38:51<1:06:52,  4.37s/it]
```

```
[Quota] LLM=4.173506 | GPU=0.0 | Used=4.173506/50.0 | Remaining=45.826494 |  
Usage=8.347012%  
36%|           | 515/1433 [38:57<1:14:23,  4.86s/it]  
  
[Quota] LLM=4.181213 | GPU=0.0 | Used=4.181213/50.0 | Remaining=45.818787 |  
Usage=8.362426%  
36%|           | 516/1433 [39:02<1:12:49,  4.76s/it]  
  
[Quota] LLM=4.188845 | GPU=0.0 | Used=4.188845/50.0 | Remaining=45.811155 |  
Usage=8.37769%  
36%|           | 517/1433 [39:05<1:07:26,  4.42s/it]  
  
[Quota] LLM=4.19663 | GPU=0.0 | Used=4.19663/50.0 | Remaining=45.80337 |  
Usage=8.39326%  
36%|           | 518/1433 [39:09<1:05:32,  4.30s/it]  
  
[Quota] LLM=4.204511 | GPU=0.0 | Used=4.204511/50.0 | Remaining=45.795489 |  
Usage=8.409022%  
36%|           | 519/1433 [39:14<1:05:17,  4.29s/it]  
  
[Quota] LLM=4.212644 | GPU=0.0 | Used=4.212644/50.0 | Remaining=45.787356 |  
Usage=8.425288%  
36%|           | 520/1433 [39:18<1:05:14,  4.29s/it]  
  
[Quota] LLM=4.220465 | GPU=0.0 | Used=4.220465/50.0 | Remaining=45.779535 |  
Usage=8.44093%  
36%|           | 521/1433 [39:22<1:03:10,  4.16s/it]  
  
[Quota] LLM=4.228088 | GPU=0.0 | Used=4.228088/50.0 | Remaining=45.771912 |  
Usage=8.456176%  
36%|           | 522/1433 [39:25<1:00:02,  3.95s/it]  
  
[Quota] LLM=4.235774 | GPU=0.0 | Used=4.235774/50.0 | Remaining=45.764226 |  
Usage=8.471548%  
36%|           | 523/1433 [39:31<1:07:05,  4.42s/it]  
  
[Quota] LLM=4.243751 | GPU=0.0 | Used=4.243751/50.0 | Remaining=45.756249 |  
Usage=8.487502%  
37%|           | 524/1433 [39:35<1:05:27,  4.32s/it]  
  
[Quota] LLM=4.251632 | GPU=0.0 | Used=4.251632/50.0 | Remaining=45.748368 |  
Usage=8.503264%  
37%|           | 525/1433 [39:39<1:02:54,  4.16s/it]  
  
[Quota] LLM=4.259567 | GPU=0.0 | Used=4.259567/50.0 | Remaining=45.740433 |  
Usage=8.519134%  
37%|           | 526/1433 [39:44<1:07:43,  4.48s/it]
```

```
[Quota] LLM=4.267664 | GPU=0.0 | Used=4.267664/50.0 |  
Remaining=45.732336000000004 | Usage=8.535328%  
37%|      | 527/1433 [39:48<1:07:40, 4.48s/it]  
  
[Quota] LLM=4.275404 | GPU=0.0 | Used=4.275404/50.0 | Remaining=45.724596 |  
Usage=8.550808%  
37%|      | 528/1433 [39:53<1:06:31, 4.41s/it]  
  
[Quota] LLM=4.283432 | GPU=0.0 | Used=4.283432/50.0 | Remaining=45.716568 |  
Usage=8.566864%  
37%|      | 529/1433 [39:56<1:03:18, 4.20s/it]  
  
[Quota] LLM=4.291004 | GPU=0.0 | Used=4.291004/50.0 | Remaining=45.708996 |  
Usage=8.582008%  
37%|      | 530/1433 [40:01<1:04:31, 4.29s/it]  
  
[Quota] LLM=4.298846 | GPU=0.0 | Used=4.298846/50.0 | Remaining=45.701154 |  
Usage=8.597692%  
37%|      | 531/1433 [40:05<1:05:44, 4.37s/it]  
  
[Quota] LLM=4.306814 | GPU=0.0 | Used=4.306814/50.0 | Remaining=45.693186 |  
Usage=8.613628%  
37%|      | 532/1433 [40:10<1:06:26, 4.42s/it]  
  
[Quota] LLM=4.314569 | GPU=0.0 | Used=4.314569/50.0 | Remaining=45.685431 |  
Usage=8.629138%  
37%|      | 533/1433 [40:14<1:05:12, 4.35s/it]  
  
[Quota] LLM=4.322339 | GPU=0.0 | Used=4.322339/50.0 | Remaining=45.677661 |  
Usage=8.644678%  
37%|      | 534/1433 [40:21<1:15:24, 5.03s/it]  
  
[Quota] LLM=4.329965 | GPU=0.0 | Used=4.329965/50.0 | Remaining=45.670035 |  
Usage=8.65993%  
37%|      | 535/1433 [40:24<1:09:27, 4.64s/it]  
  
[Quota] LLM=4.337711 | GPU=0.0 | Used=4.337711/50.0 | Remaining=45.662289 |  
Usage=8.675422%  
37%|      | 536/1433 [40:29<1:07:05, 4.49s/it]  
  
[Quota] LLM=4.345667 | GPU=0.0 | Used=4.345667/50.0 | Remaining=45.654333 |  
Usage=8.691334%  
37%|      | 537/1433 [40:33<1:07:20, 4.51s/it]  
  
[Quota] LLM=4.353503 | GPU=0.0 | Used=4.353503/50.0 | Remaining=45.646497 |  
Usage=8.707006%  
38%|      | 538/1433 [40:38<1:10:21, 4.72s/it]
```

```
[Quota] LLM=4.361393 | GPU=0.0 | Used=4.361393/50.0 | Remaining=45.638607 |  
Usage=8.722786%  
38%|           | 539/1433 [40:42<1:06:44,  4.48s/it]  
  
[Quota] LLM=4.36925 | GPU=0.0 | Used=4.36925/50.0 | Remaining=45.63075 |  
Usage=8.7385%  
38%|           | 540/1433 [40:47<1:08:42,  4.62s/it]  
  
[Quota] LLM=4.376984 | GPU=0.0 | Used=4.376984/50.0 | Remaining=45.623016 |  
Usage=8.753968%  
38%|           | 541/1433 [40:52<1:08:37,  4.62s/it]  
  
[Quota] LLM=4.385159 | GPU=0.0 | Used=4.385159/50.0 | Remaining=45.614841 |  
Usage=8.770318%  
38%|           | 542/1433 [40:56<1:06:59,  4.51s/it]  
  
[Quota] LLM=4.393037 | GPU=0.0 | Used=4.393037/50.0 | Remaining=45.606963 |  
Usage=8.786074%  
38%|           | 543/1433 [41:00<1:02:19,  4.20s/it]  
  
[Quota] LLM=4.400828 | GPU=0.0 | Used=4.400828/50.0 | Remaining=45.599172 |  
Usage=8.801656%  
38%|           | 544/1433 [41:05<1:05:32,  4.42s/it]  
  
[Quota] LLM=4.408718 | GPU=0.0 | Used=4.408718/50.0 | Remaining=45.591282 |  
Usage=8.817436%  
38%|           | 545/1433 [41:09<1:07:08,  4.54s/it]  
  
[Quota] LLM=4.416818 | GPU=0.0 | Used=4.416818/50.0 | Remaining=45.583182 |  
Usage=8.833636%  
38%|           | 546/1433 [41:13<1:02:44,  4.24s/it]  
  
[Quota] LLM=4.424381 | GPU=0.0 | Used=4.424381/50.0 | Remaining=45.575619 |  
Usage=8.848762%  
38%|           | 547/1433 [41:17<1:00:28,  4.10s/it]  
  
[Quota] LLM=4.432019 | GPU=0.0 | Used=4.432019/50.0 | Remaining=45.567981 |  
Usage=8.864038%  
38%|           | 548/1433 [41:20<57:16,  3.88s/it]  
  
[Quota] LLM=4.439669 | GPU=0.0 | Used=4.439669/50.0 | Remaining=45.560331 |  
Usage=8.879338%  
38%|           | 549/1433 [41:24<57:22,  3.89s/it]  
  
[Quota] LLM=4.447697 | GPU=0.0 | Used=4.447697/50.0 | Remaining=45.552303 |  
Usage=8.895394%  
38%|           | 550/1433 [41:30<1:06:05,  4.49s/it]
```

```
[Quota] LLM=4.455335 | GPU=0.0 | Used=4.455335/50.0 | Remaining=45.544665 |
Usage=8.91067%
38%|           | 551/1433 [41:36<1:14:36,  5.07s/it]

[Quota] LLM=4.463129 | GPU=0.0 | Used=4.463129/50.0 | Remaining=45.536871 |
Usage=8.926258%
39%|           | 552/1433 [41:39<1:05:17,  4.45s/it]

[Quota] LLM=4.470746 | GPU=0.0 | Used=4.470746/50.0 | Remaining=45.529254 |
Usage=8.941492%
39%|           | 553/1433 [41:43<1:04:11,  4.38s/it]

[Quota] LLM=4.478579 | GPU=0.0 | Used=4.478579/50.0 |
Remaining=45.521421000000004 | Usage=8.957158%
39%|           | 554/1433 [41:47<1:01:36,  4.21s/it]

[Quota] LLM=4.486433 | GPU=0.0 | Used=4.486433/50.0 | Remaining=45.513567 |
Usage=8.972866%
39%|           | 555/1433 [41:51<1:01:29,  4.20s/it]

[Quota] LLM=4.494167 | GPU=0.0 | Used=4.494167/50.0 | Remaining=45.505833 |
Usage=8.988334%
39%|           | 556/1433 [41:55<59:39,  4.08s/it]

[Quota] LLM=4.50188 | GPU=0.0 | Used=4.50188/50.0 | Remaining=45.49812 |
Usage=9.00376%
39%|           | 557/1433 [41:59<58:34,  4.01s/it]

[Quota] LLM=4.5098 | GPU=0.0 | Used=4.5098/50.0 | Remaining=45.4902 |
Usage=9.0196%
39%|           | 558/1433 [42:03<59:38,  4.09s/it]

[Quota] LLM=4.518089 | GPU=0.0 | Used=4.518089/50.0 | Remaining=45.481911 |
Usage=9.036178%
39%|           | 559/1433 [42:07<58:08,  3.99s/it]

[Quota] LLM=4.525835 | GPU=0.0 | Used=4.525835/50.0 | Remaining=45.474165 |
Usage=9.05167%
39%|           | 560/1433 [42:11<55:46,  3.83s/it]

[Quota] LLM=4.533506 | GPU=0.0 | Used=4.533506/50.0 | Remaining=45.466494 |
Usage=9.067012%
39%|           | 561/1433 [42:15<57:19,  3.94s/it]

[Quota] LLM=4.541138 | GPU=0.0 | Used=4.541138/50.0 |
Remaining=45.458861999999996 | Usage=9.082276%
39%|           | 562/1433 [42:20<1:02:41,  4.32s/it]
```

```
[Quota] LLM=4.5490070000000005 | GPU=0.0 | Used=4.5490070000000005/50.0 |  
Remaining=45.450993 | Usage=9.098014000000001%  
39%|      | 563/1433 [42:24<59:37,  4.11s/it]  
  
[Quota] LLM=4.556747 | GPU=0.0 | Used=4.556747/50.0 | Remaining=45.443253 |  
Usage=9.113494%  
39%|      | 564/1433 [42:28<1:01:53,  4.27s/it]  
  
[Quota] LLM=4.564445 | GPU=0.0 | Used=4.564445/50.0 | Remaining=45.435555 |  
Usage=9.12889%  
39%|      | 565/1433 [42:32<1:00:28,  4.18s/it]  
  
[Quota] LLM=4.572179 | GPU=0.0 | Used=4.572179/50.0 | Remaining=45.427821 |  
Usage=9.144358%  
39%|      | 566/1433 [42:36<59:27,  4.11s/it]  
  
[Quota] LLM=4.579742 | GPU=0.0 | Used=4.579742/50.0 | Remaining=45.420258 |  
Usage=9.159484%  
40%|      | 567/1433 [42:42<1:04:48,  4.49s/it]  
  
[Quota] LLM=4.587731 | GPU=0.0 | Used=4.587731/50.0 | Remaining=45.412269 |  
Usage=9.175462%  
40%|      | 568/1433 [42:45<1:02:13,  4.32s/it]  
  
[Quota] LLM=4.595645 | GPU=0.0 | Used=4.595645/50.0 | Remaining=45.404355 |  
Usage=9.19129%  
40%|      | 569/1433 [42:50<1:04:37,  4.49s/it]  
  
[Quota] LLM=4.60349 | GPU=0.0 | Used=4.60349/50.0 | Remaining=45.39651 |  
Usage=9.20698%  
40%|      | 570/1433 [42:54<1:01:59,  4.31s/it]  
  
[Quota] LLM=4.611425 | GPU=0.0 | Used=4.611425/50.0 | Remaining=45.388575 |  
Usage=9.22285%  
40%|      | 571/1433 [42:59<1:03:50,  4.44s/it]  
  
[Quota] LLM=4.619309 | GPU=0.0 | Used=4.619309/50.0 | Remaining=45.380691 |  
Usage=9.238618%  
40%|      | 572/1433 [43:02<59:10,  4.12s/it]  
  
[Quota] LLM=4.627136 | GPU=0.0 | Used=4.627136/50.0 | Remaining=45.372864 |  
Usage=9.254272%  
40%|      | 573/1433 [43:06<58:49,  4.10s/it]  
  
[Quota] LLM=4.635131 | GPU=0.0 | Used=4.635131/50.0 | Remaining=45.364869 |  
Usage=9.270262%  
40%|      | 574/1433 [43:09<54:13,  3.79s/it]
```

```
[Quota] LLM=4.642829 | GPU=0.0 | Used=4.642829/50.0 | Remaining=45.357171 |
Usage=9.285658%
40%|           | 575/1433 [43:15<1:00:17,  4.22s/it]

[Quota] LLM=4.650641 | GPU=0.0 | Used=4.650641/50.0 | Remaining=45.349359 |
Usage=9.301282%
40%|           | 576/1433 [43:19<59:58,  4.20s/it]

[Quota] LLM=4.658411 | GPU=0.0 | Used=4.658411/50.0 | Remaining=45.341589 |
Usage=9.316822%
40%|           | 577/1433 [43:23<1:00:45,  4.26s/it]

[Quota] LLM=4.666397 | GPU=0.0 | Used=4.666397/50.0 | Remaining=45.333603 |
Usage=9.332794%
40%|           | 578/1433 [43:28<1:00:49,  4.27s/it]

[Quota] LLM=4.674305 | GPU=0.0 | Used=4.674305/50.0 | Remaining=45.325694999999996 |
Usage=9.34861%
40%|           | 579/1433 [43:32<1:00:40,  4.26s/it]

[Quota] LLM=4.682132 | GPU=0.0 | Used=4.682132/50.0 | Remaining=45.317868 |
Usage=9.364264%
40%|           | 580/1433 [43:36<1:00:10,  4.23s/it]

[Quota] LLM=4.6901 | GPU=0.0 | Used=4.6901/50.0 | Remaining=45.3099 |
Usage=9.3802%
41%|           | 581/1433 [43:41<1:03:17,  4.46s/it]

[Quota] LLM=4.698188 | GPU=0.0 | Used=4.698188/50.0 | Remaining=45.301812 |
Usage=9.396376%
41%|           | 582/1433 [43:46<1:06:12,  4.67s/it]

[Quota] LLM=4.706177 | GPU=0.0 | Used=4.706177/50.0 | Remaining=45.293823 |
Usage=9.412354%
41%|           | 583/1433 [43:50<1:03:14,  4.46s/it]

[Quota] LLM=4.714004 | GPU=0.0 | Used=4.714004/50.0 | Remaining=45.285996 |
Usage=9.428008%
41%|           | 584/1433 [43:56<1:08:48,  4.86s/it]

[Quota] LLM=4.721696 | GPU=0.0 | Used=4.721696/50.0 | Remaining=45.278304 |
Usage=9.443392%
41%|           | 585/1433 [44:01<1:08:01,  4.81s/it]

[Quota] LLM=4.72958 | GPU=0.0 | Used=4.72958/50.0 | Remaining=45.27042 |
Usage=9.45916%
41%|           | 586/1433 [44:04<1:02:12,  4.41s/it]
```

```
[Quota] LLM=4.737392 | GPU=0.0 | Used=4.737392/50.0 | Remaining=45.262608 |
Usage=9.474784%
41%|           | 587/1433 [44:09<1:03:11,  4.48s/it]

[Quota] LLM=4.745027 | GPU=0.0 | Used=4.745027/50.0 | Remaining=45.254973 |
Usage=9.490054%
41%|           | 588/1433 [44:13<1:03:12,  4.49s/it]

[Quota] LLM=4.753109 | GPU=0.0 | Used=4.753109/50.0 | Remaining=45.246891 |
Usage=9.506218%
41%|           | 589/1433 [44:19<1:09:51,  4.97s/it]

[Quota] LLM=4.761419 | GPU=0.0 | Used=4.761419/50.0 |
Remaining=45.238580999999996 | Usage=9.522838%
41%|           | 590/1433 [44:26<1:17:38,  5.53s/it]

[Quota] LLM=4.769393 | GPU=0.0 | Used=4.769393/50.0 | Remaining=45.230607 |
Usage=9.538786%
41%|           | 591/1433 [44:30<1:12:33,  5.17s/it]

[Quota] LLM=4.777097 | GPU=0.0 | Used=4.777097/50.0 | Remaining=45.222903 |
Usage=9.554194%
41%|           | 592/1433 [44:35<1:11:11,  5.08s/it]

[Quota] LLM=4.785254 | GPU=0.0 | Used=4.785254/50.0 | Remaining=45.214746 |
Usage=9.570508%
41%|           | 593/1433 [44:40<1:07:51,  4.85s/it]

[Quota] LLM=4.793303 | GPU=0.0 | Used=4.793303/50.0 | Remaining=45.206697 |
Usage=9.586606%
41%|           | 594/1433 [44:46<1:12:39,  5.20s/it]

[Quota] LLM=4.801364 | GPU=0.0 | Used=4.801364/50.0 | Remaining=45.198636 |
Usage=9.602728%
42%|           | 595/1433 [44:50<1:10:51,  5.07s/it]

[Quota] LLM=4.8095930000000005 | GPU=0.0 | Used=4.8095930000000005/50.0 |
Remaining=45.190407 | Usage=9.619186000000001%
42%|           | 596/1433 [44:55<1:09:22,  4.97s/it]

[Quota] LLM=4.817363 | GPU=0.0 | Used=4.817363/50.0 | Remaining=45.182637 |
Usage=9.634726%
42%|           | 597/1433 [44:59<1:03:57,  4.59s/it]

[Quota] LLM=4.825211 | GPU=0.0 | Used=4.825211/50.0 | Remaining=45.174789 |
Usage=9.650422%
42%|           | 598/1433 [45:04<1:06:06,  4.75s/it]
```

[Quota] LLM=4.833341 | GPU=0.0 | Used=4.833341/50.0 | Remaining=45.166659 | Usage=9.666682%

42% | 599/1433 [45:09<1:06:48, 4.81s/it]

[Quota] LLM=4.841462 | GPU=0.0 | Used=4.841462/50.0 | Remaining=45.158538 | Usage=9.682924%

42% | 600/1433 [45:14<1:07:21, 4.85s/it]

[Quota] LLM=4.849445 | GPU=0.0 | Used=4.849445/50.0 | Remaining=45.150555 | Usage=9.69889%

42% | 601/1433 [45:20<1:11:24, 5.15s/it]

[Quota] LLM=4.857509 | GPU=0.0 | Used=4.857509/50.0 | Remaining=45.142491 | Usage=9.715018%

42% | 602/1433 [45:24<1:09:06, 4.99s/it]

[Quota] LLM=4.865411 | GPU=0.0 | Used=4.865411/50.0 | Remaining=45.134589 | Usage=9.730822%

42% | 603/1433 [45:29<1:06:46, 4.83s/it]

[Quota] LLM=4.87343 | GPU=0.0 | Used=4.87343/50.0 | Remaining=45.12657 | Usage=9.74686%

42% | 604/1433 [45:33<1:03:54, 4.63s/it]

[Quota] LLM=4.88156 | GPU=0.0 | Used=4.88156/50.0 | Remaining=45.11844 | Usage=9.76312%

42% | 605/1433 [45:37<1:01:54, 4.49s/it]

[Quota] LLM=4.889441 | GPU=0.0 | Used=4.889441/50.0 | Remaining=45.110559 | Usage=9.778882%

42% | 606/1433 [45:41<59:28, 4.32s/it]

[Quota] LLM=4.897388 | GPU=0.0 | Used=4.897388/50.0 | Remaining=45.102612 | Usage=9.794776%

42% | 607/1433 [45:45<59:31, 4.32s/it]

[Quota] LLM=4.905218 | GPU=0.0 | Used=4.905218/50.0 | Remaining=45.094782 | Usage=9.810436%

42% | 608/1433 [45:50<1:02:35, 4.55s/it]

[Quota] LLM=4.913048 | GPU=0.0 | Used=4.913048/50.0 | Remaining=45.086952 | Usage=9.826096%

42% | 609/1433 [45:55<1:04:08, 4.67s/it]

[Quota] LLM=4.921085 | GPU=0.0 | Used=4.921085/50.0 | Remaining=45.078915 | Usage=9.84217%

43% | 610/1433 [46:00<1:03:13, 4.61s/it]

[Quota] LLM=4.928855 | GPU=0.0 | Used=4.928855/50.0 | Remaining=45.071145 | Usage=9.85771%

43% | 611/1433 [46:06<1:10:45, 5.16s/it]

[Quota] LLM=4.936991 | GPU=0.0 | Used=4.936991/50.0 | Remaining=45.063009 | Usage=9.873982%

43% | 612/1433 [46:11<1:10:21, 5.14s/it]

[Quota] LLM=4.944839 | GPU=0.0 | Used=4.944839/50.0 | Remaining=45.055161 | Usage=9.889678%

43% | 613/1433 [46:15<1:04:36, 4.73s/it]

[Quota] LLM=4.952591 | GPU=0.0 | Used=4.952591/50.0 | Remaining=45.047409 | Usage=9.905182%

43% | 614/1433 [46:20<1:03:10, 4.63s/it]

[Quota] LLM=4.960448 | GPU=0.0 | Used=4.960448/50.0 | Remaining=45.039552 | Usage=9.920896%

43% | 615/1433 [46:23<59:06, 4.34s/it]

[Quota] LLM=4.968272 | GPU=0.0 | Used=4.968272/50.0 | Remaining=45.031728 | Usage=9.936544%

43% | 616/1433 [46:29<1:03:32, 4.67s/it]

[Quota] LLM=4.976591 | GPU=0.0 | Used=4.976591/50.0 | Remaining=45.023409 | Usage=9.953182%

43% | 617/1433 [46:34<1:05:40, 4.83s/it]

[Quota] LLM=4.984562 | GPU=0.0 | Used=4.984562/50.0 | Remaining=45.015438 | Usage=9.969124%

43% | 618/1433 [46:39<1:05:21, 4.81s/it]

[Quota] LLM=4.992809 | GPU=0.0 | Used=4.992809/50.0 | Remaining=45.007191 | Usage=9.985618%

43% | 619/1433 [46:44<1:05:59, 4.86s/it]

[Quota] LLM=5.00087 | GPU=0.0 | Used=5.00087/50.0 | Remaining=44.99913 | Usage=10.00174%

43% | 620/1433 [46:49<1:07:10, 4.96s/it]

[Quota] LLM=5.008631 | GPU=0.0 | Used=5.008631/50.0 | Remaining=44.991369 | Usage=10.017262%

43% | 621/1433 [46:53<1:02:33, 4.62s/it]

[Quota] LLM=5.016485 | GPU=0.0 | Used=5.016485/50.0 | Remaining=44.983515 | Usage=10.03297%

43% | 622/1433 [46:58<1:03:36, 4.71s/it]

[Quota] LLM=5.024567 | GPU=0.0 | Used=5.024567/50.0 | Remaining=44.975433 | Usage=10.049134%

43% | 623/1433 [47:02<1:03:20, 4.69s/it]

[Quota] LLM=5.03252 | GPU=0.0 | Used=5.03252/50.0 | Remaining=44.96748 | Usage=10.06504%

44% | 624/1433 [47:06<59:27, 4.41s/it]

[Quota] LLM=5.040086 | GPU=0.0 | Used=5.040086/50.0 | Remaining=44.959914 | Usage=10.080172%

44% | 625/1433 [47:09<54:41, 4.06s/it]

[Quota] LLM=5.047646 | GPU=0.0 | Used=5.047646/50.0 | Remaining=44.952354 | Usage=10.095292%

44% | 626/1433 [47:14<57:31, 4.28s/it]

[Quota] LLM=5.055644 | GPU=0.0 | Used=5.055644/50.0 | Remaining=44.944356 | Usage=10.111288%

44% | 627/1433 [47:18<57:01, 4.25s/it]

[Quota] LLM=5.063537 | GPU=0.0 | Used=5.063537/50.0 | Remaining=44.936463 | Usage=10.127074%

44% | 628/1433 [47:22<55:49, 4.16s/it]

[Quota] LLM=5.0714 | GPU=0.0 | Used=5.0714/50.0 | Remaining=44.9286 | Usage=10.1428%

44% | 629/1433 [47:26<54:52, 4.10s/it]

[Quota] LLM=5.079359 | GPU=0.0 | Used=5.079359/50.0 | Remaining=44.920641 | Usage=10.158718%

44% | 630/1433 [47:30<52:34, 3.93s/it]

[Quota] LLM=5.08703 | GPU=0.0 | Used=5.08703/50.0 | Remaining=44.91297 | Usage=10.17406%

44% | 631/1433 [47:34<55:26, 4.15s/it]

[Quota] LLM=5.094887 | GPU=0.0 | Used=5.094887/50.0 | Remaining=44.905113 | Usage=10.189774%

44% | 632/1433 [47:39<56:13, 4.21s/it]

[Quota] LLM=5.102798 | GPU=0.0 | Used=5.102798/50.0 | Remaining=44.897202 | Usage=10.205596%

44% | 633/1433 [47:43<56:45, 4.26s/it]

[Quota] LLM=5.110787 | GPU=0.0 | Used=5.110787/50.0 | Remaining=44.889213 | Usage=10.221574%

44% | 634/1433 [47:47<54:08, 4.07s/it]

[Quota] LLM=5.118548 | GPU=0.0 | Used=5.118548/50.0 | Remaining=44.881452 | Usage=10.237096%

44% | 635/1433 [47:51<54:40, 4.11s/it]

[Quota] LLM=5.126354 | GPU=0.0 | Used=5.126354/50.0 | Remaining=44.873646 | Usage=10.252708%

44% | 636/1433 [47:55<54:20, 4.09s/it]

[Quota] LLM=5.134214 | GPU=0.0 | Used=5.134214/50.0 | Remaining=44.865786 | Usage=10.268428%

44% | 637/1433 [48:00<57:53, 4.36s/it]

[Quota] LLM=5.142065 | GPU=0.0 | Used=5.142065/50.0 | Remaining=44.857935 | Usage=10.28413%

45% | 638/1433 [48:09<1:15:24, 5.69s/it]

[Quota] LLM=5.149889 | GPU=0.0 | Used=5.149889/50.0 | Remaining=44.850111 | Usage=10.299778%

45% | 639/1433 [48:16<1:20:51, 6.11s/it]

[Quota] LLM=5.157818 | GPU=0.0 | Used=5.157818/50.0 | Remaining=44.842182 | Usage=10.315636%

45% | 640/1433 [48:19<1:09:24, 5.25s/it]

[Quota] LLM=5.165525 | GPU=0.0 | Used=5.165525/50.0 | Remaining=44.834475 | Usage=10.33105%

45% | 641/1433 [48:25<1:11:46, 5.44s/it]

[Quota] LLM=5.174027 | GPU=0.0 | Used=5.174027/50.0 | Remaining=44.825973 | Usage=10.348054%

45% | 642/1433 [48:30<1:10:04, 5.32s/it]

[Quota] LLM=5.182037 | GPU=0.0 | Used=5.182037/50.0 | Remaining=44.817963 | Usage=10.364074%

45% | 643/1433 [48:53<2:19:38, 10.61s/it]

[Quota] LLM=5.1906740000000005 | GPU=0.0 | Used=5.1906740000000005/50.0 | Remaining=44.809326 | Usage=10.381348000000001%

45% | 644/1433 [48:56<1:51:10, 8.45s/it]

[Quota] LLM=5.198474 | GPU=0.0 | Used=5.198474/50.0 | Remaining=44.801526 | Usage=10.396948%

45% | 645/1433 [49:01<1:36:40, 7.36s/it]

[Quota] LLM=5.206382 | GPU=0.0 | Used=5.206382/50.0 | Remaining=44.793618 | Usage=10.412764%

45% | 646/1433 [49:05<1:22:35, 6.30s/it]

```
[Quota] LLM=5.214413 | GPU=0.0 | Used=5.214413/50.0 | Remaining=44.785587 |
Usage=10.428826%
45%|           | 647/1433 [49:09<1:13:43,  5.63s/it]
[Quota] LLM=5.222234 | GPU=0.0 | Used=5.222234/50.0 | Remaining=44.777766 |
Usage=10.444468%
45%|           | 648/1433 [49:13<1:08:28,  5.23s/it]
[Quota] LLM=5.23016 | GPU=0.0 | Used=5.23016/50.0 | Remaining=44.76984 |
Usage=10.46032%
45%|           | 649/1433 [49:19<1:08:53,  5.27s/it]
[Quota] LLM=5.23823 | GPU=0.0 | Used=5.23823/50.0 | Remaining=44.76177 |
Usage=10.47646%
45%|           | 650/1433 [49:22<1:02:34,  4.80s/it]
[Quota] LLM=5.24606 | GPU=0.0 | Used=5.24606/50.0 | Remaining=44.75394 |
Usage=10.49212%
45%|           | 651/1433 [49:26<58:18,  4.47s/it]
[Quota] LLM=5.253872 | GPU=0.0 | Used=5.253872/50.0 | Remaining=44.746128 |
Usage=10.507744%
45%|           | 652/1433 [49:30<58:01,  4.46s/it]
[Quota] LLM=5.261885 | GPU=0.0 | Used=5.261885/50.0 | Remaining=44.738115 |
Usage=10.52377%
46%|           | 653/1433 [49:36<1:01:32,  4.73s/it]
[Quota] LLM=5.269802 | GPU=0.0 | Used=5.269802/50.0 | Remaining=44.730198 |
Usage=10.539604%
46%|           | 654/1433 [49:41<1:03:53,  4.92s/it]
[Quota] LLM=5.277509 | GPU=0.0 | Used=5.277509/50.0 | Remaining=44.722491 |
Usage=10.555018%
46%|           | 655/1433 [49:47<1:08:24,  5.28s/it]
[Quota] LLM=5.285606 | GPU=0.0 | Used=5.285606/50.0 | Remaining=44.714394 |
Usage=10.571212%
46%|           | 656/1433 [49:52<1:04:40,  4.99s/it]
[Quota] LLM=5.293715 | GPU=0.0 | Used=5.293715/50.0 | Remaining=44.706285 |
Usage=10.58743%
46%|           | 657/1433 [49:56<1:03:41,  4.93s/it]
[Quota] LLM=5.301698 | GPU=0.0 | Used=5.301698/50.0 | Remaining=44.698302 |
Usage=10.603396%
46%|           | 658/1433 [50:01<1:02:16,  4.82s/it]
```

```
[Quota] LLM=5.30951 | GPU=0.0 | Used=5.30951/50.0 | Remaining=44.69049 |
Usage=10.61902%
46%|       | 659/1433 [50:06<1:02:15,  4.83s/it]

[Quota] LLM=5.317433 | GPU=0.0 | Used=5.317433/50.0 | Remaining=44.682567 |
Usage=10.634866%
46%|       | 660/1433 [50:09<56:41,  4.40s/it]

[Quota] LLM=5.325368 | GPU=0.0 | Used=5.325368/50.0 | Remaining=44.674632 |
Usage=10.650736%
46%|       | 661/1433 [50:13<53:34,  4.16s/it]

[Quota] LLM=5.333189 | GPU=0.0 | Used=5.333189/50.0 | Remaining=44.666811 |
Usage=10.666378%
46%|       | 662/1433 [50:18<56:03,  4.36s/it]

[Quota] LLM=5.340941 | GPU=0.0 | Used=5.340941/50.0 | Remaining=44.659059 |
Usage=10.681882%
46%|       | 663/1433 [50:22<55:33,  4.33s/it]

[Quota] LLM=5.348762 | GPU=0.0 | Used=5.348762/50.0 | Remaining=44.651238 |
Usage=10.697524%
46%|       | 664/1433 [50:27<57:18,  4.47s/it]

[Quota] LLM=5.356694 | GPU=0.0 | Used=5.356694/50.0 | Remaining=44.643306 |
Usage=10.713388%
46%|       | 665/1433 [50:32<1:00:28,  4.73s/it]

[Quota] LLM=5.364572 | GPU=0.0 | Used=5.364572/50.0 | Remaining=44.635428 |
Usage=10.729144%
46%|       | 666/1433 [50:37<1:00:45,  4.75s/it]

[Quota] LLM=5.372612 | GPU=0.0 | Used=5.372612/50.0 |
Remaining=44.627387999999996 | Usage=10.745224%
47%|       | 667/1433 [50:41<57:07,  4.48s/it]

[Quota] LLM=5.380424 | GPU=0.0 | Used=5.380424/50.0 | Remaining=44.619576 |
Usage=10.760848%
47%|       | 668/1433 [50:45<57:01,  4.47s/it]

[Quota] LLM=5.388455 | GPU=0.0 | Used=5.388455/50.0 | Remaining=44.611545 |
Usage=10.77691%
47%|       | 669/1433 [50:49<55:38,  4.37s/it]

[Quota] LLM=5.396309 | GPU=0.0 | Used=5.396309/50.0 | Remaining=44.603691 |
Usage=10.792618%
47%|       | 670/1433 [50:54<56:13,  4.42s/it]
```

[Quota] LLM=5.404109 | GPU=0.0 | Used=5.404109/50.0 | Remaining=44.595891 | Usage=10.808218%

47% | 671/1433 [50:58<55:50, 4.40s/it]

[Quota] LLM=5.411825 | GPU=0.0 | Used=5.411825/50.0 | Remaining=44.588175 | Usage=10.82365%

47% | 672/1433 [51:02<54:24, 4.29s/it]

[Quota] LLM=5.419685 | GPU=0.0 | Used=5.419685/50.0 | Remaining=44.580315 | Usage=10.83937%

47% | 673/1433 [51:08<1:00:59, 4.81s/it]

[Quota] LLM=5.427968 | GPU=0.0 | Used=5.427968/50.0 | Remaining=44.572032 | Usage=10.855936%

47% | 674/1433 [51:12<57:49, 4.57s/it]

[Quota] LLM=5.435963 | GPU=0.0 | Used=5.435963/50.0 | Remaining=44.564037 | Usage=10.871926%

47% | 675/1433 [51:18<1:01:10, 4.84s/it]

[Quota] LLM=5.44373 | GPU=0.0 | Used=5.44373/50.0 | Remaining=44.55627 | Usage=10.88746%

47% | 676/1433 [51:22<1:00:15, 4.78s/it]

[Quota] LLM=5.45177 | GPU=0.0 | Used=5.45177/50.0 | Remaining=44.548230000000004 | Usage=10.90354%

47% | 677/1433 [51:27<1:00:29, 4.80s/it]

[Quota] LLM=5.459537 | GPU=0.0 | Used=5.459537/50.0 | Remaining=44.540463 | Usage=10.919074%

47% | 678/1433 [51:33<1:04:33, 5.13s/it]

[Quota] LLM=5.46734 | GPU=0.0 | Used=5.46734/50.0 | Remaining=44.53266 | Usage=10.93468%

47% | 679/1433 [51:36<57:55, 4.61s/it]

[Quota] LLM=5.475161 | GPU=0.0 | Used=5.475161/50.0 | Remaining=44.524839 | Usage=10.950322%

47% | 680/1433 [51:40<54:39, 4.35s/it]

[Quota] LLM=5.483153 | GPU=0.0 | Used=5.483153/50.0 | Remaining=44.516847 | Usage=10.966306%

48% | 681/1433 [51:45<57:31, 4.59s/it]

[Quota] LLM=5.490959 | GPU=0.0 | Used=5.490959/50.0 | Remaining=44.509040999999996 | Usage=10.981918%

48% | 682/1433 [51:52<1:05:19, 5.22s/it]

[Quota] LLM=5.499158 | GPU=0.0 | Used=5.499158/50.0 | Remaining=44.500842 | Usage=10.998316%

48%| 683/1433 [51:57<1:03:31, 5.08s/it]

[Quota] LLM=5.507084 | GPU=0.0 | Used=5.507084/50.0 | Remaining=44.492916 | Usage=11.014168%

48%| 684/1433 [52:02<1:02:33, 5.01s/it]

[Quota] LLM=5.515082 | GPU=0.0 | Used=5.515082/50.0 | Remaining=44.484918 | Usage=11.030164%

48%| 685/1433 [52:08<1:06:16, 5.32s/it]

[Quota] LLM=5.523077 | GPU=0.0 | Used=5.523077/50.0 | Remaining=44.476923 | Usage=11.046154%

48%| 686/1433 [52:13<1:06:31, 5.34s/it]

[Quota] LLM=5.531039 | GPU=0.0 | Used=5.531039/50.0 | Remaining=44.468961 | Usage=11.062078%

48%| 687/1433 [52:18<1:03:17, 5.09s/it]

[Quota] LLM=5.53919 | GPU=0.0 | Used=5.53919/50.0 | Remaining=44.46081 | Usage=11.07838%

48%| 688/1433 [52:22<59:52, 4.82s/it]

[Quota] LLM=5.547134 | GPU=0.0 | Used=5.547134/50.0 | Remaining=44.452866 | Usage=11.094268%

48%| 689/1433 [52:27<1:01:22, 4.95s/it]

[Quota] LLM=5.554991 | GPU=0.0 | Used=5.554991/50.0 | Remaining=44.445009 | Usage=11.109982%

48%| 690/1433 [52:31<57:36, 4.65s/it]

[Quota] LLM=5.562665 | GPU=0.0 | Used=5.562665/50.0 | Remaining=44.437335 | Usage=11.12533%

48%| 691/1433 [52:36<57:41, 4.67s/it]

[Quota] LLM=5.570912 | GPU=0.0 | Used=5.570912/50.0 | Remaining=44.429088 | Usage=11.141824%

48%| 692/1433 [52:40<55:46, 4.52s/it]

[Quota] LLM=5.578805 | GPU=0.0 | Used=5.578805/50.0 | Remaining=44.421195 | Usage=11.15761%

48%| 693/1433 [52:44<55:50, 4.53s/it]

[Quota] LLM=5.586758 | GPU=0.0 | Used=5.586758/50.0 | Remaining=44.413242 | Usage=11.173516%

48%| 694/1433 [52:48<53:56, 4.38s/it]

[Quota] LLM=5.59475 | GPU=0.0 | Used=5.59475/50.0 | Remaining=44.40525 | Usage=11.1895%

48%| 695/1433 [52:53<56:09, 4.57s/it]

[Quota] LLM=5.602643 | GPU=0.0 | Used=5.602643/50.0 | Remaining=44.397357 | Usage=11.205286%

49%| 696/1433 [52:58<54:13, 4.42s/it]

[Quota] LLM=5.61056 | GPU=0.0 | Used=5.61056/50.0 | Remaining=44.38944 | Usage=11.22112%

49%| 697/1433 [53:01<51:22, 4.19s/it]

[Quota] LLM=5.61821 | GPU=0.0 | Used=5.61821/50.0 | Remaining=44.38179 | Usage=11.23642%

49%| 698/1433 [53:05<51:20, 4.19s/it]

[Quota] LLM=5.626409 | GPU=0.0 | Used=5.626409/50.0 | Remaining=44.373591 | Usage=11.252818%

49%| 699/1433 [53:10<51:49, 4.24s/it]

[Quota] LLM=5.634254 | GPU=0.0 | Used=5.634254/50.0 | Remaining=44.365746 | Usage=11.268508%

49%| 700/1433 [53:15<56:31, 4.63s/it]

[Quota] LLM=5.642315 | GPU=0.0 | Used=5.642315/50.0 | Remaining=44.35768500000004 | Usage=11.28463%

49%| 701/1433 [53:20<55:41, 4.57s/it]

[Quota] LLM=5.650187 | GPU=0.0 | Used=5.650187/50.0 | Remaining=44.349813 | Usage=11.300374%

49%| 702/1433 [53:24<55:44, 4.57s/it]

[Quota] LLM=5.658101 | GPU=0.0 | Used=5.658101/50.0 | Remaining=44.341899 | Usage=11.316202%

49%| 703/1433 [53:29<55:24, 4.55s/it]

[Quota] LLM=5.666015 | GPU=0.0 | Used=5.666015/50.0 | Remaining=44.333985 | Usage=11.33203%

49%| 704/1433 [53:36<1:03:34, 5.23s/it]

[Quota] LLM=5.674148 | GPU=0.0 | Used=5.674148/50.0 | Remaining=44.325852 | Usage=11.348296%

49%| 705/1433 [53:40<59:05, 4.87s/it]

[Quota] LLM=5.682032 | GPU=0.0 | Used=5.682032/50.0 | Remaining=44.317968 | Usage=11.364064%

49%| 706/1433 [53:45<59:57, 4.95s/it]

[Quota] LLM=5.689973 | GPU=0.0 | Used=5.689973/50.0 | Remaining=44.310027 | Usage=11.379946%

49%| 707/1433 [53:51<1:04:13, 5.31s/it]

[Quota] LLM=5.698061 | GPU=0.0 | Used=5.698061/50.0 | Remaining=44.301939 | Usage=11.396122%

49%| 708/1433 [53:55<58:21, 4.83s/it]

[Quota] LLM=5.706062 | GPU=0.0 | Used=5.706062/50.0 | Remaining=44.293938 | Usage=11.412124%

49%| 709/1433 [53:58<51:55, 4.30s/it]

[Quota] LLM=5.713769 | GPU=0.0 | Used=5.713769/50.0 | Remaining=44.286231 | Usage=11.427538%

50%| 710/1433 [54:02<51:01, 4.23s/it]

[Quota] LLM=5.721659 | GPU=0.0 | Used=5.721659/50.0 | Remaining=44.278341 | Usage=11.443318%

50%| 711/1433 [54:06<51:29, 4.28s/it]

[Quota] LLM=5.729438 | GPU=0.0 | Used=5.729438/50.0 | Remaining=44.270562 | Usage=11.458876%

50%| 712/1433 [54:11<53:51, 4.48s/it]

[Quota] LLM=5.737214 | GPU=0.0 | Used=5.737214/50.0 | Remaining=44.262786 | Usage=11.474428%

50%| 713/1433 [54:15<51:38, 4.30s/it]

[Quota] LLM=5.745008 | GPU=0.0 | Used=5.745008/50.0 | Remaining=44.254992 | Usage=11.490016%

50%| 714/1433 [54:18<48:30, 4.05s/it]

[Quota] LLM=5.752751 | GPU=0.0 | Used=5.752751/50.0 | Remaining=44.247249 | Usage=11.505502%

50%| 715/1433 [54:23<49:07, 4.10s/it]

[Quota] LLM=5.760755 | GPU=0.0 | Used=5.760755/50.0 | Remaining=44.239245 | Usage=11.52151%

50%| 716/1433 [54:28<52:03, 4.36s/it]

[Quota] LLM=5.768633 | GPU=0.0 | Used=5.768633/50.0 | Remaining=44.231367 | Usage=11.537266%

50%| 717/1433 [54:32<52:29, 4.40s/it]

[Quota] LLM=5.7765320000000004 | GPU=0.0 | Used=5.7765320000000004/50.0 | Remaining=44.223468 | Usage=11.553064000000001%

50%| 718/1433 [54:37<52:38, 4.42s/it]

```
[Quota] LLM=5.784401 | GPU=0.0 | Used=5.784401/50.0 | Remaining=44.215599 |
Usage=11.568802%
50%|      | 719/1433 [54:43<58:09,  4.89s/it]

[Quota] LLM=5.792534 | GPU=0.0 | Used=5.792534/50.0 | Remaining=44.207466 |
Usage=11.585068%
50%|      | 720/1433 [54:47<56:49,  4.78s/it]

[Quota] LLM=5.80052 | GPU=0.0 | Used=5.80052/50.0 | Remaining=44.19948 |
Usage=11.60104%
50%|      | 721/1433 [54:52<55:32,  4.68s/it]

[Quota] LLM=5.808323 | GPU=0.0 | Used=5.808323/50.0 | Remaining=44.191677 |
Usage=11.616646%
50%|      | 722/1433 [54:56<53:52,  4.55s/it]

[Quota] LLM=5.816252 | GPU=0.0 | Used=5.816252/50.0 | Remaining=44.183748 |
Usage=11.632504%
50%|      | 723/1433 [55:01<56:36,  4.78s/it]

[Quota] LLM=5.824142 | GPU=0.0 | Used=5.824142/50.0 | Remaining=44.175858 |
Usage=11.648284%
51%|      | 724/1433 [55:05<53:56,  4.56s/it]

[Quota] LLM=5.832266 | GPU=0.0 | Used=5.832266/50.0 | Remaining=44.167734 |
Usage=11.664532%
51%|      | 725/1433 [55:10<53:19,  4.52s/it]

[Quota] LLM=5.840255 | GPU=0.0 | Used=5.840255/50.0 | Remaining=44.159745 |
Usage=11.68051%
51%|      | 726/1433 [55:14<52:41,  4.47s/it]

[Quota] LLM=5.84834 | GPU=0.0 | Used=5.84834/50.0 | Remaining=44.15166 |
Usage=11.69668%
51%|      | 727/1433 [55:18<51:47,  4.40s/it]

[Quota] LLM=5.856446 | GPU=0.0 | Used=5.856446/50.0 | Remaining=44.143554 |
Usage=11.712892%
51%|      | 728/1433 [55:22<51:09,  4.35s/it]

[Quota] LLM=5.86445 | GPU=0.0 | Used=5.86445/50.0 | Remaining=44.13555 |
Usage=11.7289%
51%|      | 729/1433 [55:26<47:38,  4.06s/it]

[Quota] LLM=5.872181 | GPU=0.0 | Used=5.872181/50.0 | Remaining=44.127819 |
Usage=11.744362%
51%|      | 730/1433 [55:30<48:50,  4.17s/it]
```

```
[Quota] LLM=5.880047 | GPU=0.0 | Used=5.880047/50.0 | Remaining=44.119953 |
Usage=11.760094%
51%|      | 731/1433 [55:35<50:18,  4.30s/it]

[Quota] LLM=5.887937 | GPU=0.0 | Used=5.887937/50.0 | Remaining=44.112063 |
Usage=11.775874%
51%|      | 732/1433 [55:41<58:03,  4.97s/it]

[Quota] LLM=5.895911 | GPU=0.0 | Used=5.895911/50.0 | Remaining=44.104089 |
Usage=11.791822%
51%|      | 733/1433 [55:46<55:42,  4.78s/it]

[Quota] LLM=5.9038640000000004 | GPU=0.0 | Used=5.9038640000000004/50.0 |
Remaining=44.096136 | Usage=11.807728000000001%
51%|      | 734/1433 [55:50<53:55,  4.63s/it]

[Quota] LLM=5.911649 | GPU=0.0 | Used=5.911649/50.0 | Remaining=44.088351 |
Usage=11.823298%
51%|      | 735/1433 [55:55<54:16,  4.67s/it]

[Quota] LLM=5.919398 | GPU=0.0 | Used=5.919398/50.0 | Remaining=44.080602 |
Usage=11.838796%
51%|      | 736/1433 [55:59<53:00,  4.56s/it]

[Quota] LLM=5.92733 | GPU=0.0 | Used=5.92733/50.0 | Remaining=44.07267 |
Usage=11.85466%
51%|      | 737/1433 [56:04<53:31,  4.61s/it]

[Quota] LLM=5.935445 | GPU=0.0 | Used=5.935445/50.0 | Remaining=44.064555 |
Usage=11.87089%
52%|      | 738/1433 [56:07<49:56,  4.31s/it]

[Quota] LLM=5.94335 | GPU=0.0 | Used=5.94335/50.0 | Remaining=44.05665 |
Usage=11.8867%
52%|      | 739/1433 [56:12<49:53,  4.31s/it]

[Quota] LLM=5.951375 | GPU=0.0 | Used=5.951375/50.0 | Remaining=44.048625 |
Usage=11.90275%
52%|      | 740/1433 [56:17<52:01,  4.50s/it]

[Quota] LLM=5.959271 | GPU=0.0 | Used=5.959271/50.0 | Remaining=44.040729 |
Usage=11.918542%
52%|      | 741/1433 [56:21<51:08,  4.43s/it]

[Quota] LLM=5.967032 | GPU=0.0 | Used=5.967032/50.0 | Remaining=44.032968 |
Usage=11.934064%
52%|      | 742/1433 [56:24<47:29,  4.12s/it]
```

```
[Quota] LLM=5.974772 | GPU=0.0 | Used=5.974772/50.0 | Remaining=44.025228 |  
Usage=11.949544%  
52%|      | 743/1433 [56:29<48:15,  4.20s/it]  
  
[Quota] LLM=5.982869 | GPU=0.0 | Used=5.982869/50.0 | Remaining=44.017131 |  
Usage=11.965738%  
52%|      | 744/1433 [56:34<52:27,  4.57s/it]  
  
[Quota] LLM=5.990909 | GPU=0.0 | Used=5.990909/50.0 | Remaining=44.009091 |  
Usage=11.981818%  
52%|      | 745/1433 [56:38<50:09,  4.37s/it]  
  
[Quota] LLM=5.998757 | GPU=0.0 | Used=5.998757/50.0 | Remaining=44.001243 |  
Usage=11.997514%  
52%|      | 746/1433 [56:44<54:06,  4.73s/it]  
  
[Quota] LLM=6.006545 | GPU=0.0 | Used=6.006545/50.0 | Remaining=43.993455 |  
Usage=12.01309%  
52%|      | 747/1433 [56:48<53:10,  4.65s/it]  
  
[Quota] LLM=6.0143510000000004 | GPU=0.0 | Used=6.0143510000000004/50.0 |  
Remaining=43.985649 | Usage=12.028702000000001%  
52%|      | 748/1433 [56:53<54:33,  4.78s/it]  
  
[Quota] LLM=6.022208 | GPU=0.0 | Used=6.022208/50.0 | Remaining=43.977792 |  
Usage=12.044416%  
52%|      | 749/1433 [56:59<56:44,  4.98s/it]  
  
[Quota] LLM=6.030035 | GPU=0.0 | Used=6.030035/50.0 | Remaining=43.969965 |  
Usage=12.06007%  
52%|      | 750/1433 [57:02<52:34,  4.62s/it]  
  
[Quota] LLM=6.037844 | GPU=0.0 | Used=6.037844/50.0 | Remaining=43.962156 |  
Usage=12.075688%  
52%|      | 751/1433 [57:06<49:41,  4.37s/it]  
  
[Quota] LLM=6.045833 | GPU=0.0 | Used=6.045833/50.0 | Remaining=43.954167 |  
Usage=12.091666%  
52%|      | 752/1433 [57:11<51:01,  4.50s/it]  
  
[Quota] LLM=6.0539060000000005 | GPU=0.0 | Used=6.0539060000000005/50.0 |  
Remaining=43.946094 | Usage=12.107812000000001%  
53%|      | 753/1433 [57:14<47:17,  4.17s/it]  
  
[Quota] LLM=6.061778 | GPU=0.0 | Used=6.061778/50.0 |  
Remaining=43.938221999999996 | Usage=12.123556%  
53%|      | 754/1433 [57:19<48:36,  4.29s/it]
```

```
[Quota] LLM=6.06995 | GPU=0.0 | Used=6.06995/50.0 | Remaining=43.93005 |
Usage=12.1399%
53%|      | 755/1433 [57:24<52:41,  4.66s/it]

[Quota] LLM=6.077897 | GPU=0.0 | Used=6.077897/50.0 | Remaining=43.922103 |
Usage=12.155794%
53%|      | 756/1433 [57:29<51:38,  4.58s/it]

[Quota] LLM=6.08618 | GPU=0.0 | Used=6.08618/50.0 | Remaining=43.91382 |
Usage=12.17236%
53%|      | 757/1433 [57:34<52:43,  4.68s/it]

[Quota] LLM=6.094202 | GPU=0.0 | Used=6.094202/50.0 | Remaining=43.905798 |
Usage=12.188404%
53%|      | 758/1433 [57:40<57:46,  5.14s/it]

[Quota] LLM=6.102473 | GPU=0.0 | Used=6.102473/50.0 | Remaining=43.897527 |
Usage=12.204946%
53%|      | 759/1433 [57:45<56:08,  5.00s/it]

[Quota] LLM=6.110798 | GPU=0.0 | Used=6.110798/50.0 | Remaining=43.889202 |
Usage=12.221596%
53%|      | 760/1433 [57:49<55:24,  4.94s/it]

[Quota] LLM=6.118853 | GPU=0.0 | Used=6.118853/50.0 | Remaining=43.881147 |
Usage=12.237706%
53%|      | 761/1433 [57:56<59:38,  5.32s/it]

[Quota] LLM=6.1268270000000005 | GPU=0.0 | Used=6.1268270000000005/50.0 |
Remaining=43.873173 | Usage=12.253654000000001%
53%|      | 762/1433 [58:00<54:42,  4.89s/it]

[Quota] LLM=6.134576 | GPU=0.0 | Used=6.134576/50.0 | Remaining=43.865424 |
Usage=12.269152%
53%|      | 763/1433 [58:04<53:37,  4.80s/it]

[Quota] LLM=6.142481 | GPU=0.0 | Used=6.142481/50.0 |
Remaining=43.857518999999996 | Usage=12.284962%
53%|      | 764/1433 [58:08<49:55,  4.48s/it]

[Quota] LLM=6.150383 | GPU=0.0 | Used=6.150383/50.0 | Remaining=43.849617 |
Usage=12.300766%
53%|      | 765/1433 [58:13<51:47,  4.65s/it]

[Quota] LLM=6.158171 | GPU=0.0 | Used=6.158171/50.0 | Remaining=43.841829 |
Usage=12.316342%
53%|      | 766/1433 [58:19<55:29,  4.99s/it]
```

[Quota] LLM=6.165929 | GPU=0.0 | Used=6.165929/50.0 | Remaining=43.834071 | Usage=12.331858%

54%| 767/1433 [58:23<54:13, 4.88s/it]

[Quota] LLM=6.173756 | GPU=0.0 | Used=6.173756/50.0 | Remaining=43.826244 | Usage=12.347512%

54%| 768/1433 [58:28<52:32, 4.74s/it]

[Quota] LLM=6.181505 | GPU=0.0 | Used=6.181505/50.0 | Remaining=43.818495 | Usage=12.36301%

54%| 769/1433 [58:31<48:36, 4.39s/it]

[Quota] LLM=6.189356 | GPU=0.0 | Used=6.189356/50.0 | Remaining=43.810643999999996 | Usage=12.378712%

54%| 770/1433 [58:36<49:02, 4.44s/it]

[Quota] LLM=6.197153 | GPU=0.0 | Used=6.197153/50.0 | Remaining=43.802847 | Usage=12.394306%

54%| 771/1433 [58:40<46:38, 4.23s/it]

[Quota] LLM=6.204935 | GPU=0.0 | Used=6.204935/50.0 | Remaining=43.795065 | Usage=12.40987%

54%| 772/1433 [58:45<49:02, 4.45s/it]

[Quota] LLM=6.213233 | GPU=0.0 | Used=6.213233/50.0 | Remaining=43.786767 | Usage=12.426466%

54%| 773/1433 [58:50<50:41, 4.61s/it]

[Quota] LLM=6.221195 | GPU=0.0 | Used=6.221195/50.0 | Remaining=43.778805 | Usage=12.44239%

54%| 774/1433 [58:54<51:20, 4.68s/it]

[Quota] LLM=6.228749 | GPU=0.0 | Used=6.228749/50.0 | Remaining=43.771251 | Usage=12.457498%

54%| 775/1433 [58:59<49:31, 4.52s/it]

[Quota] LLM=6.23651 | GPU=0.0 | Used=6.23651/50.0 | Remaining=43.76349 | Usage=12.47302%

54%| 776/1433 [59:03<48:04, 4.39s/it]

[Quota] LLM=6.244244 | GPU=0.0 | Used=6.244244/50.0 | Remaining=43.755756 | Usage=12.488488%

54%| 777/1433 [59:06<45:41, 4.18s/it]

[Quota] LLM=6.252038 | GPU=0.0 | Used=6.252038/50.0 | Remaining=43.747962 | Usage=12.504076%

54%| 778/1433 [59:12<50:03, 4.59s/it]

[Quota] LLM=6.259904 | GPU=0.0 | Used=6.259904/50.0 | Remaining=43.740096 | Usage=12.519808%

54%| 779/1433 [59:18<54:21, 4.99s/it]

[Quota] LLM=6.268262 | GPU=0.0 | Used=6.268262/50.0 | Remaining=43.731738 | Usage=12.536523999999998%

54%| 780/1433 [59:22<53:22, 4.90s/it]

[Quota] LLM=6.276254 | GPU=0.0 | Used=6.276254/50.0 | Remaining=43.723746 | Usage=12.552507999999998%

55%| 781/1433 [59:27<51:24, 4.73s/it]

[Quota] LLM=6.28406 | GPU=0.0 | Used=6.28406/50.0 | Remaining=43.71594 | Usage=12.56811999999999%

55%| 782/1433 [59:32<53:11, 4.90s/it]

[Quota] LLM=6.291929 | GPU=0.0 | Used=6.291929/50.0 | Remaining=43.708071000000004 | Usage=12.583858000000001%

55%| 783/1433 [59:37<51:59, 4.80s/it]

[Quota] LLM=6.299978 | GPU=0.0 | Used=6.299978/50.0 | Remaining=43.700022 | Usage=12.599956%

55%| 784/1433 [59:42<53:02, 4.90s/it]

[Quota] LLM=6.307967 | GPU=0.0 | Used=6.307967/50.0 | Remaining=43.692033 | Usage=12.615933999999998%

55%| 785/1433 [59:46<50:00, 4.63s/it]

[Quota] LLM=6.315821 | GPU=0.0 | Used=6.315821/50.0 | Remaining=43.684179 | Usage=12.631642%

55%| 786/1433 [59:50<47:02, 4.36s/it]

[Quota] LLM=6.323828 | GPU=0.0 | Used=6.323828/50.0 | Remaining=43.676172 | Usage=12.647656%

55%| 787/1433 [59:56<52:46, 4.90s/it]

[Quota] LLM=6.331805 | GPU=0.0 | Used=6.331805/50.0 | Remaining=43.668195 | Usage=12.66361%

55%| 788/1433 [1:00:01<52:34, 4.89s/it]

[Quota] LLM=6.339995 | GPU=0.0 | Used=6.339995/50.0 | Remaining=43.660005 | Usage=12.67999%

55%| 789/1433 [1:00:07<57:13, 5.33s/it]

[Quota] LLM=6.34796 | GPU=0.0 | Used=6.34796/50.0 | Remaining=43.65204 | Usage=12.69592%

55%| 790/1433 [1:00:12<56:38, 5.28s/it]

[Quota] LLM=6.355892 | GPU=0.0 | Used=6.355892/50.0 | Remaining=43.644108 | Usage=12.711784000000002%

55%| 791/1433 [1:00:16<53:03, 4.96s/it]

[Quota] LLM=6.363665 | GPU=0.0 | Used=6.363665/50.0 | Remaining=43.636335 | Usage=12.72733%

55%| 792/1433 [1:00:21<50:41, 4.74s/it]

[Quota] LLM=6.371423 | GPU=0.0 | Used=6.371423/50.0 | Remaining=43.628577 | Usage=12.742846%

55%| 793/1433 [1:00:24<47:47, 4.48s/it]

[Quota] LLM=6.379256 | GPU=0.0 | Used=6.379256/50.0 | Remaining=43.620744 | Usage=12.758512%

55%| 794/1433 [1:00:29<48:50, 4.59s/it]

[Quota] LLM=6.387218 | GPU=0.0 | Used=6.387218/50.0 | Remaining=43.612782 | Usage=12.774436%

55%| 795/1433 [1:00:34<48:37, 4.57s/it]

[Quota] LLM=6.395147 | GPU=0.0 | Used=6.395147/50.0 | Remaining=43.604853 | Usage=12.790294%

56%| 796/1433 [1:00:38<47:06, 4.44s/it]

[Quota] LLM=6.403025 | GPU=0.0 | Used=6.403025/50.0 | Remaining=43.596975 | Usage=12.80605000000003%

56%| 797/1433 [1:00:42<45:55, 4.33s/it]

[Quota] LLM=6.411089 | GPU=0.0 | Used=6.411089/50.0 | Remaining=43.588911 | Usage=12.82217799999997%

56%| 798/1433 [1:00:46<45:01, 4.25s/it]

[Quota] LLM=6.418784 | GPU=0.0 | Used=6.418784/50.0 | Remaining=43.581216 | Usage=12.837568%

56%| 799/1433 [1:00:49<41:02, 3.88s/it]

[Quota] LLM=6.426446 | GPU=0.0 | Used=6.426446/50.0 | Remaining=43.573554 | Usage=12.85289200000002%

56%| 800/1433 [1:00:53<41:25, 3.93s/it]

[Quota] LLM=6.434687 | GPU=0.0 | Used=6.434687/50.0 | Remaining=43.565313 | Usage=12.869374%

56%| 801/1433 [1:00:57<40:15, 3.82s/it]

[Quota] LLM=6.442757 | GPU=0.0 | Used=6.442757/50.0 | Remaining=43.557243 | Usage=12.885514%

56%| 802/1433 [1:01:01<42:22, 4.03s/it]

```
[Quota] LLM=6.450557 | GPU=0.0 | Used=6.450557/50.0 | Remaining=43.549443 |
Usage=12.901114%
56%|      | 803/1433 [1:01:05<42:59,  4.09s/it]

[Quota] LLM=6.45848 | GPU=0.0 | Used=6.45848/50.0 | Remaining=43.54152 |
Usage=12.91696%
56%|      | 804/1433 [1:01:10<44:03,  4.20s/it]

[Quota] LLM=6.46646 | GPU=0.0 | Used=6.46646/50.0 | Remaining=43.53354 |
Usage=12.932920000000001%
56%|      | 805/1433 [1:01:14<44:44,  4.27s/it]

[Quota] LLM=6.474362 | GPU=0.0 | Used=6.474362/50.0 | Remaining=43.525638 |
Usage=12.948724%
56%|      | 806/1433 [1:01:19<44:52,  4.29s/it]

[Quota] LLM=6.482282 | GPU=0.0 | Used=6.482282/50.0 | Remaining=43.517718 |
Usage=12.964564000000001%
56%|      | 807/1433 [1:01:23<44:34,  4.27s/it]

[Quota] LLM=6.490253 | GPU=0.0 | Used=6.490253/50.0 | Remaining=43.509747 |
Usage=12.980506%
56%|      | 808/1433 [1:01:27<44:36,  4.28s/it]

[Quota] LLM=6.498266 | GPU=0.0 | Used=6.498266/50.0 | Remaining=43.501734 |
Usage=12.996532%
56%|      | 809/1433 [1:01:31<44:13,  4.25s/it]

[Quota] LLM=6.50633 | GPU=0.0 | Used=6.50633/50.0 | Remaining=43.49367 |
Usage=13.01266%
57%|      | 810/1433 [1:01:35<43:33,  4.20s/it]

[Quota] LLM=6.514004 | GPU=0.0 | Used=6.514004/50.0 | Remaining=43.485996 |
Usage=13.028008%
57%|      | 811/1433 [1:01:40<44:38,  4.31s/it]

[Quota] LLM=6.521939 | GPU=0.0 | Used=6.521939/50.0 | Remaining=43.478061 |
Usage=13.043878000000001%
57%|      | 812/1433 [1:01:45<45:10,  4.37s/it]

[Quota] LLM=6.529988 | GPU=0.0 | Used=6.529988/50.0 | Remaining=43.470012 |
Usage=13.059976%
57%|      | 813/1433 [1:01:50<47:52,  4.63s/it]

[Quota] LLM=6.538094 | GPU=0.0 | Used=6.538094/50.0 | Remaining=43.461906 |
Usage=13.076188%
57%|      | 814/1433 [1:01:54<47:14,  4.58s/it]
```

[Quota] LLM=6.546017 | GPU=0.0 | Used=6.546017/50.0 | Remaining=43.453983 | Usage=13.092034%

57%| 815/1433 [1:01:58<43:38, 4.24s/it]

[Quota] LLM=6.553694 | GPU=0.0 | Used=6.553694/50.0 | Remaining=43.446306 | Usage=13.107388%

57%| 816/1433 [1:02:01<41:51, 4.07s/it]

[Quota] LLM=6.561464 | GPU=0.0 | Used=6.561464/50.0 | Remaining=43.438536 | Usage=13.122928%

57%| 817/1433 [1:02:06<42:27, 4.14s/it]

[Quota] LLM=6.569387 | GPU=0.0 | Used=6.569387/50.0 | Remaining=43.430613 | Usage=13.138774%

57%| 818/1433 [1:02:10<43:14, 4.22s/it]

[Quota] LLM=6.5774 | GPU=0.0 | Used=6.5774/50.0 | Remaining=43.4226 | Usage=13.1548%

57%| 819/1433 [1:02:14<42:01, 4.11s/it]

[Quota] LLM=6.585158 | GPU=0.0 | Used=6.585158/50.0 | Remaining=43.414842 | Usage=13.170315999999998%

57%| 820/1433 [1:02:20<47:12, 4.62s/it]

[Quota] LLM=6.593675 | GPU=0.0 | Used=6.593675/50.0 | Remaining=43.406325 | Usage=13.18735%

57%| 821/1433 [1:02:24<47:17, 4.64s/it]

[Quota] LLM=6.601424 | GPU=0.0 | Used=6.601424/50.0 | Remaining=43.398576 | Usage=13.202848%

57%| 822/1433 [1:02:28<43:42, 4.29s/it]

[Quota] LLM=6.609026 | GPU=0.0 | Used=6.609026/50.0 | Remaining=43.390974 | Usage=13.218052%

57%| 823/1433 [1:02:33<45:17, 4.45s/it]

[Quota] LLM=6.616763 | GPU=0.0 | Used=6.616763/50.0 | Remaining=43.383237 | Usage=13.233525999999998%

58%| 824/1433 [1:02:38<48:35, 4.79s/it]

[Quota] LLM=6.624557 | GPU=0.0 | Used=6.624557/50.0 | Remaining=43.375443 | Usage=13.249114%

58%| 825/1433 [1:02:47<1:01:19, 6.05s/it]

[Quota] LLM=6.63275 | GPU=0.0 | Used=6.63275/50.0 | Remaining=43.36725 | Usage=13.2655%

58%| 826/1433 [1:02:52<56:15, 5.56s/it]

[Quota] LLM=6.640598 | GPU=0.0 | Used=6.640598/50.0 | Remaining=43.359402 | Usage=13.281196000000001%

58%| 827/1433 [1:02:56<51:01, 5.05s/it]

[Quota] LLM=6.648563 | GPU=0.0 | Used=6.648563/50.0 | Remaining=43.351437 | Usage=13.297126%

58%| 828/1433 [1:03:00<49:07, 4.87s/it]

[Quota] LLM=6.656693 | GPU=0.0 | Used=6.656693/50.0 | Remaining=43.343307 | Usage=13.313386%

58%| 829/1433 [1:03:06<53:07, 5.28s/it]

[Quota] LLM=6.664856 | GPU=0.0 | Used=6.664856/50.0 | Remaining=43.335144 | Usage=13.329712000000002%

58%| 830/1433 [1:03:10<48:53, 4.86s/it]

[Quota] LLM=6.672665 | GPU=0.0 | Used=6.672665/50.0 | Remaining=43.327335 | Usage=13.345533%

58%| 831/1433 [1:03:14<46:12, 4.61s/it]

[Quota] LLM=6.680411 | GPU=0.0 | Used=6.680411/50.0 | Remaining=43.319589 | Usage=13.360822%

58%| 832/1433 [1:03:18<43:09, 4.31s/it]

[Quota] LLM=6.688166 | GPU=0.0 | Used=6.688166/50.0 | Remaining=43.311834 | Usage=13.376332%

58%| 833/1433 [1:03:23<46:00, 4.60s/it]

[Quota] LLM=6.69638 | GPU=0.0 | Used=6.69638/50.0 | Remaining=43.30362 | Usage=13.39276%

58%| 834/1433 [1:03:28<46:41, 4.68s/it]

[Quota] LLM=6.704303 | GPU=0.0 | Used=6.704303/50.0 | Remaining=43.295697 | Usage=13.408606%

58%| 835/1433 [1:03:32<44:18, 4.45s/it]

[Quota] LLM=6.712001 | GPU=0.0 | Used=6.712001/50.0 | Remaining=43.287999 | Usage=13.424001999999998%

58%| 836/1433 [1:03:35<40:45, 4.10s/it]

[Quota] LLM=6.719795 | GPU=0.0 | Used=6.719795/50.0 | Remaining=43.280205 | Usage=13.43959%

58%| 837/1433 [1:03:39<40:34, 4.09s/it]

[Quota] LLM=6.727673 | GPU=0.0 | Used=6.727673/50.0 | Remaining=43.272327 | Usage=13.455346000000002%

58%| 838/1433 [1:03:44<42:34, 4.29s/it]

[Quota] LLM=6.735416 | GPU=0.0 | Used=6.735416/50.0 | Remaining=43.264584 | Usage=13.470832%

59%| 839/1433 [1:03:49<45:34, 4.60s/it]

[Quota] LLM=6.743348 | GPU=0.0 | Used=6.743348/50.0 | Remaining=43.256652 | Usage=13.486696%

59%| 840/1433 [1:03:54<46:53, 4.74s/it]

[Quota] LLM=6.7513250000000005 | GPU=0.0 | Used=6.7513250000000005/50.0 | Remaining=43.248675 | Usage=13.502650000000003%

59%| 841/1433 [1:03:59<45:56, 4.66s/it]

[Quota] LLM=6.759437 | GPU=0.0 | Used=6.759437/50.0 | Remaining=43.240563 | Usage=13.518874%

59%| 842/1433 [1:04:04<47:08, 4.79s/it]

[Quota] LLM=6.767414 | GPU=0.0 | Used=6.767414/50.0 | Remaining=43.232586 | Usage=13.534828%

59%| 843/1433 [1:04:08<46:15, 4.71s/it]

[Quota] LLM=6.775307 | GPU=0.0 | Used=6.775307/50.0 | Remaining=43.224693 | Usage=13.550614%

59%| 844/1433 [1:04:14<48:29, 4.94s/it]

[Quota] LLM=6.783341 | GPU=0.0 | Used=6.783341/50.0 | Remaining=43.216659 | Usage=13.566682%

59%| 845/1433 [1:04:19<49:39, 5.07s/it]

[Quota] LLM=6.791213 | GPU=0.0 | Used=6.791213/50.0 | Remaining=43.208787 | Usage=13.582426%

59%| 846/1433 [1:04:24<49:04, 5.02s/it]

[Quota] LLM=6.799163 | GPU=0.0 | Used=6.799163/50.0 | Remaining=43.200837 | Usage=13.598326%

59%| 847/1433 [1:04:28<45:08, 4.62s/it]

[Quota] LLM=6.806969 | GPU=0.0 | Used=6.806969/50.0 | Remaining=43.193031 | Usage=13.613938000000001%

59%| 848/1433 [1:04:34<48:37, 4.99s/it]

[Quota] LLM=6.8147210000000005 | GPU=0.0 | Used=6.8147210000000005/50.0 | Remaining=43.185279 | Usage=13.629442%

59%| 849/1433 [1:04:37<44:04, 4.53s/it]

[Quota] LLM=6.822491 | GPU=0.0 | Used=6.822491/50.0 | Remaining=43.177509 | Usage=13.644982%

59%| 850/1433 [1:04:43<48:17, 4.97s/it]

[Quota] LLM=6.830303 | GPU=0.0 | Used=6.830303/50.0 | Remaining=43.169697 | Usage=13.660606%

59% | 851/1433 [1:04:47<45:39, 4.71s/it]

[Quota] LLM=6.838091 | GPU=0.0 | Used=6.838091/50.0 | Remaining=43.161909 | Usage=13.676182%

59% | 852/1433 [1:04:52<44:59, 4.65s/it]

[Quota] LLM=6.846095 | GPU=0.0 | Used=6.846095/50.0 | Remaining=43.153905 | Usage=13.692190000000002%

60% | 853/1433 [1:04:56<42:48, 4.43s/it]

[Quota] LLM=6.854039 | GPU=0.0 | Used=6.854039/50.0 | Remaining=43.145961 | Usage=13.708078%

60% | 854/1433 [1:05:00<42:02, 4.36s/it]

[Quota] LLM=6.86171 | GPU=0.0 | Used=6.86171/50.0 | Remaining=43.13829 | Usage=13.72342%

60% | 855/1433 [1:05:04<39:58, 4.15s/it]

[Quota] LLM=6.869366 | GPU=0.0 | Used=6.869366/50.0 | Remaining=43.130634 | Usage=13.738732%

60% | 856/1433 [1:05:08<40:17, 4.19s/it]

[Quota] LLM=6.877628 | GPU=0.0 | Used=6.877628/50.0 | Remaining=43.122372 | Usage=13.755256%

60% | 857/1433 [1:05:11<37:42, 3.93s/it]

[Quota] LLM=6.885317 | GPU=0.0 | Used=6.885317/50.0 | Remaining=43.114683 | Usage=13.77063399999998%

60% | 858/1433 [1:05:15<38:21, 4.00s/it]

[Quota] LLM=6.893369 | GPU=0.0 | Used=6.893369/50.0 | Remaining=43.106631 | Usage=13.78673799999998%

60% | 859/1433 [1:05:23<47:31, 4.97s/it]

[Quota] LLM=6.901277 | GPU=0.0 | Used=6.901277/50.0 | Remaining=43.098723 | Usage=13.802554%

60% | 860/1433 [1:05:27<45:47, 4.79s/it]

[Quota] LLM=6.9092 | GPU=0.0 | Used=6.9092/50.0 | Remaining=43.0908 | Usage=13.8184%

60% | 861/1433 [1:05:38<1:02:34, 6.56s/it]

[Quota] LLM=6.917123 | GPU=0.0 | Used=6.917123/50.0 | Remaining=43.08287699999996 | Usage=13.834246%

60% | 862/1433 [1:05:42<56:41, 5.96s/it]

[Quota] LLM=6.924995 | GPU=0.0 | Used=6.924995/50.0 | Remaining=43.075005 | Usage=13.849990000000002%

60%| 863/1433 [1:05:47<53:53, 5.67s/it]

[Quota] LLM=6.93305 | GPU=0.0 | Used=6.93305/50.0 | Remaining=43.06695 | Usage=13.866100000000001%

60%| 864/1433 [1:05:50<46:49, 4.94s/it]

[Quota] LLM=6.940829 | GPU=0.0 | Used=6.940829/50.0 | Remaining=43.059171 | Usage=13.881658%

60%| 865/1433 [1:05:56<47:31, 5.02s/it]

[Quota] LLM=6.948764 | GPU=0.0 | Used=6.948764/50.0 | Remaining=43.051236 | Usage=13.897528000000001%

60%| 866/1433 [1:06:00<45:37, 4.83s/it]

[Quota] LLM=6.956642 | GPU=0.0 | Used=6.956642/50.0 | Remaining=43.043358 | Usage=13.913284%

61%| 867/1433 [1:06:05<45:17, 4.80s/it]

[Quota] LLM=6.964541 | GPU=0.0 | Used=6.964541/50.0 | Remaining=43.035459 | Usage=13.929081999999998%

61%| 868/1433 [1:06:09<43:26, 4.61s/it]

[Quota] LLM=6.97223 | GPU=0.0 | Used=6.97223/50.0 | Remaining=43.027770000000004 | Usage=13.94446%

61%| 869/1433 [1:06:15<47:53, 5.09s/it]

[Quota] LLM=6.9806420000000005 | GPU=0.0 | Used=6.9806420000000005/50.0 | Remaining=43.019358 | Usage=13.961284000000001%

61%| 870/1433 [1:06:19<44:18, 4.72s/it]

[Quota] LLM=6.988409 | GPU=0.0 | Used=6.988409/50.0 | Remaining=43.011591 | Usage=13.976818%

61%| 871/1433 [1:06:23<42:27, 4.53s/it]

[Quota] LLM=6.996251 | GPU=0.0 | Used=6.996251/50.0 | Remaining=43.003749 | Usage=13.992502000000002%

61%| 872/1433 [1:06:27<42:10, 4.51s/it]

[Quota] LLM=7.004102 | GPU=0.0 | Used=7.004102/50.0 | Remaining=42.995898 | Usage=14.008204%

61%| 873/1433 [1:06:32<40:52, 4.38s/it]

[Quota] LLM=7.011722 | GPU=0.0 | Used=7.011722/50.0 | Remaining=42.988278 | Usage=14.023444%

61%| 874/1433 [1:06:36<41:15, 4.43s/it]

[Quota] LLM=7.019495 | GPU=0.0 | Used=7.019495/50.0 | Remaining=42.980505 | Usage=14.038990000000002%

61% | 875/1433 [1:06:40<39:20, 4.23s/it]

[Quota] LLM=7.027196 | GPU=0.0 | Used=7.027196/50.0 | Remaining=42.972804 | Usage=14.054391999999998%

61% | 876/1433 [1:06:45<40:26, 4.36s/it]

[Quota] LLM=7.03508 | GPU=0.0 | Used=7.03508/50.0 | Remaining=42.96492 | Usage=14.070159999999998%

61% | 877/1433 [1:06:50<42:37, 4.60s/it]

[Quota] LLM=7.043 | GPU=0.0 | Used=7.043/50.0 | Remaining=42.957 | Usage=14.086000000000002%

61% | 878/1433 [1:06:55<43:13, 4.67s/it]

[Quota] LLM=7.051025 | GPU=0.0 | Used=7.051025/50.0 | Remaining=42.948975 | Usage=14.102049999999998%

61% | 879/1433 [1:06:59<43:03, 4.66s/it]

[Quota] LLM=7.058996 | GPU=0.0 | Used=7.058996/50.0 | Remaining=42.941004 | Usage=14.117992%

61% | 880/1433 [1:07:05<46:49, 5.08s/it]

[Quota] LLM=7.066988 | GPU=0.0 | Used=7.066988/50.0 | Remaining=42.933012 | Usage=14.133976%

61% | 881/1433 [1:07:11<47:57, 5.21s/it]

[Quota] LLM=7.074908 | GPU=0.0 | Used=7.074908/50.0 | Remaining=42.925092 | Usage=14.149815999999998%

62% | 882/1433 [1:07:16<46:57, 5.11s/it]

[Quota] LLM=7.083128 | GPU=0.0 | Used=7.083128/50.0 | Remaining=42.916872 | Usage=14.16625599999999%

62% | 883/1433 [1:07:20<44:40, 4.87s/it]

[Quota] LLM=7.090931 | GPU=0.0 | Used=7.090931/50.0 | Remaining=42.909069 | Usage=14.181862%

62% | 884/1433 [1:07:26<48:59, 5.35s/it]

[Quota] LLM=7.098704 | GPU=0.0 | Used=7.098704/50.0 | Remaining=42.901296 | Usage=14.197408%

62% | 885/1433 [1:07:33<53:19, 5.84s/it]

[Quota] LLM=7.106648 | GPU=0.0 | Used=7.106648/50.0 | Remaining=42.893352 | Usage=14.213296%

62% | 886/1433 [1:07:39<52:17, 5.74s/it]

[Quota] LLM=7.114832 | GPU=0.0 | Used=7.114832/50.0 | Remaining=42.885168 | Usage=14.229664%

62%| 887/1433 [1:07:43<48:57, 5.38s/it]

[Quota] LLM=7.122713 | GPU=0.0 | Used=7.122713/50.0 | Remaining=42.877287 | Usage=14.245426%

62%| 888/1433 [1:07:48<45:57, 5.06s/it]

[Quota] LLM=7.130453 | GPU=0.0 | Used=7.130453/50.0 | Remaining=42.869547 | Usage=14.260906%

62%| 889/1433 [1:07:53<46:40, 5.15s/it]

[Quota] LLM=7.138751 | GPU=0.0 | Used=7.138751/50.0 | Remaining=42.861249 | Usage=14.277502%

62%| 890/1433 [1:07:57<43:03, 4.76s/it]

[Quota] LLM=7.146407 | GPU=0.0 | Used=7.146407/50.0 | Remaining=42.853593000000004 | Usage=14.292814000000002%

62%| 891/1433 [1:08:01<40:55, 4.53s/it]

[Quota] LLM=7.154441 | GPU=0.0 | Used=7.154441/50.0 | Remaining=42.845559 | Usage=14.308882%

62%| 892/1433 [1:08:05<38:13, 4.24s/it]

[Quota] LLM=7.162358 | GPU=0.0 | Used=7.162358/50.0 | Remaining=42.837642 | Usage=14.324716%

62%| 893/1433 [1:08:09<39:24, 4.38s/it]

[Quota] LLM=7.170146 | GPU=0.0 | Used=7.170146/50.0 | Remaining=42.829854 | Usage=14.34029199999998%

62%| 894/1433 [1:08:14<39:46, 4.43s/it]

[Quota] LLM=7.177913 | GPU=0.0 | Used=7.177913/50.0 | Remaining=42.822086999999996 | Usage=14.355825999999999%

62%| 895/1433 [1:08:18<39:32, 4.41s/it]

[Quota] LLM=7.185977 | GPU=0.0 | Used=7.185977/50.0 | Remaining=42.814023 | Usage=14.371954%

63%| 896/1433 [1:08:25<46:41, 5.22s/it]

[Quota] LLM=7.193684 | GPU=0.0 | Used=7.193684/50.0 | Remaining=42.806316 | Usage=14.387368%

63%| 897/1433 [1:08:30<45:55, 5.14s/it]

[Quota] LLM=7.201784 | GPU=0.0 | Used=7.201784/50.0 | Remaining=42.798216 | Usage=14.403568%

63%| 898/1433 [1:08:34<43:09, 4.84s/it]

```
[Quota] LLM=7.209752 | GPU=0.0 | Used=7.209752/50.0 | Remaining=42.790248 |
Usage=14.419504%
63%|      | 899/1433 [1:08:40<44:14,  4.97s/it]

[Quota] LLM=7.217636 | GPU=0.0 | Used=7.217636/50.0 | Remaining=42.782364 |
Usage=14.435272%
63%|      | 900/1433 [1:08:44<42:38,  4.80s/it]

[Quota] LLM=7.225388 | GPU=0.0 | Used=7.225388/50.0 | Remaining=42.774612 |
Usage=14.45077599999998%
63%|      | 901/1433 [1:08:48<40:15,  4.54s/it]

[Quota] LLM=7.233161 | GPU=0.0 | Used=7.233161/50.0 | Remaining=42.766839 |
Usage=14.46632200000002%
63%|      | 902/1433 [1:08:53<42:22,  4.79s/it]

[Quota] LLM=7.241105 | GPU=0.0 | Used=7.241105/50.0 | Remaining=42.758895 |
Usage=14.48221%
63%|      | 903/1433 [1:08:58<42:17,  4.79s/it]

[Quota] LLM=7.248812 | GPU=0.0 | Used=7.248812/50.0 | Remaining=42.751188 |
Usage=14.497624%
63%|      | 904/1433 [1:09:02<39:58,  4.53s/it]

[Quota] LLM=7.256552 | GPU=0.0 | Used=7.256552/50.0 | Remaining=42.743448 |
Usage=14.51310399999998%
63%|      | 905/1433 [1:09:07<42:12,  4.80s/it]

[Quota] LLM=7.264772 | GPU=0.0 | Used=7.264772/50.0 | Remaining=42.735228 |
Usage=14.529544%
63%|      | 906/1433 [1:09:13<43:34,  4.96s/it]

[Quota] LLM=7.272701 | GPU=0.0 | Used=7.272701/50.0 | Remaining=42.727299 |
Usage=14.545402%
63%|      | 907/1433 [1:09:20<50:27,  5.75s/it]

[Quota] LLM=7.280792 | GPU=0.0 | Used=7.280792/50.0 | Remaining=42.719208 |
Usage=14.561584%
63%|      | 908/1433 [1:09:24<45:25,  5.19s/it]

[Quota] LLM=7.288997 | GPU=0.0 | Used=7.288997/50.0 | Remaining=42.711003 |
Usage=14.577994%
63%|      | 909/1433 [1:09:31<48:42,  5.58s/it]

[Quota] LLM=7.297112 | GPU=0.0 | Used=7.297112/50.0 | Remaining=42.702888 |
Usage=14.594224%
64%|      | 910/1433 [1:09:35<45:10,  5.18s/it]
```

```
[Quota] LLM=7.30505 | GPU=0.0 | Used=7.30505/50.0 | Remaining=42.69495 |
Usage=14.61009999999997%
64%|      | 911/1433 [1:09:39<41:53,  4.82s/it]

[Quota] LLM=7.312832 | GPU=0.0 | Used=7.312832/50.0 | Remaining=42.687168 |
Usage=14.62566399999999%
64%|      | 912/1433 [1:09:44<42:11,  4.86s/it]

[Quota] LLM=7.321112 | GPU=0.0 | Used=7.321112/50.0 | Remaining=42.678888 |
Usage=14.642224%
64%|      | 913/1433 [1:09:49<42:50,  4.94s/it]

[Quota] LLM=7.328876 | GPU=0.0 | Used=7.328876/50.0 | Remaining=42.671124 |
Usage=14.657752000000002%
64%|      | 914/1433 [1:09:53<40:00,  4.63s/it]

[Quota] LLM=7.336616 | GPU=0.0 | Used=7.336616/50.0 | Remaining=42.663384 |
Usage=14.673232%
64%|      | 915/1433 [1:09:59<44:30,  5.16s/it]

[Quota] LLM=7.344485 | GPU=0.0 | Used=7.344485/50.0 | Remaining=42.655515 |
Usage=14.688969999999998%
64%|      | 916/1433 [1:10:07<52:05,  6.04s/it]

[Quota] LLM=7.352498 | GPU=0.0 | Used=7.352498/50.0 | Remaining=42.647502 |
Usage=14.704996000000001%
64%|      | 917/1433 [1:10:12<47:24,  5.51s/it]

[Quota] LLM=7.3602170000000005 | GPU=0.0 | Used=7.3602170000000005/50.0 |
Remaining=42.639783 | Usage=14.720434000000001%
64%|      | 918/1433 [1:10:16<44:20,  5.17s/it]

[Quota] LLM=7.367927 | GPU=0.0 | Used=7.367927/50.0 | Remaining=42.632073 |
Usage=14.735854000000002%
64%|      | 919/1433 [1:10:20<42:04,  4.91s/it]

[Quota] LLM=7.37582 | GPU=0.0 | Used=7.37582/50.0 | Remaining=42.62418 |
Usage=14.75163999999998%
64%|      | 920/1433 [1:10:27<45:24,  5.31s/it]

[Quota] LLM=7.383713 | GPU=0.0 | Used=7.383713/50.0 | Remaining=42.616287 |
Usage=14.767426%
64%|      | 921/1433 [1:10:35<54:14,  6.36s/it]

[Quota] LLM=7.391696 | GPU=0.0 | Used=7.391696/50.0 |
Remaining=42.608304000000004 | Usage=14.783391999999997%
64%|      | 922/1433 [1:10:40<48:21,  5.68s/it]
```

[Quota] LLM=7.399562 | GPU=0.0 | Used=7.399562/50.0 | Remaining=42.600438 | Usage=14.79912399999999%

64%| 923/1433 [1:10:43<43:30, 5.12s/it]

[Quota] LLM=7.407428 | GPU=0.0 | Used=7.407428/50.0 | Remaining=42.592572 | Usage=14.814856%

64%| 924/1433 [1:10:48<42:38, 5.03s/it]

[Quota] LLM=7.415222 | GPU=0.0 | Used=7.415222/50.0 | Remaining=42.584778 | Usage=14.830444000000002%

65%| 925/1433 [1:10:57<51:03, 6.03s/it]

[Quota] LLM=7.423268 | GPU=0.0 | Used=7.423268/50.0 | Remaining=42.576732 | Usage=14.84653599999999%

65%| 926/1433 [1:11:02<49:12, 5.82s/it]

[Quota] LLM=7.431122 | GPU=0.0 | Used=7.431122/50.0 | Remaining=42.568878 | Usage=14.86224399999999%

65%| 927/1433 [1:11:09<51:47, 6.14s/it]

[Quota] LLM=7.43912 | GPU=0.0 | Used=7.43912/50.0 | Remaining=42.56088 | Usage=14.878240000000002%

65%| 928/1433 [1:11:13<46:18, 5.50s/it]

[Quota] LLM=7.447058 | GPU=0.0 | Used=7.447058/50.0 | Remaining=42.552942 | Usage=14.894116000000002%

65%| 929/1433 [1:11:18<44:34, 5.31s/it]

[Quota] LLM=7.454855 | GPU=0.0 | Used=7.454855/50.0 | Remaining=42.545145 | Usage=14.90971%

65%| 930/1433 [1:11:23<44:36, 5.32s/it]

[Quota] LLM=7.4628440000000005 | GPU=0.0 | Used=7.4628440000000005/50.0 | Remaining=42.537155999999996 | Usage=14.925688000000001%

65%| 931/1433 [1:11:27<41:22, 4.95s/it]

[Quota] LLM=7.47071 | GPU=0.0 | Used=7.47071/50.0 | Remaining=42.52929 | Usage=14.94141999999999%

65%| 932/1433 [1:11:33<43:00, 5.15s/it]

[Quota] LLM=7.478585 | GPU=0.0 | Used=7.478585/50.0 | Remaining=42.521415 | Usage=14.95717%

65%| 933/1433 [1:11:37<41:22, 4.97s/it]

[Quota] LLM=7.486253 | GPU=0.0 | Used=7.486253/50.0 | Remaining=42.513747 | Usage=14.972506%

65%| 934/1433 [1:11:42<40:47, 4.91s/it]

[Quota] LLM=7.494248 | GPU=0.0 | Used=7.494248/50.0 | Remaining=42.505752 | Usage=14.98849599999999%

65%| 935/1433 [1:11:47<40:14, 4.85s/it]

[Quota] LLM=7.502198 | GPU=0.0 | Used=7.502198/50.0 | Remaining=42.497802 | Usage=15.004396%

65%| 936/1433 [1:11:51<39:01, 4.71s/it]

[Quota] LLM=7.510349 | GPU=0.0 | Used=7.510349/50.0 | Remaining=42.489651 | Usage=15.020698%

65%| 937/1433 [1:11:56<38:25, 4.65s/it]

[Quota] LLM=7.518194 | GPU=0.0 | Used=7.518194/50.0 | Remaining=42.481806 | Usage=15.036388%

65%| 938/1433 [1:12:00<37:19, 4.52s/it]

[Quota] LLM=7.526318 | GPU=0.0 | Used=7.526318/50.0 | Remaining=42.473682 | Usage=15.052636%

66%| 939/1433 [1:12:05<38:03, 4.62s/it]

[Quota] LLM=7.534322 | GPU=0.0 | Used=7.534322/50.0 | Remaining=42.465678 | Usage=15.068644%

66%| 940/1433 [1:12:08<35:53, 4.37s/it]

[Quota] LLM=7.542056 | GPU=0.0 | Used=7.542056/50.0 | Remaining=42.457944 | Usage=15.084112%

66%| 941/1433 [1:12:12<34:45, 4.24s/it]

[Quota] LLM=7.549754 | GPU=0.0 | Used=7.549754/50.0 | Remaining=42.450246 | Usage=15.099508%

66%| 942/1433 [1:12:16<34:19, 4.19s/it]

[Quota] LLM=7.557536 | GPU=0.0 | Used=7.557536/50.0 | Remaining=42.442464 | Usage=15.115072%

66%| 943/1433 [1:12:21<34:47, 4.26s/it]

[Quota] LLM=7.565231 | GPU=0.0 | Used=7.565231/50.0 | Remaining=42.434769 | Usage=15.130462%

66%| 944/1433 [1:12:25<34:43, 4.26s/it]

[Quota] LLM=7.573196 | GPU=0.0 | Used=7.573196/50.0 | Remaining=42.426804 | Usage=15.146392%

66%| 945/1433 [1:12:29<32:47, 4.03s/it]

[Quota] LLM=7.580933 | GPU=0.0 | Used=7.580933/50.0 | Remaining=42.419067 | Usage=15.16186599999999%

66%| 946/1433 [1:12:33<33:24, 4.12s/it]

[Quota] LLM=7.588919 | GPU=0.0 | Used=7.588919/50.0 | Remaining=42.411081 | Usage=15.177838%

66% | 947/1433 [1:12:37<34:13, 4.22s/it]

[Quota] LLM=7.596776 | GPU=0.0 | Used=7.596776/50.0 | Remaining=42.403224 | Usage=15.19355199999999%

66% | 948/1433 [1:12:41<32:42, 4.05s/it]

[Quota] LLM=7.604594 | GPU=0.0 | Used=7.604594/50.0 | Remaining=42.395406 | Usage=15.209188%

66% | 949/1433 [1:12:45<31:53, 3.95s/it]

[Quota] LLM=7.612481 | GPU=0.0 | Used=7.612481/50.0 | Remaining=42.387519 | Usage=15.224962%

66% | 950/1433 [1:12:49<31:53, 3.96s/it]

[Quota] LLM=7.620176 | GPU=0.0 | Used=7.620176/50.0 | Remaining=42.379824 | Usage=15.24035199999998%

66% | 951/1433 [1:12:53<31:20, 3.90s/it]

[Quota] LLM=7.627922 | GPU=0.0 | Used=7.627922/50.0 | Remaining=42.372078 | Usage=15.255844%

66% | 952/1433 [1:12:57<32:01, 4.00s/it]

[Quota] LLM=7.635725 | GPU=0.0 | Used=7.635725/50.0 | Remaining=42.364275 | Usage=15.27145%

67% | 953/1433 [1:13:02<35:20, 4.42s/it]

[Quota] LLM=7.643594 | GPU=0.0 | Used=7.643594/50.0 | Remaining=42.356406 | Usage=15.287188000000002%

67% | 954/1433 [1:13:07<35:04, 4.39s/it]

[Quota] LLM=7.651319 | GPU=0.0 | Used=7.651319/50.0 | Remaining=42.348681 | Usage=15.30263799999998%

67% | 955/1433 [1:13:11<35:09, 4.41s/it]

[Quota] LLM=7.6590620000000005 | GPU=0.0 | Used=7.6590620000000005/50.0 | Remaining=42.340938 | Usage=15.318124%

67% | 956/1433 [1:13:15<34:49, 4.38s/it]

[Quota] LLM=7.666871 | GPU=0.0 | Used=7.666871/50.0 | Remaining=42.333129 | Usage=15.333742%

67% | 957/1433 [1:13:20<36:18, 4.58s/it]

[Quota] LLM=7.67486 | GPU=0.0 | Used=7.67486/50.0 | Remaining=42.32514 | Usage=15.34972%

67% | 958/1433 [1:13:24<34:43, 4.39s/it]

```
[Quota] LLM=7.682489 | GPU=0.0 | Used=7.682489/50.0 |  
Remaining=42.317510999999996 | Usage=15.364978%  
67%|      | 959/1433 [1:13:29<36:04,  4.57s/it]  
  
[Quota] LLM=7.690355 | GPU=0.0 | Used=7.690355/50.0 | Remaining=42.309645 |  
Usage=15.38071%  
67%|      | 960/1433 [1:13:35<39:48,  5.05s/it]  
  
[Quota] LLM=7.6982990000000004 | GPU=0.0 | Used=7.6982990000000004/50.0 |  
Remaining=42.301701 | Usage=15.396598000000001%  
67%|      | 961/1433 [1:13:40<39:30,  5.02s/it]  
  
[Quota] LLM=7.706402 | GPU=0.0 | Used=7.706402/50.0 | Remaining=42.293598 |  
Usage=15.412804%  
67%|      | 962/1433 [1:13:44<36:16,  4.62s/it]  
  
[Quota] LLM=7.714148 | GPU=0.0 | Used=7.714148/50.0 | Remaining=42.285852 |  
Usage=15.428296%  
67%|      | 963/1433 [1:13:48<33:36,  4.29s/it]  
  
[Quota] LLM=7.721795 | GPU=0.0 | Used=7.721795/50.0 | Remaining=42.278205 |  
Usage=15.443590000000002%  
67%|      | 964/1433 [1:13:54<38:19,  4.90s/it]  
  
[Quota] LLM=7.729754 | GPU=0.0 | Used=7.729754/50.0 | Remaining=42.270246 |  
Usage=15.459508%  
67%|      | 965/1433 [1:13:59<39:43,  5.09s/it]  
  
[Quota] LLM=7.73759 | GPU=0.0 | Used=7.73759/50.0 | Remaining=42.26241 |  
Usage=15.47518%  
67%|      | 966/1433 [1:14:03<36:54,  4.74s/it]  
  
[Quota] LLM=7.745327 | GPU=0.0 | Used=7.745327/50.0 | Remaining=42.254673 |  
Usage=15.49065399999997%  
67%|      | 967/1433 [1:14:08<36:56,  4.76s/it]  
  
[Quota] LLM=7.753439 | GPU=0.0 | Used=7.753439/50.0 | Remaining=42.246561 |  
Usage=15.506878000000002%  
68%|      | 968/1433 [1:14:14<39:15,  5.07s/it]  
  
[Quota] LLM=7.761467 | GPU=0.0 | Used=7.761467/50.0 |  
Remaining=42.238533000000004 | Usage=15.522934%  
68%|      | 969/1433 [1:14:18<37:02,  4.79s/it]  
  
[Quota] LLM=7.769171 | GPU=0.0 | Used=7.769171/50.0 | Remaining=42.230829 |  
Usage=15.538342%  
68%|      | 970/1433 [1:14:23<36:20,  4.71s/it]
```

```
[Quota] LLM=7.776899 | GPU=0.0 | Used=7.776899/50.0 | Remaining=42.223101 |
Usage=15.553797999999999%
68%|      | 971/1433 [1:14:28<37:33,  4.88s/it]

[Quota] LLM=7.785281 | GPU=0.0 | Used=7.785281/50.0 | Remaining=42.214719 |
Usage=15.570562000000002%
68%|      | 972/1433 [1:14:33<38:09,  4.97s/it]

[Quota] LLM=7.793354 | GPU=0.0 | Used=7.793354/50.0 | Remaining=42.206646 |
Usage=15.586708%
68%|      | 973/1433 [1:14:38<37:05,  4.84s/it]

[Quota] LLM=7.8013580000000005 | GPU=0.0 | Used=7.8013580000000005/50.0 |
Remaining=42.198642 | Usage=15.602716%
68%|      | 974/1433 [1:14:41<34:36,  4.52s/it]

[Quota] LLM=7.809569 | GPU=0.0 | Used=7.809569/50.0 | Remaining=42.19043100000004 |
Usage=15.619138%
68%|      | 975/1433 [1:14:47<37:31,  4.92s/it]

[Quota] LLM=7.817417 | GPU=0.0 | Used=7.817417/50.0 | Remaining=42.182583 |
Usage=15.634834%
68%|      | 976/1433 [1:14:51<35:09,  4.62s/it]

[Quota] LLM=7.825241 | GPU=0.0 | Used=7.825241/50.0 | Remaining=42.174759 |
Usage=15.65048199999998%
68%|      | 977/1433 [1:14:55<34:13,  4.50s/it]

[Quota] LLM=7.833221 | GPU=0.0 | Used=7.833221/50.0 | Remaining=42.166779 |
Usage=15.666442%
68%|      | 978/1433 [1:14:59<32:35,  4.30s/it]

[Quota] LLM=7.841192 | GPU=0.0 | Used=7.841192/50.0 | Remaining=42.158808 |
Usage=15.682384000000003%
68%|      | 979/1433 [1:15:03<30:37,  4.05s/it]

[Quota] LLM=7.848767 | GPU=0.0 | Used=7.848767/50.0 | Remaining=42.151233 |
Usage=15.697534%
68%|      | 980/1433 [1:15:07<30:12,  4.00s/it]

[Quota] LLM=7.856453 | GPU=0.0 | Used=7.856453/50.0 | Remaining=42.143547 |
Usage=15.712906000000002%
68%|      | 981/1433 [1:15:11<31:36,  4.20s/it]

[Quota] LLM=7.864787 | GPU=0.0 | Used=7.864787/50.0 | Remaining=42.135213 |
Usage=15.729574%
69%|      | 982/1433 [1:15:17<34:20,  4.57s/it]
```

```
[Quota] LLM=7.872578 | GPU=0.0 | Used=7.872578/50.0 | Remaining=42.127422 |
Usage=15.745156%
69%|      | 983/1433 [1:15:22<36:04,  4.81s/it]

[Quota] LLM=7.880666 | GPU=0.0 | Used=7.880666/50.0 | Remaining=42.119334 |
Usage=15.761332%
69%|      | 984/1433 [1:15:27<37:24,  5.00s/it]

[Quota] LLM=7.888556 | GPU=0.0 | Used=7.888556/50.0 | Remaining=42.111444 |
Usage=15.777112%
69%|      | 985/1433 [1:15:32<36:01,  4.83s/it]

[Quota] LLM=7.896668 | GPU=0.0 | Used=7.896668/50.0 | Remaining=42.103332 |
Usage=15.793336%
69%|      | 986/1433 [1:15:38<38:03,  5.11s/it]

[Quota] LLM=7.905125 | GPU=0.0 | Used=7.905125/50.0 | Remaining=42.094875 |
Usage=15.81025%
69%|      | 987/1433 [1:15:42<36:22,  4.89s/it]

[Quota] LLM=7.913048 | GPU=0.0 | Used=7.913048/50.0 | Remaining=42.086952 |
Usage=15.826096%
69%|      | 988/1433 [1:15:46<34:47,  4.69s/it]

[Quota] LLM=7.921148 | GPU=0.0 | Used=7.921148/50.0 | Remaining=42.078852 |
Usage=15.842296000000001%
69%|      | 989/1433 [1:15:51<35:50,  4.84s/it]

[Quota] LLM=7.929719 | GPU=0.0 | Used=7.929719/50.0 | Remaining=42.070281 |
Usage=15.859438%
69%|      | 990/1433 [1:15:57<37:53,  5.13s/it]

[Quota] LLM=7.937726 | GPU=0.0 | Used=7.937726/50.0 | Remaining=42.062274 |
Usage=15.875451999999997%
69%|      | 991/1433 [1:16:02<37:58,  5.15s/it]

[Quota] LLM=7.946021 | GPU=0.0 | Used=7.946021/50.0 | Remaining=42.053979 |
Usage=15.892042%
69%|      | 992/1433 [1:16:08<38:07,  5.19s/it]

[Quota] LLM=7.954052 | GPU=0.0 | Used=7.954052/50.0 | Remaining=42.045948 |
Usage=15.908104000000002%
69%|      | 993/1433 [1:16:13<39:10,  5.34s/it]

[Quota] LLM=7.962206 | GPU=0.0 | Used=7.962206/50.0 | Remaining=42.037794 |
Usage=15.924411999999998%
69%|      | 994/1433 [1:16:18<36:42,  5.02s/it]
```

```
[Quota] LLM=7.970129 | GPU=0.0 | Used=7.970129/50.0 | Remaining=42.029871 |
Usage=15.940257999999998%
69%|      | 995/1433 [1:16:22<35:17,  4.83s/it]

[Quota] LLM=7.978502 | GPU=0.0 | Used=7.978502/50.0 | Remaining=42.021498 |
Usage=15.957004%
70%|      | 996/1433 [1:16:26<34:15,  4.70s/it]

[Quota] LLM=7.986542 | GPU=0.0 | Used=7.986542/50.0 | Remaining=42.013458 |
Usage=15.973084000000002%
70%|      | 997/1433 [1:16:31<34:08,  4.70s/it]

[Quota] LLM=7.994489 | GPU=0.0 | Used=7.994489/50.0 | Remaining=42.005511 |
Usage=15.988977999999998%
70%|      | 998/1433 [1:16:37<36:26,  5.03s/it]

[Quota] LLM=8.002304 | GPU=0.0 | Used=8.002304/50.0 | Remaining=41.997696 |
Usage=16.004608%
70%|      | 999/1433 [1:16:42<37:04,  5.13s/it]

[Quota] LLM=8.010248 | GPU=0.0 | Used=8.010248/50.0 |
Remaining=41.989751999999996 | Usage=16.020496%
70%|      | 1000/1433 [1:16:47<35:26,  4.91s/it]

[Quota] LLM=8.018489 | GPU=0.0 | Used=8.018489/50.0 | Remaining=41.981511 |
Usage=16.036978%
70%|      | 1001/1433 [1:16:51<34:37,  4.81s/it]

[Quota] LLM=8.026409 | GPU=0.0 | Used=8.026409/50.0 | Remaining=41.973591 |
Usage=16.052818%
70%|      | 1002/1433 [1:16:57<35:34,  4.95s/it]

[Quota] LLM=8.034491 | GPU=0.0 | Used=8.034491/50.0 | Remaining=41.965509 |
Usage=16.068982%
70%|      | 1003/1433 [1:17:01<33:36,  4.69s/it]

[Quota] LLM=8.042273 | GPU=0.0 | Used=8.042273/50.0 | Remaining=41.957727 |
Usage=16.084546%
70%|      | 1004/1433 [1:17:07<36:53,  5.16s/it]

[Quota] LLM=8.050658 | GPU=0.0 | Used=8.050658/50.0 | Remaining=41.949342 |
Usage=16.101316%
70%|      | 1005/1433 [1:17:11<35:31,  4.98s/it]

[Quota] LLM=8.058488 | GPU=0.0 | Used=8.058488/50.0 | Remaining=41.941512 |
Usage=16.116976%
70%|      | 1006/1433 [1:17:17<36:22,  5.11s/it]
```

[Quota] LLM=8.066291 | GPU=0.0 | Used=8.066291/50.0 | Remaining=41.933709 | Usage=16.132582%

70%| 1007/1433 [1:17:20<32:46, 4.62s/it]

[Quota] LLM=8.073869 | GPU=0.0 | Used=8.073869/50.0 | Remaining=41.926131 | Usage=16.147738%

70%| 1008/1433 [1:17:23<29:28, 4.16s/it]

[Quota] LLM=8.081633 | GPU=0.0 | Used=8.081633/50.0 | Remaining=41.918367 | Usage=16.163266%

70%| 1009/1433 [1:17:31<35:43, 5.06s/it]

[Quota] LLM=8.089655 | GPU=0.0 | Used=8.089655/50.0 | Remaining=41.910345 | Usage=16.17931%

70%| 1010/1433 [1:17:35<33:17, 4.72s/it]

[Quota] LLM=8.097395 | GPU=0.0 | Used=8.097395/50.0 | Remaining=41.902605 | Usage=16.19479%

71%| 1011/1433 [1:17:38<30:37, 4.35s/it]

[Quota] LLM=8.10509 | GPU=0.0 | Used=8.10509/50.0 | Remaining=41.8949099999999996 | Usage=16.21018%

71%| 1012/1433 [1:17:42<30:12, 4.31s/it]

[Quota] LLM=8.112791 | GPU=0.0 | Used=8.112791/50.0 | Remaining=41.887209 | Usage=16.225582%

71%| 1013/1433 [1:17:46<28:55, 4.13s/it]

[Quota] LLM=8.120474 | GPU=0.0 | Used=8.120474/50.0 | Remaining=41.879526 | Usage=16.240948%

71%| 1014/1433 [1:17:50<29:39, 4.25s/it]

[Quota] LLM=8.128511 | GPU=0.0 | Used=8.128511/50.0 | Remaining=41.871489 | Usage=16.257022%

71%| 1015/1433 [1:17:54<27:52, 4.00s/it]

[Quota] LLM=8.136317 | GPU=0.0 | Used=8.136317/50.0 | Remaining=41.863683 | Usage=16.272634%

71%| 1016/1433 [1:17:59<29:03, 4.18s/it]

[Quota] LLM=8.144132 | GPU=0.0 | Used=8.144132/50.0 | Remaining=41.855868 | Usage=16.288264%

71%| 1017/1433 [1:18:03<28:51, 4.16s/it]

[Quota] LLM=8.151914 | GPU=0.0 | Used=8.151914/50.0 | Remaining=41.848086 | Usage=16.303828%

71%| 1018/1433 [1:18:07<29:29, 4.26s/it]

```
[Quota] LLM=8.159792 | GPU=0.0 | Used=8.159792/50.0 |  
Remaining=41.840208000000004 | Usage=16.319584%  
71%|       | 1019/1433 [1:18:13<32:11,  4.66s/it]  
  
[Quota] LLM=8.167769 | GPU=0.0 | Used=8.167769/50.0 | Remaining=41.832231 |  
Usage=16.335538%  
71%|       | 1020/1433 [1:18:17<32:03,  4.66s/it]  
  
[Quota] LLM=8.175968 | GPU=0.0 | Used=8.175968/50.0 | Remaining=41.824032 |  
Usage=16.351936%  
71%|       | 1021/1433 [1:18:22<32:05,  4.67s/it]  
  
[Quota] LLM=8.184224 | GPU=0.0 | Used=8.184224/50.0 | Remaining=41.815776 |  
Usage=16.368448%  
71%|       | 1022/1433 [1:18:26<30:38,  4.47s/it]  
  
[Quota] LLM=8.191973 | GPU=0.0 | Used=8.191973/50.0 |  
Remaining=41.808026999999996 | Usage=16.383946%  
71%|       | 1023/1433 [1:18:30<29:26,  4.31s/it]  
  
[Quota] LLM=8.199911 | GPU=0.0 | Used=8.199911/50.0 | Remaining=41.800089 |  
Usage=16.399822%  
71%|       | 1024/1433 [1:18:35<29:44,  4.36s/it]  
  
[Quota] LLM=8.20796 | GPU=0.0 | Used=8.20796/50.0 | Remaining=41.79204 |  
Usage=16.41592%  
72%|       | 1025/1433 [1:18:39<29:15,  4.30s/it]  
  
[Quota] LLM=8.215838 | GPU=0.0 | Used=8.215838/50.0 | Remaining=41.784162 |  
Usage=16.431676%  
72%|       | 1026/1433 [1:18:44<31:29,  4.64s/it]  
  
[Quota] LLM=8.223845 | GPU=0.0 | Used=8.223845/50.0 | Remaining=41.776155 |  
Usage=16.44769%  
72%|       | 1027/1433 [1:18:49<31:49,  4.70s/it]  
  
[Quota] LLM=8.231948 | GPU=0.0 | Used=8.231948/50.0 | Remaining=41.768052 |  
Usage=16.463896%  
72%|       | 1028/1433 [1:18:55<33:39,  4.99s/it]  
  
[Quota] LLM=8.240198 | GPU=0.0 | Used=8.240198/50.0 | Remaining=41.759802 |  
Usage=16.480396%  
72%|       | 1029/1433 [1:18:59<32:05,  4.77s/it]  
  
[Quota] LLM=8.248271 | GPU=0.0 | Used=8.248271/50.0 | Remaining=41.751729 |  
Usage=16.496542%  
72%|       | 1030/1433 [1:19:04<33:46,  5.03s/it]
```

[Quota] LLM=8.256575 | GPU=0.0 | Used=8.256575/50.0 | Remaining=41.743425 | Usage=16.51315%

72%| 1031/1433 [1:19:08<30:51, 4.61s/it]

[Quota] LLM=8.264366 | GPU=0.0 | Used=8.264366/50.0 | Remaining=41.735634 | Usage=16.528732%

72%| 1032/1433 [1:19:12<30:19, 4.54s/it]

[Quota] LLM=8.272478 | GPU=0.0 | Used=8.272478/50.0 | Remaining=41.727522 | Usage=16.544956%

72%| 1033/1433 [1:19:16<29:01, 4.35s/it]

[Quota] LLM=8.280374 | GPU=0.0 | Used=8.280374/50.0 | Remaining=41.719626 | Usage=16.560748%

72%| 1034/1433 [1:19:20<27:57, 4.20s/it]

[Quota] LLM=8.28824 | GPU=0.0 | Used=8.28824/50.0 | Remaining=41.71176 | Usage=16.57648%

72%| 1035/1433 [1:19:26<30:23, 4.58s/it]

[Quota] LLM=8.296289 | GPU=0.0 | Used=8.296289/50.0 | Remaining=41.703711 | Usage=16.592578%

72%| 1036/1433 [1:19:31<31:06, 4.70s/it]

[Quota] LLM=8.304482 | GPU=0.0 | Used=8.304482/50.0 | Remaining=41.695518 | Usage=16.608964%

72%| 1037/1433 [1:19:35<30:37, 4.64s/it]

[Quota] LLM=8.312456 | GPU=0.0 | Used=8.312456/50.0 | Remaining=41.687544 | Usage=16.624912%

72%| 1038/1433 [1:19:45<39:58, 6.07s/it]

[Quota] LLM=8.320403 | GPU=0.0 | Used=8.320403/50.0 | Remaining=41.679597 | Usage=16.640806%

73%| 1039/1433 [1:19:48<35:26, 5.40s/it]

[Quota] LLM=8.328257 | GPU=0.0 | Used=8.328257/50.0 | Remaining=41.671743 | Usage=16.656514%

73%| 1040/1433 [1:19:53<33:54, 5.18s/it]

[Quota] LLM=8.336357 | GPU=0.0 | Used=8.336357/50.0 | Remaining=41.663643 | Usage=16.672714%

73%| 1041/1433 [1:19:59<35:04, 5.37s/it]

[Quota] LLM=8.344328 | GPU=0.0 | Used=8.344328/50.0 | Remaining=41.655671999999996 | Usage=16.688656%

73%| 1042/1433 [1:20:02<30:41, 4.71s/it]

```
[Quota] LLM=8.352137 | GPU=0.0 | Used=8.352137/50.0 | Remaining=41.647863 |
Usage=16.704274%
73%|      | 1043/1433 [1:20:08<33:54,  5.22s/it]

[Quota] LLM=8.360135 | GPU=0.0 | Used=8.360135/50.0 | Remaining=41.639865 |
Usage=16.72027%
73%|      | 1044/1433 [1:20:14<34:05,  5.26s/it]

[Quota] LLM=8.368091 | GPU=0.0 | Used=8.368091/50.0 | Remaining=41.631909 |
Usage=16.736182%
73%|      | 1045/1433 [1:20:19<32:53,  5.09s/it]

[Quota] LLM=8.376203 | GPU=0.0 | Used=8.376203/50.0 |
Remaining=41.623796999999996 | Usage=16.752406%
73%|      | 1046/1433 [1:20:23<32:05,  4.98s/it]

[Quota] LLM=8.384213 | GPU=0.0 | Used=8.384213/50.0 | Remaining=41.615787 |
Usage=16.768426%
73%|      | 1047/1433 [1:20:28<31:12,  4.85s/it]

[Quota] LLM=8.392445 | GPU=0.0 | Used=8.392445/50.0 | Remaining=41.607555 |
Usage=16.78489%
73%|      | 1048/1433 [1:20:32<29:25,  4.58s/it]

[Quota] LLM=8.400275 | GPU=0.0 | Used=8.400275/50.0 | Remaining=41.599725 |
Usage=16.80055%
73%|      | 1049/1433 [1:20:38<32:09,  5.03s/it]

[Quota] LLM=8.408615 | GPU=0.0 | Used=8.408615/50.0 | Remaining=41.591385 |
Usage=16.81723%
73%|      | 1050/1433 [1:20:42<30:51,  4.83s/it]

[Quota] LLM=8.416724 | GPU=0.0 | Used=8.416724/50.0 | Remaining=41.583276 |
Usage=16.833448%
73%|      | 1051/1433 [1:20:48<31:56,  5.02s/it]

[Quota] LLM=8.424704 | GPU=0.0 | Used=8.424704/50.0 | Remaining=41.575296 |
Usage=16.849408%
73%|      | 1052/1433 [1:20:52<30:33,  4.81s/it]

[Quota] LLM=8.432786 | GPU=0.0 | Used=8.432786/50.0 | Remaining=41.567214 |
Usage=16.865572%
73%|      | 1053/1433 [1:20:56<29:45,  4.70s/it]

[Quota] LLM=8.440658 | GPU=0.0 | Used=8.440658/50.0 | Remaining=41.559342 |
Usage=16.881316%
74%|      | 1054/1433 [1:21:02<31:55,  5.05s/it]
```

```
[Quota] LLM=8.449136 | GPU=0.0 | Used=8.449136/50.0 |  
Remaining=41.550864000000004 | Usage=16.898272%  
74%|      | 1055/1433 [1:21:07<31:46,  5.04s/it]  
  
[Quota] LLM=8.457329 | GPU=0.0 | Used=8.457329/50.0 | Remaining=41.542671 |  
Usage=16.914658%  
74%|      | 1056/1433 [1:21:13<32:39,  5.20s/it]  
  
[Quota] LLM=8.465429 | GPU=0.0 | Used=8.465429/50.0 | Remaining=41.534571 |  
Usage=16.930858%  
74%|      | 1057/1433 [1:21:18<32:39,  5.21s/it]  
  
[Quota] LLM=8.473421 | GPU=0.0 | Used=8.473421/50.0 | Remaining=41.526579 |  
Usage=16.946842%  
74%|      | 1058/1433 [1:21:23<31:32,  5.05s/it]  
  
[Quota] LLM=8.481716 | GPU=0.0 | Used=8.481716/50.0 | Remaining=41.518284 |  
Usage=16.963432%  
74%|      | 1059/1433 [1:21:28<31:10,  5.00s/it]  
  
[Quota] LLM=8.489681000000001 | GPU=0.0 | Used=8.489681000000001/50.0 |  
Remaining=41.510318999999996 | Usage=16.979362000000002%  
74%|      | 1060/1433 [1:21:35<35:31,  5.72s/it]  
  
[Quota] LLM=8.498339 | GPU=0.0 | Used=8.498339/50.0 | Remaining=41.501661 |  
Usage=16.996678%  
74%|      | 1061/1433 [1:21:41<35:32,  5.73s/it]  
  
[Quota] LLM=8.506247 | GPU=0.0 | Used=8.506247/50.0 | Remaining=41.493753 |  
Usage=17.012494%  
74%|      | 1062/1433 [1:21:45<32:34,  5.27s/it]  
  
[Quota] LLM=8.514278000000001 | GPU=0.0 | Used=8.514278000000001/50.0 |  
Remaining=41.485721999999996 | Usage=17.028556000000002%  
74%|      | 1063/1433 [1:21:50<32:01,  5.19s/it]  
  
[Quota] LLM=8.521916 | GPU=0.0 | Used=8.521916/50.0 | Remaining=41.478084 |  
Usage=17.043832%  
74%|      | 1064/1433 [1:21:55<31:46,  5.17s/it]  
  
[Quota] LLM=8.529956 | GPU=0.0 | Used=8.529956/50.0 | Remaining=41.470044 |  
Usage=17.059912%  
74%|      | 1065/1433 [1:22:00<31:51,  5.19s/it]  
  
[Quota] LLM=8.537918 | GPU=0.0 | Used=8.537918/50.0 | Remaining=41.462082 |  
Usage=17.075836%  
74%|      | 1066/1433 [1:22:05<31:33,  5.16s/it]
```

[Quota] LLM=8.546231 | GPU=0.0 | Used=8.546231/50.0 | Remaining=41.453769 | Usage=17.092462%

74%| 1067/1433 [1:22:09<29:21, 4.81s/it]

[Quota] LLM=8.554073 | GPU=0.0 | Used=8.554073/50.0 | Remaining=41.445927 | Usage=17.108146%

75%| 1068/1433 [1:22:16<33:11, 5.46s/it]

[Quota] LLM=8.561942 | GPU=0.0 | Used=8.561942/50.0 | Remaining=41.438058 | Usage=17.123884%

75%| 1069/1433 [1:22:22<32:54, 5.42s/it]

[Quota] LLM=8.56988 | GPU=0.0 | Used=8.56988/50.0 | Remaining=41.43012 | Usage=17.13976%

75%| 1070/1433 [1:22:25<29:33, 4.89s/it]

[Quota] LLM=8.577692 | GPU=0.0 | Used=8.577692/50.0 | Remaining=41.422308 | Usage=17.155384%

75%| 1071/1433 [1:22:30<29:30, 4.89s/it]

[Quota] LLM=8.585609 | GPU=0.0 | Used=8.585609/50.0 | Remaining=41.414391 | Usage=17.171218%

75%| 1072/1433 [1:22:35<28:29, 4.74s/it]

[Quota] LLM=8.593733 | GPU=0.0 | Used=8.593733/50.0 | Remaining=41.406267 | Usage=17.187466%

75%| 1073/1433 [1:22:39<27:29, 4.58s/it]

[Quota] LLM=8.601923 | GPU=0.0 | Used=8.601923/50.0 | Remaining=41.398077 | Usage=17.203846%

75%| 1074/1433 [1:22:45<30:56, 5.17s/it]

[Quota] LLM=8.60996 | GPU=0.0 | Used=8.60996/50.0 | Remaining=41.39004 | Usage=17.21992%

75%| 1075/1433 [1:22:49<28:14, 4.73s/it]

[Quota] LLM=8.617835 | GPU=0.0 | Used=8.617835/50.0 | Remaining=41.382165 | Usage=17.23567%

75%| 1076/1433 [1:22:53<26:11, 4.40s/it]

[Quota] LLM=8.625731 | GPU=0.0 | Used=8.625731/50.0 | Remaining=41.374269 | Usage=17.251462%

75%| 1077/1433 [1:23:00<31:12, 5.26s/it]

[Quota] LLM=8.633513 | GPU=0.0 | Used=8.633513/50.0 | Remaining=41.366487 | Usage=17.267026%

75%| 1078/1433 [1:23:04<29:15, 4.95s/it]

```
[Quota] LLM=8.64128 | GPU=0.0 | Used=8.64128/50.0 | Remaining=41.35872 |  
Usage=17.28256%  
75%|      | 1079/1433 [1:23:10<30:28,  5.17s/it]  
  
[Quota] LLM=8.649302 | GPU=0.0 | Used=8.649302/50.0 | Remaining=41.350698 |  
Usage=17.298604%  
75%|      | 1080/1433 [1:23:16<31:19,  5.32s/it]  
  
[Quota] LLM=8.65745 | GPU=0.0 | Used=8.65745/50.0 | Remaining=41.34255 |  
Usage=17.3149%  
75%|      | 1081/1433 [1:23:19<28:33,  4.87s/it]  
  
[Quota] LLM=8.665505 | GPU=0.0 | Used=8.665505/50.0 |  
Remaining=41.334495000000004 | Usage=17.33101%  
76%|      | 1082/1433 [1:23:24<28:13,  4.82s/it]  
  
[Quota] LLM=8.673581 | GPU=0.0 | Used=8.673581/50.0 | Remaining=41.326419 |  
Usage=17.347162%  
76%|      | 1083/1433 [1:23:29<27:19,  4.68s/it]  
  
[Quota] LLM=8.681615 | GPU=0.0 | Used=8.681615/50.0 | Remaining=41.318385 |  
Usage=17.36323%  
76%|      | 1084/1433 [1:23:33<27:33,  4.74s/it]  
  
[Quota] LLM=8.689604 | GPU=0.0 | Used=8.689604/50.0 | Remaining=41.310396 |  
Usage=17.379208%  
76%|      | 1085/1433 [1:23:37<25:54,  4.47s/it]  
  
[Quota] LLM=8.697404 | GPU=0.0 | Used=8.697404/50.0 | Remaining=41.302596 |  
Usage=17.394808%  
76%|      | 1086/1433 [1:23:42<25:41,  4.44s/it]  
  
[Quota] LLM=8.705435 | GPU=0.0 | Used=8.705435/50.0 | Remaining=41.294565 |  
Usage=17.41087%  
76%|      | 1087/1433 [1:23:46<25:32,  4.43s/it]  
  
[Quota] LLM=8.713721 | GPU=0.0 | Used=8.713721/50.0 | Remaining=41.286279 |  
Usage=17.427442%  
76%|      | 1088/1433 [1:23:50<25:19,  4.41s/it]  
  
[Quota] LLM=8.721869 | GPU=0.0 | Used=8.721869/50.0 | Remaining=41.278131 |  
Usage=17.443738%  
76%|      | 1089/1433 [1:23:55<25:54,  4.52s/it]  
  
[Quota] LLM=8.729663 | GPU=0.0 | Used=8.729663/50.0 | Remaining=41.270337 |  
Usage=17.459326%  
76%|      | 1090/1433 [1:24:01<27:45,  4.86s/it]
```

[Quota] LLM=8.737472 | GPU=0.0 | Used=8.737472/50.0 | Remaining=41.262528 | Usage=17.474944%

76% | 1091/1433 [1:24:05<27:02, 4.74s/it]

[Quota] LLM=8.745221 | GPU=0.0 | Used=8.745221/50.0 | Remaining=41.254779 | Usage=17.490442%

76% | 1092/1433 [1:24:09<24:55, 4.38s/it]

[Quota] LLM=8.752964 | GPU=0.0 | Used=8.752964/50.0 | Remaining=41.247036 | Usage=17.505928%

76% | 1093/1433 [1:24:13<24:46, 4.37s/it]

[Quota] LLM=8.76086 | GPU=0.0 | Used=8.76086/50.0 | Remaining=41.23914 | Usage=17.52172%

76% | 1094/1433 [1:24:18<25:18, 4.48s/it]

[Quota] LLM=8.768681 | GPU=0.0 | Used=8.768681/50.0 | Remaining=41.231319 | Usage=17.537362%

76% | 1095/1433 [1:24:23<25:44, 4.57s/it]

[Quota] LLM=8.776745 | GPU=0.0 | Used=8.776745/50.0 | Remaining=41.223255 | Usage=17.55349%

76% | 1096/1433 [1:24:28<27:22, 4.87s/it]

[Quota] LLM=8.784575 | GPU=0.0 | Used=8.784575/50.0 | Remaining=41.215424999999996 | Usage=17.56915%

77% | 1097/1433 [1:24:32<25:24, 4.54s/it]

[Quota] LLM=8.792207 | GPU=0.0 | Used=8.792207/50.0 | Remaining=41.207793 | Usage=17.584414%

77% | 1098/1433 [1:24:36<23:51, 4.27s/it]

[Quota] LLM=8.800241 | GPU=0.0 | Used=8.800241/50.0 | Remaining=41.199759 | Usage=17.600482%

77% | 1099/1433 [1:24:40<23:40, 4.25s/it]

[Quota] LLM=8.808149 | GPU=0.0 | Used=8.808149/50.0 | Remaining=41.191851 | Usage=17.616298%

77% | 1100/1433 [1:24:44<22:36, 4.07s/it]

[Quota] LLM=8.816033000000001 | GPU=0.0 | Used=8.816033000000001/50.0 | Remaining=41.183966999999996 | Usage=17.632066000000002%

77% | 1101/1433 [1:24:47<21:38, 3.91s/it]

[Quota] LLM=8.823677 | GPU=0.0 | Used=8.823677/50.0 | Remaining=41.176323 | Usage=17.647354%

77% | 1102/1433 [1:24:52<22:39, 4.11s/it]

[Quota] LLM=8.832065 | GPU=0.0 | Used=8.832065/50.0 | Remaining=41.167935 | Usage=17.66413%

77% | 1103/1433 [1:24:57<24:33, 4.47s/it]

[Quota] LLM=8.840405 | GPU=0.0 | Used=8.840405/50.0 | Remaining=41.159594999999996 | Usage=17.68081%

77% | 1104/1433 [1:25:01<23:59, 4.38s/it]

[Quota] LLM=8.848442 | GPU=0.0 | Used=8.848442/50.0 | Remaining=41.151558 | Usage=17.696884%

77% | 1105/1433 [1:25:05<23:10, 4.24s/it]

[Quota] LLM=8.85698 | GPU=0.0 | Used=8.85698/50.0 | Remaining=41.14302 | Usage=17.71396%

77% | 1106/1433 [1:25:12<26:49, 4.92s/it]

[Quota] LLM=8.864984 | GPU=0.0 | Used=8.864984/50.0 | Remaining=41.135016 | Usage=17.729968%

77% | 1107/1433 [1:25:15<24:41, 4.54s/it]

[Quota] LLM=8.87297 | GPU=0.0 | Used=8.87297/50.0 | Remaining=41.12703 | Usage=17.74594%

77% | 1108/1433 [1:25:20<24:24, 4.51s/it]

[Quota] LLM=8.880782 | GPU=0.0 | Used=8.880782/50.0 | Remaining=41.11921800000004 | Usage=17.761564%

77% | 1109/1433 [1:25:25<26:12, 4.85s/it]

[Quota] LLM=8.889077 | GPU=0.0 | Used=8.889077/50.0 | Remaining=41.110923 | Usage=17.778154%

77% | 1110/1433 [1:25:30<26:05, 4.85s/it]

[Quota] LLM=8.89685 | GPU=0.0 | Used=8.89685/50.0 | Remaining=41.10315 | Usage=17.7937%

78% | 1111/1433 [1:25:34<24:33, 4.58s/it]

[Quota] LLM=8.904635 | GPU=0.0 | Used=8.904635/50.0 | Remaining=41.095365 | Usage=17.80927%

78% | 1112/1433 [1:25:39<25:13, 4.71s/it]

[Quota] LLM=8.912462 | GPU=0.0 | Used=8.912462/50.0 | Remaining=41.087538 | Usage=17.824924%

78% | 1113/1433 [1:25:44<25:05, 4.70s/it]

[Quota] LLM=8.920436 | GPU=0.0 | Used=8.920436/50.0 | Remaining=41.079564 | Usage=17.840872%

78% | 1114/1433 [1:25:48<24:48, 4.67s/it]

[Quota] LLM=8.928362 | GPU=0.0 | Used=8.928362/50.0 | Remaining=41.071638 | Usage=17.856724%

78% | 1115/1433 [1:25:52<23:46, 4.48s/it]

[Quota] LLM=8.936006 | GPU=0.0 | Used=8.936006/50.0 | Remaining=41.063994 | Usage=17.872012%

78% | 1116/1433 [1:25:56<22:27, 4.25s/it]

[Quota] LLM=8.943818 | GPU=0.0 | Used=8.943818/50.0 | Remaining=41.056182 | Usage=17.887636%

78% | 1117/1433 [1:26:01<23:14, 4.41s/it]

[Quota] LLM=8.951546 | GPU=0.0 | Used=8.951546/50.0 | Remaining=41.048454 | Usage=17.903092%

78% | 1118/1433 [1:26:05<22:04, 4.21s/it]

[Quota] LLM=8.959334 | GPU=0.0 | Used=8.959334/50.0 | Remaining=41.040666 | Usage=17.918668%

78% | 1119/1433 [1:26:09<21:57, 4.19s/it]

[Quota] LLM=8.966984 | GPU=0.0 | Used=8.966984/50.0 | Remaining=41.033016 | Usage=17.933968%

78% | 1120/1433 [1:26:13<21:53, 4.20s/it]

[Quota] LLM=8.974601 | GPU=0.0 | Used=8.974601/50.0 | Remaining=41.025399 | Usage=17.949202%

78% | 1121/1433 [1:26:17<21:59, 4.23s/it]

[Quota] LLM=8.982545 | GPU=0.0 | Used=8.982545/50.0 | Remaining=41.017455 | Usage=17.96509%

78% | 1122/1433 [1:26:22<22:02, 4.25s/it]

[Quota] LLM=8.990408 | GPU=0.0 | Used=8.990408/50.0 | Remaining=41.009592 | Usage=17.980816%

78% | 1123/1433 [1:26:26<21:32, 4.17s/it]

[Quota] LLM=8.998208 | GPU=0.0 | Used=8.998208/50.0 | Remaining=41.001792 | Usage=17.996416%

78% | 1124/1433 [1:26:30<21:31, 4.18s/it]

[Quota] LLM=9.006146 | GPU=0.0 | Used=9.006146/50.0 | Remaining=40.993854 | Usage=18.012292%

79% | 1125/1433 [1:26:35<22:41, 4.42s/it]

[Quota] LLM=9.014165 | GPU=0.0 | Used=9.014165/50.0 | Remaining=40.985835 | Usage=18.02833%

79% | 1126/1433 [1:26:38<21:29, 4.20s/it]

[Quota] LLM=9.022145 | GPU=0.0 | Used=9.022145/50.0 | Remaining=40.977855 | Usage=18.04429%

79%| 1127/1433 [1:26:42<20:35, 4.04s/it]

[Quota] LLM=9.02981 | GPU=0.0 | Used=9.02981/50.0 | Remaining=40.97019 | Usage=18.05962%

79%| 1128/1433 [1:26:47<22:11, 4.37s/it]

[Quota] LLM=9.037673 | GPU=0.0 | Used=9.037673/50.0 | Remaining=40.962327 | Usage=18.075346%

79%| 1129/1433 [1:26:51<21:37, 4.27s/it]

[Quota] LLM=9.045851 | GPU=0.0 | Used=9.045851/50.0 | Remaining=40.954149 | Usage=18.091702%

79%| 1130/1433 [1:26:56<21:50, 4.33s/it]

[Quota] LLM=9.054101 | GPU=0.0 | Used=9.054101/50.0 | Remaining=40.945899 | Usage=18.108202%

79%| 1131/1433 [1:27:01<22:43, 4.52s/it]

[Quota] LLM=9.062039 | GPU=0.0 | Used=9.062039/50.0 | Remaining=40.937961 | Usage=18.124078%

79%| 1132/1433 [1:27:06<24:17, 4.84s/it]

[Quota] LLM=9.070148 | GPU=0.0 | Used=9.070148/50.0 | Remaining=40.929852 | Usage=18.140296%

79%| 1133/1433 [1:27:12<25:06, 5.02s/it]

[Quota] LLM=9.07853 | GPU=0.0 | Used=9.07853/50.0 | Remaining=40.92147 | Usage=18.15706%

79%| 1134/1433 [1:27:16<23:32, 4.73s/it]

[Quota] LLM=9.08669 | GPU=0.0 | Used=9.08669/50.0 | Remaining=40.9133099999999996 | Usage=18.17338%

79%| 1135/1433 [1:27:22<25:40, 5.17s/it]

[Quota] LLM=9.095249 | GPU=0.0 | Used=9.095249/50.0 | Remaining=40.904751 | Usage=18.190498%

79%| 1136/1433 [1:27:27<25:39, 5.18s/it]

[Quota] LLM=9.10316 | GPU=0.0 | Used=9.10316/50.0 | Remaining=40.89684 | Usage=18.20632%

79%| 1137/1433 [1:27:31<24:09, 4.90s/it]

[Quota] LLM=9.111224 | GPU=0.0 | Used=9.111224/50.0 | Remaining=40.888776 | Usage=18.222448%

79%| 1138/1433 [1:27:37<25:23, 5.16s/it]

[Quota] LLM=9.119384 | GPU=0.0 | Used=9.119384/50.0 | Remaining=40.880616 | Usage=18.238768%

79% | 1139/1433 [1:27:42<24:01, 4.90s/it]

[Quota] LLM=9.127514 | GPU=0.0 | Used=9.127514/50.0 | Remaining=40.872486 | Usage=18.255028%

80% | 1140/1433 [1:27:48<25:35, 5.24s/it]

[Quota] LLM=9.135584 | GPU=0.0 | Used=9.135584/50.0 | Remaining=40.864416 | Usage=18.271168%

80% | 1141/1433 [1:27:52<24:46, 5.09s/it]

[Quota] LLM=9.14333 | GPU=0.0 | Used=9.14333/50.0 | Remaining=40.85667 | Usage=18.28666%

80% | 1142/1433 [1:27:58<25:55, 5.35s/it]

[Quota] LLM=9.151187 | GPU=0.0 | Used=9.151187/50.0 | Remaining=40.848813 | Usage=18.302374%

80% | 1143/1433 [1:28:03<24:31, 5.07s/it]

[Quota] LLM=9.159155 | GPU=0.0 | Used=9.159155/50.0 | Remaining=40.840845 | Usage=18.31831%

80% | 1144/1433 [1:28:07<23:01, 4.78s/it]

[Quota] LLM=9.166799 | GPU=0.0 | Used=9.166799/50.0 | Remaining=40.833201 | Usage=18.333598%

80% | 1145/1433 [1:28:15<27:29, 5.73s/it]

[Quota] LLM=9.174911 | GPU=0.0 | Used=9.174911/50.0 | Remaining=40.825089 | Usage=18.349822%

80% | 1146/1433 [1:28:19<25:47, 5.39s/it]

[Quota] LLM=9.182768 | GPU=0.0 | Used=9.182768/50.0 | Remaining=40.817232000000004 | Usage=18.365536%

80% | 1147/1433 [1:28:24<25:17, 5.30s/it]

[Quota] LLM=9.191162 | GPU=0.0 | Used=9.191162/50.0 | Remaining=40.808838 | Usage=18.382324%

80% | 1148/1433 [1:28:30<25:08, 5.29s/it]

[Quota] LLM=9.199442 | GPU=0.0 | Used=9.199442/50.0 | Remaining=40.800558 | Usage=18.398884%

80% | 1149/1433 [1:28:35<25:09, 5.32s/it]

[Quota] LLM=9.207287 | GPU=0.0 | Used=9.207287/50.0 | Remaining=40.792713 | Usage=18.414574%

80% | 1150/1433 [1:28:39<23:17, 4.94s/it]

[Quota] LLM=9.215246 | GPU=0.0 | Used=9.215246/50.0 | Remaining=40.784754 | Usage=18.430492%

80%| 1151/1433 [1:28:43<21:43, 4.62s/it]

[Quota] LLM=9.223178 | GPU=0.0 | Used=9.223178/50.0 | Remaining=40.776821999999996 | Usage=18.446356%

80%| 1152/1433 [1:28:47<21:05, 4.50s/it]

[Quota] LLM=9.231263 | GPU=0.0 | Used=9.231263/50.0 | Remaining=40.768737 | Usage=18.462526%

80%| 1153/1433 [1:28:51<20:06, 4.31s/it]

[Quota] LLM=9.239231 | GPU=0.0 | Used=9.239231/50.0 | Remaining=40.760768999999996 | Usage=18.478462%

81%| 1154/1433 [1:28:56<21:07, 4.54s/it]

[Quota] LLM=9.247127 | GPU=0.0 | Used=9.247127/50.0 | Remaining=40.752873 | Usage=18.494254%

81%| 1155/1433 [1:29:01<21:11, 4.57s/it]

[Quota] LLM=9.254966 | GPU=0.0 | Used=9.254966/50.0 | Remaining=40.745034000000004 | Usage=18.509932%

81%| 1156/1433 [1:29:05<20:09, 4.37s/it]

[Quota] LLM=9.262946 | GPU=0.0 | Used=9.262946/50.0 | Remaining=40.737054 | Usage=18.525892%

81%| 1157/1433 [1:29:09<20:22, 4.43s/it]

[Quota] LLM=9.270962 | GPU=0.0 | Used=9.270962/50.0 | Remaining=40.729038 | Usage=18.541924%

81%| 1158/1433 [1:29:13<19:54, 4.34s/it]

[Quota] LLM=9.279152 | GPU=0.0 | Used=9.279152/50.0 | Remaining=40.720848000000004 | Usage=18.558304%

81%| 1159/1433 [1:29:20<23:02, 5.04s/it]

[Quota] LLM=9.287249 | GPU=0.0 | Used=9.287249/50.0 | Remaining=40.712751 | Usage=18.574498%

81%| 1160/1433 [1:29:25<22:27, 4.93s/it]

[Quota] LLM=9.295145 | GPU=0.0 | Used=9.295145/50.0 | Remaining=40.704855 | Usage=18.59029%

81%| 1161/1433 [1:29:29<20:55, 4.61s/it]

[Quota] LLM=9.302867 | GPU=0.0 | Used=9.302867/50.0 | Remaining=40.697133 | Usage=18.605734%

81%| 1162/1433 [1:29:35<22:36, 5.01s/it]

```
[Quota] LLM=9.310883 | GPU=0.0 | Used=9.310883/50.0 |  
Remaining=40.689116999999996 | Usage=18.621766%  
81%|     | 1163/1433 [1:29:40<22:54,  5.09s/it]  
  
[Quota] LLM=9.318983 | GPU=0.0 | Used=9.318983/50.0 | Remaining=40.681017 |  
Usage=18.637966%  
81%|     | 1164/1433 [1:29:44<21:40,  4.83s/it]  
  
[Quota] LLM=9.327059 | GPU=0.0 | Used=9.327059/50.0 | Remaining=40.672941 |  
Usage=18.654118%  
81%|     | 1165/1433 [1:29:48<20:26,  4.58s/it]  
  
[Quota] LLM=9.335171 | GPU=0.0 | Used=9.335171/50.0 | Remaining=40.664829 |  
Usage=18.670342%  
81%|     | 1166/1433 [1:29:53<20:40,  4.65s/it]  
  
[Quota] LLM=9.343013000000001 | GPU=0.0 | Used=9.343013000000001/50.0 |  
Remaining=40.656987 | Usage=18.686026000000002%  
81%|     | 1167/1433 [1:29:57<19:48,  4.47s/it]  
  
[Quota] LLM=9.351149 | GPU=0.0 | Used=9.351149/50.0 | Remaining=40.648851 |  
Usage=18.702298%  
82%|     | 1168/1433 [1:30:03<21:33,  4.88s/it]  
  
[Quota] LLM=9.359138 | GPU=0.0 | Used=9.359138/50.0 | Remaining=40.640862 |  
Usage=18.718276%  
82%|     | 1169/1433 [1:30:08<21:29,  4.88s/it]  
  
[Quota] LLM=9.367202 | GPU=0.0 | Used=9.367202/50.0 | Remaining=40.632798 |  
Usage=18.734404%  
82%|     | 1170/1433 [1:30:13<21:41,  4.95s/it]  
  
[Quota] LLM=9.375509 | GPU=0.0 | Used=9.375509/50.0 | Remaining=40.624491 |  
Usage=18.751018%  
82%|     | 1171/1433 [1:30:18<21:23,  4.90s/it]  
  
[Quota] LLM=9.383672 | GPU=0.0 | Used=9.383672/50.0 |  
Remaining=40.61632799999996 | Usage=18.767344%  
82%|     | 1172/1433 [1:30:24<23:39,  5.44s/it]  
  
[Quota] LLM=9.391835 | GPU=0.0 | Used=9.391835/50.0 | Remaining=40.608165 |  
Usage=18.78367%  
82%|     | 1173/1433 [1:30:29<22:56,  5.30s/it]  
  
[Quota] LLM=9.400061 | GPU=0.0 | Used=9.400061/50.0 | Remaining=40.599939 |  
Usage=18.800122%  
82%|     | 1174/1433 [1:30:34<21:57,  5.09s/it]
```

[Quota] LLM=9.408026 | GPU=0.0 | Used=9.408026/50.0 | Remaining=40.591974 | Usage=18.816052%

82% | 1175/1433 [1:30:38<20:18, 4.72s/it]

[Quota] LLM=9.415805 | GPU=0.0 | Used=9.415805/50.0 | Remaining=40.584195 | Usage=18.83161%

82% | 1176/1433 [1:30:44<22:25, 5.24s/it]

[Quota] LLM=9.423788 | GPU=0.0 | Used=9.423788/50.0 | Remaining=40.576212 | Usage=18.847576%

82% | 1177/1433 [1:30:48<20:24, 4.78s/it]

[Quota] LLM=9.431705000000001 | GPU=0.0 | Used=9.431705000000001/50.0 | Remaining=40.568295 | Usage=18.863410000000002%

82% | 1178/1433 [1:30:51<18:49, 4.43s/it]

[Quota] LLM=9.439658 | GPU=0.0 | Used=9.439658/50.0 | Remaining=40.560342 | Usage=18.879316%

82% | 1179/1433 [1:30:56<19:28, 4.60s/it]

[Quota] LLM=9.447662 | GPU=0.0 | Used=9.447662/50.0 | Remaining=40.552338 | Usage=18.895324%

82% | 1180/1433 [1:31:00<18:27, 4.38s/it]

[Quota] LLM=9.455774 | GPU=0.0 | Used=9.455774/50.0 | Remaining=40.544226 | Usage=18.911548%

82% | 1181/1433 [1:31:04<18:01, 4.29s/it]

[Quota] LLM=9.463877 | GPU=0.0 | Used=9.463877/50.0 | Remaining=40.536123 | Usage=18.927754%

82% | 1182/1433 [1:31:10<20:08, 4.81s/it]

[Quota] LLM=9.472346 | GPU=0.0 | Used=9.472346/50.0 | Remaining=40.527654 | Usage=18.944692%

83% | 1183/1433 [1:31:15<19:48, 4.75s/it]

[Quota] LLM=9.48026 | GPU=0.0 | Used=9.48026/50.0 | Remaining=40.51974 | Usage=18.96052%

83% | 1184/1433 [1:31:20<20:24, 4.92s/it]

[Quota] LLM=9.488276 | GPU=0.0 | Used=9.488276/50.0 | Remaining=40.511724 | Usage=18.976552%

83% | 1185/1433 [1:31:25<20:12, 4.89s/it]

[Quota] LLM=9.49637 | GPU=0.0 | Used=9.49637/50.0 | Remaining=40.50363 | Usage=18.99274%

83% | 1186/1433 [1:31:30<19:52, 4.83s/it]

```
[Quota] LLM=9.50453 | GPU=0.0 | Used=9.50453/50.0 | Remaining=40.49547 |  
Usage=19.00906%  
83%|     | 1187/1433 [1:31:34<19:02,  4.64s/it]  
  
[Quota] LLM=9.512756 | GPU=0.0 | Used=9.512756/50.0 |  
Remaining=40.487244000000004 | Usage=19.025512%  
83%|     | 1188/1433 [1:31:39<19:19,  4.73s/it]  
  
[Quota] LLM=9.521015 | GPU=0.0 | Used=9.521015/50.0 | Remaining=40.478985 |  
Usage=19.04203%  
83%|     | 1189/1433 [1:31:43<18:52,  4.64s/it]  
  
[Quota] LLM=9.528854 | GPU=0.0 | Used=9.528854/50.0 | Remaining=40.471146 |  
Usage=19.057708%  
83%|     | 1190/1433 [1:31:48<18:22,  4.54s/it]  
  
[Quota] LLM=9.536675 | GPU=0.0 | Used=9.536675/50.0 | Remaining=40.463325 |  
Usage=19.07335%  
83%|     | 1191/1433 [1:31:52<18:31,  4.59s/it]  
  
[Quota] LLM=9.544694 | GPU=0.0 | Used=9.544694/50.0 | Remaining=40.455306 |  
Usage=19.089388%  
83%|     | 1192/1433 [1:31:57<18:00,  4.49s/it]  
  
[Quota] LLM=9.552665 | GPU=0.0 | Used=9.552665/50.0 | Remaining=40.447335 |  
Usage=19.10533%  
83%|     | 1193/1433 [1:32:01<17:43,  4.43s/it]  
  
[Quota] LLM=9.560621 | GPU=0.0 | Used=9.560621/50.0 | Remaining=40.439379 |  
Usage=19.121242%  
83%|     | 1194/1433 [1:32:04<16:19,  4.10s/it]  
  
[Quota] LLM=9.568514 | GPU=0.0 | Used=9.568514/50.0 | Remaining=40.431486 |  
Usage=19.137028%  
83%|     | 1195/1433 [1:32:08<15:56,  4.02s/it]  
  
[Quota] LLM=9.576296 | GPU=0.0 | Used=9.576296/50.0 | Remaining=40.423704 |  
Usage=19.152592%  
83%|     | 1196/1433 [1:32:13<16:29,  4.18s/it]  
  
[Quota] LLM=9.584177 | GPU=0.0 | Used=9.584177/50.0 | Remaining=40.415823 |  
Usage=19.168354%  
84%|     | 1197/1433 [1:32:18<17:16,  4.39s/it]  
  
[Quota] LLM=9.592031 | GPU=0.0 | Used=9.592031/50.0 | Remaining=40.407969 |  
Usage=19.184062%  
84%|     | 1198/1433 [1:32:21<16:17,  4.16s/it]
```

[Quota] LLM=9.599819 | GPU=0.0 | Used=9.599819/50.0 | Remaining=40.400181 | Usage=19.199638%

84%| 1199/1433 [1:32:26<16:50, 4.32s/it]

[Quota] LLM=9.607769 | GPU=0.0 | Used=9.607769/50.0 | Remaining=40.392231 | Usage=19.215538%

84%| 1200/1433 [1:32:31<17:11, 4.43s/it]

[Quota] LLM=9.615866 | GPU=0.0 | Used=9.615866/50.0 | Remaining=40.384134 | Usage=19.231732%

84%| 1201/1433 [1:32:37<19:03, 4.93s/it]

[Quota] LLM=9.624002 | GPU=0.0 | Used=9.624002/50.0 | Remaining=40.375997999999996 | Usage=19.248004%

84%| 1202/1433 [1:32:40<17:32, 4.56s/it]

[Quota] LLM=9.632132 | GPU=0.0 | Used=9.632132/50.0 | Remaining=40.367868 | Usage=19.264264%

84%| 1203/1433 [1:32:47<19:29, 5.08s/it]

[Quota] LLM=9.640163 | GPU=0.0 | Used=9.640163/50.0 | Remaining=40.359837 | Usage=19.280326%

84%| 1204/1433 [1:32:50<17:42, 4.64s/it]

[Quota] LLM=9.648038 | GPU=0.0 | Used=9.648038/50.0 | Remaining=40.351962 | Usage=19.296076%

84%| 1205/1433 [1:32:54<16:58, 4.47s/it]

[Quota] LLM=9.655952 | GPU=0.0 | Used=9.655952/50.0 | Remaining=40.344048 | Usage=19.311904%

84%| 1206/1433 [1:32:59<17:31, 4.63s/it]

[Quota] LLM=9.663983 | GPU=0.0 | Used=9.663983/50.0 | Remaining=40.336017 | Usage=19.327966%

84%| 1207/1433 [1:33:03<16:03, 4.26s/it]

[Quota] LLM=9.671696 | GPU=0.0 | Used=9.671696/50.0 | Remaining=40.328304 | Usage=19.343392%

84%| 1208/1433 [1:33:07<16:31, 4.41s/it]

[Quota] LLM=9.679853 | GPU=0.0 | Used=9.679853/50.0 | Remaining=40.320147 | Usage=19.359706%

84%| 1209/1433 [1:33:12<17:01, 4.56s/it]

[Quota] LLM=9.688241 | GPU=0.0 | Used=9.688241/50.0 | Remaining=40.311759 | Usage=19.376482%

84%| 1210/1433 [1:33:16<16:12, 4.36s/it]

```
[Quota] LLM=9.696017 | GPU=0.0 | Used=9.696017/50.0 | Remaining=40.303983 |  
Usage=19.392034%  
85%|      | 1211/1433 [1:33:22<17:57,  4.85s/it]  
  
[Quota] LLM=9.704249 | GPU=0.0 | Used=9.704249/50.0 |  
Remaining=40.295750999999996 | Usage=19.408498%  
85%|      | 1212/1433 [1:33:27<17:47,  4.83s/it]  
  
[Quota] LLM=9.712457 | GPU=0.0 | Used=9.712457/50.0 | Remaining=40.287543 |  
Usage=19.424914%  
85%|      | 1213/1433 [1:33:32<18:01,  4.92s/it]  
  
[Quota] LLM=9.720266 | GPU=0.0 | Used=9.720266/50.0 | Remaining=40.279734 |  
Usage=19.440532%  
85%|      | 1214/1433 [1:33:36<17:09,  4.70s/it]  
  
[Quota] LLM=9.728507 | GPU=0.0 | Used=9.728507/50.0 | Remaining=40.271493 |  
Usage=19.457014%  
85%|      | 1215/1433 [1:33:41<16:59,  4.68s/it]  
  
[Quota] LLM=9.736505 | GPU=0.0 | Used=9.736505/50.0 | Remaining=40.263495 |  
Usage=19.47301%  
85%|      | 1216/1433 [1:33:45<16:06,  4.45s/it]  
  
[Quota] LLM=9.744284 | GPU=0.0 | Used=9.744284/50.0 | Remaining=40.255716 |  
Usage=19.488568%  
85%|      | 1217/1433 [1:33:49<15:54,  4.42s/it]  
  
[Quota] LLM=9.752459 | GPU=0.0 | Used=9.752459/50.0 | Remaining=40.247541 |  
Usage=19.504918%  
85%|      | 1218/1433 [1:33:54<15:47,  4.41s/it]  
  
[Quota] LLM=9.760169 | GPU=0.0 | Used=9.760169/50.0 | Remaining=40.239831 |  
Usage=19.520338%  
85%|      | 1219/1433 [1:33:58<15:10,  4.26s/it]  
  
[Quota] LLM=9.767813 | GPU=0.0 | Used=9.767813/50.0 |  
Remaining=40.232186999999996 | Usage=19.535626%  
85%|      | 1220/1433 [1:34:01<14:41,  4.14s/it]  
  
[Quota] LLM=9.775601 | GPU=0.0 | Used=9.775601/50.0 | Remaining=40.224399 |  
Usage=19.551202%  
85%|      | 1221/1433 [1:34:06<15:00,  4.25s/it]  
  
[Quota] LLM=9.783623 | GPU=0.0 | Used=9.783623/50.0 | Remaining=40.216377 |  
Usage=19.567246%  
85%|      | 1222/1433 [1:34:11<15:42,  4.47s/it]
```

```
[Quota] LLM=9.791792 | GPU=0.0 | Used=9.791792/50.0 | Remaining=40.208208 |
Usage=19.583584%
85%|      | 1223/1433 [1:34:15<15:21,  4.39s/it]
[Quota] LLM=9.799655 | GPU=0.0 | Used=9.799655/50.0 | Remaining=40.200345 |
Usage=19.59931%
85%|      | 1224/1433 [1:34:21<17:05,  4.91s/it]
[Quota] LLM=9.807623 | GPU=0.0 | Used=9.807623/50.0 | Remaining=40.192377 |
Usage=19.615246%
85%|      | 1225/1433 [1:34:26<16:59,  4.90s/it]
[Quota] LLM=9.815603 | GPU=0.0 | Used=9.815603/50.0 |
Remaining=40.184397000000004 | Usage=19.631206%
86%|      | 1226/1433 [1:34:32<17:31,  5.08s/it]
[Quota] LLM=9.82358 | GPU=0.0 | Used=9.82358/50.0 | Remaining=40.17642 |
Usage=19.64716%
86%|      | 1227/1433 [1:34:38<18:20,  5.34s/it]
[Quota] LLM=9.831515 | GPU=0.0 | Used=9.831515/50.0 |
Remaining=40.168485000000004 | Usage=19.66303%
86%|      | 1228/1433 [1:34:42<17:04,  5.00s/it]
[Quota] LLM=9.839393 | GPU=0.0 | Used=9.839393/50.0 | Remaining=40.160607 |
Usage=19.678786%
86%|      | 1229/1433 [1:34:46<16:42,  4.91s/it]
[Quota] LLM=9.847421 | GPU=0.0 | Used=9.847421/50.0 | Remaining=40.152579 |
Usage=19.694842%
86%|      | 1230/1433 [1:34:51<16:31,  4.88s/it]
[Quota] LLM=9.855656 | GPU=0.0 | Used=9.855656/50.0 |
Remaining=40.144344000000004 | Usage=19.711312%
86%|      | 1231/1433 [1:34:55<15:17,  4.54s/it]
[Quota] LLM=9.863588 | GPU=0.0 | Used=9.863588/50.0 | Remaining=40.136412 |
Usage=19.727176%
86%|      | 1232/1433 [1:35:01<16:35,  4.95s/it]
[Quota] LLM=9.871619 | GPU=0.0 | Used=9.871619/50.0 | Remaining=40.128381 |
Usage=19.743238%
86%|      | 1233/1433 [1:35:05<15:26,  4.63s/it]
[Quota] LLM=9.879422 | GPU=0.0 | Used=9.879422/50.0 | Remaining=40.120578 |
Usage=19.758844%
86%|      | 1234/1433 [1:35:11<16:31,  4.98s/it]
```

[Quota] LLM=9.88736 | GPU=0.0 | Used=9.88736/50.0 | Remaining=40.11264 | Usage=19.77472%  
86%| 1235/1433 [1:35:16<16:24, 4.97s/it]

[Quota] LLM=9.895652 | GPU=0.0 | Used=9.895652/50.0 | Remaining=40.104348 | Usage=19.791304%  
86%| 1236/1433 [1:35:20<16:06, 4.91s/it]

[Quota] LLM=9.903812 | GPU=0.0 | Used=9.903812/50.0 | Remaining=40.096188 | Usage=19.807624%  
86%| 1237/1433 [1:35:27<17:31, 5.37s/it]

[Quota] LLM=9.911969000000001 | GPU=0.0 | Used=9.911969000000001/50.0 | Remaining=40.088031 | Usage=19.823938000000002%

86%| 1238/1433 [1:35:33<17:58, 5.53s/it]

[Quota] LLM=9.920072 | GPU=0.0 | Used=9.920072/50.0 | Remaining=40.079928 | Usage=19.840144%  
86%| 1239/1433 [1:35:37<16:44, 5.18s/it]

[Quota] LLM=9.928022 | GPU=0.0 | Used=9.928022/50.0 | Remaining=40.071978 | Usage=19.856044%  
87%| 1240/1433 [1:35:42<16:14, 5.05s/it]

[Quota] LLM=9.93599 | GPU=0.0 | Used=9.93599/50.0 | Remaining=40.064009999999996 | Usage=19.87198%

87%| 1241/1433 [1:35:46<14:51, 4.64s/it]

[Quota] LLM=9.944009 | GPU=0.0 | Used=9.944009/50.0 | Remaining=40.055991 | Usage=19.888018%  
87%| 1242/1433 [1:35:51<15:12, 4.78s/it]

[Quota] LLM=9.952034 | GPU=0.0 | Used=9.952034/50.0 | Remaining=40.047966 | Usage=19.904068%  
87%| 1243/1433 [1:35:55<14:30, 4.58s/it]

[Quota] LLM=9.959873 | GPU=0.0 | Used=9.959873/50.0 | Remaining=40.040127 | Usage=19.919746%  
87%| 1244/1433 [1:35:59<14:32, 4.62s/it]

[Quota] LLM=9.967979 | GPU=0.0 | Used=9.967979/50.0 | Remaining=40.032021 | Usage=19.935958%  
87%| 1245/1433 [1:36:04<14:26, 4.61s/it]

[Quota] LLM=9.976415 | GPU=0.0 | Used=9.976415/50.0 | Remaining=40.023585 | Usage=19.95283%  
87%| 1246/1433 [1:36:08<13:34, 4.36s/it]

[Quota] LLM=9.984071 | GPU=0.0 | Used=9.984071/50.0 | Remaining=40.015929 | Usage=19.968142%

87% | 1247/1433 [1:36:12<13:21, 4.31s/it]

[Quota] LLM=9.991982 | GPU=0.0 | Used=9.991982/50.0 | Remaining=40.008018 | Usage=19.983964%

87% | 1248/1433 [1:36:16<12:37, 4.10s/it]

[Quota] LLM=9.999938 | GPU=0.0 | Used=9.999938/50.0 | Remaining=40.000062 | Usage=19.999876%

87% | 1249/1433 [1:36:20<13:08, 4.28s/it]

[Quota] LLM=10.008074 | GPU=0.0 | Used=10.008074/50.0 | Remaining=39.991926 | Usage=20.016148%

87% | 1250/1433 [1:36:25<13:02, 4.28s/it]

[Quota] LLM=10.016057 | GPU=0.0 | Used=10.016057/50.0 | Remaining=39.983943 | Usage=20.032114%

87% | 1251/1433 [1:36:29<12:59, 4.28s/it]

[Quota] LLM=10.024031 | GPU=0.0 | Used=10.024031/50.0 | Remaining=39.975969 | Usage=20.048062%

87% | 1252/1433 [1:36:34<13:22, 4.44s/it]

[Quota] LLM=10.03196 | GPU=0.0 | Used=10.03196/50.0 | Remaining=39.96804 | Usage=20.06392%

87% | 1253/1433 [1:36:38<13:35, 4.53s/it]

[Quota] LLM=10.039883 | GPU=0.0 | Used=10.039883/50.0 | Remaining=39.960117 | Usage=20.079766%

88% | 1254/1433 [1:36:42<12:43, 4.27s/it]

[Quota] LLM=10.047545 | GPU=0.0 | Used=10.047545/50.0 | Remaining=39.952455 | Usage=20.09509%

88% | 1255/1433 [1:36:46<12:35, 4.25s/it]

[Quota] LLM=10.055459 | GPU=0.0 | Used=10.055459/50.0 | Remaining=39.944541 | Usage=20.110918%

88% | 1256/1433 [1:36:51<12:42, 4.31s/it]

[Quota] LLM=10.063013 | GPU=0.0 | Used=10.063013/50.0 | Remaining=39.936987 | Usage=20.126026%

88% | 1257/1433 [1:36:55<12:41, 4.32s/it]

[Quota] LLM=10.070723 | GPU=0.0 | Used=10.070723/50.0 | Remaining=39.929277 | Usage=20.141446%

88% | 1258/1433 [1:37:00<12:53, 4.42s/it]

[Quota] LLM=10.0784 | GPU=0.0 | Used=10.0784/50.0 | Remaining=39.9216 | Usage=20.1568%

88%| 1259/1433 [1:37:05<13:56, 4.81s/it]

[Quota] LLM=10.086551 | GPU=0.0 | Used=10.086551/50.0 | Remaining=39.913449 | Usage=20.173102%

88%| 1260/1433 [1:37:10<13:45, 4.77s/it]

[Quota] LLM=10.094339 | GPU=0.0 | Used=10.094339/50.0 | Remaining=39.905661 | Usage=20.188678%

88%| 1261/1433 [1:37:14<13:04, 4.56s/it]

[Quota] LLM=10.102421 | GPU=0.0 | Used=10.102421/50.0 | Remaining=39.897579 | Usage=20.204842%

88%| 1262/1433 [1:37:18<12:31, 4.39s/it]

[Quota] LLM=10.110143 | GPU=0.0 | Used=10.110143/50.0 | Remaining=39.889857 | Usage=20.220286%

88%| 1263/1433 [1:37:23<13:07, 4.63s/it]

[Quota] LLM=10.117952 | GPU=0.0 | Used=10.117952/50.0 | Remaining=39.882048 | Usage=20.235904%

88%| 1264/1433 [1:37:27<12:36, 4.47s/it]

[Quota] LLM=10.125812 | GPU=0.0 | Used=10.125812/50.0 | Remaining=39.87418800000004 | Usage=20.251624%

88%| 1265/1433 [1:37:32<12:43, 4.54s/it]

[Quota] LLM=10.133492 | GPU=0.0 | Used=10.133492/50.0 | Remaining=39.86650799999996 | Usage=20.266984%

88%| 1266/1433 [1:37:37<12:59, 4.67s/it]

[Quota] LLM=10.141229 | GPU=0.0 | Used=10.141229/50.0 | Remaining=39.85877100000004 | Usage=20.282458%

88%| 1267/1433 [1:37:41<12:31, 4.53s/it]

[Quota] LLM=10.149077 | GPU=0.0 | Used=10.149077/50.0 | Remaining=39.850923 | Usage=20.298154%

88%| 1268/1433 [1:37:45<12:04, 4.39s/it]

[Quota] LLM=10.156928 | GPU=0.0 | Used=10.156928/50.0 | Remaining=39.843072 | Usage=20.313856%

89%| 1269/1433 [1:37:51<12:38, 4.62s/it]

[Quota] LLM=10.16504 | GPU=0.0 | Used=10.16504/50.0 | Remaining=39.83496 | Usage=20.33008%

89%| 1270/1433 [1:37:56<13:27, 4.96s/it]

[Quota] LLM=10.172852 | GPU=0.0 | Used=10.172852/50.0 | Remaining=39.827148 | Usage=20.345704%

89%| 1271/1433 [1:38:01<12:47, 4.74s/it]

[Quota] LLM=10.180805 | GPU=0.0 | Used=10.180805/50.0 | Remaining=39.819195 | Usage=20.36161%

89%| 1272/1433 [1:38:05<12:15, 4.57s/it]

[Quota] LLM=10.188968000000001 | GPU=0.0 | Used=10.188968000000001/50.0 | Remaining=39.811032 | Usage=20.377936000000002%

89%| 1273/1433 [1:38:09<12:07, 4.55s/it]

[Quota] LLM=10.19693 | GPU=0.0 | Used=10.19693/50.0 | Remaining=39.80307 | Usage=20.39386%

89%| 1274/1433 [1:38:14<12:36, 4.76s/it]

[Quota] LLM=10.205018 | GPU=0.0 | Used=10.205018/50.0 | Remaining=39.794982 | Usage=20.410036%

89%| 1275/1433 [1:38:20<12:54, 4.90s/it]

[Quota] LLM=10.213007 | GPU=0.0 | Used=10.213007/50.0 | Remaining=39.786993 | Usage=20.426014%

89%| 1276/1433 [1:38:24<12:40, 4.84s/it]

[Quota] LLM=10.220804 | GPU=0.0 | Used=10.220804/50.0 | Remaining=39.779196 | Usage=20.441608%

89%| 1277/1433 [1:38:29<11:59, 4.62s/it]

[Quota] LLM=10.228736 | GPU=0.0 | Used=10.228736/50.0 | Remaining=39.771264 | Usage=20.457472%

89%| 1278/1433 [1:38:32<11:01, 4.27s/it]

[Quota] LLM=10.236425 | GPU=0.0 | Used=10.236425/50.0 | Remaining=39.763575 | Usage=20.47285%

89%| 1279/1433 [1:38:36<11:02, 4.30s/it]

[Quota] LLM=10.244225 | GPU=0.0 | Used=10.244225/50.0 | Remaining=39.755775 | Usage=20.48845%

89%| 1280/1433 [1:38:40<10:42, 4.20s/it]

[Quota] LLM=10.244225 | GPU=0.0 | Used=10.244225/50.0 | Remaining=39.755775 | Usage=20.48845%

89%| 1281/1433 [1:38:45<10:51, 4.29s/it]

[Quota] LLM=10.259942 | GPU=0.0 | Used=10.259942/50.0 | Remaining=39.740058 | Usage=20.519884%

89%| 1282/1433 [1:38:49<10:49, 4.30s/it]

```
[Quota] LLM=10.267754 | GPU=0.0 | Used=10.267754/50.0 | Remaining=39.732246 | Usage=20.535508%
90%|      | 1283/1433 [1:38:54<11:14,  4.49s/it]

[Quota] LLM=10.275671 | GPU=0.0 | Used=10.275671/50.0 | Remaining=39.724329 | Usage=20.551342%
90%|      | 1284/1433 [1:38:57<10:10,  4.10s/it]

[Quota] LLM=10.283498 | GPU=0.0 | Used=10.283498/50.0 | Remaining=39.716502 | Usage=20.566996%
90%|      | 1285/1433 [1:39:02<10:20,  4.19s/it]

[Quota] LLM=10.291331 | GPU=0.0 | Used=10.291331/50.0 | Remaining=39.708669 | Usage=20.582662%
90%|      | 1286/1433 [1:39:06<10:04,  4.11s/it]

[Quota] LLM=10.299032 | GPU=0.0 | Used=10.299032/50.0 | Remaining=39.700968 | Usage=20.598064%
90%|      | 1287/1433 [1:39:10<09:56,  4.08s/it]

[Quota] LLM=10.30682 | GPU=0.0 | Used=10.30682/50.0 | Remaining=39.69318 | Usage=20.61364%
90%|      | 1288/1433 [1:39:15<10:41,  4.43s/it]

[Quota] LLM=10.31465 | GPU=0.0 | Used=10.31465/50.0 | Remaining=39.68535 | Usage=20.6293%
90%|      | 1289/1433 [1:39:20<10:51,  4.53s/it]

[Quota] LLM=10.322411 | GPU=0.0 | Used=10.322411/50.0 | Remaining=39.677589 | Usage=20.644822%
90%|      | 1290/1433 [1:39:23<10:05,  4.24s/it]

[Quota] LLM=10.330154 | GPU=0.0 | Used=10.330154/50.0 | Remaining=39.669846 | Usage=20.660308%
90%|      | 1291/1433 [1:39:27<10:06,  4.27s/it]

[Quota] LLM=10.337924 | GPU=0.0 | Used=10.337924/50.0 | Remaining=39.662076 | Usage=20.675848%
90%|      | 1292/1433 [1:39:32<09:53,  4.21s/it]

[Quota] LLM=10.345739 | GPU=0.0 | Used=10.345739/50.0 | Remaining=39.654261 | Usage=20.691478%
90%|      | 1293/1433 [1:39:36<09:51,  4.22s/it]

[Quota] LLM=10.353512 | GPU=0.0 | Used=10.353512/50.0 | Remaining=39.646488 | Usage=20.707024%
90%|      | 1294/1433 [1:39:40<09:25,  4.07s/it]
```

```
[Quota] LLM=10.361357 | GPU=0.0 | Used=10.361357/50.0 | Remaining=39.638643 | Usage=20.722714%
90%|      | 1295/1433 [1:39:45<10:36,  4.61s/it]

[Quota] LLM=10.369163 | GPU=0.0 | Used=10.369163/50.0 | Remaining=39.630837 | Usage=20.738326%
90%|      | 1296/1433 [1:39:50<10:44,  4.70s/it]

[Quota] LLM=10.37699 | GPU=0.0 | Used=10.37699/50.0 | Remaining=39.62301 | Usage=20.75398%
91%|      | 1297/1433 [1:39:55<10:40,  4.71s/it]

[Quota] LLM=10.384967 | GPU=0.0 | Used=10.384967/50.0 | Remaining=39.615033 | Usage=20.769934%
91%|      | 1298/1433 [1:40:00<10:42,  4.76s/it]

[Quota] LLM=10.392899 | GPU=0.0 | Used=10.392899/50.0 | Remaining=39.607101 | Usage=20.785798%
91%|      | 1299/1433 [1:40:05<10:55,  4.89s/it]

[Quota] LLM=10.401098 | GPU=0.0 | Used=10.401098/50.0 | Remaining=39.598902 | Usage=20.802196%
91%|      | 1300/1433 [1:40:09<10:10,  4.59s/it]

[Quota] LLM=10.408943 | GPU=0.0 | Used=10.408943/50.0 | Remaining=39.591057 | Usage=20.817886%
91%|      | 1301/1433 [1:40:14<10:31,  4.78s/it]

[Quota] LLM=10.416725 | GPU=0.0 | Used=10.416725/50.0 | Remaining=39.583275 | Usage=20.83345%
91%|      | 1302/1433 [1:40:18<09:50,  4.51s/it]

[Quota] LLM=10.424552 | GPU=0.0 | Used=10.424552/50.0 | Remaining=39.575448 | Usage=20.849104%
91%|      | 1303/1433 [1:40:23<09:54,  4.57s/it]

[Quota] LLM=10.432319 | GPU=0.0 | Used=10.432319/50.0 | Remaining=39.567681 | Usage=20.864638%
91%|      | 1304/1433 [1:40:27<09:16,  4.32s/it]

[Quota] LLM=10.440095 | GPU=0.0 | Used=10.440095/50.0 | Remaining=39.559905 | Usage=20.88019%
91%|      | 1305/1433 [1:40:31<09:29,  4.45s/it]

[Quota] LLM=10.448015 | GPU=0.0 | Used=10.448015/50.0 | Remaining=39.551985 | Usage=20.89603%
91%|      | 1306/1433 [1:40:36<09:45,  4.61s/it]
```

```
[Quota] LLM=10.455908 | GPU=0.0 | Used=10.455908/50.0 | Remaining=39.544092 |
Usage=20.911816%
91% | 1307/1433 [1:40:41<09:33, 4.55s/it]

[Quota] LLM=10.463588 | GPU=0.0 | Used=10.463588/50.0 | Remaining=39.536412 |
Usage=20.927176%
91% | 1308/1433 [1:40:46<09:44, 4.67s/it]

[Quota] LLM=10.471478 | GPU=0.0 | Used=10.471478/50.0 | Remaining=39.528522 |
Usage=20.942956%
91% | 1309/1433 [1:40:50<09:39, 4.67s/it]

[Quota] LLM=10.479521 | GPU=0.0 | Used=10.479521/50.0 | Remaining=39.520479 |
Usage=20.959042%
91% | 1310/1433 [1:40:57<11:05, 5.41s/it]

[Quota] LLM=10.487354 | GPU=0.0 | Used=10.487354/50.0 |
Remaining=39.512646000000004 | Usage=20.974708%
91% | 1311/1433 [1:41:04<11:35, 5.70s/it]

[Quota] LLM=10.495412 | GPU=0.0 | Used=10.495412/50.0 | Remaining=39.504588 |
Usage=20.990824%
92% | 1312/1433 [1:41:08<10:50, 5.37s/it]

[Quota] LLM=10.503404 | GPU=0.0 | Used=10.503404/50.0 | Remaining=39.496596 |
Usage=21.006808%
92% | 1313/1433 [1:41:12<09:53, 4.95s/it]

[Quota] LLM=10.511297 | GPU=0.0 | Used=10.511297/50.0 | Remaining=39.488703 |
Usage=21.022594%
92% | 1314/1433 [1:41:20<11:20, 5.72s/it]

[Quota] LLM=10.519103 | GPU=0.0 | Used=10.519103/50.0 | Remaining=39.480897 |
Usage=21.038206%
92% | 1315/1433 [1:41:24<10:20, 5.26s/it]

[Quota] LLM=10.527077 | GPU=0.0 | Used=10.527077/50.0 | Remaining=39.472923 |
Usage=21.054154%
92% | 1316/1433 [1:41:28<09:29, 4.86s/it]

[Quota] LLM=10.53485 | GPU=0.0 | Used=10.53485/50.0 | Remaining=39.46515 |
Usage=21.0697%
92% | 1317/1433 [1:41:32<08:54, 4.61s/it]

[Quota] LLM=10.542683 | GPU=0.0 | Used=10.542683/50.0 | Remaining=39.457317 |
Usage=21.085366%
92% | 1318/1433 [1:41:36<08:31, 4.44s/it]
```

```
[Quota] LLM=10.550552 | GPU=0.0 | Used=10.550552/50.0 |  
Remaining=39.449448000000004 | Usage=21.101104%  
92%|      | 1319/1433 [1:41:41<08:28,  4.46s/it]  
  
[Quota] LLM=10.55837 | GPU=0.0 | Used=10.55837/50.0 | Remaining=39.44163 |  
Usage=21.11674%  
92%|      | 1320/1433 [1:41:45<08:34,  4.55s/it]  
  
[Quota] LLM=10.566236 | GPU=0.0 | Used=10.566236/50.0 | Remaining=39.433764 |  
Usage=21.132472%  
92%|      | 1321/1433 [1:41:50<08:18,  4.45s/it]  
  
[Quota] LLM=10.573991 | GPU=0.0 | Used=10.573991/50.0 | Remaining=39.426009 |  
Usage=21.147982%  
92%|      | 1322/1433 [1:41:54<08:27,  4.57s/it]  
  
[Quota] LLM=10.581827 | GPU=0.0 | Used=10.581827/50.0 | Remaining=39.418172999999996 | Usage=21.163654%  
92%|      | 1323/1433 [1:42:00<08:45,  4.78s/it]  
  
[Quota] LLM=10.58966 | GPU=0.0 | Used=10.58966/50.0 | Remaining=39.41034 |  
Usage=21.17932%  
92%|      | 1324/1433 [1:42:04<08:25,  4.63s/it]  
  
[Quota] LLM=10.59746 | GPU=0.0 | Used=10.59746/50.0 | Remaining=39.40254 |  
Usage=21.19492%  
92%|      | 1325/1433 [1:42:09<08:27,  4.70s/it]  
  
[Quota] LLM=10.605263 | GPU=0.0 | Used=10.605263/50.0 | Remaining=39.394737 |  
Usage=21.210526%  
93%|      | 1326/1433 [1:42:14<08:36,  4.82s/it]  
  
[Quota] LLM=10.613162 | GPU=0.0 | Used=10.613162/50.0 | Remaining=39.386838 |  
Usage=21.226324%  
93%|      | 1327/1433 [1:42:19<08:50,  5.00s/it]  
  
[Quota] LLM=10.621067 | GPU=0.0 | Used=10.621067/50.0 | Remaining=39.378933 |  
Usage=21.242134%  
93%|      | 1328/1433 [1:42:24<08:23,  4.80s/it]  
  
[Quota] LLM=10.628768 | GPU=0.0 | Used=10.628768/50.0 | Remaining=39.371232 |  
Usage=21.257536%  
93%|      | 1329/1433 [1:42:27<07:38,  4.41s/it]  
  
[Quota] LLM=10.636622000000001 | GPU=0.0 | Used=10.636622000000001/50.0 |  
Remaining=39.363378 | Usage=21.273244000000002%  
93%|      | 1330/1433 [1:42:31<07:16,  4.24s/it]
```

```
[Quota] LLM=10.644473 | GPU=0.0 | Used=10.644473/50.0 | Remaining=39.355527 |
Usage=21.288946%
93%|      | 1331/1433 [1:42:38<08:20,  4.90s/it]

[Quota] LLM=10.65236 | GPU=0.0 | Used=10.65236/50.0 | Remaining=39.34764 |
Usage=21.30472%
93%|      | 1332/1433 [1:42:42<07:55,  4.71s/it]

[Quota] LLM=10.660409 | GPU=0.0 | Used=10.660409/50.0 | Remaining=39.339591 |
Usage=21.320818%
93%|      | 1333/1433 [1:42:46<07:29,  4.49s/it]

[Quota] LLM=10.668182 | GPU=0.0 | Used=10.668182/50.0 | Remaining=39.331818 |
Usage=21.336364%
93%|      | 1334/1433 [1:42:50<07:29,  4.54s/it]

[Quota] LLM=10.676015 | GPU=0.0 | Used=10.676015/50.0 | Remaining=39.323985 |
Usage=21.35203%
93%|      | 1335/1433 [1:42:55<07:33,  4.63s/it]

[Quota] LLM=10.683602 | GPU=0.0 | Used=10.683602/50.0 | Remaining=39.316398 |
Usage=21.367204%
93%|      | 1336/1433 [1:42:59<07:05,  4.39s/it]

[Quota] LLM=10.691495 | GPU=0.0 | Used=10.691495/50.0 | Remaining=39.308505 |
Usage=21.38299%
93%|      | 1337/1433 [1:43:03<07:02,  4.40s/it]

[Quota] LLM=10.699463 | GPU=0.0 | Used=10.699463/50.0 | Remaining=39.300537 |
Usage=21.398926%
93%|      | 1338/1433 [1:43:08<06:47,  4.29s/it]

[Quota] LLM=10.707221 | GPU=0.0 | Used=10.707221/50.0 |
Remaining=39.292778999999996 | Usage=21.414442%
93%|      | 1339/1433 [1:43:12<06:56,  4.43s/it]

[Quota] LLM=10.715162 | GPU=0.0 | Used=10.715162/50.0 | Remaining=39.284838 |
Usage=21.430324%
94%|      | 1340/1433 [1:43:17<06:54,  4.45s/it]

[Quota] LLM=10.722977 | GPU=0.0 | Used=10.722977/50.0 | Remaining=39.277023 |
Usage=21.445954%
94%|      | 1341/1433 [1:43:21<06:29,  4.23s/it]

[Quota] LLM=10.730708 | GPU=0.0 | Used=10.730708/50.0 | Remaining=39.269292 |
Usage=21.461416%
94%|      | 1342/1433 [1:43:26<07:05,  4.67s/it]
```

[Quota] LLM=10.738922 | GPU=0.0 | Used=10.738922/50.0 | Remaining=39.261078 | Usage=21.477844%

94%| 1343/1433 [1:43:30<06:48, 4.54s/it]

[Quota] LLM=10.746707 | GPU=0.0 | Used=10.746707/50.0 | Remaining=39.253293 | Usage=21.493414%

94%| 1344/1433 [1:43:35<06:48, 4.59s/it]

[Quota] LLM=10.754855 | GPU=0.0 | Used=10.754855/50.0 | Remaining=39.245145 | Usage=21.50971%

94%| 1345/1433 [1:43:39<06:30, 4.44s/it]

[Quota] LLM=10.762895 | GPU=0.0 | Used=10.762895/50.0 | Remaining=39.237105 | Usage=21.52579%

94%| 1346/1433 [1:43:43<06:17, 4.34s/it]

[Quota] LLM=10.77089 | GPU=0.0 | Used=10.77089/50.0 | Remaining=39.22911 | Usage=21.54178%

94%| 1347/1433 [1:43:49<06:36, 4.61s/it]

[Quota] LLM=10.779065 | GPU=0.0 | Used=10.779065/50.0 | Remaining=39.220935 | Usage=21.55813%

94%| 1348/1433 [1:43:53<06:31, 4.60s/it]

[Quota] LLM=10.787135 | GPU=0.0 | Used=10.787135/50.0 | Remaining=39.212865 | Usage=21.57427%

94%| 1349/1433 [1:43:58<06:36, 4.72s/it]

[Quota] LLM=10.795184 | GPU=0.0 | Used=10.795184/50.0 | Remaining=39.204816 | Usage=21.590368%

94%| 1350/1433 [1:44:03<06:26, 4.66s/it]

[Quota] LLM=10.803176 | GPU=0.0 | Used=10.803176/50.0 | Remaining=39.196824 | Usage=21.606352%

94%| 1351/1433 [1:44:07<06:10, 4.52s/it]

[Quota] LLM=10.810901 | GPU=0.0 | Used=10.810901/50.0 | Remaining=39.189099 | Usage=21.621802%

94%| 1352/1433 [1:44:10<05:38, 4.18s/it]

[Quota] LLM=10.818944 | GPU=0.0 | Used=10.818944/50.0 | Remaining=39.181056 | Usage=21.637888%

94%| 1353/1433 [1:44:14<05:22, 4.03s/it]

[Quota] LLM=10.827008 | GPU=0.0 | Used=10.827008/50.0 | Remaining=39.172992 | Usage=21.654016%

94%| 1354/1433 [1:44:19<05:41, 4.33s/it]

```
[Quota] LLM=10.834991 | GPU=0.0 | Used=10.834991/50.0 | Remaining=39.165009 |
Usage=21.669982%  
95%|      | 1355/1433 [1:44:24<05:56,  4.57s/it]  
  
[Quota] LLM=10.842935 | GPU=0.0 | Used=10.842935/50.0 | Remaining=39.157065 |
Usage=21.68587%  
95%|      | 1356/1433 [1:44:28<05:45,  4.49s/it]  
  
[Quota] LLM=10.850657 | GPU=0.0 | Used=10.850657/50.0 | Remaining=39.149343 |
Usage=21.701314%  
95%|      | 1357/1433 [1:44:33<05:32,  4.38s/it]  
  
[Quota] LLM=10.858397 | GPU=0.0 | Used=10.858397/50.0 | Remaining=39.141603 |
Usage=21.716794%  
95%|      | 1358/1433 [1:44:37<05:30,  4.40s/it]  
  
[Quota] LLM=10.866212 | GPU=0.0 | Used=10.866212/50.0 | Remaining=39.133787999999996 |
Usage=21.732424%  
95%|      | 1359/1433 [1:44:41<05:28,  4.43s/it]  
  
[Quota] LLM=10.874492 | GPU=0.0 | Used=10.874492/50.0 | Remaining=39.125507999999996 |
Usage=21.748984%  
95%|      | 1360/1433 [1:44:46<05:31,  4.54s/it]  
  
[Quota] LLM=10.882688 | GPU=0.0 | Used=10.882688/50.0 | Remaining=39.117312 |
Usage=21.765376%  
95%|      | 1361/1433 [1:44:51<05:25,  4.52s/it]  
  
[Quota] LLM=10.890854000000001 | GPU=0.0 | Used=10.890854000000001/50.0 |
Remaining=39.109145999999996 | Usage=21.781708000000002%  
95%|      | 1362/1433 [1:44:55<05:22,  4.54s/it]  
  
[Quota] LLM=10.89926 | GPU=0.0 | Used=10.89926/50.0 | Remaining=39.10074 |
Usage=21.79852%  
95%|      | 1363/1433 [1:44:59<05:05,  4.36s/it]  
  
[Quota] LLM=10.907018 | GPU=0.0 | Used=10.907018/50.0 | Remaining=39.092982 |
Usage=21.814036%  
95%|      | 1364/1433 [1:45:03<04:56,  4.30s/it]  
  
[Quota] LLM=10.915127 | GPU=0.0 | Used=10.915127/50.0 | Remaining=39.084873 |
Usage=21.830254%  
95%|      | 1365/1433 [1:45:08<04:51,  4.29s/it]  
  
[Quota] LLM=10.922861 | GPU=0.0 | Used=10.922861/50.0 | Remaining=39.077139 |
Usage=21.845722%  
95%|      | 1366/1433 [1:45:13<04:59,  4.48s/it]
```

[Quota] LLM=10.931174 | GPU=0.0 | Used=10.931174/50.0 | Remaining=39.068826 | Usage=21.862348%

95%| 1367/1433 [1:45:18<05:15, 4.79s/it]

[Quota] LLM=10.939316 | GPU=0.0 | Used=10.939316/50.0 | Remaining=39.060684 | Usage=21.878632%

95%| 1368/1433 [1:45:23<05:07, 4.72s/it]

[Quota] LLM=10.94762 | GPU=0.0 | Used=10.94762/50.0 | Remaining=39.05238 | Usage=21.89524%

96%| 1369/1433 [1:45:28<05:09, 4.83s/it]

[Quota] LLM=10.956239 | GPU=0.0 | Used=10.956239/50.0 | Remaining=39.043761 | Usage=21.912478%

96%| 1370/1433 [1:45:32<05:00, 4.77s/it]

[Quota] LLM=10.964807 | GPU=0.0 | Used=10.964807/50.0 | Remaining=39.035193 | Usage=21.929614%

96%| 1371/1433 [1:45:37<04:55, 4.76s/it]

[Quota] LLM=10.973081 | GPU=0.0 | Used=10.973081/50.0 | Remaining=39.026919 | Usage=21.946162%

96%| 1372/1433 [1:45:42<04:54, 4.84s/it]

[Quota] LLM=10.981685 | GPU=0.0 | Used=10.981685/50.0 | Remaining=39.018315 | Usage=21.96337%

96%| 1373/1433 [1:45:48<05:02, 5.05s/it]

[Quota] LLM=10.989866 | GPU=0.0 | Used=10.989866/50.0 | Remaining=39.010134 | Usage=21.979732%

96%| 1374/1433 [1:45:52<04:49, 4.90s/it]

[Quota] LLM=10.998281 | GPU=0.0 | Used=10.998281/50.0 | Remaining=39.001719 | Usage=21.996562%

96%| 1375/1433 [1:45:57<04:47, 4.96s/it]

[Quota] LLM=11.006741 | GPU=0.0 | Used=11.006741/50.0 | Remaining=38.993259 | Usage=22.013482%

96%| 1376/1433 [1:46:02<04:29, 4.73s/it]

[Quota] LLM=11.014862 | GPU=0.0 | Used=11.014862/50.0 | Remaining=38.985138 | Usage=22.029724%

96%| 1377/1433 [1:46:07<04:40, 5.00s/it]

[Quota] LLM=11.02364 | GPU=0.0 | Used=11.02364/50.0 | Remaining=38.97636 | Usage=22.04728%

96%| 1378/1433 [1:46:12<04:26, 4.85s/it]

[Quota] LLM=11.031815 | GPU=0.0 | Used=11.031815/50.0 | Remaining=38.968185 | Usage=22.06363%

96% | 1379/1433 [1:46:16<04:06, 4.56s/it]

[Quota] LLM=11.040065 | GPU=0.0 | Used=11.040065/50.0 | Remaining=38.959935 | Usage=22.08013%

96% | 1380/1433 [1:46:21<04:07, 4.67s/it]

[Quota] LLM=11.048045 | GPU=0.0 | Used=11.048045/50.0 | Remaining=38.951955 | Usage=22.09609%

96% | 1381/1433 [1:46:25<04:03, 4.69s/it]

[Quota] LLM=11.056292000000001 | GPU=0.0 | Used=11.056292000000001/50.0 | Remaining=38.943708 | Usage=22.112584000000002%

96% | 1382/1433 [1:46:30<04:01, 4.74s/it]

[Quota] LLM=11.064512 | GPU=0.0 | Used=11.064512/50.0 | Remaining=38.935488 | Usage=22.129024%

97% | 1383/1433 [1:46:35<03:58, 4.78s/it]

[Quota] LLM=11.07278 | GPU=0.0 | Used=11.07278/50.0 | Remaining=38.92722 | Usage=22.14556%

97% | 1384/1433 [1:46:39<03:44, 4.57s/it]

[Quota] LLM=11.080505 | GPU=0.0 | Used=11.080505/50.0 | Remaining=38.919495 | Usage=22.16101%

97% | 1385/1433 [1:46:43<03:33, 4.45s/it]

[Quota] LLM=11.088449 | GPU=0.0 | Used=11.088449/50.0 | Remaining=38.911551 | Usage=22.176898%

97% | 1386/1433 [1:46:47<03:23, 4.32s/it]

[Quota] LLM=11.09618 | GPU=0.0 | Used=11.09618/50.0 | Remaining=38.903819999999996 | Usage=22.19236%

97% | 1387/1433 [1:46:52<03:19, 4.34s/it]

[Quota] LLM=11.104058 | GPU=0.0 | Used=11.104058/50.0 | Remaining=38.895942 | Usage=22.208116%

97% | 1388/1433 [1:46:56<03:12, 4.27s/it]

[Quota] LLM=11.111873 | GPU=0.0 | Used=11.111873/50.0 | Remaining=38.888127 | Usage=22.223746%

97% | 1389/1433 [1:47:02<03:28, 4.74s/it]

[Quota] LLM=11.119715 | GPU=0.0 | Used=11.119715/50.0 | Remaining=38.880285 | Usage=22.23943%

97% | 1390/1433 [1:47:06<03:22, 4.71s/it]

[Quota] LLM=11.127758 | GPU=0.0 | Used=11.127758/50.0 | Remaining=38.872242 | Usage=22.255516%

97%| 1391/1433 [1:47:11<03:20, 4.77s/it]

[Quota] LLM=11.135696 | GPU=0.0 | Used=11.135696/50.0 | Remaining=38.864304000000004 | Usage=22.271392%

97%| 1392/1433 [1:47:15<03:08, 4.60s/it]

[Quota] LLM=11.143562 | GPU=0.0 | Used=11.143562/50.0 | Remaining=38.856438 | Usage=22.287124%

97%| 1393/1433 [1:47:20<03:07, 4.68s/it]

[Quota] LLM=11.151503 | GPU=0.0 | Used=11.151503/50.0 | Remaining=38.848497 | Usage=22.303006%

97%| 1394/1433 [1:47:25<03:08, 4.84s/it]

[Quota] LLM=11.159282 | GPU=0.0 | Used=11.159282/50.0 | Remaining=38.840718 | Usage=22.318564%

97%| 1395/1433 [1:47:30<03:02, 4.81s/it]

[Quota] LLM=11.167436 | GPU=0.0 | Used=11.167436/50.0 | Remaining=38.832564 | Usage=22.334872%

97%| 1396/1433 [1:47:34<02:52, 4.66s/it]

[Quota] LLM=11.175608 | GPU=0.0 | Used=11.175608/50.0 | Remaining=38.824392 | Usage=22.351216%

97%| 1397/1433 [1:47:40<02:55, 4.89s/it]

[Quota] LLM=11.183873 | GPU=0.0 | Used=11.183873/50.0 | Remaining=38.816127 | Usage=22.367746%

98%| 1398/1433 [1:47:45<02:52, 4.94s/it]

[Quota] LLM=11.191871 | GPU=0.0 | Used=11.191871/50.0 | Remaining=38.808129 | Usage=22.383742%

98%| 1399/1433 [1:47:49<02:37, 4.63s/it]

[Quota] LLM=11.199869 | GPU=0.0 | Used=11.199869/50.0 | Remaining=38.800131 | Usage=22.399738%

98%| 1400/1433 [1:47:54<02:35, 4.71s/it]

[Quota] LLM=11.208077 | GPU=0.0 | Used=11.208077/50.0 | Remaining=38.791923 | Usage=22.416154%

98%| 1401/1433 [1:47:58<02:30, 4.70s/it]

[Quota] LLM=11.216312 | GPU=0.0 | Used=11.216312/50.0 | Remaining=38.783688 | Usage=22.432624%

98%| 1402/1433 [1:48:02<02:13, 4.32s/it]

[Quota] LLM=11.224058 | GPU=0.0 | Used=11.224058/50.0 | Remaining=38.775942 | Usage=22.448116%

98%| 1403/1433 [1:48:09<02:34, 5.16s/it]

[Quota] LLM=11.231888 | GPU=0.0 | Used=11.231888/50.0 | Remaining=38.768112 | Usage=22.463776%

98%| 1404/1433 [1:48:13<02:16, 4.71s/it]

[Quota] LLM=11.239772 | GPU=0.0 | Used=11.239772/50.0 | Remaining=38.760228 | Usage=22.479544%

98%| 1405/1433 [1:48:16<02:03, 4.43s/it]

[Quota] LLM=11.239772 | GPU=0.0 | Used=11.239772/50.0 | Remaining=38.760228 | Usage=22.479544%

98%| 1406/1433 [1:48:20<01:56, 4.31s/it]

[Quota] LLM=11.255687 | GPU=0.0 | Used=11.255687/50.0 | Remaining=38.744313 | Usage=22.511374%

98%| 1407/1433 [1:48:25<01:57, 4.51s/it]

[Quota] LLM=11.26385 | GPU=0.0 | Used=11.26385/50.0 | Remaining=38.73615 | Usage=22.5277%

98%| 1408/1433 [1:48:30<01:55, 4.63s/it]

[Quota] LLM=11.272256 | GPU=0.0 | Used=11.272256/50.0 | Remaining=38.727744 | Usage=22.544512%

98%| 1409/1433 [1:48:34<01:46, 4.42s/it]

[Quota] LLM=11.280305 | GPU=0.0 | Used=11.280305/50.0 | Remaining=38.719695 | Usage=22.56061%

98%| 1410/1433 [1:48:39<01:45, 4.57s/it]

[Quota] LLM=11.288387 | GPU=0.0 | Used=11.288387/50.0 | Remaining=38.711613 | Usage=22.576774%

98%| 1411/1433 [1:48:44<01:44, 4.75s/it]

[Quota] LLM=11.296589 | GPU=0.0 | Used=11.296589/50.0 | Remaining=38.703411 | Usage=22.593178%

99%| 1412/1433 [1:48:50<01:46, 5.05s/it]

[Quota] LLM=11.304791 | GPU=0.0 | Used=11.304791/50.0 | Remaining=38.695209 | Usage=22.609582%

99%| 1413/1433 [1:48:54<01:32, 4.64s/it]

[Quota] LLM=11.312906 | GPU=0.0 | Used=11.312906/50.0 | Remaining=38.687094 | Usage=22.625812%

99%| 1414/1433 [1:49:00<01:38, 5.19s/it]

```
[Quota] LLM=11.321276 | GPU=0.0 | Used=11.321276/50.0 | Remaining=38.678724 |
Usage=22.642552%
99%|      | 1415/1433 [1:49:06<01:37,  5.43s/it]

[Quota] LLM=11.329421 | GPU=0.0 | Used=11.329421/50.0 |
Remaining=38.670579000000004 | Usage=22.658842%
99%|      | 1416/1433 [1:49:10<01:25,  5.05s/it]

[Quota] LLM=11.337593 | GPU=0.0 | Used=11.337593/50.0 | Remaining=38.662407 |
Usage=22.675186%
99%|      | 1417/1433 [1:49:15<01:17,  4.84s/it]

[Quota] LLM=11.345507 | GPU=0.0 | Used=11.345507/50.0 | Remaining=38.654493 |
Usage=22.691014%
99%|      | 1418/1433 [1:49:19<01:11,  4.75s/it]

[Quota] LLM=11.353541 | GPU=0.0 | Used=11.353541/50.0 | Remaining=38.646459 |
Usage=22.707082%
99%|      | 1419/1433 [1:49:24<01:04,  4.59s/it]

[Quota] LLM=11.361725 | GPU=0.0 | Used=11.361725/50.0 | Remaining=38.638275 |
Usage=22.72345%
99%|      | 1420/1433 [1:49:28<01:00,  4.66s/it]

[Quota] LLM=11.369975 | GPU=0.0 | Used=11.369975/50.0 | Remaining=38.630025 |
Usage=22.73995%
99%|      | 1421/1433 [1:49:34<00:58,  4.86s/it]

[Quota] LLM=11.37803 | GPU=0.0 | Used=11.37803/50.0 | Remaining=38.62197 |
Usage=22.75606%
99%|      | 1422/1433 [1:49:39<00:53,  4.86s/it]

[Quota] LLM=11.386166 | GPU=0.0 | Used=11.386166/50.0 | Remaining=38.613834 |
Usage=22.772332%
99%|      | 1423/1433 [1:49:43<00:46,  4.66s/it]

[Quota] LLM=11.39447 | GPU=0.0 | Used=11.39447/50.0 | Remaining=38.60553 |
Usage=22.78894%
99%|      | 1424/1433 [1:49:47<00:40,  4.55s/it]

[Quota] LLM=11.402417 | GPU=0.0 | Used=11.402417/50.0 | Remaining=38.597583 |
Usage=22.804834%
99%|      | 1425/1433 [1:49:52<00:36,  4.55s/it]

[Quota] LLM=11.410769 | GPU=0.0 | Used=11.410769/50.0 | Remaining=38.589231 |
Usage=22.821538%
100%|     | 1426/1433 [1:49:57<00:33,  4.82s/it]
```

```

[Quota] LLM=11.419166 | GPU=0.0 | Used=11.419166/50.0 | Remaining=38.580833999999996 | Usage=22.838332%
100%|      | 1427/1433 [1:50:01<00:27,  4.51s/it]

[Quota] LLM=11.427122 | GPU=0.0 | Used=11.427122/50.0 | Remaining=38.572878 | Usage=22.854244%
100%|      | 1428/1433 [1:50:07<00:24,  4.97s/it]

[Quota] LLM=11.435267 | GPU=0.0 | Used=11.435267/50.0 | Remaining=38.564733000000004 | Usage=22.870534%
100%|      | 1429/1433 [1:50:12<00:20,  5.11s/it]

[Quota] LLM=11.443661 | GPU=0.0 | Used=11.443661/50.0 | Remaining=38.556339 | Usage=22.887322%
100%|      | 1430/1433 [1:50:17<00:14,  4.95s/it]

[Quota] LLM=11.451599 | GPU=0.0 | Used=11.451599/50.0 | Remaining=38.548401 | Usage=22.903198%
100%|      | 1431/1433 [1:50:21<00:09,  4.65s/it]

[Quota] LLM=11.459693 | GPU=0.0 | Used=11.459693/50.0 | Remaining=38.540307 | Usage=22.919386%
100%|      | 1432/1433 [1:50:25<00:04,  4.40s/it]

[Quota] LLM=11.467607 | GPU=0.0 | Used=11.467607/50.0 | Remaining=38.532393 | Usage=22.935214%
100%|      | 1433/1433 [1:50:29<00:00,  4.63s/it]

```

```
[ ]: # Build final augmented DataFrame
rubist_aug = pd.DataFrame(augmented_rows)
```

```

Final shape: (4299, 4)
label_level
stereotype    1433
neutral       1433
unrelated     1433
Name: count, dtype: int64

```

```
[ ]: # Save final file
rubist_aug.to_csv("COMP0173_Temp_Data/rubist_aug.csv", index=False)
```

## 4 References

- [1] Theo King, Zekun Wu, Adriano Koshiyama, Emre Kazim, and Philip Treleaven. 2024. HEARTS: A holistic framework for explainable, sustainable and robust text stereotype detection. arXiv preprint arXiv:2409.11579. Available at: <https://arxiv.org/abs/2409.11579> (Accessed: 4 December 2025). <https://doi.org/10.48550/arXiv.2409.11579>
- [2] Theo King, Zekun Wu, Adriano Koshiyama, Emre Kazim, and Philip Treleaven. 2024. HEARTS-Text-Stereotype-Detection (GitHub Repository). Available at: <https://github.com/holistic-ai/HEARTS-Text-Stereotype-Detection> (Accessed: 4 December 2025).
- [3] Theo King, Zekun Wu, Adriano Koshiyama, Emre Kazim, and Philip Treleaven. Holistic AI. 2024. EMGSD: Expanded Multi-Group Stereotype Dataset (HuggingFace Dataset). Available at: <https://huggingface.co/datasets/holistic-ai/EMGSD> (Accessed: 4 December 2025).
- [4] University College London Technical Support Group (TSG). 2025. GPU Access and Usage Documentation. Available at: <https://tsg.cs.ucl.ac.uk/gpus/> (Accessed: 6 December 2025).
- [5] United Nations. 2025. The 2030 Agenda for Sustainable Development. Available at: <https://sdgs.un.org/2030agenda> (Accessed: 6 December 2025).
- [6] Veronika Grigoreva, Anastasiia Ivanova, Ilseyar Alimova, and Ekaterina Artemova. 2024. RuBia: A Russian Language Bias Detection Dataset. Available at: <https://arxiv.org/abs/2403.17553> (Accessed: 9 December 2025).
- [7] Veronika Grigoreva, Anastasiia Ivanova, Ilseyar Alimova, and Ekaterina Artemova. 2024. RuBia-Dataset (GitHub Repository). Available at: <https://github.com/vergrig/RuBia-Dataset> (Accessed: 9 December 2025).
- [8] Sismetanin. 2020. Toxic Comments Detection in Russian (GitHub Repository). Available at: <https://github.com/sismetanin/toxic-comments-detection-in-russian> (Accessed: 9 December 2025).
- [9] DeepPavlov. 2019. RuBERT-base-cased (Hugging Face Model). Available at: <https://huggingface.co/DeepPavlov/rubert-base-cased> (Accessed: 9 December 2025).
- [10] AI-Forever. 2023. RuBERT-base (Hugging Face Model). Available at: <https://huggingface.co/ai-forever/ruBert-base> (Accessed: 9 December 2025).
- [11] Hugging Face. 2024. XLM-RoBERTa: Model Documentation. Available at: [https://huggingface.co/docs/transformers/en/model\\_doc/xlm-roberta](https://huggingface.co/docs/transformers/en/model_doc/xlm-roberta) (Accessed: 9 December 2025).
- [12] DeepPavlov. 2020. ruBERT-base-cased-sentence (Hugging Face Model). Available at: <https://huggingface.co/DeepPavlov/rubert-base-cased-sentence> (Accessed: 9 December 2025).

### 4.1 References: Manually Collected Russian Stereotype Sources

#### 4.1.1 Gender Stereotypes

- [1] AdMe Media. 2023. 11 . Available at: <https://adme.media/articles/11-stereotipov-o-muzhchinah-i-zhenschinah-kotorye-davno-ustareli-no-mnogie-s-nimi-tak-i-zhivu> (Accessed: 9 December 2025).

- [2] T—Journal. Gender stereotypes: cases and examples. Available at: <https://t-j.ru/gender-stereotypes-cases/> (Accessed: 9 December 2025).
- [3] Klinika Expert. . Available at: <https://klinikaexpert.ru/articles/v-yarlykah-lozhnye-stereotipy-o-muzhchinah-i-zhenschinah> (Accessed: 9 December 2025).
- [4] Zerkalo News. 2024. . Available at: <https://news.zerkalo.io/life/54154.html> (Accessed: 9 December 2025).
- [5] Siberian Branch of the Russian Academy of Sciences (SB RAS). . Available at: <https://www.sbras.info/articles/editors/gendernye-stereotipy-v-kotorye-pora-perestat-verit> (Accessed: 9 December 2025).
- [6] RBC Ukraine. 2024. . Available at: <https://www.rbc.ua/ukr/styler/rozpovsyudzheni-ta-zastarili-stereotipi-cholovikiv-1709482914.html> (Accessed: 9 December 2025).
- [7] Burning Hut. . Available at: <https://burninghut.ru/stereotipy-o-muzhchinakh-i-zhenshhinakh/> (Accessed: 9 December 2025).
- [8] Sravni.ru. 5 . Available at: <https://www.sravni.ru/text/5-stereotipov-o-muzhchinakh-i-zhenschinhakh-kotorye-plokhovlijajut-na-finansy-semi/> (Accessed: 9 December 2025).

## 4.2 References: Manually Collected Russian Stereotype Sources

### 4.2.1 Profession Stereotypes

- [1] 1 September / . Available at: <https://psy.1sept.ru/article.php?ID=200301712> (Accessed: 9 December 2025).
- [2] Peopletalk. 10 , . Available at: <https://peopletalk.ru/article/10-samyh-populyarnyh-stereotipov-o-professiyah-kotorye-besyat/> (Accessed: 9 December 2025).
- [3] Adukar. . Available at: <https://adukar.com/by/news/abiturientu/stereotipy-o-professiyah> (Accessed: 9 December 2025).
- [4] KemGMU. : . Available at: [https://kemgmu.ru/about\\_the\\_university/news/11672/](https://kemgmu.ru/about_the_university/news/11672/) (Accessed: 9 December 2025).
- [5] Dzen. . Available at: <https://dzen.ru/a/YZ810X7HcALiVr1k> (Accessed: 9 December 2025).
- [6] NAFI Analytics Center. IT- . Available at: <https://nafi.ru/en/analytics/samye-rasprostranennye-stereotipy-rossiyan-ob-it-professiyakh-/> (Accessed: 9 December 2025).
- [7] MosDigitals. : . Available at: <https://mosdigitals.ru/blog/pochemu-ne-lyubyat-yuristov-osnovnye-prichiny-i-stereotipy> (Accessed: 9 December 2025).

- [8] Medium (Juris Prudence). . Available at: <https://medium.com/juris-prudence/- - - - -cc3de98cdfec> (Accessed: 9 December 2025).
- [9] Tilda. . Available at: [https://stereotypes\\_actors.tilda.ws/](https://stereotypes_actors.tilda.ws/) (Accessed: 9 December 2025).
- [10] VK. . Available at: [https://m.vk.com/wall-181816199\\_1644](https://m.vk.com/wall-181816199_1644) (Accessed: 9 December 2025).
- [11] ECVDO (Kovrov). . Available at: <https://kovrov.ecvdo.ru/states/stereotipy-o-professii-ekonomista-chto-pravda-a-chto-vymysel> (Accessed: 9 December 2025).
- [12] ECVDO (Urzhum). . Available at: <https://urzhum.ecvdo.ru/states/razvenchivaem-mify-o-professii-finansista> (Accessed: 9 December 2025).
- [13] ECVDO (Azov). . Available at: <https://azov.ecvdo.ru/states/mify-o-populyarnyh-professiyah> (Accessed: 9 December 2025).
- [14] ECVDO (Igrim). . Available at: <https://igrim.ecvdo.ru/states/mify-i-pravda-o-rabote-v-sfere-obrazovaniya> (Accessed: 9 December 2025).
- [15] BroDude. 8 . Available at: <https://brodude.ru/8-stereotipov-o-rabote-shef-povara/> (Accessed: 9 December 2025).
- [16] Maxim Online. 7 . Available at: <https://www.maximonline.ru/lifestyle/7-glavnnykh-mifov-o-rabote-bortprovodnikov-id6443401/> (Accessed: 9 December 2025).
- [17] Sports.ru. . Available at: <https://www.sports.ru/football/blogs/477244.html> (Accessed: 9 December 2025).
- [18] DonNMU. . Available at: <https://dnmu.edu.ua/old/pro-meditsinu/3633-o-hirurgah> (Accessed: 9 December 2025).
- [19] VC.ru. . Available at: <https://vc.ru/id3158218/1127046-mify-o-pilotah-ty-tozhe-tak-dumal> (Accessed: 9 December 2025).
- [20] Psychologies. 5 . Available at: <https://www.psychologies.ru/wellbeing/5-strashnyih-mifov-o-detskom-balete-v-kotoryie-pora-perestat-verit/> (Accessed: 9 December 2025).
- [21] Masterok LiveJournal. . Available at: <https://masterok.livejournal.com/11781279.html> (Accessed: 9 December 2025).
- [22] AdMe Media. 10 , . Available at: <https://adme.media/articles/10-mifov-o-balete-kotorye-kinoshniki-pridumali-radi-vau-effekta-a-my-i-kupilis-2514646/> (Accessed: 9 December 2025).
- [23] Mir24. 8 . Available at: <https://mir24.tv/articles/16379683/8-glavnnyh-stereotipov-o-rabote-advokata-merkantilnye-ciniki-ili-professionaly> (Accessed: 9 December 2025).
- [24] ECVDO (Omsk). . Available at: [https://omsk.ecvdo.ru/states/mify-o-dizajnerskoj-professii-razbiraemsya-chто-pravda-a-chто-net](https://omsk.ecvdo.ru/states/mify-o-dizajnerskoj-professii-razbiraemsya-chto-pravda-a-chto-net) (Accessed: 9 December 2025).

- [25] Contented Media. 5 . Available at: <https://media.contented.ru/vdohnovenie/kofebrejk/5-mifov-o-dizaynerah/> (Accessed: 9 December 2025).
- [26] Ashley Home. . Available at: <https://ashleyhome.am/ru/blogs/news/steriotipy-o-dizainerax> (Accessed: 9 December 2025).
- [27] Lifehacker. 7 . Available at: <https://lifehacker.ru/7-stereotipov-o-rabote-barmena/> (Accessed: 9 December 2025).
- [28] LiveJournal (Klepachsv). . Available at: <https://klepachsv.livejournal.com/24127.html> (Accessed: 9 December 2025).
- [29] Championat. -9 . Available at: <https://www.championat.com/lifestyle/article-4793635-razoblachaem-top-9-stereotipov-o-sportsmenah-pravda-ili-vymysel.html> (Accessed: 9 December 2025).
- [30] Krasotuli. . Available at: <https://krasotuli.com/25199-stereotipy-o-parikmaherskom-dele-razvenchivaem-mify.html> (Accessed: 9 December 2025).
- [31] Maycenter. . Available at: <https://maycenter.ru/blog/mifi-o-professii-parikmakhera> (Accessed: 9 December 2025).
- [32] OK.ru. . Available at: <https://m.ok.ru/group/70000000389722/topic/157405964165210> (Accessed: 9 December 2025).
- [33] Nacasting. . Available at: <https://nacasting.ru/statii/mify-o-rezhisserakh> (Accessed: 9 December 2025).
- [34] ECVDO (Omsk). . Available at: <https://omsk.ecvdo.ru/states/mify-o-professii-arkhitektora-chto-na-samom-dele-vazhno-dlya-uspeha-v-etoj-sfere> (Accessed: 9 December 2025).
- [35] LookAtMe. . Available at: <http://www.lookatme.ru/flow/posts/fashion-radar/181029-mify-i-realii-o-modelnom-biznese> (Accessed: 9 December 2025).
- [36] Rylskova. . Available at: <https://rylskova.com/> / - / (Accessed: 9 December 2025).

### 4.3 References: Manually Collected Russian Stereotype Sources

#### 4.3.1 Nationality Stereotypes

- [1] Kuban24. 10 . Available at: <https://kuban24.tv/item/10-samyh-rasprostranennyh-stereotipov-o-raznyh-natsionalnostyah> (Accessed: 9 December 2025).
- [2] Mir24. 10 . Available at: <https://mir24.tv/articles/16626782/10-stereotipov-ob-indejcah-razoblachenie-mifov-i-udivitelnye-fakty> (Accessed: 9 December 2025).
- [3] Tandem. Russian stereotypes: fact or fiction. Available at: <https://tandem.net/ru/blog/russian-stereotypes-fact-fiction> (Accessed: 9 December 2025).

- [4] Tandem. British stereotypes: fact or myth. Available at: <https://tandem.net/ru/blog/british-stereotypes-fact-or-myth> (Accessed: 9 December 2025).
- [5] Meschool. -7 . Available at: <https://meschool.ru/poleznoe/top-7-stereotipov-o-britancakh/> (Accessed: 9 December 2025).
- [6] Linguacats. . Available at: <https://linguacats.com/ru/stati/o-chjom-molchat-frantsuzhenki> (Accessed: 9 December 2025).
- [7] Smapse LiveJournal. . Available at: <https://smapse.livejournal.com/699411.html> (Accessed: 9 December 2025).
- [8] Français Club. . Available at: <https://francaisclub.ru/> -- / (Accessed: 9 December 2025).
- [9] BestPrivateGuides. . Available at: <https://www.bestprivateguides.com/articles/stereotipi-ob-italyantsah-art-69.php> (Accessed: 9 December 2025).
- [10] Mamin Klub. : . Available at: <https://maminklub.lv/rebionok/stereotipy-ob-ispantsakh-pravda-i-lozh-623074/> (Accessed: 9 December 2025).
- [11] Pikabu. : . Available at: [https://pikabu.ru/story/stereotipyi-ob\\_ispantsakh\\_\\_pravda\\_i\\_vyimyisel\\_6306872](https://pikabu.ru/story/stereotipyi-ob_ispantsakh__pravda_i_vyimyisel_6306872) (Accessed: 9 December 2025).
- [12] Chetyre-Jelania LiveJournal. . Available at: <https://chetyre-jelania.livejournal.com/197804.html> (Accessed: 9 December 2025).
- [13] StayPoland. . Available at: <https://www.staypoland.com/ru/poland/stereotipy-o-polshe/> (Accessed: 9 December 2025).
- [14] ABEA Ukraine. -10 . Available at: <https://abea.com.ua/ru/top-10-pravdyvykh-stereotypov-ob-ukrayne-y-ukrayntsakh> (Accessed: 9 December 2025).
- [15] VancouverOK. 15 , . Available at: <https://vancouverok.com/15-stereotipov-o-kanadtsah-kotorye-yavlyayutsya-pravdoj/> (Accessed: 9 December 2025).
- [16] NashVancouver. 6 . Available at: <https://nashvancouver.com/6-lozhnih-stereotipov-o-kanadcax/> (Accessed: 9 December 2025).
- [17] AMI Visa. : . Available at: <https://amivisa.ru/blog/usa/stereotipy-ob-amerikancax-pravda-vymysel/> (Accessed: 9 December 2025).
- [18] English-Language.ru. . Available at: <https://www.english-language.ru/articles/informative/stereotipyi-ob-amerikanczax/> (Accessed: 9 December 2025).
- [19] Dzen. . Available at: <https://dzen.ru/a/ZfHtqQPJj3LVcVW> (Accessed: 9 December 2025).
- [20] Reddit (AskCentralAsia). Stereotypes of different Central Asian nations. Available at: [https://www.reddit.com/r/AskCentralAsia/comments/ahes8h/what\\_are\\_the\\_stereotypes\\_of\\_the\\_different\\_central/?tl=ru](https://www.reddit.com/r/AskCentralAsia/comments/ahes8h/what_are_the_stereotypes_of_the_different_central/?tl=ru) (Accessed: 9 December 2025).
- [21] Tsargrad. 5 . Available at: [https://am.tsargrad.tv/articles/5-stereotipov-pro-armjan-i-armeniju\\_395939](https://am.tsargrad.tv/articles/5-stereotipov-pro-armjan-i-armeniju_395939) (Accessed: 9 December 2025).
- [22] AdMe Media. : 10 , Available at: <https://adme.media/articles/>

ya-zhivu-v-irane-i-hochu-rasskazat-o-10-veschah-kotorye-otkroyut-etu-stranu-s-drugoj-storony-  
(Accessed: 9 December 2025).

[23] ChillTravel. . Available at: <https://chilltravel.ru/iindiastereotipi> (Accessed: 9 December 2025).

[24] ChaoChay. 9 . Available at: <https://www.chaochay.ru/blog/9-mifov-o-kitae-i-kitajcah> (Accessed: 9 December 2025).

[25] Smapse. 7 . Available at: <https://smapse.ru/7-banalnyh-stereotipov-o-zhitelyah-yuzhnoj-korei/> (Accessed: 9 December 2025).

[26] Lifehacker. . Available at: <https://lifehacker.ru/stereotipy-o-severnoi-koree/> (Accessed: 9 December 2025).

[27] Moya Planeta. : . Available at: [https://moya-planeta.ru/reports/view/yaponcy\\_lomka\\_stereotipov\\_35074](https://moya-planeta.ru/reports/view/yaponcy_lomka_stereotipov_35074) (Accessed: 9 December 2025).

[28] Smapse. 15 , . Available at: <https://smapse.ru/15-stereotipov-o-yaponcah-kotorye-oni-nenavidiyat/> (Accessed: 9 December 2025).

## 4.4 References: Manually Collected Russian Stereotype Sources

### 4.4.1 LGBTQ+ Stereotypes

[1] Denis Balin LiveJournal. . Available at: <https://denis-balin.livejournal.com/329915.html> (Accessed: 9 December 2025).

[2] Sojka.io. LGBT+: , . Available at: <https://sojka.io/ru/guides/lgbt> (Accessed: 9 December 2025).

[3] Raznoobrasije. (PDF). Available at: [http://raznoobrasije.org/wp-content/uploads/2020/07/2020\\_Raznoobrasije\\_1\\_Mythen-und-Fakten-über-LGB.pdf](http://raznoobrasije.org/wp-content/uploads/2020/07/2020_Raznoobrasije_1_Mythen-und-Fakten-über-LGB.pdf) (Accessed: 9 December 2025).

[4] Sphere Queer. Bisexual Week 2023: . Available at: <https://spherequeer.org/bisexual-week-2023/> (Accessed: 9 December 2025).

[5] GPress. . Available at: <https://gpress.info/2020/03/13/stereotipy-o-lgbt-1/> (Accessed: 9 December 2025).

[6] ParniPlus. . Available at: <https://parniplus.com/lgbt-movement/myths-about-bisexuality/> (Accessed: 9 December 2025).

[7] Kok.team. . Available at: <https://www.kok.team/ru/2018-04-26/stereotipy-o-lesbiyankah> (Accessed: 9 December 2025).

[8] Yvision. : . Available at: <https://yvision.kz/post/gei-i-lesbiyanki-mify-i-fakty-seksualnaya-patologiya-ili-estestvennyy-process-298823> (Accessed: 9 December 2025).

[9] WhatIsGood. - . Available at: <https://whatisgood.ru/theory/analytics/ulovki-lgbt-propagand/> (Accessed: 9 December 2025).

[10] Holod Media. - . Available at: <https://holod.media/2023/05/15/myths-about-trans-people/> (Accessed: 9 December 2025).

[11] VK OvsYanArt. - : . Available at: <https://vk.com/@ovsyanart-trans-people> (Accessed: 9 December 2025).

[12] RostovGazeta. . Available at: <https://rostovgazeta.ru/news/2017-02-17/samye-rasprostranennye-mify-o-transgenderah-1353439> (Accessed: 9 December 2025).